cONCEPTuAL

The **N**etwork **C**orrectness and **P**erformance **T**esting **L**anguage

Scott Pakin, pakin@lanl.gov

This document describes coNCePTuaL version 1.5.1b.

Table of Contents

1 Introduction

This document presents a simple, special-purpose language called coNCePTuaL. coN-CePTuaL is intended for rapidly generating programs that measure the performance and/or test the correctness of networks and network protocol layers. A few lines of coN-CePTuaL code can produce programs that would take significantly more effort to write in a conventional programming language.

coNCePTuaL is not merely a language specification. The coNCePTuaL toolset includes a compiler, run-time library, and associated utility programs that enable users to analyze network behavior quickly, conveniently, and accurately.

1.1 Motivation

A frequently reinvented wheel among network researchers is a suite of programs that test a network's performance. A problem with having umpteen versions of performance tests is that it leads to a variety in the way results are reported; colloquially, apples are often compared to oranges. Consider a bandwidth test. Does a bandwidth test run for a fixed number of iterations or a fixed length of time? Is bandwidth measured as ping-pong bandwidth (i.e., $2 \times$ message length $\div$ round-trip time) or unidirectional throughput (N messages in one direction followed by a single acknowledgement message)? Is the acknowledgement message of minimal length or as long as the entire message? Does its length contribute to the total bandwidth? Is data sent unidirectionally or in both directions at once? How many warmup messages (if any) are sent before the timing loop? Is there a delay after the warmup messages (to give the network a chance to reclaim any scarce resources)? Are receives nonblocking (possibly allowing overlap in the NIC) or blocking?

The motivation behind creating coNCePTuaL, a simple specification language designed for describing network benchmarks, is that it enables a benchmark to be described sufficiently tersely as to fit easily in a report or research paper, facilitating peer review of the experimental setup and timing measurements. Because coNCePTuaL code is simple to write, network tests can be developed and deployed with low turnaround times—useful when the results of one test suggest a following test that should be written. Because coN-CePTuaL is special-purpose its run-time system can perform the following functions, which benchmark writers often neglect to implement:

- logging information about the environment under which the benchmark ran: operating system, CPU architecture and clock speed, timer type and resolution, etc.
- aborting a program if it takes longer than a predetermined length of time to complete
- writing measurement data and descriptive statistics to a variety of output formats, including the input formats of various graph-plotting programs

coNCePTuaL is not limited to network peformance tests, however. It can also be used for network verification. That is, coNCePTuaL programs can be used to locate failed links or to determine the frequency of bit errors—even those that may sneak past the network's CRC hardware.

In addition, because coNCePTuaL is a very high-level language, the coNCePTuaL compiler's backend has a great deal of potential. It would be possible for the backend to produce a variety of target formats such as Fortran + MPI, Perl + sockets, C + a network

vendor's low-level messaging layer, and so forth. It could directly manipulate a network simulator. It could feed into a graphics program to produce a space-time diagram of a CONCEPTUAL program. The possibilities are endless.

1.2 Limitations

Although CONCEPTUAL can express a wide variety of race-free communication patterns it cannot currently express data-dependent communication. For example, CONCEPTUAL canot express a master-worker pattern in which a master task sends a message to a worker task as a reaction to that particular worker's sending of a message to the master. Such a communication pattern is not independent of the order in which messages happen to arrive from the workers. Similarly, CONCEPTUAL cannot use run-time performance data to guide its operations. It is therefore not currently possible to express a communication benchmark that repeats until the standard error of some performance metric drops below a given threshold. These limitations may be lifted in a future release of the system.

1.3 Typesetting conventions

The following table showcases the typesetting conventions used in this manual to attribute various meanings to text. Note that not all of the conventions are typographically distinct.

`-a` `--abcdef`	command-line options (e.g., `-C` or `--help`)
`ABCDEF`	environment variables (e.g., `PATH`)
`<abcdef>`	nonterminals in the CONCEPTUAL grammar (e.g., `<ident>`)
`abcdef`	commands to enter on the keyboard (e.g., `make install`)
`abcdef`	file and directory names (e.g., `conceptual.pdf`)
`ABCDEF`	CONCEPTUAL keywords (e.g., `RECEIVE`)
`abcdef`	variables, constants, functions, and types in any language (e.g., `bit_errors` or `gettimeofday()`)
abcdef	metasyntactic variables and formal function parameters (e.g., *fan-out*)
`'abcdef'`	snippets of code, command lines, files, etc. (e.g., '`10 MOD 3`')

2 Installation

cONCEPTuaL uses the GNU Autotools (Autoconf, Automake, and Libtool) to increase portability, to automate compilation, and to facilitate installation. As of this writing, cONCEPTuaL has passed *make check* (see) on the following platforms:

Architecture	OS	Compiler	
IA-32	Linux	gcc	(GNU)
		icc	(Intel)
		opencc	(Open64)
	FreeBSD	gcc	(GNU)
	OpenBSD	gcc	(GNU)
	NetBSD	gcc	(GNU)
	Solaris	gcc	(GNU)
		cc	(Sun)
	Syllable	gcc	(GNU)
	Windows	gcc	(GNU)
	(via Cygwin)		
x86-64	Linux	gcc	(GNU)
		pgcc	(PGI)
		pathcc	(PathScale)
		llvm-gcc	(LLVM)
		clang	(LLVM)
	Catamount	gcc	(GNU)
		pgcc	(PGI)
IA-64	Linux	gcc	(GNU)
		ecc	(Intel)
PowerPC	Linux	gcc	(GNU)
		xlc	(IBM)
	AIX	gcc	(GNU)
		xlc	(IBM)
	MacOS X	gcc	(GNU)
	BLRTS	xlc	(IBM)
Cell (Power)	Linux	gcc	(GNU)
Cray X1	UNICOS/mp	cc	(Cray)
UltraSPARC	Solaris	gcc	(GNU)
		cc	(Sun)
MIPS	IRIX	gcc	(GNU)

		cc	(MIPSpro)
Alpha	Linux	gcc	(GNU)
		ccc	(HP)
	Tru64	gcc	(GNU)
		cc	(HP)
ARM	Linux	gcc	(GNU)

In its simplest form, cONCEpTuAL installation works by executing the following commands at the operating-system prompt:

```
./configure
make
make install
```

(`configure` is normally run as *./configure* to force it to run from the current directory on the assumption that `.` is not in the executable search path.) We now describe those three installation steps in detail, listing a variety of customization options for each step.

2.1 *configure*

`configure` is a Bourne-shell script that analyzes your system's capabilities (compiler features, library and header-file availability, function and datatype availability, linker flags for various options, etc.) and custom-generates a `Makefile` and miscellaneous other files. `configure` accepts a variety of command-line options. *./configure --help* lists all of the options. The following are some of the more useful ones:

--disable-shared

> cONCEpTuAL normally installs both static and dynamic libraries. While dynamic libraries have a number of advantages they do need to be installed on all nodes that run the compiled cONCEpTuAL programs. If global installation is not convenient/feasible, *--disable-shared* can be used to force static linking of executables. Note, however, that `libncptlmodule.so`, the Python interface to the cONCEpTuAL run-time library, needs to be built as a shared object so that it can be loaded dynamically into a running Python interpreter. *--disable-shared* inhibits the compilation and installation of `libncptlmodule.so`.

--prefix=directory

> *make install* normally installs cONCEpTuAL into the `/usr/local` directory. The *--prefix* option instructs `configure` to write a `Makefile` with a different installation directory. For example, *--prefix=/local/encap/conceptual-1.5.1b* will cause cONCEpTuAL's files to be installed in `/local/encap/conceptual-1.5.1b/bin`, `/local/encap/conceptual-1.5.1b/include`, etc.

`--with-ignored-libs=lib1,lib2,...`

> In some circumstances it may be necessary to prevent CoNCePTuaL from using certain libraries even when `./configure` detects them and believes them to be usable. The `--with-ignored-libs` configuration option forces `./configure` to ignore one or more specified libraries. Only the base name of each library should be used; omit directory names, the 'lib' prefix (on Unix-like systems), and the file suffix. For example, to disable the use of `/usr/local/lib/libpapi.a` you should specify `--with-ignored-libs=papi`.

`--without-fork`

> `./configure` detects automatically if your system provides a working `fork()` function. However, it cannot detect if `fork()` correctly spawns a child process but corrupts the parent's memory map while doing so, as is the case when using some InfiniBand software stacks. The `--without-fork` option inhibits the use of `fork()` and well as functions that implicitly invoke `fork()` such as `system()` and `popen()`.

`--with-gettimeofday`

> The CoNCePTuaL run-time library is able to use any of a variety of platform-specific microsecond timers to take timing measurements. (See Section 6.3.5 [Time-related functions], page 170, for a complete list.) The `--with-gettimeofday` option forces the run-time library to utilize instead the generic C `gettimeofday()` function. This can be useful in the rare, but not impossible, case that a quirk in some particular platform misleads one of CoNCePTuaL's other timers. The `validatetimer` utility (see [Validating the coNCePTuaL timer], page 9) can help determine whether `--with-gettimeofday` is necessary.

`--with-mpi-wtime`

> On some systems the most accurate timer available is provided by the `MPI_Wtime()` function in the MPI library. The `--with-mpi-wtime` option forces the run-time library to measure elapsed time using `MPI_Wtime()` instead of any of the other available timers. (See Section 6.3.5 [Time-related functions], page 170, for a complete list). The ramifications of `--with-mpi-wtime` are threefold:
>
> 1. The option requires that you link all CoNCePTuaL programs against an MPI library and run them like any other MPI program. (You may need to set *CPPFLAGS*, *LIBS*, *LDFLAGS*, or some of the other command-line variables described below.)
>
> 2. `MPI_Wtime()` may be *less* accurate than some of the other timers available to CoNCePTuaL. In many MPI implementations, `MPI_Wtime()` simply invokes `gettimeofday()`, for instance.
>
> 3. Although this is a rare problem, it may not be safe to invoke `MPI_Wtime()` without first invoking `MPI_Init()`. Fortunately, proper juxtaposition of the two functions is not a concern for the CoNCePTuaL C+MPI backend (see Section 3.3.2 [The c_mpi backend], page 22), which ensures that `MPI_Init()` is invoked before `MPI_Wtime()`.

In short, you should specify `--with-mpi-wtime` only if you have good reason to believe that `MPI_Wtime()` is likely to produce the most accurate timing measurements on your system.

`CC=C compiler`

        `configure` automatically searches for a C compiler to use. To override its selection, assign a value to `CC` on the command line. For example, `./configure CC=ecc` will cause CONCEPTUAL to be built with `ecc`.

`CFLAGS=C compiler flags`
`LDFLAGS=linker flags`
`CPPFLAGS=C preprocessor flags`
`LIBS=extra libraries`

        Like `CC`, these variables override the values determined automatically by `configure`. As an illustration, `./configure CPPFLAGS="-DSPECIAL -I/home/pakin/include/special -I." CFLAGS="-O3 -g -Wall -W" LDFLAGS=--static LIBS="-lz /usr/lib/libstuff.a"` assigns values to all four variables.

`MPICC=C compiler`
`MPICPPFLAGS=C preprocessor flags`
`MPILDFLAGS=extra linker flags`
`MPILIBS=extra libraries`

        These variables are analagous to `CC`, `CPPFLAGS`, `LDFLAGS`, and `LIBS`, respectively. The difference is that they are not used to build the CONCEPTUAL run-time library but rather to build user programs targeted to the C+MPI compiler backend. For example, if your MPI installation lacks an `mpicc` script, you may need to specify extra header files and libraries explicitly: `./configure MPICPPFLAGS="-I/usr/lib/mpi/include" MPILIBS="-lmpich"`.

As a rather complex illustration of how some of the preceding options (as well as a few mentioned by `./configure --help`) might be combined, the following is how CON-CEPTUAL was once configured to cross-compile from a Linux/PowerPC build machine to a prototype of the BlueGene/L supercomputer (containing, at the time, 2048 embedded PowerPC processors, each executing a minimal run-time system, BLRTS). IBM's `xlc` compiler was accessed via a wrapper script called `mpcc`.

    `./configure CFLAGS="-g -O -qmaxmem=64000" CC=/bgl/local/bin/mpcc`
    `CPP="gcc -E" --host=powerpc-ibm-linux-gnu --build=powerpc-unknown-linux-gnu`
    `--with-alignment=8 --with-gettimeofday --prefix=/bgl/bgguest/LANL/ncptl`
    `MPICC=/bgl/local/bin/mpcc CPPFLAGS=-I/BlueLight/floor/bglsys/include`

It's always best to specify environment variables as arguments to `./configure` because the `configure` script writes its entire command line as a comment to `config.log` and as a shell command to `config.status` to enable re-running `./configure` with exactly the same parameters.

When `./configure` finishes running it outputs a list of the warning messages that were issued during the run. If no warnings were issued, `./configure` will output 'Configuration completed without any errors or warnings.'. Warnings are also written to `config.log`

and can therefore be redisplayed at any time by executing a shell command such as `grep WARNING config.log`.

2.2 `make`

Running `make` by itself will compile the CONCEPTUAL run-time library. However, the `Makefile` generated by `configure` can perform a variety of other actions, as well:

`make check`

> Perform a series of regression tests on the CONCEPTUAL run-time library. This is a good thing to do after a `make` to ensure that the run-time library built properly on your system. When `make check` finishes it summarizes the test results. The following output signifies a successful completion of `make check`:

```
===================
All 21 tests passed
===================
```

> The total number of tests performed depends upon the way that CONCEPTUAL was configured. CONCEPTUAL components that could not be built are not tested.

> If any tests behave unexpectedly it may be possible to gain more information about the source of the problem by examining the corresponding `.log` file left behind in the `tests` directory.

> Tests can also be run individually:

```
cd tests
./runtime_random
```

`make clean`
`make distclean`
`make maintainer-clean`

> `make clean` deletes all files generated by a preceding `make` command. `make distclean` deletes all files generated by a preceding `./configure` command. `make maintainer-clean` delete all generated files. Run `make maintainer-clean` only if you have not-too-ancient versions of the GNU Autotools (Autoconf 2.53, Automake 1.6, and Libtool 1.4) because those are needed to regenerate some of the generated files. The sequence of operations to regenerate all of the configuration files needed by CONCEPTUAL is shown below.

```
libtoolize --force --copy
aclocal
autoheader
automake --add-missing --copy
autoconf
```

make install

> Install CONCEPTUAL, including the compiler, run-time library, header files, and tools. *make install* is described in detail in Section 2.3 [make install], page 10.

make uninstall

> Remove all of the files that *make install* installed. Most of the top-level directories are retained, however, as `make` cannot guarantee that these are not needed by other applications.

make info

make pdf

make docbook

> Produce the CONCEPTUAL user's guide (this document) in, respectively, Emacs info format, PDF format, or DocBook format. The resulting documentation (`conceptual.info*`, `conceptual.pdf`, or `conceptual.xml`) is created in the `doc` subdirectory.

make ncptl-logextract.html

> CONCEPTUAL comes with a postprocessor called `ncptl-logextract` that facilitates extracting information from CONCEPTUAL-produced log files. The complete `ncptl-logextract` documentation is presented in Section 3.5.2 [ncptl-logextract], page 53. As is readily apparent from that documentation, `ncptl-logextract` supports an overwhelming number of command-line options. To make the `ncptl-logextract` documentation more approachable, the *make ncptl-logextract.html* command creates a dynamic HTML version of it (and stores in the `doc` subdirectory). The result, `ncptl-logextract.html`, initially presents only the top level of the `ncptl-logextract` option hierarchy. Users can then click on the name of a command-line option to expand or contract the list of subobtions. This interactive behavior makes it easy for a user to get more information on some options without being distracted by the documentation for the others.

make empty.log

> Create an empty log file called `empty.log` that contains a complete prologue and epilogue but no data. This is convenient for validating that the CONCEPTUAL run-time library was built using your preferred build options.

make stylesheets

> CONCEPTUAL can automatically produce stylesheets for a variety of programs. These stylesheets make keywords, comments, strings, and other terms in the language visually distinct from each other for a more aesthetically appealing appearance. Currently, *make stylesheets* produces a LaTeX 2_ε package (`ncptl.sty`), an a2ps style sheet (`ncptl.ssh`), an Emacs major mode (`ncptl-mode.el` and `ncptl-mode.elc`), a Vim syntax file (`ncptl.vim`), a Source-highlight language definition (`ncptl.lang`), and a GeSHi language file (`ncptl.php`). Each of these can be built individually if desired. (For example, *make ncptl-mode.vim* will create only `ncptl-mode.vim`.) Note that the `Makefile` currently lacks provisions for installing these files so

whichever stylesheets are desired will need to be installed manually. Stylesheet installation is detailed in [Installing stylesheets], page 10.

`make modulefile`

> The Environment Modules package facilitates configuring the operating-system shell for a given application. The `make modulefile` command creates a `conceptual_1.5.1b` modulefile that checks for conflicts with previously loaded CoNCePTuaL modulefiles then sets the `PATH`, `MANPATH`, and `LD_LIBRARY_PATH` environment variables to values appropriate values as determined by `configure` (see Section 2.1 [configure], page 4).
>
> Normally, `conceptual_1.5.1b` should be installed in the system's module path (as described by the `MODULEPATH` environment variable). However, users without administrator access can still use the CoNCePTuaL modulefile as a convenient mechanism for properly setting all of the environment variables needed by CoNCePTuaL:

```
make modulefile
module load ./conceptual_1.5.1b
```

> See the `module` man page for more information about modules.

`make dist`

> Package together all of the files needed to rebuild CoNCePTuaL. The resulting file is called `conceptual-1.5.1b.tar.gz` (for this version of CoNCePTuaL).

`make all`   Although `all` is the default target it can also be specified explicitly. Doing so is convenient when performing multiple actions at once, e.g., `make clean all`.

`make tags`

> Produce/update a `TAGS` file that the Emacs text editor can use to find function declarations, macro definitions, variable definitions, `typedef`s, etc. in the CoN-CePTuaL run-time library source code. This is useful primarily for developers wishing to interface with the CoNCePTuaL run-time library. Read the Emacs documentation for `M-x find-tag` for more information.

`make gui`   Compile the CoNCePTuaL GUI, producing `ncptlGUI-1.5.1b.jar`. Note that compilation requires both a Java compiler (e.g., `javac`) and the Jython Python-to-Java compiler (`jythonc`). Unfortunately, at the time of this writing (January 2009), `jythonc`'s future is uncertain (cf. http://www.jython.org/Project/jythonc.html). Hence, `make gui` has been tested only with `jythonc` version 2.2.x, not any later versions.

Validating the coNCePTuaL timer

`make` automatically builds a program called `validatetimer`. `validatetimer` helps validate that the real-time clock used by the CoNCePTuaL run-time library accurately measures wall-clock time. The idea is to compare CoNCePTuaL's timer to an external clock (i.e., one not associated with the computer). Simply follow the program's prompts:

```
% validatetimer
Press <Enter> to start the clock ...
Press <Enter> again in exactly 60 seconds ...


coNCePTuaL measured 60.005103 seconds.
coNCePTuaL timer error = 0.008505%
```

If the difference between coNCePTuaL's timer and an external clock is significant, then performance results from coNCePTuaL—and possibly from other programs, as well—should not be trusted. Note that only extreme differences in timings are significant; there will always be *some* error caused by human response time and by system I/O speed. In the case that there *is* an extreme performance difference,[1] the `--with-gettimeofday` option to `configure` (see Section 2.1 [configure], page 4) may be a viable workaround.

`validatetimer` takes an optional command-line argument, which is the number of seconds of wall-clock time to expect. The default is '60'. Larger numbers help amortize error; smaller numbers enable the program to finish sooner.

2.3 `make install`

The coNCePTuaL compiler and run-time library are installed with `make install`. Although `configure` can specify the default installation directory (see Section 2.1 [configure], page 4), this can be overridden at `make install` time in one of two ways. `make DESTDIR=prefix install` prepends *prefix* to every directory when installing. However, the files are installed believing that `DESTDIR` was not specified. For example, `make DESTDIR=/mnt install` would cause executables to be installed into `/mnt/usr/local/bin`, but if any of these are symbolic links, the link will omit the `/mnt` prefix.

The second technique for overriding installation directories is to specify a new value for 'prefix' on the command line. That is, `make prefix=/opt/ncptl install` will install into `/opt/ncptl/bin`, `/opt/ncptl/include`, `/opt/ncptl/man`, etc., regardless of the `--prefix` value given to `configure`. coNCePTuaL's `Makefile` provides even finer-grained control than that. Instead of—or in addition to—specifying a *prefix* option on the command line, individual installation directories can be named explicitly. These include `bindir`, `datadir`, `libdir`, `includedir`, `infodir`, `mandir`, `pkgdatadir`, `pythondir`, and many others. Scrutinize the `Makefile` to find a particular directory that should be overridden.

The remainder of this section presents a number of optional installation steps that add coNCePTuaL support to a variety of third-party software packages.

Installing stylesheets

The `make stylesheets` command (see Section 2.2 [make], page 7) produces a variety of stylesheets for presenting coNCePTuaL code in a more pleasant format than ordinary, monochromatic text. Stylesheets must currently be installed manually as per the following instructions:

---

[1] To date, extreme performance differences have been observed primarily on PowerPC-based systems. The PowerPC cycle counter is clocked at a different rate from the CPU speed, which may confuse coNCePTuaL. The run-time library compensates for this behavior on all tested platforms (see Chapter 2 [Installation], page 3), but the user should nevertheless make sure to run `validatetimer` to verify that coNCePTuaL's timer is sufficiently accurate.

`ncptl.sty`

> `ncptl.sty` is typically installed in *texmf*`/tex/latex/misc`, where *texmf*
> is likely to be `/usr/local/share/texmf`. On a Web2c version of TEX the
> command *kpsewhich -expand-var='$TEXMFLOCAL'* should output the correct
> value of *texmf*. In most TEX distributions the filename database needs to
> be refreshed after a new package is installed. See http://www.tex.ac.
> uk/cgi-bin/texfaq2html?label=instpackages for more information.
> `ncptl.sty` is merely a customization of the `listings` package that defines a
> new language called '`ncptl`'. See the `listings` documentation for instructions
> on typesetting source code.

`ncptl.ssh`

> Running *a2ps --list=defaults* outputs (among other things) the a2ps library
> path. `ncptl.ssh` should be installed in one of the `sheets` directories listed
> there, typically `/usr/share/a2ps/sheets`.

`ncptl-mode.el`
`ncptl-mode.elc`

> `ncptl-mode.el` and `ncptl-mode.elc` belong in a local Elisp directory
> that is part of the Emacs `load-path`, e.g., `/usr/share/emacs/site-lisp`.
> The following Elisp code, which belongs in `~/.emacs` for GNU Emacs or
> `~/.xemacs/init.el` for XEmacs, makes Emacs set `ncptl-mode` whenever
> opening a file with extension `.ncptl`:

```
(autoload 'ncptl-mode "ncptl-mode"
  "Major mode for editing coNCePTuaL programs." t)
(add-to-list 'auto-mode-alist '("\\.ncptl$" . ncptl-mode))
```

> Syntax highlighting should be enabled by default. If it isn't, the Emacs com-
> mand *M-x font-lock-mode* should enable it for the current buffer.

`ncptl.vim`

> Vim's syntax-file directory may be named after the Vim version, e.g.,
> `/usr/share/vim/vim61/syntax` for Vim 6.1. Put `ncptl.vim` there. To
> associate `.ncptl` files with CONCEPTUAL code, the following lines need
> to be added to Vim's `filetype.vim` file somewhere between the '`augroup`
> `filetypedetect`' line and the '`augroup END`' line:

```
" coNCePTuaL
au BufNewFile,BufRead *.ncptl          setf ncptl
```

`ncptl.lang`

> Source-highlight stores all of its helper files in a single directory, typically
> `/usr/share/source-highlight`. Put `ncptl.lang` there. To associate `.ncptl`
> files with CONCEPTUAL code you will also need to add the following line to
> the `lang.map` file in the same directory:

```
ncptl = ncptl.lang
```

ncptl.php

GeSHi stores all of its language files in a `geshi` subdirectory. Simply put `ncptl.php` there, and it should be ready for use. A quick way to convert CONCEPTUAL programs to color-coded HTML is to use the command-line version of the PHP interpreter with the following shell script.

```
#!/usr/bin/env php

<html>
<head>
<title><?php echo basename($argv[1]); ?></title>
</head>
<body>
<?php
include("GeSHi_directory/geshi.php");
$ncptl_code = file_get_contents($argv[1]);
geshi_highlight($ncptl_code, "ncptl");
?>
</body>
</html>
```

Simply replace *GeSHi_directory* with the name of your top-level GeSHi directory and run the script with the name of a CONCEPTUAL file as its sole argument.

A GeSHi extension for Mediawiki introduces a '`<syntaxhighlight>`' tag for presenting color-coded program listings on a wiki. `ncptl.php` makes it easy to use CONCEPTUAL code in this context.

SLOCCount

SLOCCount is a utility that counts the number of lines of code in a file, excluding blank lines and comments. SLOCCount supports a variety of programming languages and it is straightforward to get it to support CONCEPTUAL, as well. The procedure follows the "Adding support for new languages" section of the SLOCCount manual:

1. Create an `ncptl_count` script with the following contents:

```
#! /bin/sh

generic_count "#" $@
```

2. Mark the script executable and install it somewhere in your executable search path.

3. Edit SLOCCount's `break_filelist` Perl script to include the following association in the `%file_extensions` hash:

```
"ncptl" => "ncptl",    # coNCePTuaL
```

pkg-config

The pkg-config utility helps ensure that programs are given appropriate compiler and linker
flags to use a particular package's C header files and libraries. cONCEPTuAL's `configure`
script (see Section 2.1 [configure], page 4) automatically produces a pkg-config configuration
file for the cONCEPTuAL header file (`ncptl.h`) and run-time library (`libncptl`). This
configuration file, `ncptl.pc`, should be installed in one of the directories searched by pkg-
config (`/usr/lib/pkgconfig` on some systems). Once `ncptl.pc` is installed, pkg-config
can be used to compile C programs that require the cONCEPTuAL header file and link
programs that require the cONCEPTuAL run-time library, as is shown in the following
example:

```
cc `pkg-config --cflags ncptl` -c myprog.c
cc -o myprog myprog.o `pkg-config --libs ncptl`
```

3 Usage

CoNCePTuaL is more than just a language; it is a complete toolset that consists of the following components:

- the CoNCePTuaL language (see Chapter 4 [Grammar], page 84)
- a compiler and run-time library for CoNCePTuaL programs (see Section 3.2 [Compiling coNCePTuaL programs], page 18)
- a set of compiler backends that can generate code for a variety of languages and communication layers (see Section 3.3 [Supplied backends], page 20)
- utilities to help analyze the results (see Section 3.5.2 [ncptl-logextract], page 53)

This chapter explains how to compile and run CoNCePTuaL programs and how to interpret the log files they output.

3.1 The coNCePTuaL GUI

The CoNCePTuaL graphical user interface (GUI) is the easiest way to get started with CoNCePTuaL. Instead of writing code in the CoNCePTuaL language (documented in detail in Chapter 4 [Grammar], page 84), a user merely *draws* a communication pattern using the mouse, and the CoNCePTuaL GUI automatically produces a CoNCePTuaL program from that illustration. The generated program can then be compiled just like a hand-coded CoNCePTuaL program as per the instructions in Section 3.2 [Compiling coNCePTuaL programs], page 18.

The CoNCePTuaL GUI is written in Java and therefore requires a Java virtual machine (JVM) to run. However, the CoNCePTuaL GUI is quite portable and should run identically on every platform for which a JVM exists. The following is a typical command for launching the CoNCePTuaL GUI from the command line:

```
java -jar ncptlGUI-1.5.1b.jar
```

The CoNCePTuaL GUI is split into two main panels. The *program panel* displays the components that make up the program and how they interact with each other. The *dialog panel* displays fields for setting the options of selected components. Furthermore, a *menu bar* provides access to various GUI-wide operations, and a *command bar* includes buttons for creating new components in the program panel.
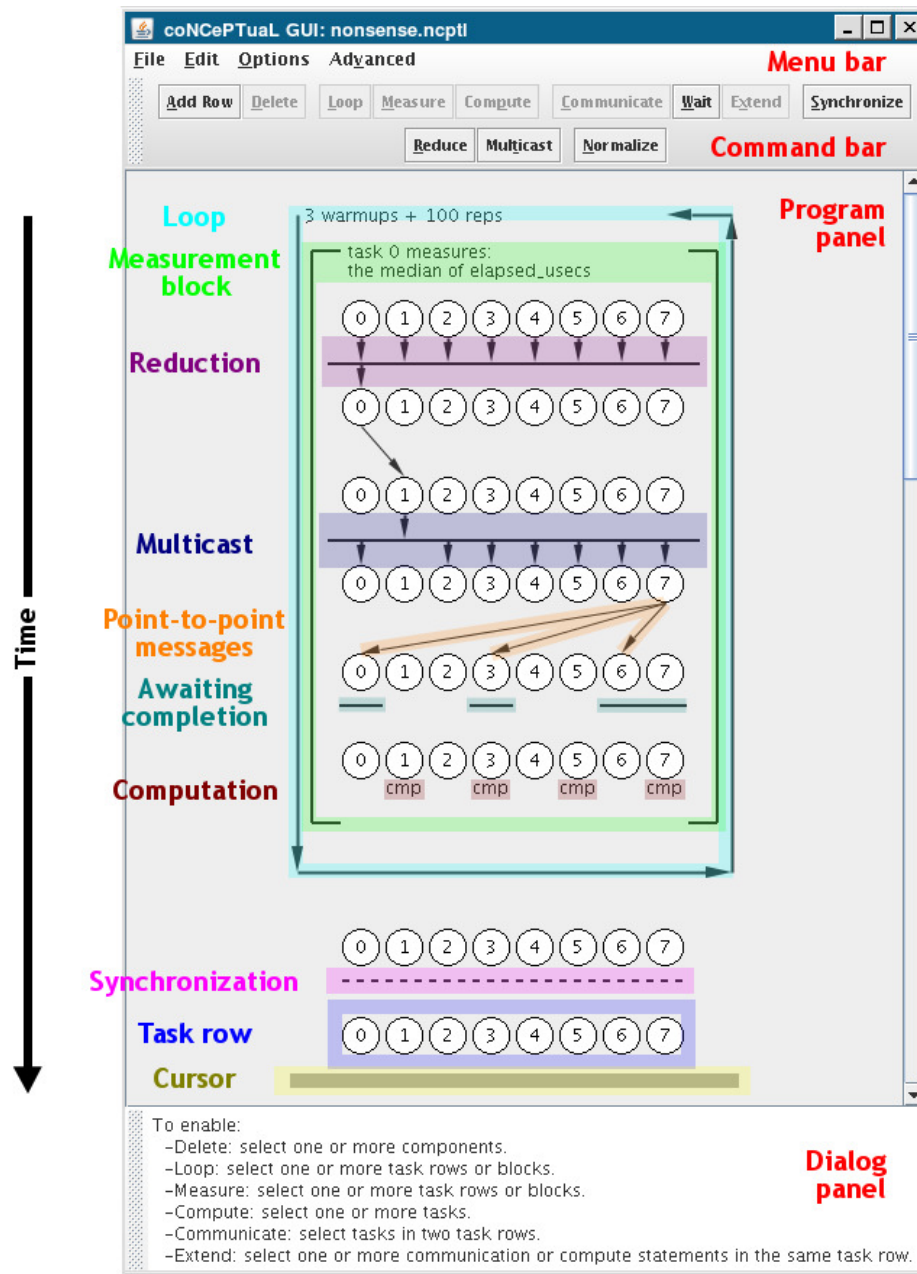
3.1.1 Components

Components are graphical representations of CONCEPTUAL objects and operations used by the CONCEPTUAL GUI to specify programs. Components are added to the program panel by clicking on their corresponding buttons in the command bar. To select a component in the program panel, simply left-click on it. If the component has parameters that can be edited, a dialog will appear in the dialog panel. Multiple components can be selected by dragging the mouse over target components or holding down Ctrl as you left-click on several components.

The CONCEPTUAL GUI lets a user manipulate the following components:

task/task row

> Tasks represent operational units in CONCEPTUAL programs and are analogous to a process or thread in a parallel program. For example, a CONCEPTUAL operation like point-to-point communication has a source task that specifies how the message is sent and a target task that specifies how the message is received. Tasks are displayed graphically as numbered circles in the CONCEPTUAL GUI. A task row represents the total number of tasks available for operations at each step of a CONCEPTUAL program. See Section 4.3 [Task descriptions], page 106, for information on how to specify subsets of a program's tasks.

communication message

> Messages between tasks are displayed as directed edges (arrows) from source task to target task in the CONCEPTUAL GUI. Messages are added to the program panel via the command bar or by dragging the mouse from a source task to a target task. Communication messages in the CONCEPTUAL GUI correspond to the SEND and RECEIVE statements in the CONCEPTUAL language (see Section 4.4.2 [Sending], page 113, and Section 4.4.3 [Receiving], page 116).

awaiting completion

> Messages that are sent/received asynchronously must eventually be waited on. Awaits message completion is displayed as a solid line under the associated tasks in the CONCEPTUAL GUI. Awaiting completion in the CONCEPTUAL GUI correspond to the AWAIT COMPLETION statement in the CONCEPTUAL language (see Section 4.4.4 [Awaiting completion], page 117).

loop   One can add a loop around selected components to repeat the corresponding CONCEPTUAL operations. Loops in the CONCEPTUAL GUI correspond to the FOR statement in the CONCEPTUAL language (see Section 4.7.2 [Iterating], page 126).

measurement block

> One can log timing or other measurements of CONCEPTUAL operations by placing them in a measurement block. Measurement blocks in the CONCEPTUAL GUI correspond to the LOGS statement in the CONCEPTUAL language (see Section 4.5.3 [Writing to a log file], page 120).

computation/sleeping

> Artificial computation (really a spin loop) and sleeping, both of which delay the program for a given length of time, can be performed on a set of tasks. Computation is shown with 'cmp' under a task in the CONCEPTUAL GUI, and sleeping is shown with 'slp'. Computation/sleeping in the CONCEPTUAL GUI corresponds to the COMPUTE and SLEEP statements in the CONCEPTUAL language (see Section 4.8.2 [Delaying execution], page 134).

multicasting

> A multicast operation sends a message from a source task to multiple target tasks. Multicasting in the CONCEPTUAL GUI corresponds to the MULTICAST statement in the CONCEPTUAL language (see Section 4.4.5 [Multicasting], page 117).

reduction   A reduction operation combines messages from multiple source tasks to a single target task. Reduction in the CONCEPTUAL GUI corresponds to the `REDUCE` statement in the CONCEPTUAL language (see Section 4.4.6 [Reducing], page 118).

synchronization

   Barrier synchronization forces a set of tasks to wait until each task in the set reaches the synchronization point before any task in the set proceeds past the synchronization point. Synchronization is displayed as a dotted line under the associated tasks in the CONCEPTUAL GUI. Synchronization in the CONCEP-TUAL GUI corresponds to the `SYNCHRONIZE` statement in the CONCEPTUAL language (see Section 4.4.7 [Synchronizing], page 119).

3.1.2 Menu bar

The *File* menu, which appears only when the CONCEPTUAL GUI is granted access to the filesystem, contains *New*, *Open*, *Save*, *Save As*, *Print*, and *Quit* commands that exhibit the expected behavior. Programs are saved as CONCEPTUAL source code that can then be compiled with the CONCEPTUAL compiler as per Section 3.2 [Compiling coNCePTuaL programs], page 18. *Print* prints a graphical view of the program as it appears on screen. (See Section 3.3.10 [The latex_vis backend], page 37, for a more sophisticated way to produce graphical views of CONCEPTUAL programs.)

   The *Edit* menu provides the usual *Cut*, *Copy*, *Paste*, and *Undo* commands.

   Options→Settings opens a dialog in the dialog panel for setting the number of tasks in a task row. The default number of tasks in a task row is 16.

   Advanced→Add conditional adds a conditional statement to a program at the current cursor position in the program panel. A dialog in the dialog panel will open for entering the conditional expression. A placeholder expression '`1 = 1`' is set by default.

   Advanced→Command line options opens a dialog in the dialog panel for adding command-line options to a program. A placeholder '`reps`' ('`number of repetitions`') option is set by default when this command is selected.

3.1.3 Command bar

The following buttons appear on the command bar:

Add Row   Insert a new, empty task row at the cursor position.

Delete   Delete the selected components.

Loop   Add a loop around the selected components.

Measure   Add a measurement block around the selected components.

Compute   Make the selected tasks "compute" or sleep for a length of time.

Communicate

   Add point-to-point communication between selected tasks (different task rows). This can also be achieved by dragging an arrow from a source task to a target task.

Wait   Make the selected tasks (same task row) or all tasks in the row above the cursor wait for all outstanding messages sent or received asynchronously to complete.

Extend   Extend a communication or computation pattern across an entire task row.

Synchronize

Synchronize the selected tasks (same task row) or all tasks in the task row above the cursor.

Reduce   Reduce data from the selected tasks in one task row to the selected tasks in the next task row. With no selection, reduce data from all tasks above the cursor to task 0 below the cursor.

Multicast   Multicast data from the selected tasks in one task row to the selected tasks in the next task row. With no selection, multicast data from task 0 above the cursor to all tasks below the cursor.

Normalize

Put the program into standard form as it will appear when translated into CONCEPTUAL code. After normalization, components are drawn as early in time as possible without changing the program's semantics.

3.2 Compiling coNCePTuaL programs

The CONCEPTUAL compiler is called `ncptl` and is, by default, installed into `/usr/local/bin`. Executing *ncptl --help* produces a brief usage string:

```
Usage: ncptl [--backend=<string>] [--quiet] [--no-link | --no-compile]
        [--keep-ints] [--lenient] [--filter=<sed expr>] [--output=<file>]
        <file.ncptl> | --program=<program>
        [<backend-specific options>]

     ncptl --help

     ncptl [--backend=<string>] --help-backend
```

The usage string is followed by a list of installed backends.

The following list describes each compiler option in turn:

--backend (abbreviation: *-b*)

Specify the module that CONCEPTUAL should use as the compiler backend. *ncptl* must be told which backend to use with either *--backend=backend* or by setting the environment variable *NCPTL_BACKEND* to the desired backend. Running *ncptl --help* lists the available backends. Most CONCEPTUAL backends are code generators. For example, `c_mpi` causes `ncptl` to compile CONCEPTUAL programs into C using MPI as the communication library. However, a backend need not generate code directly—or at all. The `c_trace` backend (see Section 3.3.4 [The c_trace backend], page 24), for instance, supplements the code generated by another backend by adding tracing output to it.

ncptl searches for backends first using *NCPTL_PATH*, an environment variable containing a colon-separated list of directories (default: empty); then, in the directory in which CONCEPTUAL installed all of its Python files; and finally,

in the default Python search path. Non-directories (e.g., the `.zip` archives used in newer versions of Python) are not searched.

If no backend is specified, `ncptl` runs the given CONCEPTUAL program through the lexer, parser, and semantic analyzer but does not generate an output file.

`--quiet` (abbreviation: `-q`)

The `--quiet` option tells `ncptl` and the chosen backend to output minimal status information.

`--no-link` (abbreviation: `-c`)

By default, `ncptl` instructs the backend to compile and link the user's CON-CEPTUAL program into an executable file. `--no-link` tells the backend to skip the linking step and produce only an object file.

`--no-compile` (abbreviation: `-E`)

By default, `ncptl` instructs the backend to compile and link the user's CON-CEPTUAL program into an executable file. `--no-compile` tells the backend to skip both the compilation and the linking step and to produce only a source file in the target language.

`--keep-ints` (abbreviation: `-K`)

CONCEPTUAL backends normally delete any files created as part of the compiling or linking process. `--keep-ints` tells `ncptl` and the chosen backend to preserve their intermediate files.

`--lenient` (abbreviation: `-L`)

The `--lenient` option tells the compiler to permit certain constructs that would otherwise result in a compilation error. First, using the same command-line option (either the long or short variant) for two different variables normally generates an '`Option opt is multiply defined`' error. (See Section 4.9.2 [Command-line arguments], page 139, for a description of how to declare command-line options in CONCEPTUAL.) `--lenient` tells the CONCEPTUAL compiler to automatically rename duplicate options to avoid conflicts. Only the option strings can be renamed; the programmer must still ensure that the option variables are unique. Second, using a variable without declaring it normally produces an error message at compile time. Passing `--lenient` to `ncptl` tells the compiler to automatically generate a command-line option for each missing variable. This is convenient when entering brief programs on the command line with `--program` (described below) as it can save a significant amount of typing.

`--filter` (abbreviation: `-f`)

The `--filter` option applies a `sed`-style substitution expression to the backend-translated code (e.g., a `.c` file output by the `c_udgram` backend or a `.tex` file output by the `latex_vis` backend) before the backend compiles it. The `--filter` option can be used multiple times on the command line; filters are applied in the order specified. Substitution expressions must be of the form '`s/pattern/replacement/flags`', although the '`/`' characters can be replaced

by any other character. *pattern* is a regular expression; *replacement* is an optional replacement string; and, *flags* is a sequence of zero or more modifiers from the set {'`i`', '`l`', '`m`', '`s`', '`u`', '`x`'}, as described in the Python Library Reference. For example, '`i`' means to perform a case-insensitive substitution. In addition, the '`g`' flag performs a global search-and-replace instead of replacing only the first occurrence of *pattern*. An important difference between `--filter` and `sed` is that omitting the '`g`' flag instructs `--filter` to make at most one substitution *total* while it instructs `sed` to make at most one substitution *per line*.

`--output` (abbreviation: `-o`)

> `ncptl` normally writes its output to a file with the same base name as the input file (or `a.out` if the program was specified on the command line using `--program`). `--output` lets the user specify a file to which to write the generated code.

`--program` (abbreviation: `-p`)

> Because CoNCePTuaL programs can be quite short `--program` enables a program to be specified in its entirety on the command line. The alternative to using `--program` is to specify the name of a file containing a CoNCePTuaL program. By convention, CoNCePTuaL programs have a `.ncptl` file extension.

`--help-backend` (abbreviation: `-H`)

> Describe additional options that are meaningful only to the specified backend. The `--backend` option must be used in conjunction with `--help-backend`.

The following—to be entered without line breaks—is a sample command line:

```
ncptl --backend=c_mpi --output=sample.c --program='Task 0 sends a 0
   byte message to task 1 then task 1 sends a 0 byte message to task 0
   then task 0 logs elapsed_usecs/2 as "Startup latency (usecs)".'
```

`ncptl` stops processing the command line at the first unrecognized option it encounters. That option and all subsequent options—including those which `ncptl` would otherwise process—are passed to the backend without interpretation by `ncptl`. Furthermore, the `--` (i.e., empty) option tells `ncptl` explicitly to stop processing the command line at that point. For example, in the command `ncptl --backend=some_backend --lenient myprogram.ncptl -- --program`, ncptl will process the `--backend` and `--lenient` options but will pass `--program` to the `some_backend` backend even though `ncptl` has its own `--program` option.[1]

## 3.3 Supplied backends

The CoNCePTuaL 1.5.1b distribution includes the following compiler backends:

`c_seq`   Generate ANSI C code with no communication support.

`c_mpi`   Generate ANSI C code with calls to the MPI library for communication.

---

[1] As an aside, `--help-backend` is essentially equivalent to '`-- --help`'; the `--help-backend` synonym is provided merely for convenience.

c_udgram   Generate ANSI C code that communicates using Unix-domain (i.e., local to a single machine) datagram sockets.

c_trace   Instrument a C-based backend either to include a call to `fprintf()` before every program event or to utilize the `curses` library to display graphically the execution of a selected task.

c_profile
            Instrument a C-based backend either to write event timings and tallies to each log file or to the standard error device if the CONCEPTUAL program doesn't use a log file.

interpret
            Interpret a CONCEPTUAL program, simulating any number of processors and checking for common problems such as deadlocks and mismatched sends and receives.

stats   Output statistics of a program's execution—message tallies, byte counts, communication peers, network bisection crossings, event tallies, etc.

picl   Output in the PICL trace format a logical-time trace of a CONCEPTUAL program's communication pattern.

paraver   Output in the Paraver trace format a logical-time trace of a CONCEPTUAL program's communication pattern.

latex_vis
            Use LaTeX to produce an Encapsulated PostScript visualization of a program's communication pattern.

dot_ast   Output a program's parse tree in the Graphviz dot format.

libsea_ast
            Output a program's parse tree in the LibSea graph file format.

CONCEPTUAL employs a highly modular software structure for its backends. Many of the backends listed above are built atop other backends. The following figure illustrates the current set of dependencies:
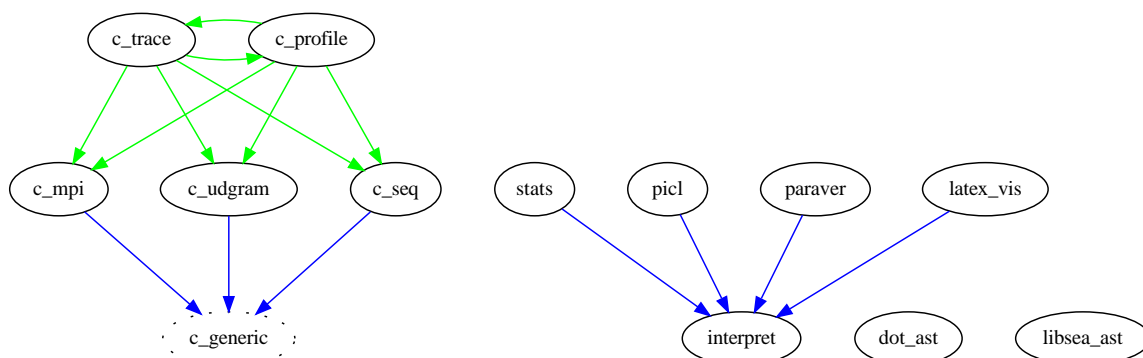
Some dependencies are defined as a static characteristic of a backend. For example, the `c_mpi` backend is hardwired to derive some of its functionality from `c_generic`. Other dependencies are determined dynamically. For example, the `c_trace` backend must be instructed to derive its functionality from one of `c_profile`, `c_mpi`, `c_udgram`, or `c_seq`. (See Section 3.3.4 [The c_trace backend], page 24, for more information on the `c_trace` backend.)

Except for `c_generic`, all of the compiler backends are described in turn in the following sections. The `c_generic` backend is unique because it is used exclusively to construct C-based backends; it does nothing by itself. See Section 6.2 [Backend creation], page 157, for a detailed description of `c_generic`.

Each backend accepts a `--help` option that explains the backend's command-line options. The easiest way to request help from a specific backend is with 'ncptl `--backend=backend -- --help`'. The empty `--` option, mentioned in Section 3.2 [Compiling coNCePTuaL programs], page 18, prevents the compiler from intercepting `--help` and providing the standard, backend-independent information.

3.3.1 The `c_seq` backend

The `c_seq` backend is intended primarily to provide backend developers with a minimal C-based backend that can be used as a starting point for creating new backends. See Section 6.2 [Backend creation], page 157, explains how to write backends.

3.3.2 The `c_mpi` backend

The `c_mpi` backend is coNCePTuaL's workhorse. It generates parallel programs written in ANSI C that communicate using the industry-standard MPI messaging library.

By default, `c_mpi` produces an executable program that can be run with mpirun, prun, pdsh, or whatever other job-launching program is normally used to run MPI programs. When ncptl is run with the `--no-link` option, `c_mpi` produces an object file that needs to be linked with the appropriate MPI library. When ncptl is run with the `--no-compile` option, `c_mpi` outputs ANSI C code that must be both compiled and linked.

`c_mpi` honors the following environment variables when compiling and linking C+MPI programs: *MPICC*, *MPICPPFLAGS*, *MPICFLAGS*, *MPILDFLAGS*, *MPILIBS*. If any of these variables is not found in the environment, `c_mpi` will use the value specified/discovered at configuration time (see Section 2.1 [configure], page 4). *MPICC* defaults to the value of *CC*; *MPICFLAGS* defaults to the value of *CFLAGS*; the remaining variables are appended respectively to *CPPFLAGS*, *LDFLAGS*, and *LIBS*.

The following is a complete list of MPI functions employed by the `c_mpi` backend: MPI_Allreduce(), MPI_Alltoallv(), MPI_Barrier(), MPI_Bcast(), MPI_Comm_group(), MPI_Comm_rank(), MPI_Comm_size(), MPI_Comm_split(), MPI_Errhandler_create(), MPI_Errhandler_set(), MPI_Finalize(), MPI_Group_translate_ranks(), MPI_Init(), MPI_Irecv(), MPI_Isend(), MPI_Recv(), MPI_Reduce(), MPI_Send(), MPI_Ssend(), MPI_Waitall(). In addition, if *./configure* is passed the `--with-mpi-wtime` option as described in Section 2.1 [configure], page 4, then *all* backends that utilize the coNCePTuaL run-time library, including `c_mpi`, will use MPI_Wtime() for taking performance measurements.

Command-line options for `c_mpi`

When `ncptl` is passed `--backend=c_mpi` as a command-line option, `c_mpi` processes the following backend-specific command-line options:

`--ssend`   In the generated code, `MPI_Isend()` and `MPI_Irecv()` are used for asynchronous communication and `MPI_Send()` and `MPI_Recv()` are normally used for synchronous communication. However, the `c_mpi`-specific compiler option `--ssend` instructs `c_mpi` to replace all calls to `MPI_Send()` in the generated code with calls to `MPI_Ssend()`, MPI's synchronizing send function. A program's log files indicate whether the program was built to use `MPI_Send()` or `MPI_Ssend()`.

`--reduce=MPI_Op`
             The default reduction operation is `MPI_SUM` but a different operation can be specified using the `c_mpi`-specific `--reduce` compiler option. A program's log files indicate the reduction operation that was used.

Implementation of reductions

The `c_mpi` backend implements the CoNCePTuaL `REDUCE` statement (see Section 4.4.6 [Reducing], page 118) as follows. Many-to-one reductions are implemented with a single call to `MPI_Reduce()`. Many-to-many reductions in which the sources exactly match the targets are implemented with a single call to `MPI_Allreduce()`. All other reductions are implemented by calling `MPI_Reduce()` to reduce the data to the first target task then `MPI_Bcast()` to distribute the reduced data to the remaining targets.

Because MPI requires that the source and target buffers in an `MPI_Reduce()` or `MPI_Allreduce()` be different, the `c_mpi` backend utilizes *two* buffers when `UNIQUE` (see [Unique messages], page 109) is specified. It utilizes two *adjacent* buffers when `FROM BUFFER` or `INTO BUFFER` (see [Buffer control], page 113) is specified.

Implementation of multicasts

The `MULTICAST` statement (see Section 4.4.5 [Multicasting], page 117) enables any set of tasks to transmit data to any other set of tasks. The `c_mpi` backend implements `MULTICAST` using the simpler `MPI_Bcast()` when there is a single initiator (determined dynamically at run time) and `MPI_Alltoallv()` otherwise. In all cases, `c_mpi` creates an MPI communicator that contains exactly the tasks involved in the multicast (as senders and/or receivers).

MPI requires that all tasks involved in an `MPI_Bcast()` belong to the same MPI communicator. Hence, in the CoNCePTuaL statement

```
TASK 4 MULTICASTS A 64 KILOBYTE MESSAGE TO TASKS r SUCH THAT r
IS IN {1, 3, 5, 7, 9}
```

the generated communicator contains tasks 1, 3, 4, 5, 7, and 9, even though task 4 transmits but does not receive any data.

Rather than repeatedly call `MPI_Bcast()`, many-to-one and many-to-many multicasts such as

```
TASKS from_tasks SUCH THAT from_tasks < 8 MULTICAST 3 2-KILOBYTE
MESSAGES to TASKS to_tasks SUCH THAT 3 DIVIDES to_tasks
```

perform a single `MPI_Alltoallv()`. Two message buffers are used in this case, one for all outgoing data and one for all incoming data.

3.3.3 The `c_udgram` backend

CoNCePTuaL program development on a workstation is facilitated by the `c_udgram` backend. `c_udgram` runs on only a single machine but, unlike `c_seq`, supports all of CoNCePTuaL's communication statements. Communication is performed over Unix-domain datagram sockets. Unix-domain datagrams are reliable and guarantee order (unlike UDP/IP datagrams) but have a maximum packet size. `c_udgram` backend write this maximum to every log file and automatically packetizes larger messages.

By default, `c_udgram` produces an executable program that can be run directly from the command line. When `ncptl` is run with the `--no-link` option, `c_udgram` produces an object file that needs to be linked with the appropriate sockets library (on systems that require a separate library for socket calls). When `ncptl` is run with the `--no-compile` option, `c_udgram` outputs ANSI C code that must be both compiled and linked. Like all C-based backends, `c_udgram` honors the `CC`, `CPPFLAGS`, `LDFLAGS`, and `LIBS` environment variables when compiling and linking. Values not found in the environment are taken from those specified/discovered at configuration time (see Section 2.1 [configure], page 4).

In addition to supporting the default set of command-line options, programs generated using the `c_udgram` backend further support a `--tasks` option that designates the number of tasks to use:

```
-T, --tasks=<number>          Number of tasks to use [default: 1]
```

`c_udgram` programs spawn one OS-level process for each task in the program. They also create a number of sockets in the current directory named '`c_udgram_<tag>`'. These are automatically deleted if the program exits cleanly but will need to be removed manually in the case that the program is killed by an non-trappable signal.

Implementation of collectives

The `c_udgram` normally uses logarithmic-time algorithms for multicasts (see Section 4.4.5 [Multicasting], page 117), reductions (see Section 4.4.6 [Reducing], page 118), and barriers (see Section 4.4.7 [Synchronizing], page 119). One-to-many multicasts disperse data in a binary-tree pattern. Reductions reduce data in a binary-tree pattern to the root, which then disperses the data in a binary-tree pattern in the many-to-many case. Barriers synchronize in a butterfly pattern. The only non-logarithmic-time algorithm is for many-to-many multicasts, which are implemented as one logarithmic-time, one-to-many multicast for each sender.

3.3.4 The `c_trace` backend

While most CoNCePTuaL backends are code generators, the `c_trace` backend adds tracing code to the code produced by a code-generating backend. `c_trace` is useful as a debugging aid and as a means to help understand the control flow of a CoNCePTuaL program.

Command-line options for `c_trace`

When `ncptl` is passed `--backend=c_trace` as a command-line option, `c_trace` processes the following backend-specific command-line options:

--*trace=backend*

> Specify a backend that will produce C code for `c_trace` to trace. The --*trace* option is required to use `c_trace`; `c_trace` will issue an error message if --*trace* is not specified. The restrictions on *backend* are that it must produce C code and must be derived from the `c_generic` backend (see Section 6.2 [Backend creation], page 157). Improper backends cause `c_trace` to abort abnormally.

--*curses*   Instead of injecting `fprintf()` statements into the generated C code, inject calls to the `curses` (or `ncurses`) library to show graphically the line of code currently executing on a given processor.

Default `c_trace` tracing

Without --*curses*, `c_trace` alters the generated C code to write data like the following to the standard error device:

```
[TRACE] phys: 1 | virt: 1 | action: RECV | event: 1 / 44001 | lines: 18 - 18
[TRACE] phys: 0 | virt: 0 | action: RESET | event: 1 / 88023 | lines: 17 - 17
[TRACE] phys: 0 | virt: 0 | action: SEND | event: 2 / 88023 | lines: 18 - 18
[TRACE] phys: 0 | virt: 0 | action: RECV | event: 3 / 88023 | lines: 19 - 19
[TRACE] phys: 1 | virt: 1 | action: SEND | event: 2 / 44001 | lines: 19 - 19
[TRACE] phys: 1 | virt: 1 | action: RECV | event: 3 / 44001 | lines: 18 - 18
[TRACE] phys: 0 | virt: 0 | action: CODE | event: 4 / 88023 | lines: 20 - 21
[TRACE] phys: 0 | virt: 0 | action: RESET | event: 5 / 88023 | lines: 17 - 17
[TRACE] phys: 0 | virt: 0 | action: SEND | event: 6 / 88023 | lines: 18 - 18
[TRACE] phys: 0 | virt: 0 | action: RECV | event: 7 / 88023 | lines: 19 - 19
```

$$\vdots$$

The format is designed to be easy to read and easy for a program to parse. Each line of trace data begins with the string '`[TRACE]`' and lists the (physical) processor number, the (virtual) task ID, the action (a.k.a., event type) that is about to be performed, the current event number and total number of events that will execute on the given processor, and the range of lines of source code to which the current event corresponds. An "event" corresponds more-or-less to a statement in the CONCEPTUAL language.[2] Loops are unrolled at initialization time and therefore produce no events. Section B.3 [Event types], page 215, lists and briefly describes the various event types.

`c_trace` tracing with `curses`

The --*curses* option enables a more interactive tracing environment. Generated programs must be linked with the `curses` (or compatible, such as `ncurses`) library. The resulting executable supports the following additional command-line options:

-*D*, --*delay=number*

> delay in milliseconds after each screen update ('0'=no delay)

---

[2] A more precise correspondence is to a <*simple_stmt*> in the formal grammar presented in Chapter 4 [Grammar], page 84.

-M, --monitor=number

          processor number to monitor

-B, --breakpoint=number

          source line at which to enter single-stepping mode ('-1'=none; '0'=first event)

When the program is run it brings up a screen like the following:

```
 1.   # Determine computational "noise"
 2.
 3.   Require language version "1.5".
 4.
 5.   accesses is "Number of data accesses to perform" and comes from
 6.     "--accesses" or "-a" with default 500000.
 7.
 8.   trials is "Number of timings to take" and comes from "--timings" or
 9.     "-t" with default 1000.
10.
11.   For trials repetitions {
12.     all tasks reset their counters then
13.     all tasks touch a 1 word memory region accesses times with stride 0 w
14.     all tasks log a histogram of elapsed_usecs as "Actual time (usecs)"
15.   }



Phys: 0  Virt: 0  Action: RESET     Event:    1/3001
```

The program displays its source code (truncated vertically if too tall and truncated horizontally if too wide) at the top of the screen and a status bar at the bottom of the screen. As the program executes, a cursor indicates the line of source code that is currently executing. Likewise, the status bar updates dynamically to indicate the processor's current task ID, action, and event number. In curses mode, the program's standard output (see Section 4.5.2 [Writing to standard output], page 120) is suppressed so as not to disrupt the trace display.

Programs traced with c_trace and the --curses option are made interactive and support the following (case-insensitive) keyboard commands:

'S'   Enable single-stepping mode. While single-stepping mode is enabled the traced processor will execute only one event per keystroke from the user.

'space'   Disable single-stepping mode. The program executes without further user intervention.

'D'   Delete the breakpoint.

'Q'   Quit the program. The log file will indicate that the program did not run to completion.

All other keystrokes cause the program to advance to the next event immediately.

A single breakpoint can be set using the program's `-B` or `--breakpoint` command-line option. Whenever the monitored processor reaches the source-code line at which a breakpoint has been set, it enters single-stepping mode exactly as if `S` were pressed. Setting a breakpoint at line 0 tells the program to begin single-stepping as soon as the program begins. Note that only lines corresponding to CoNCePTuaL events can support breakpoints.

Offline tracing with `curses`

The `c_trace` backend can be told to trace by writing messages to the standard error device (see [Default c_trace tracing], page 25) or by employing an interactive display (see [c_trace tracing with curses], page 25). These two alternatives can be combined to support offline tracing of a CoNCePTuaL program. The idea is to compile the program without the `--curses` option. When running the program, the standard-error output should be redirected to a file. The `ncptl-replaytrace` utility, which comes with CoNCePTuaL, can then be used to play back the program's execution by reading and displaying the file of redirected trace data.

`ncptl-replaytrace` accepts the following command-line options, which correspond closely to those accepted by a program compiled with the `--curses` option to `c_trace` (see [c_trace tracing with curses], page 25):

`--trace =file`
> Specify a file containing redirected trace data. *file* defaults to the standard input device.

`--delay =number`
> Specify the delay in milliseconds after each screen update ('0'=no delay).

`--monitor =processor`
> Specify the processor number to monitor. *processor* defaults to '0'.

`--breakpoint =line`
> Specify a line of source code at which to enter single-stepping mode ('`-1`'=none; '0'=first event).

In addition, `ncptl-replaytrace` requires that the CoNCePTuaL source-code file be specified on the command line, as the source code is not included in the trace data.

The interactive display presented by the offline `ncptl-replaytrace` tool is nearly identical to that presented by a program compiled with the `--curses` option to `c_trace`. See [c_trace tracing with curses], page 25, for a usage description.

3.3.5 The `c_profile` backend

Profiling communication events can help explain surprising performance measurements. A profiling tool that automatically instruments CoNCePTuaL programs is simpler than manually adding `LOGS` or `OUTPUTS` statements to an existing program.

Like the `c_trace` backend (see Section 3.3.4 [The c_trace backend], page 24), the `c_profile` backend adds code to that produced by other backends. The `c_profile` backend accepts a single—mandatory—command-line argument, `--profile=backend`, which designates a target backend to use. The restrictions on *backend* are that it must produce C code and must be derived from the `c_generic` backend (see Section 6.2 [Backend creation], page 157). Improper backends cause `c_profile` to abort abnormally.

If the profiled CoNCePTuaL program produces log files, the log files will include profiling information in the epilogue, one line for each event type that was used at least once by the corresponding process. (Section B.3 [Event types], page 215, lists and briefly describes the various event types.) Each line names an event and presents the total number of microseconds spent processing that event, a tally of the number of times that event type was executed, and the quotient of those two numbers. In addition, the amount of memory used to store the event list is also logged.

For example, the following extract from a log file indicates that the process that wrote it spent a total of 6.7 s processing 22,000 'RECV' events (which includes waiting time) for an average of $303.5\,\mu$s per 'RECV' event:

```
# Profile of SEND (microseconds, count, average): 5267469 22001 239.4
# Profile of RECV (microseconds, count, average): 6676521 22000 303.5
# Profile of REPEAT (microseconds, count, average): 11985167 1 11985167.0
# Profile of NEWSTMT (microseconds, count, average): 43 1 43.0
# Profile of CODE (microseconds, count, average): 5516 1 5516.0
# Profile of event memory: 528 bytes (6 events * 88 bytes/event)
```

Each 'REPEAT' event includes the time for all of the events that it repeats.

Although the preceding log-file excerpt indicates that a total of $22001+22000+1+1+1 = 44004$ events were executed, the 'event memory' line clarifies that the event list contained only 6 unique events and therefore required only 528 bytes of memory.

Profiled programs that do not produce log files write profiling information to the standard error device. Because all processes may share a single standard error device, each line of output is preceded by a processor ID as in the following example:

```
1 SEND 5527125 22000 251.2
1 RECV 6322699 22001 287.4
1 REPEAT 11894523 1 11894523.0
1 event-memory 352 4 88
0 SEND 5267469 22001 239.4
0 RECV 6676521 22000 303.5
0 REPEAT 11985167 1 11985167.0
0 event-memory 352 4 88
```

The columns are *processor ID*, *event*, *total microseconds*, *tally*, and *average microseconds* except when *event* is 'event-memory' in which case the columns are *processor ID*, 'event-memory', *total bytes*, *number of events*, and *bytes per event*. The intention is for the output to be easily parseable using tools such as awk.

3.3.6 The interpret backend

Like the c_udgram backend (see Section 3.3.3 [The c_udgram backend], page 24), the interpret backend is designed to help programmers ensure the correctness of CoNCePTuaL code. The interpret backend does not output code. As its name implies, interpret is an *interpreter* of CoNCePTuaL programs rather than a compiler. interpret exhibits the following salient features:

1. Some programs run faster than with a compiler because the interpreter does not actually send messages. `interpret` merely simulates communication. It also skips over statements such as `COMPUTES`/`SLEEPS` (see Section 4.8.2 [Delaying execution], page 134) and `TOUCHES` (see Section 4.8.3 [Touching memory], page 134).

2. `interpret` can simulate massively parallel computer systems from a single process.

3. As `interpret` runs it checks for common communication errors such as deadlocks, asynchronous sends and receives that are never completed, and blocking operations left over at the end of the program (which would likely cause hung tasks under a real messaging layer).

The drawbacks are that `interpret` is slow when interpreting control-intensive programs and that timing measurements are not indicative of any real network. (The `interpret` backend utilizes logical time rather than physical time.) `interpret` is intended primarily as a development tool for helping ensure the correctness of CoNCePTuaL programs.

The `interpret` backend accepts all of the command-line options described in Section 3.4 [Running coNCePTuaL programs], page 46, plus the following four options:

```
-H, --hierarchy=<string>      Latency hierarchy as a comma-separated list
                              of task_factor:latency_delta pairs [default:
                              "tasks:1"]
-K, --kill-reps=<number>      If nonzero, perform FOR...REPETITIONS loop
                              bodies exactly once [default: 0]
-M, --mcastsync=<number>      Perform an implicit synchronization after a
                              multicast (0=no; 1=yes) [default: 0]
-T, --tasks=<number>          Number of tasks to use [default: 1]
```

Normally, the `interpret` backend assigns unit latency to every communication operation. The `--hierarchy` option can make communication with distant tasks observe more latency than communication with nearby tasks. An explanation of the argument to `--hierarchy` is presented in [Task latency hierarchies], page 31.

To save execution time, the `--kill-reps` option alters the behavior of all `FOR...REPETITIONS` statements in the program, treating them as if they read, 'FOR 1 REPETITION'. That is, it ignores warmup repetitions, synchronizations, and the specified number of repetitions, always using '1' instead.

A multicast operation (see Section 4.4.5 [Multicasting], page 117) is normally treated as multiple point-to-point operations with the same send time. The `--mcastsync` option instructs the `interpret` backend to perform an implicit barrier synchronization at the end of the multicast.

The `--tasks` option specifies the number of tasks to simulate. Because this number can be quite large the `NCPTL_LOG_ONLY` environment variable (see Appendix C [Environment Variables], page 217) may be used to limit the set of processors that are allowed to create log files. That way, if task 0 is the only task out of thousands that logs any data, `NCPTL_LOG_ONLY` can specify that only one log file will be produced, not thousands. By default, all processors create a log file.

All other command-line arguments are passed to the program being interpreted.

The `--output` option described in Section 3.2 [Compiling coNCePTuaL programs], page 18, has special meaning to the `interpret` backend. When `--output` is used, `interpret` dumps a list of events to the specified file after a successful run. For example, the CoNCePTuaL program 'ALL TASKS t ASYNCHRONOUSLY SEND A 384 BYTE MESSAGE TO TASK t XOR 2 THEN ALL TASKS AWAIT COMPLETION' results in the following event dump:

```
Task 0 posted a NEWSTMT at time 0 and completed it at time 0
Task 0 posted a RECEIVE at time 0 and completed it at time 0
Task 0 posted a SEND at time 1 and completed it at time 1
Task 0 posted a WAIT_ALL at time 2 and completed it at time 2
Task 1 posted a NEWSTMT at time 0 and completed it at time 0
Task 1 posted a RECEIVE at time 0 and completed it at time 0
Task 1 posted a SEND at time 1 and completed it at time 1
Task 1 posted a WAIT_ALL at time 2 and completed it at time 2
Task 2 posted a NEWSTMT at time 0 and completed it at time 0
Task 2 posted a RECEIVE at time 0 and completed it at time 0
Task 2 posted a SEND at time 1 and completed it at time 1
Task 2 posted a WAIT_ALL at time 2 and completed it at time 2
Task 3 posted a NEWSTMT at time 0 and completed it at time 0
Task 3 posted a RECEIVE at time 0 and completed it at time 0
Task 3 posted a SEND at time 1 and completed it at time 1
Task 3 posted a WAIT_ALL at time 2 and completed it at time 2
```

As an example of the `interpret` backend's usage, here's how to simulate 100,000 processors communicating in a simple ring pattern:

```
% ncptl --backend=interpret --lenient --program='All tasks t send
    nummsgs 1024 gigabyte messages to task t+1 then task num_tasks-1
    sends nummsgs 1024 gigabyte messages to task 0.' --tasks=100000
    --nummsgs=5
```

The preceding command ran to completion in under 5 minutes on a 1.5 GHz Xeon uniprocessor workstation—not too bad considering that 488 petabytes of data are transmitted on the program's critical path.

The `interpret` backend is especially useful for finding communication-related program errors:

```
% ncptl --backend=interpret --quiet --program='All tasks t send
    a 10 doubleword message to task (t+1) mod num_tasks.' --tasks=3
<command line>: The following tasks have deadlocked: 0 --> 2 --> 1
    --> 0
```

Deadlocked tasks are shown with '-->' signifying "is blocked waiting for". In the preceding example, all receives are posted before all sends. Hence, task 0 is blocked waiting for task 2 to send it a message. Task 2, in turn, is blocked waiting for task 1 to sent it a message. Finally, task 1 is blocked waiting for task 0 to send it a message, which creates a cycle of dependencies.

The `interpret` backend can find other errors, as well:

```
% ncptl --backend=interpret --quiet --program='All tasks t
    asynchronously send a 10 doubleword message to task (t+1) mod
    num_tasks.' --tasks=4
```

```
<command line>: The program ended with the following leftover-event
    errors:
  * Task 0 posted an asynchronous RECEIVE that was never waited for
  * Task 0 posted an asynchronous SEND that was never waited for
  * Task 0 sent a message to task 1 that was never received
  * Task 1 posted an asynchronous RECEIVE that was never waited for
  * Task 1 posted an asynchronous SEND that was never waited for
  * Task 1 sent a message to task 2 that was never received
  * Task 2 posted an asynchronous RECEIVE that was never waited for
  * Task 2 posted an asynchronous SEND that was never waited for
  * Task 2 sent a message to task 3 that was never received
  * Task 3 posted an asynchronous RECEIVE that was never waited for
  * Task 3 posted an asynchronous SEND that was never waited for
  * Task 3 sent a message to task 0 that was never received
```

(A message received `ASYNCHRONOUSLY` is not considered received until after the corresponding `AWAITS COMPLETION`; hence, all of the 'was never received' messages listed above.)

```
% ncptl --backend=interpret --quiet --program='Task 0 sends a 40
    kilobyte message to unsuspecting task 1 then task 0 receives a 40
    kilobyte message from task 1.' --tasks=2
<command line>: The program ended with the following leftover-event
    errors:
  * Task 0 sent a message to task 1 that was never received
  * Task 1 terminated before satisfying task 0's RECEIVE operation
```

In short, it is well worth testing the correctness of new coNCePTuaL programs with `interpret` before performing timing runs with one of the message-passing backends.

Task latency hierarchies

The `interpret` backend normally simulates a flat network in which communication between any two tasks takes unit latency. However, the `--hierarchy` option lets one specify a hierarchy of latencies: latency l_1 within a set of t_1 tasks, latency $l_1 + l_2$ within a set of $t_1 \cdot t_2$ tasks, latency $l_1 + l_2 + l_3$ within a set of $t_1 \cdot t_2 \cdot t_3$ tasks, and so forth.

The argument to `--hierarchy` is a comma-separated list of ⟨*task factor*:*latency delta*⟩ pairs. The *latency delta* component is optional and defaults to '`1`'. If the list ends with '`...`', the final ⟨*task factor*:*latency delta*⟩ pair is repeatedly indefinitely. All *task factor* values must be positive integers, and all *latency delta* values must be nonnegative integers.

As an example, `--hierarchy=4` (or `--hierarchy=4:1`) partitions the program's tasks into sets of four with one unit of additional latency to communicate with another set. Hence, tasks 0, 1, 2, and 3 can communicate with each other in unit time, as can tasks 4, 5, 6, and 7. However, communication between any task in the first set and any task in the second set (e.g., between tasks 3 and 4) takes an additional unit of time for a total latency of two units.

As a more complex example, consider a cluster of symmetric multiprocessors (SMPs) in which each SMP comprises two processor sockets with each socket containing a quad-core processor (four CPUs). Further assume that the SMPs are networked together via a fat-tree network consisting of 12-port switches. In this example,

let's say that it takes unit latency to communicate within a socket, two units of latency to communicate with another socket in the same SMP, and an additional three units of latency to traverse each switch. This configuration can be specified to the `interpret` backend as `--hierarchy="4:1, 2:1, 12:3, 12:6, ..."` (or simply `--hierarchy=4,2,12:3,12:6, ...`). With this setting, task 0, for instance, communicates with tasks 1–3 in 1 time unit, with tasks 4–7 in 2 time units, with tasks 8–95 in 5 time units (one switch crossing), with tasks 96–1151 in 11 time units (three switch crossings—a level 0 switch, a level 1 switch, and another level 0 switch), with tasks 1152–13,823 in 17 time units (five switch crossings—of switch levels 0, 1, 2, 1, and 0), and so forth.

3.3.7 The `stats` backend

The `stats` backend outputs various statistics about a CONCEPTUAL program. It can be used to help verify a program's correctness or merely to search for interesting patterns within a complex communication pattern. The following is some sample output from a program compiled with `stats`:

```
Execution parameters
--------------------
  Number of processors:        16
  Random-number seed:    -1727114895
  Command line:          ncptl --backend=stats hycom.ncptl --xdim=4 --ydim=4 --tasks=16 --iter=10
  Timestamp:             Wed Jan  4 17:32:04 2006

Message traffic
---------------
  Total messages sent:         710
  Total bytes sent:         1454080
  Unique message sizes sent: 2048

Per-processor message traffic
-----------------------------
  Processors sending a total of 61440 bytes:    0
  Processors sending a total of 81920 bytes:    1-3, 12-15
  Processors sending a total of 102400 bytes:   4-11
  Processors receiving a total of 61440 bytes:  1-3, 13-15
  Processors receiving a total of 81920 bytes:  5-7, 9-11
  Processors receiving a total of 122880 bytes: 12
  Processors receiving a total of 143360 bytes: 4, 8
  Processors receiving a total of 184320 bytes: 0
  Processors sending a total of 30 messages:    0
  Processors sending a total of 40 messages:    1-3, 12-15
  Processors sending a total of 50 messages:    4-11
  Processors receiving a total of 30 messages:  1-3, 13-15
  Processors receiving a total of 40 messages:  5-7, 9-11
  Processors receiving a total of 60 messages:  12
  Processors receiving a total of 70 messages:  4, 8
  Processors receiving a total of 90 messages:  0

Processor SEND-event peers
--------------------------
  Processors posting SEND events to offsets {-12, -4, +1, +3}:    12
  Processors posting SEND events to offsets {-8, -4, +1, +3, +4}: 8
  Processors posting SEND events to offsets {-4, -3, -1}:         15
  Processors posting SEND events to offsets {-4, -3, -1, +4}:     7, 11
  Processors posting SEND events to offsets {-4, -2, -1, +1}:     14
  Processors posting SEND events to offsets {-4, -2, -1, +1, +4}: 6, 10
```

```
        Processors posting SEND events to offsets {-4, -1, +1}:        13
        Processors posting SEND events to offsets {-4, -1, +1, +4}:    5, 9
        Processors posting SEND events to offsets {-4, +1, +3, +4}:    4
        Processors posting SEND events to offsets {-3, -1, +4}:        3
        Processors posting SEND events to offsets {-2, -1, +1, +4}:    2
        Processors posting SEND events to offsets {-1, +1, +4}:        1
        Processors posting SEND events to offsets {+1, +3, +4}:        0

    Network bisection crossings
    ---------------------------
      Bisection messages:    100
      Bisection bytes:    204800

    Event tallies
    -------------
      Total number of LOG events:       10
      Total number of NEWSTMT events:   48
      Total number of RECEIVE events:  710
      Total number of RESET events:     10
      Total number of SEND events:     710
      Total number of WAIT_ALL events: 480

    Per-processor event sets
    ------------------------
      Processors executing only {LOG, NEWSTMT, RECEIVE, RESET, SEND, WAIT_ALL}: 0
      Processors executing only {NEWSTMT, RECEIVE, SEND, WAIT_ALL}:             1-15

    Per-processor event tallies
    ---------------------------
      Processors executing 10 LOG events:      0
      Processors executing 3 NEWSTMT events:   0-15
      Processors executing 30 RECEIVE events:  1-3, 13-15
      Processors executing 40 RECEIVE events:  5-7, 9-11
      Processors executing 60 RECEIVE events:  12
      Processors executing 70 RECEIVE events:  4, 8
      Processors executing 90 RECEIVE events:  0
      Processors executing 10 RESET events:    0
      Processors executing 30 SEND events:     0
      Processors executing 40 SEND events:     1-3, 12-15
      Processors executing 50 SEND events:     4-11
      Processors executing 30 WAIT_ALL events: 0-15
```

stats is derived from interpret (see Section 3.3.6 [The interpret backend], page 28) and therefore supports the interpret backend's options as well as the standard options described in Section 3.4 [Running coNCePTuaL programs], page 46. However, because stats does not produce log files, the --logfile option is absent. stats additionally supports the following three command-line options:

```
  -E, --expand-lists=<number>     0=collapse lists of numbers into ranges;
                                  1=show all numbers [default: 0]
  -F, --format=<string>           Output format, either "text", "excelcsv",
                                  or "sep:<string>" [default: "text"]
  -X, --exclude=<string>          Name of a category or individual field to
                                  exclude from output [default: ""]
```

The `--expand-lists` option tells the `stats` backend whether it should output sets of numbers as comma-separated ranges (e.g., '`1-3, 12-15`') or as individual numbers (e.g., '`1, 2, 3, 12, 13, 14, 15`'). The `--format` option specifies the output format. By default, it outputs human-readable text, as shown in the above example. However, `--format=excelcsv` outputs the data in comma-separated value format suitable for loading directly into Microsoft Excel. A more general form of computer-parseable output is specified with the '`sep`' sub-option, which takes a separator string as a sub-option. For example, `--format=sep:"#"` produces output like the following:

```
"Event tallies"#"Total number of LOG events"#10
#"Total number of NEWSTMT events"#48
#"Total number of RECEIVE events"#710
#"Total number of RESET events"#10
#"Total number of SEND events"#710
#"Total number of WAIT_ALL events"#480
```

The difference between `--format=excelcsv` and `--format=sep:","` is that the former employs a number of bizarre special cases when quoting strings so as to coerce Excel into properly processing the file. The latter simply inserts a comma between columns with no special string processing other than preceding the characters '"' and '\' with '\'.

The `stats` backend issues a '`Column separator sep appears within a field`' warning if the separator string appears in any of the output fields. (Note that it will always appear in the '`Command line`' line because the `--format` line is specified on the command line.) The idea is to warn the user that automatic parsing of the output (e.g., using `awk`) may need to be careful when processing lines containing such fields.

As can be seen from the sample output, `stats` outputs a lot of information. The `--exclude` option—which can be specified repeatedly on the command line—provides the user with fine-grained control over which output to suppress. For example, `--exclude="Total number of NEWSTMT events"` tells `stats` not to output the line with that key. Similarly, `--exclude="Processor SEND-event peers"` eliminates an entire category of information.

3.3.8 The `picl` backend

PICL, the Portable Instrumented Communication Library, defines a trace file format that records an execution trace of a message-passing application. In particular, MPICL is an MPI-specific implementation of PICL that facilitates instrumenting MPI programs. Normally, one passes a PICL file to a performance-visualization tool such as ParaGraph, which renders the trace data using any of a number of graphical views.

Although MPICL is compatible with the `c_mpi` backend (see Section 3.3.2 [The c_mpi backend], page 22), cONCEPTUAL provides a `picl` backend that produces PICL files directly. The intention is to use the PICL format to represent an idealized view of a communication pattern rather than to store a time-sensitive execution trace. For example, timing events recorded by `picl` occur at logical times instead of physical times and statements such as `COMPUTES`/`SLEEPS` (see Section 4.8.2 [Delaying execution], page 134) and `TOUCHES` (see Section 4.8.3 [Touching memory], page 134) take either zero time or unit time, depending upon a command-line option. By providing an idealized view of a communication pattern, `picl` abstracts away timing artifacts so events that one would expect to occur simultaneously are written to the trace file as being exactly simultaneous.

As an example, the following is the PICL output produced by passing `picl` the CON-CEPTuaL program 'TASK 0 SENDS A 32 KILOBYTE MESSAGE TO TASK 1 THEN TASK 1 SENDS A 256 BYTE MESSAGE TO TASK 0.':

```
-3 -901 0.00000 0 0 0
-3 -901 0.00000 1 1 0
-3 -21 0.00001 0 0 4 2 32768 1 1 1
-4 -21 0.00001 0 0 0
-3 -52 0.00001 1 1 3 2 1 0 0
-3 -52 0.00002 0 0 3 2 1 1 1
-4 -52 0.00002 1 1 4 2 32768 1 0 0
-3 -21 0.00002 1 1 4 2 256 1 0 0
-4 -21 0.00002 1 1 0
-4 -52 0.00003 0 0 4 2 256 1 1 1
-4 -901 0.00003 1 1 0
-4 -901 0.00004 0 0 0
```

See the PICL manual, *A new PICL trace file format* (ORNL/TM-12125), for a detailed description of the various fields used in the preceding trace file.

ParaGraph is a visualization tool that reads PICL trace files and graphically displays/animates the corresponding execution in a variety of formats. Running ParaGraph on output from `picl` makes it easy for a CONCEPTuaL programmer to explain complex communication patterns to other people.

`picl` is derived from `interpret` (see Section 3.3.6 [The interpret backend], page 28) and therefore supports the `interpret` backend's `--tasks` and `--mcastsync` options as well as the standard options described in Section 3.4 [Running coNCePTuaL programs], page 46. However, because `picl` does not produce log files, the `--logfile` option is absent. `picl` additionally supports the following two command-line options:

```
-A, --all-events=<number>      0=include only communication events;
                               1=include all events  [default: 0]
-F, --frequency=<number>       PICL event frequency (Hz) [default: 100000]
```

By default, only communication events are written to the trace file. If `--all-events` is set to '1' then all events are written to the trace file. Events not related to communication are defined to take unit time. Regardless of the setting of `--all-events`, the OUTPUTS statement (see Section 4.5.2 [Writing to standard output], page 120) writes a PICL 'tracemsg' event to the trace file. A user can tell ParaGraph to pause visualization at 'tracemsg' events, making it possible to isolate key components of a CONCEPTuaL program's execution.

Because PICL events are marked with physical time (a floating-point number of seconds) but `picl` uses exclusively logical time, `picl` needs to associate a (fabricated) physical time with each logical time. The `--frequency` option specifies that mapping. By default, `picl` pretends that each unit of logical time corresponds to 1/100000 of a second (i.e, 10 microseconds) of physical time. As an example, the default time step in ParaGraph is 1 microsecond; this can be specified to `picl` with `--frequency=1E6`.

One of the goals of CONCEPTuaL is to facilitate the explanation of network performance tests. The `picl` backend aids in the explanation by making it easy to show graphically how tasks communicate with each other in an arbitrary CONCEPTuaL program.

3.3.9 The `paraver` backend

The Paraver trace visualizer accepts a trace file format that records an execution trace of a message-passing (or other) application and uses that to display application timelines and statistical data. Normally, one uses an application-instrumentation tool such as Extrae to generate Paraver traces, and Extrae is, in fact, compatible with programs generated using the coNCePTuaL `c_mpi` backend (see Section 3.3.2 [The c_mpi backend], page 22). However, the `paraver` backend produces Paraver files directly. The intention is to use the Paraver format to represent an idealized view of a communication pattern rather than to store a time-sensitive execution trace. For example, timing events recorded by `paraver` occur at logical times instead of physical times and statements such as COMPUTES/SLEEPS (see Section 4.8.2 [Delaying execution], page 134) and TOUCHES (see Section 4.8.3 [Touching memory], page 134) take either zero time or unit time, depending upon a command-line option. By providing an idealized view of a communication pattern, `paraver` abstracts away timing artifacts so events that one would expect to occur simultaneously are written to the trace file as being exactly simultaneous.

As an example, the following is the Paraver output produced by passing the `paraver` backend the coNCePTuaL program, 'TASK 0 SENDS A 32 KILOBYTE MESSAGE TO TASK 1 THEN TASK 1 SENDS A 256 BYTE MESSAGE TO TASK 0.':

```
#Paraver (06/01/2011 at 11:22):2001_ns:2(1,1):1:2(1:1,1:2),0
3:1:1:1:1:0:0:2:1:2:1:1000:1000:32768:0
2:1:1:1:1:0:1000000:6
2:1:1:1:1:0:1000000:13
2:2:1:2:1:0:1000000:6
2:2:1:2:1:0:1000000:9
1:1:1:1:1:0:1000:4
1:2:1:2:1:0:1000:3
3:2:1:2:1:1000:1000:1:1:1:1:2000:2000:256:0
2:1:1:1:1:1000:1000000:9
2:2:1:2:1:1000:1000000:13
1:1:1:1:1:1000:2000:3
1:2:1:2:1:1000:2000:4
```

The `paraver` backend also generates an associated configuration file that maps program state and event numbers to names and colors. For example, state 4 represents "blocking send" and is colored a neon fuchsia in Paraver's graphical interface; event 1000000 represents "coNCePTuaL," and event 1000000, value 13 represents a coNCePTuaL 'SEND' event. See the manual, *Paraver Parallel Program Visualization and Analysis Tool: Tracefile Description*, for a detailed description of the various fields used in the preceding trace file.

`paraver` is derived from `interpret` (see Section 3.3.6 [The interpret backend], page 28) and therefore supports the `interpret` backend's `--tasks` and `--mcastsync` options as well as the standard options described in Section 3.4 [Running coNCePTuaL programs], page 46. However, because `paraver` does not produce log files, the `--logfile` option is absent. `paraver` additionally supports the following two command-line options:

```
 -D, --dimemas-events=<number> 0=no Dimemas events; 1=extra events for
                               Dimemas simulator [default: 0]
 -E, --conc-events=<number>    0=exclude names of coNCePTuaL event types;
                               1=include them [default: 0]
 -O, --comp-time=<number>      Time spent in each non-communication event
                               (ns) [default: 0]
 -P, --event-time=<number>     Paraver event time (ns) [default: 1000]
 -R, --conc-source=<number>    0=exclude references to coNCePTuaL source
                               lines; 1=include them [default: 1]
```

By default, only CONCEPTUAL communication events are written to the trace file. If
--*comp-time* is set to a positive integer then all events are written to the trace file, and all
non-communication events take the specified amount of time. Because Paraver events are
marked with physical time (a integral number of nanoseconds) but `paraver` uses exclusively
logical time, `paraver` needs to associate a (fabricated) physical time with each logical time.
The --*event-time* option specifies that mapping. By default, `paraver` pretends that each
unit of logical time corresponds to 1 microsecond (i.e, 1000 nanoseconds) of physical time.
In addition to CONCEPTUAL, Paraver also has a notion of an "event." The --*conc-source*
and --*conc-events* options tell the `paraver` backend which Paraver events to include in the
trace file. If --*conc-source* is '1' (the default), Paraver will display text like the following
in its Info panel:

        User Event at 529,000 ns myfile.ncptl source line Unknown value 203

which indicates that the most recently clicked time in the Paraver timeline represents logical
time 529 and corresponds to line 203 of `myfile.ncptl`. If --*conc-events* is '1', Paraver
will display text like the following in its Info panel:

        User Event at 529,000 ns coNCePTuaL event RECEIVE

which indicates that the most recently clicked time in the Paraver timeline represents logical
time 529 and represents a CONCEPTUAL interpreter 'RECEIVE' event (see Section 3.3.6
[The interpret backend], page 28). The Dimemas network simulator accepts the same traces
as the Paraver visualizer but expects to see a number of additional events in the trace file.
If --*dimemas-events* is '1', the `paraver` backend will generate those extra events.

    One of the goals of CONCEPTUAL is to facilitate the explanation of network performance
tests. The `paraver` backend aids in the explanation by making it easy to show graphically
how tasks communicate with each other in an arbitrary CONCEPTUAL program.

3.3.10 The `latex_vis` backend

The `latex_vis` backend visualizes a program's communication pattern as an Encapsulated
PostScript time-space diagram. The backend thereby provides a static counterpart to the
animated visualizations made possible by the `picl` and `paraver` backends and the Para-
Graph and Paraver trace visualizers (see Section 3.3.8 [The picl backend], page 34, and
Section 3.3.9 [The paraver backend], page 36). `latex_vis` outputs LaTeX code with calls to
the PSTricks package (available from http://www.tug.org/applications/PSTricks/ but
shipped with most LaTeX distributions) to draw the figure, then runs *latex* (or whatever
the *LATEX* environment variable is set to) to convert the .tex file to DVI format and *dvips*
(or whatever the *DVIPS* environment variable is set to) to convert the .dvi file into EPS
format. Finally, `latex_vis` runs the graphic through Ghostscript (*gs* or whatever the *GS*

environment variable is set to) to tighten the graphic's bounding box.[3] `latex_vis` requires that LATEX and PSTricks be installed in order to produce its visualizations. However, it can still run with *--no-compile* (see Section 3.2 [Compiling coNCePTuaL programs], page 18) to produce a `.tex` file that can later be run manually through *latex*.

As an example, the following is the default `latex_vis` output from the coNCePTuaL program 'TASK 0 SENDS A 1 MEGABYTE MESSAGE TO TASK 1 THEN ALL TASKS SYNCHRONIZE THEN TASK 1 SENDS A 3 KILOBYTE MESSAGE TO TASK 0' when run with three tasks:

Currently, the output diagram does not indicate message size but this may change in a future release of `latex_vis`.

The `latex_vis` backend has a number of uses. First, it can be used to illustrate a nontraditional communication pattern for a presentation, research paper, or technical report. Second, it can be used as a teaching aid to demonstrate common communication patterns (e.g., a butterfly pattern) to students. Third, it can be used as a debugging aid to ensure that a coNCePTuaL program is, in fact, performing the expected communication operations.

`latex_vis` is derived from `interpret` (see Section 3.3.6 [The interpret backend], page 28) and therefore supports the `interpret` backend's *--tasks* and *--mcastsync* options as well as the standard options described in Section 3.4 [Running coNCePTuaL programs], page 46. (The `interpret` backend's *--kill-reps* option can be handy for reducing the picture's size by showing only a single iteration of each loop.) Because `latex_vis` does not produce log files, the *--logfile* option is absent. `latex_vis` additionally supports the following command-line options:

---

[3] If Ghostscript is not installed or fails to run, `latex_vis` issues a warning message, not an error message. The figure is still usable without a tight bounding box and a loose bounding box can be corrected manually by editing the '%%BoundingBox:' line in the generated EPS file.

```
-A, --annotate=<string>        Annotation level (0=no annotations;
                               1=annotate communication events;
                               2=annotate all events;
                               "<event>..."=annotate only the
                               specified events) [default: "0"]
-B, --binary-tasks=<number>    Display task numbers in binary rather
                               than decimal (0=decimal; 1=binary)
                               [default: 0]
-E, --every-event=<number>     Events requiring nonzero time to complete
                               (0=only communication events;
                               1=every event) [default: 0]
-G, --stagger=<number>         Number of points by which to stagger
                               overlapping arrows [default: 2]
-L, --source-lines=<number>    Associate source-code line numbers with
                               each event annotation (0=no; 1=yes)
                               [default: 0]
-R, --arrow-width=<string>     Python expression to map m, representing
                               a message size in bytes, to an arrow
                               width in points [default: "1"]
-Z, --zero-latency=<number>    Depict communication as having zero
                               latency (0=unit latency; 1=zero
                               latency) [default: 0]
```

In the preceding list a "point" refers to a PostScript point. TEX calls these "big points" and defines $1\,\mathrm{bp} \equiv 1/72$")

The `--annotate` option places adjacent to appropriate nodes in the output diagram a list of textual annotations that indicate the communication operations that were posted or completed and the non-communication operations that were executed by the corresponding task at the corresponding time. Some sample annotations are 'Post SEND', 'Complete SEND', and 'Execute OUTPUT'. With `--annotate=1`, only communication events are annotated; with `--annotate=2`, all events are annotated;[4]; otherwise, a list of specific events can be annotated. For example, `--annotate=SYNC,MCAST` causes only the 'SYNC' and 'MCAST' events to be annotated. Compiling a coNCePTuaL program with the `latex_vis` backend and the `--keep-ints` options (see Section 3.2 [Compiling coNCePTuaL programs], page 18) produces a `.tex` file that contains in the prologue comments a list of all events used by a program that were and were not annotated.

Task numbers are normally shown in base ten. The `--binary-tasks` option causes task numbers to be output in base two and using the same number of bits for each task. For example, a 3-task visualization with `--binary-tasks` would number the tasks '00', '01', and '11'.

Because `latex_vis` is intended primarily for visualizing communication patterns, by default only communication operations take nonzero time to complete. The `--every-event` option indicates that all events—including local events such as 'SLEEP', 'COMPUTE', 'LOG', and 'OUTPUT'—should be deemed to complete in unit time.

---

[4] To avoid the confusion of annotating what is essentially a "do nothing" event, `latex_vis` does not annotate the 'NEWSTMT' event, which is injected at the beginning of each top-level statement.

`latex_vis` output distinguishes coincident arrows (consider the phrase 'TASK 0 ASYNCHRONOUSLY SENDS 5 MESSAGES TO TASK 1') by staggering them slightly. The `--stagger` option specifies the number of PostScript points by which to stagger each overlapping arrow with a default of 2 points. Large stagger values are more visually distinctive while small stagger values allow arrows to drift less from their associated nodes.

The `--source-lines` option, when used with `--annotate`, augments each event annotation with the corresponding line(s) of source code that produced that event. This feature improves the utility of `latex_vis` as a debugging aid for coNCePTuaL programs.

By default, all arrows that indicate message transmissions are drawn with equally thick line widths. The `--arrow-width` option lets you specify a Python function to map a message size in bytes, m, to a line width in PostScript points, with the default being 1 point. As a simple example, `--arrow-width=2` doubles the arrow width for all message (i.e., it represents the constant function $f : m \mapsto 2$). A more typical example would be `--arrow-width="log10(m)"` (representing $f : m \mapsto \log_{10}(m)$), which causes a tenfold increase in message size to yield a unit increase in line width. The argument to `--arrow-width` can be any Python expression using either built-in functions or functions from the Python `math` module. See the Python library reference for details.

Normally, messages are considered to require one time unit to travel from source to destination. The `--zero-latency` option shows messages as being received in the same timestep as they were sent. Some communication patterns are more aesthetically pleasing when drawn this way.

Further customizations

In addition to the options described above, the front end's `--filter` option (see Section 3.2 [Compiling coNCePTuaL programs], page 18) is a useful mechanism for customizing the formatting of the communication diagram. For example, specifying `--filter="s/rowsep=30bp/rowsep=1.5in/g"` increases the separation between rows from 30 bp (where $1\,\text{bp} \equiv 1/72"$) to 1.5". See the PSTricks documentation for more information about the PSTricks commands used in the generated `.tex` files.

To facilitate the use of `--filter`, the `latex_vis` backend uses a helper macro (`\viscolor`) to define colors. `\viscolor` takes an argument of the form '*name=color*' and defines a macro `\name` `color` that expands to *color*. To further facilitate the use of `--filter`, the LaTeX code generated by the `latex_vis` backend contains a number of strategically placed placeholder comments of the form '`% PLACEHOLDER: tag`'. A `--filter` command can thereby insert code into the document by replacing an appropriate '`PLACEHOLDER`' line. In alphabetical order, the currently defined placeholder tags are `ANNOTATIONS`, `COLORS`, `COMMUNICATION`, `DEADLOCK`, `DOCUMENT`, `END`, `NODESHAPE`, `PACKAGES`, `PSMATRIX`, `TEXTOEPS`, and `TIMELINE`. Look through any `latex_vis`-generated `.tex` file to see where these are situated. As an example, the following command-line options define a "chartreuse" color then change the color used to indicate point-to-point messages from blue to chartreuse:

```
--filter="s/% PLACEHOLDER: COLORS/\\newrgbcolor{chartreuse}{0.5 1 0}/"
--filter="s/sendrecv=blue/sendrecv=chartreuse/"
```

(According to the PSTricks documentation, the predefined colors are `red`, `green`, `blue`, `cyan`, `magenta`, and `yellow`, and the predefined grayscales are `black`, `darkgray`, `gray`, `lightgray`, and `white`.)

Tasks that are active at a given time are drawn using the `\task` macro, which takes the task number as an argument. Tasks that are idle at a given time are drawn using the `\idle` macro, which is initially defined to be the same as `\task`. The following "cookbook" examples showcase some of the ways that the power of LaTeX and PSTricks can be exploited to display idle tasks in a variety of different styles (best used with `--annotate=2`):

omitting idle tasks

```
--filter="s/\\let\\idle=\\task/\\newcommand*{\\idle}[1]{[mnode=R]}/"
```

showing each idle task as a gray dot

```
--filter="s/\\let\\idle=\\task/\\newcommand*{\\idle}[1]{[mnode=dot,
linecolor=gray]}/"
```

drawing idle tasks with a dotted circle instead of a solid circle

```
--filter="s/\\let\\idle=\\task/\\newcommand*{\\idle}[1]
{[linestyle=dotted]\\task{#1}}/"
```

As an alternative to replacing the '`\let\idle=\task`' binding, the preceding substitutions can be expressed as the insertion of a LaTeX `\renewcommand`. That is, idle tasks can also be omitted by specifying '`--filter="s/% PLACEHOLDER: NODESHAPE/\\renewcommand*{\\idle}[1]{[mnode=R]}/"`'.

In short, the `latex_vis` backend produces highly customizable illustrations of communication patterns. Because `latex_vis` produces commented LaTeX code, any customization not provided through the use of `--filter` or one of the backend-specific command-line options is easily performed by compiling with `--keep-ints` (see Section 3.2 [Compiling coNCePTuaL programs], page 18) and editing the generated LaTeX code.

3.3.11 The `dot_ast` backend

dot is a format for describing graphs in terms of their edges and vertices. The tools in the Graphviz suite typeset dot files in a variety of output formats and using a variety of graph-layout algorithms. CoNCePTuaL's `dot_ast` backend outputs in dot format the abstract-syntax tree corresponding to a given CoNCePTuaL program. As an example, `dot_ast` renders the one-line CoNCePTuaL program 'TASK 0 SLEEPS FOR 10 SECONDS' as follows:

```
┌──────────────┐
│   program    │
├──────────────┤
│    line 1    │
└──────────────┘
        │
        ▼
┌────────────────────┬─────┐
│ top_level_stmt_list │ 1L │
├────────────────────┴─────┤
│          line 1          │
└──────────────────────────┘
        │
        ▼
┌──────────────────┐
│ top_level_stmt   │
├──────────────────┤
│      line 1      │
└──────────────────┘
        │
        ▼
┌─────────────────┬─────┐
│ simple_stmt_list │ 1L │
├─────────────────┴─────┤
│        line 1         │
└───────────────────────┘
        │
        ▼
┌──────────────┐
│ simple_stmt  │
├──────────────┤
│    line 1    │
└──────────────┘
        │
        ▼
┌──────────────┐
│  sleeps_for  │
├──────────────┤
│    line 1    │
└──────────────┘
```

1)   TASK 0 SLEEPS FOR 10 SECONDS

sleeps_for → source_task, expr, time_unit 'seconds'

source_task (line 1) → task_expr 'expr' (line 1) → expr (line 1) → ifelse_expr (line 1) → add_expr (line 1) → mult_expr (line 1) → unary_expr (line 1) → power_expr (line 1) → primary_expr (line 1) → integer 0L (line 1)

expr (line 1) → ifelse_expr (line 1) → add_expr (line 1) → mult_expr (line 1) → unary_expr (line 1) → power_expr (line 1) → primary_expr (line 1) → integer 10L (line 1)

time_unit 'seconds' (line 1)

dot_ast is expected to be of particular use to backend developers, who can use it to help prioritize the methods that need to be implemented (i.e., implementing first the AST node types needed by in a trivial program, then those needed by successively more complex programs).

The dot_ast backend accepts the following options from the ncptl command line:

`--format=dot_format`

> The programs in the Graphviz suite can output graphs in a variety of formats such as PostScript, SVG, and PNG. By default, the `dot_ast` backend outputs PostScript. The `--format` option specifies an alternate format to use. At the time of this writing, the Graphviz programs support the following formats: `canon`, `cmap`, `dot`, `fig`, `gd`, `gd2`, `gif`, `hpgl`, `imap`, `ismap`, `jpeg`, `jpg`, `mif`, `mp`, `pcl`, `pic`, `plain`, `plain-ext`, `png`, `ps`, `ps2`, `svg`, `svgz`, `vrml`, `vtx`, `wbmp`, and `xdot`. See the Graphviz documentation for more information about these formats.

`--node-code=characters`

> To facilitate associating nodes in the AST with fragments of the CONCEPTUAL program being graphed, the `dot_ast` backend provides a `--node-code` option that labels each node with the fragment of code to which it corresponds. The argument to `--node-code` is a number of characters at which to truncate the code fragment or '`-1`' to inhibit truncation. The purpose of truncation is to prevent excessively large nodes from disturbing the graph layout. The <*program*> node (see Section 4.10 [Complete programs], page 142), for example, includes the complete program source code if not truncated.

`--extra-dot=dot_code`

> The `dot_ast` backend's `--extra-dot` option enables the user to inject arbitrary dot code into the generated file. For example, specifying `--extra-dot="node [shape=Mrecord]"`[5] tells dot to use draw nodes as rounded rectangles and specifying `--extra-dot='edge [color="green"]'` colors all edges green. `--extra-dot` can be specified repeatedly on the command line; `dot_ast` concatenates all of the extra dot code with intervening semicolons.

`--compress`

> Normally, every node of the AST is drawn. The `--compress` option makes the resulting dot graph more readable by eliding chains of single-child nodes. For example, `--compress` reduces the previous graph of '`TASK 0 SLEEPS FOR 10 SECONDS`' to the following:

---

[5] `dot_ast` automatically places a semicolon after the extra dot code.
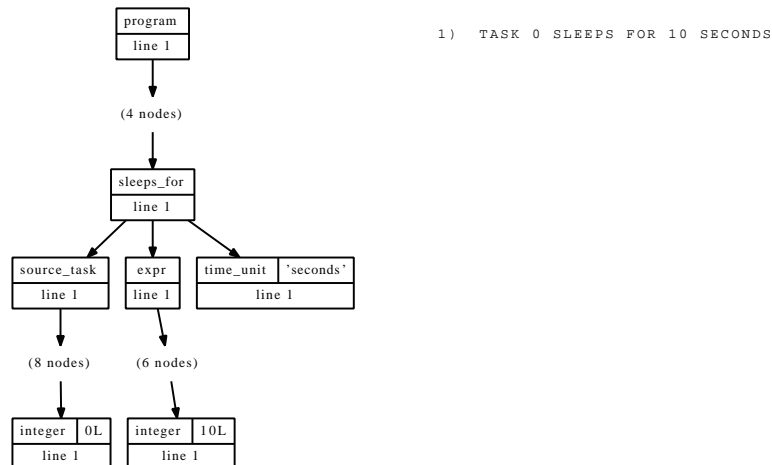
`--no-lines`

> By default, each AST node indicates the lines in the program's source code to which it corresponds. The `--no-lines` option suppresses the outputting of source-code line numbers.

`--no-attrs`

> Every node in the AST has a type. Some nodes additionally have an attribute. `dot_ast` normally outputs attributes but `--no-attrs` prevents `dot_ast` from doing so.

`--no-source`

> The complete source code corresponding to the AST is included in the generated dot graph unless `--no-source` is specified on the command line.

The *DOT* environment variable names the Graphviz program that `dot_ast` should run on the generated code. If *DOT* is not set, `dot_ast` uses whatever value was specified/discovered at configuration time (see Section 2.1 [configure], page 4), with the default being `dot`. By default, `dot_ast` produces dot code and runs this through the designated Graphviz program to produce a PostScript file (or whatever format is named by the `--format` option). If `ncptl` is run with either the `--no-link` or `--no-compile` options, it produces a dot file that should be run manually through `dot` or another Graphviz tool.

3.3.12 The `libsea_ast` backend

Like the `dot_ast` backend (see Section 3.3.11 [The dot_ast backend], page 41), the `libsea_ast` backend visualizes a CONCEPTUAL parse tree. This capability is primarily useful to backend developers to help prioritize the methods that need to be implemented (i.e., implementing first the AST node types needed by in a trivial program, then those needed by successively more complex programs).

The `libsea_ast` backend specifies a CONCEPTUAL program's parse tree in the LibSea graph file format, which can then be visualized and manipulated interactively using the Walrus visualization tool. The primary advantage of `libsea_ast` over `dot_ast` is that `libsea_ast` can handle much larger parse trees than `dot_ast`. For example, the following image illustrates the Walrus GUI displaying a 702,985-node CONCEPTUAL parse tree:[6]

---

[6] The original source program contained over 6000 lines of CONCEPTUAL code.
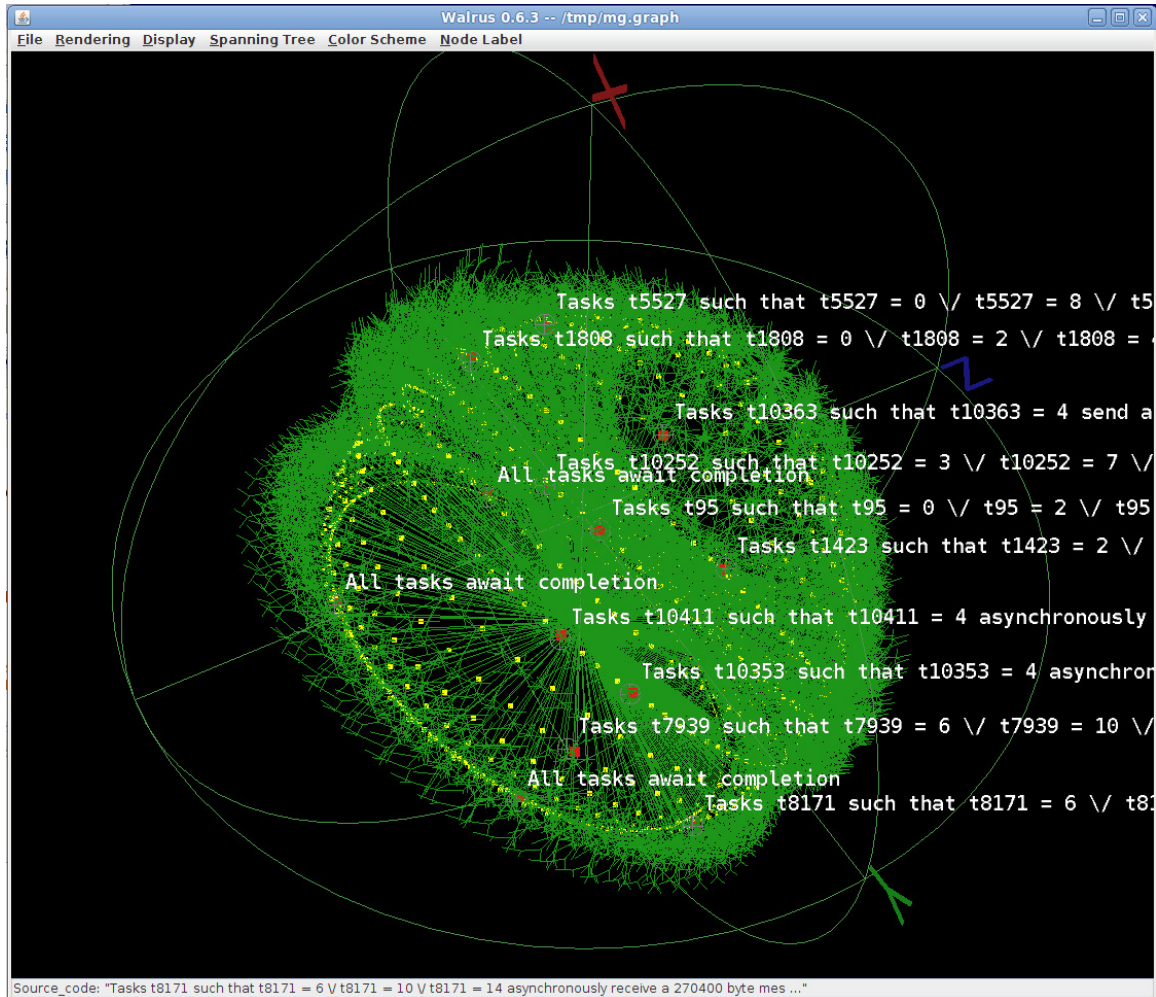
A secondary advantage of `libsea_ast` is that Walrus lets a user selectively view information. For example, in the preceding image, only <*simple_stmt*> nodes (see Section 4.7 [Complex statements], page 125) are highlighted; the source code for only certain nodes is displayed; and all node type and node attribute information is hidden.

The `libsea_ast` backend accepts the following option from the `ncptl` command line:

`--node-code=characters`

        Every node in the generated LibSea graph lists the fragment of CONCEPTUAL code to which it corresponds. For large CONCEPTUAL programs, these strings can get large. The <*program*> node (see Section 4.10 [Complete programs], page 142), for example, includes the complete program source code. Because the Walrus visualizer does not currently wrap long lines, most of a long string is never even presented to the user and therefore serves no practical purpose. The argument to `--node-code` is a number of characters at which to truncate the code fragment or '`-1`' to inhibit truncation.

3.4 Running coNCePTuaL programs

CoNCePTuaL programs can be run like any other program built with the same compiler and communication library. For example if a program `myprog` was built with CoNCePTuaL's C+MPI backend, the program might be run with a command like *mpirun -np nodes myprog* or *prun -Nnodes myprog* or *pdsh -w node_list myprog*. The important point is that job launching is external to CoNCePTuaL. A CoNCePTuaL program is oblivious to whether it is being run with a single thread on each multiprocessor node or with one thread on each CPU, for example. However, CoNCePTuaL log files do include the host name in the prologue comments (see Section 3.5.1 [Log-file format], page 49) so job-launching parameters can potentially be inferred from those.

CoNCePTuaL programs automatically support a "help" option. This is usually specified as *--help* or *-?*, depending on which option-parsing library `configure` configured in. (See Section 2.1 [configure], page 4.) The output of running *myprog --help* most likely looks something like this:

```
Usage: myprog [OPTION...]
  -C, --comment=<string>       Additional commentary to write to the log
                               file, @FILE to import commentary from FILE,
                               or !COMMAND to import commentary from COMMAND
                               (may be specified repeatedly)
  -L, --logfile=<string>       Log-file template [default: "a.out-%p.log"]
  -N, --no-trap=<string>       List of signals that should not be trapped
                               [default: ""]
  -S, --seed=<number>          Seed for the random-number generator
                               [default: 0]


Help options:
  -?, --help                   Show this help message
  --usage                      Display brief usage message
```

Although a CoNCePTuaL program can specify its own command-line options (see Section 4.9.2 [Command-line arguments], page 139), a few are provided by default. In addition to *--help* these include *--comment*, *--logfile*, *--no-trap*, and *--seed*.

--comment

> *--comment* makes it possible to add arbitrary commentary to a log file. This is useful for incorporating information that CoNCePTuaL would be unable to (or simply does not currently) determine on its own, for example, *--comment="Last experiment before upgrading the network device driver"*. Two special cases are supported:
>
> 1. If the comment string begins with '@' then the remainder of the string is treated as a filename. Each line of the corresponding file is treated as a separate comment string. Hence, if the file `sysdesc.txt` contains the lines 'Using FooBarNet' and 'Quux is enabled', then specifying *--comment=sysdesc.txt* is similar to specifying both *--comment="Using FooBarNet"* and *--comment="Quux is enabled"*.

2. If the comment string begins with '!' then the remainder of the string
   is treated as a shell command. The command is executed and each line
   of its output is treated as a separate comment string. For example,
   --comment='!lspci | grep -i net' executes `lspci`, extracts only
   those lines containing the string 'net', and makes log-file comments
   out of the result. Note that *--comment='!command'* differs from
   --comment="'command'" in that the former causes *command* to be
   executed individually by each process in the program while the latter
   executes *command* only once and only before launching the program.
   Also note that '!' must be escaped in `csh` and derivitive shells (i.e.,
   --comment='\!command').

As an example, the command line arguments *--comment="This is
a simple comment."* *--comment=!lspci* *--comment=@/proc/version*
--comment="This is another simple comment." produce log-file lines like
the following:

```
# User comment 1: This is a simple comment.
# Output of 'lspci', line 1: 00:00.0 Host bridge: Intel Corp. 82860 860 (Wombat) Chipset Host
# Output of 'lspci', line 2: 00:01.0 PCI bridge: Intel Corp. 82850 850 (Tehama) Chipset AGP E
# Output of 'lspci', line 3: 00:02.0 PCI bridge: Intel Corp. 82860 860 (Wombat) Chipset AGP E
# Output of 'lspci', line 4: 00:1e.0 PCI bridge: Intel Corp. 82801BA/CA/DB PCI Bridge (rev 04
# Output of 'lspci', line 5: 00:1f.0 ISA bridge: Intel Corp. 82801BA ISA Bridge (LPC) (rev 04
# Output of 'lspci', line 6: 00:1f.1 IDE interface: Intel Corp. 82801BA IDE U100 (rev 04)
# Output of 'lspci', line 7: 00:1f.2 USB Controller: Intel Corp. 82801BA/BAM USB (Hub  (rev 0
# Output of 'lspci', line 8: 00:1f.3 SMBus: Intel Corp. 82801BA/BAM SMBus (rev 04)
# Output of 'lspci', line 9: 00:1f.4 USB Controller: Intel Corp. 82801BA/BAM USB (Hub  (rev 0
# Output of 'lspci', line 10: 00:1f.5 Multimedia audio controller: Intel Corp. 82801BA/BAM AC
# Output of 'lspci', line 11: 01:00.0 VGA compatible controller: nVidia Corporation NV11 [GeF
# Output of 'lspci', line 12: 02:1f.0 PCI bridge: Intel Corp. 82806AA PCI64 Hub PCI Bridge (r
# Output of 'lspci', line 13: 03:00.0 PIC: Intel Corp. 82806AA PCI64 Hub Advanced Programmabl
# Output of 'lspci', line 14: 04:0b.0 Ethernet controller: 3Com Corporation 3c905C-TX/TX-M [T
# Output of 'lspci', line 15: 04:0d.0 Multimedia audio controller: Creative Labs SB Live! EMU
# Output of 'lspci', line 16: 04:0d.1 Input device controller: Creative Labs SB Live! MIDI/Ga
# Contents of /proc/version, line 1: Linux version 2.4.20-28.7 (bhcompile@porky.devel.redhat.
# User comment 2: This is another simple comment.
```

To facilitate log-file parsing, all colons in the name of an '@'-file or '!'-command
are written as periods. This affects only the display, not the file to read or
command to execute.

Be careful when using shell backquotes with *--comment* (e.g., as in
--comment="'who'"). Different shells have different ways of handling
newlines output by a backquoted command. In some shells (e.g., `bash`), the
CONCEPTUAL program sees the entire output with embedded newlines as
a single argument; in others (e.g., `tcsh`), each line of output constitutes a
separate command-line argument. Because CONCEPTUAL programs currently
ignore arguments that do not begin with '--', the comment written to the log
file in the latter case terminates at the first newline. In virtually all shells,
the double quotes around the backquoted command are needed to prevent
the shell from splitting arguments at word boundaries. As a consequence,
--comment='who' logs only the first word output by the *who* command.

`--logfile`

> `--logfile` specifies a template for naming log files. Each task maintains a log file based on the template name but with '`%p`' replaced with the processor number, '`%r`' replaced with the run number (the smallest nonnegative integer that produces a filename that does not already exist), and '`%%`' replaced with a literal "%" character. The program outputs an error message and aborts if the log-file template does not contain at least one '`%p`'. The only exception is that an empty template (i.e., `--logfile=""`) inhibits the production of log files entirely.
>
> Like C's `printf()` function, a numeric field width can occur between the '`%`' and the conversion specifier (the '`p`' or '`r`' in the case of `--logfile`). The field is padded on the left with spaces to the given width. More practically, if the field width begins with the number '`0`' the field is padded on the left with zeroes to the given width. For example, specifying `--logfile="mydata-%03p.log"` on the command line produces log files named `mydata-000.log`, `mydata-001.log`, `mydata-002.log`, `mydata-003.log`, and so forth.

`--no-trap`

> `--no-trap` specifies a list of signals or ranges of signals that should not be trapped. For example, `--no-trap=10-12,17` prevents signals 10, 11, 12, and 17 from being trapped.[7] Signals can also be referred to by name, with or without a '`SIG`' prefix. Also, names and numbers can be freely mixed. Hence, `--no-trap=10-12,INT,17,SIGSTOP,SIGCONT` is a valid argument to a CoN-CePTuaL program. Because signal reception can adversely affect performance, coNCePTuaL's default behavior is to terminate the program on receipt of a signal. However, some signals may be necessary for the underlying communication layer's proper operation. `--no-trap` enables such signals to pass through coNCePTuaL untouched. (Some signals, however, are needed by coNCeP-TuaL or by a particular backend and are always trapped.)

`--seed`   `--seed` (which selects a different default value on each run) is used in any program that utilizes the `RANDOM TASK` construct (see Section 4.7.3 [Binding variables], page 130) or that sends message `WITH VERIFICATION` (see Section 4.4.1 [Message specifications], page 108).

If the `LOGS` statement (see Section 4.5.3 [Writing to a log file], page 120) is used anywhere in a coNCePTuaL program, then *all* processes write a log file. This is done because coNCePTuaL log files—even those that contain no measurement data—include a wealth of important information in prologue comments, as described Section 3.5 [Interpreting coN-CePTuaL log files], page 49. As a consequence, a program run with thousands of processes produces thousands of log files. If process 0 is the only process that logs actual data, the `ncptl-logmerge` script (see Section 3.5.3 [ncptl-logmerge], page 76) can merge these log files into a single, more manageable, file. For situations in which it is unreasonable for every process to write a log file (e.g., if the filesystem is unable to handle large numbers of simultaneous file creations, the `NCPTL_LOG_ONLY` environment variable lets the user limit the set of processes that produce log files. `NCPTL_LOG_ONLY` accepts a comma-separated

---

[7] On some platforms, these signals correspond to `SIGUSR1`, `SIGSEGV`, `SIGUSR2`, and `SIGCHLD`, respectively.

list of dash-separated process ranges such as '0-3,12-16,24,25,32-48'. Only processes included in the list produce log files.

3.5 Interpreting coNCePTuaL log files

Any CoNCePTuaL program that uses the `LOGS` keyword (see Section 4.5.3 [Writing to a log file], page 120) will produce a log file as it runs. The CoNCePTuaL run-time library writes log files in a simple, plain-text format. In addition to measurement data, a wealth of information is stored within log-file comments. CoNCePTuaL comes with a tool, `ncptl-logextract`, which can extract data and other information from a log file and convert it into any of a variety of other formats.

3.5.1 Log-file format

The CoNCePTuaL run-time library writes log files in the following (textual) format:

- Lines beginning with '#' are comments.

- Columns are separated by commas.

- Strings are output between double quotes. Literal double-quotes are output as '\"' and literal backslashes are output as '\\'.

A sample log file is listed below. The log file is presented in its entirety.

```
##############################################################################
# ===================
# coNCePTuaL log file
# ===================
# coNCePTuaL version: 0.6.4a
# coNCePTuaL backend: c_mpi (C + MPI)
# Executable name: /home/pakin/src/coNCePTuaL/example
# Working directory: /home/pakin/src/coNCePTuaL
# Command line: ./example
# Number of tasks: 2
# Processor (0<=P<tasks): 0
# Host name: a1
# Operating system: Linux version 2.4.21-3.5qsnet (root@a31) (gcc version 2.96 20000731 (Red Hat Linux 7.
# CPU vendor: GenuineIntel
# CPU architecture: ia64
# CPU count: 2
# CPU frequency: 1300000000 Hz (1.3 GHz)
# Cycle-counter frequency: 1300000000 Hz (1.3 GHz)
# OS page size: 16384 bytes
# Physical memory: 2047901696 bytes (1.9 GB)
# Elan capability: [1965771c.6213bf5d.3a26411c.5b4c5d58] Version 10002 Type a001 Context 640.640.640 Node
# coNCePTuaL configuration: ./configure '--prefix=/tmp/ncptl' 'MPICPPFLAGS=-I/usr/local/include' 'CFLAGS=
# Library compiler+linker: /opt/intel-7.1.033/compiler70/ia64/bin/ecc
# Library compiler version: Intel(R) C++ gcc 3.0 mode [7.1]
# Library compiler options: -g -O3 -ansi_alias -ansi
# Library linker options: -lrmscall -lelan -lpopt
# Library compiler mode: LP64
# Dynamic libraries used: /usr/lib/qsnet/elan3/lib/librmscall.so.1 /usr/lib/qsnet/elan3/lib/libelan.so.1
# Microsecond timer type: gettimeofday()
# Average microsecond timer overhead: <1 microsecond
# Microsecond timer increment: 1.00466 +/- 0.123576 microseconds (ideal: 1 +/- 0)
# Minimum sleep time: 1946.6 +/- 31.5872 microseconds (ideal: 1 +/- 0)
# WARNING: Sleeping exhibits poor granularity (not a serious problem).
# WARNING: Sleeping has a large error component (not a serious problem).
```

```
# Process CPU timer: getrusage()
# Process CPU-time increment: 976.59 +/- 0.494311 microseconds (ideal: 1 +/- 0)
# WARNING: Process timer exhibits poor granularity (not a serious problem).
# Log-file template: example-%p.log
# Number of minutes after which to kill the job (-1=never): -1
# List of signals which should not be trapped: 14
# Log-file checkpointing interval: 60 seconds (i.e., 1 minute)
# MPI send routine: MPI_Send()
# MPI error checking: off
# Front-end compilation command line: ncptl --backend=c_mpi example.ncptl
# Back-end compilation command line: /usr/lib/mpi/mpi_intel/bin/mpicc -I/tmp/ncptl/include  -I/usr/local/
# Log creator: Scott Pakin
# Log creation time: Mon Dec 19 12:02:18 2005
#
# Environment variables
# ---------------------
# CVS_RSH: /usr/bin/ssh
# DISPLAY: localhost:19.0
# DYNINSTAPI_RT_LIB: /home/pakin/dyninstAPI-3.0/lib/i386-unknown-linux2.2/libdyninstAPI_RT.so.1
# DYNINST_ROOT: /home/pakin/dyninstAPI-3.0
# EDITOR: /usr/bin/emacs
# GROUP: CCS3
# HOME: /home/pakin
# HOST: a0
# HOSTNAME: a0
# HOSTTYPE: unknown
# KDEDIR: /usr
# LANG: en_US
# LD_LIBRARY_PATH: /opt/intel-7.1.033/compiler70/ia64/lib:/users/pakin/lib:/usr/lib:/usr/ccs/lib:/opt/SUN
# LESSOPEN: |/usr/bin/lesspipe.sh %s
# LOGNAME: pakin
# LPDEST: lwy
# LS_COLORS: no=00:fi=00:di=01;34:ln=01;36:pi=40;33:so=01;35:bd=40;33;01:cd=40;33;01:or=01;05;37;41:mi=01
# MACHTYPE: unknown
# MAIL: /var/mail/pakin
# MANPATH: /opt/intel-7.1.033/compiler70/man:/home/pakin/man:/usr/man:/opt/SUNWspro/man:/usr/dt/man:/usr/
# MOZILLA_HOME: /usr/local/netscape/java
# NAME: Scott Pakin
# ORGANIZATION: Los Alamos National Lab
# OSTYPE: linux
# PATH: /opt/intel-7.1.033/compiler70/ia64/bin:.:/home/pakin/bin:/usr/local/bin:/usr/dt/bin:/usr/openwin/
# PRINTER: lwy
# PVM_ROOT: /usr/share/pvm3
# PVM_RSH: /usr/bin/rsh
# PWD: /home/pakin/src/coNCePTuaL
# QTDIR: /usr/lib/qt-2.3.1
# REMOTEHOST: antero.c3.lanl.gov
# RMS_JOBID: 308
# RMS_MACHINE: a
# RMS_NNODES: 2
# RMS_NODEID: 0
# RMS_NPROCS: 2
# RMS_PROCID: 0
# RMS_RANK: 0
# RMS_RESOURCEID: parallel.318
# RMS_STOPONELANINIT: 0
# SHELL: /bin/tcsh
# SHLVL: 2
# SSH_AGENT_PID: 8586
```

```
# SSH_ASKPASS: /usr/libexec/openssh/gnome-ssh-askpass
# SSH_AUTH_SOCK: /tmp/ssh-XXfG457q/agent.8562
# SSH_CLIENT: 128.165.20.177 47456 22
# SSH_TTY: /dev/pts/10
# SUPPORTED: en_US:en
# TERM: xterm
# TZ: MST7MDT
# USER: pakin
# VENDOR: unknown
#
# coNCePTuaL source code
# ---------------------
#     FOR 10 REPETITIONS {
#        TASK 0 RESETS ITS COUNTERS THEN
#        TASK 0 SENDS A 0 BYTE MESSAGE TO TASK 1 THEN
#        TASK 1 SENDS A 0 BYTE MESSAGE TO TASK 0 THEN
#        TASK 0 LOGS EACH elapsed_usecs/2 AS "Latency (usecs)"
#     }
#
################################################################################
"Latency (usecs)"
"(all data)"
192.5
7
5.5
5.5
5
5
5
5.5
5
5.5
################################################################################
# Program exited normally.
# Log completion time: Mon Dec 19 12:02:18 2005
# Elapsed time: 0 seconds
# Process CPU usage (user+system): 0 seconds
# Number of interrupts received (all CPUs): 22
# Peak memory allocation: 1072987 bytes (1.0 MB)
################################################################################
```

As the preceding example indicates, a log file's comment block can be divided into multiple stanzas:

- a list of $\langle key{:}value \rangle$ pairs that describe various characteristics of the run-time environment, including hardware and software identification, timer quality, values of command-line arguments, and a timestamp
- a dump of the environment variables active when the program ran
- the complete coNCePTuaL source program

Two rows of column headers and the measurement data follow the prologue comment block. A set of epilogue comments completes the log file.

The motivation for writing so much information to the log file is to facilitate reproduction of the experiment. The ideal situation is for a third party to be able to look at a coNCePTuaL log file and from that, recreate the original experiment and get identical results.

Some of the comments that may benefit from additional explanation include the following:

'Library compiler mode'

'LP64' means "**L**ong integers and **P**ointers contain exactly **64** bits while ordinary integers contain exactly 32 bits". 'ILP32' means "ordinary **I**ntegers, **L**ong integers, and **P**ointers all contain exactly **32** bits". The library compiler mode will be 'nonstandard' for any other combination of datatype sizes.

'Average timer overhead'

During initialization, the coNCePTuaL run-time library performs some calibration routines. Among these routines is a measurement of the quality of whatever mechanisms the library is using to measure elapsed time. In the sample log file presented above, the mechanism used was 'inline assembly code', meaning the run-time library reads the hardware cycle counter without going through a standard library or system call. The complete set of timer mechanisms and selection criteria is presented in Section 6.3.5 [Time-related functions], page 170, in the documentation for the ncptl_time() function.

The log file then reports "average timer overhead" as the mean time between back-to-back invocations of whichever timer routine is being used. Ideally, the mean should be '<1 microsecond' but this is not the case on all systems. Large values indicate that a performance penalty is charged every time a coNCePTuaL program reads the timer.

'Timer increment'

In addition to measuring the overhead of reading the timer, the coNCePTuaL run-time library also measures timer accuracy. The library expects to be able to read the timer with microsecond accuracy. That is, the time reported should not increase by more than a microsecond between successive readings of the timer. To gauge timer accuracy, the library's initialization routine performs a number of back-to-back invocations of the timer routine and reports the mean and standard deviation of the number of microseconds that elapsed between readings, discarding any deltas of zero microseconds. Ideally, the microsecond timer, when read multiple times in rapid succession, should report nonzero increments of exactly one microsecond with no variation. The log file will contain warning messages if the increment or standard deviation are excessively large, as this may indicate a large margin of error in the measurement data.

'Process CPU-time increment'
'Process CPU usage (user+system)'

Log files end with an epilogue section that includes 'Process CPU usage (user+system)', which indicates the subset of total wall-clock time for which the program was running ('user') or for which the operating system was running on the program's behalf ('system'). The log-file prologue reports as 'Process CPU-time increment' the resolution of the timer user to report process CPU time. Note that process CPU time is not exported to coNCePTuaL programs; it is therefore much less critical than the wall-clock timer and is reported primarily for informational purposes.

'`Number of interrupts received (all CPUs)`'

On certain platforms, CoNCePTuaL can tally the number of CPU interrupts that were processed during the run of the program. Because a multiprocessor may migrate tasks among CPUs during their execution, a per-CPU interrupt count may have little merit. Consequently, the number reported represents the sum across all CPUs in the same node (but not across nodes). CoNCePTuaL attempts to read interrupt information from the `/proc/interrupts` file if no alternative mechanism is available. In case having a large number of processes accessing `/proc/interrupts` poses a problem the `--disable-proc-interrupts` configuration option prevents programs from accessing that file.

'`Peak memory allocation`'

CoNCePTuaL programs heap-allocate memory for a variety of purposes: message buffers, event lists (see Section 6.2.3 [Generated code], page 162), unaggregated performance data (see [Computing aggregates], page 122), etc. Allocated memory that is no longer needed is returned to the heap. The total amount of allocated memory the program is holding therefore increases and decreases over time. The '`Peak memory allocation`' comment reports the maximum amount of memory the program held at any given time. If this value nears or exceeds the value reported for '`Physical memory`', it is possible that paging overheads may be negatively impacting some of the program's timing measurements. On systems without demand paging, exceeding '`Physical memory`' is likely to crash the program; a large '`Peak memory allocation`' on a smaller run can therefore help explain why a larger run is crashing. '`Peak memory allocation`' includes only memory allocated using the memory-allocation functions described in Section 6.3.3 [Memory-allocation functions], page 168.

3.5.2 `ncptl-logextract`

To facilitate converting CoNCePTuaL log files into input data for other applications, CoNCePTuaL provides a Perl script called `ncptl-logextract`. `ncptl-logextract` can extract the data from a log file—as well as various information that appears in the log file's comments—into a variety of formats suitable for graphing or typesetting.

Running `ncptl-logextract --usage` causes `ncptl-logextract` to list a synopsis of its core command-line options to the standard output device; running `ncptl-logextract --help` produces basic usage information; and, running `ncptl-logextract --man` outputs a complete manual page. See [ncptl-logextract manual page], page 54, shows the `ncptl-logextract` documentation as produced by `ncptl-logextract --man`.

NAME

ncptl-logextract - Extract various bits of information from a coNCePTuaL log file

SYNOPSIS

ncptl-logextract *--usage* | *--help* | *--man*

ncptl-logextract [*--extract*=[data|params|env|source|warnings]] [*--format=format*] [*format-specific options...*] [*--before=string*] [*--after=string*] [*--force-merge*[=*number*]] [*--procs=string*] [*--quiet*] [*--verbose*] [*--output=filename*] [*filename...*]

DESCRIPTION

Background   coNCePTuaL is a domain-specific programming language designed to facilitate writing networking benchmarks and validation suites. coNCePTuaL programs can log data to a file but in only a single file format. `ncptl-logextract` extracts this log data and outputs it in a variety of formats for use with other applications.

The coNCePTuaL-generated log files that serve as input to `ncptl-logextract` are plain ASCII files. Syntactically, they contain a number of newline-separated tables. Each table contains a number of newline-separated rows of comma-separated columns. This is known generically as *comma-separated value* or *CSV* format. Each table begins with two rows of header text followed by one or more rows of numbers. Text is written within double quotes. Double-quote characters and backslashes within text are escaped with a backslash. No other escaped characters are recognized. Lines that begin with `#` are considered comments.

Semantically, there are four types of data present in every coNCePTuaL-generated log file:

1. The complete source code of the coNCePTuaL program that produced the log file

2. Characteristics of the run-time environment and the values of all command-line parameters

3. A list of warning messages that coNCePTuaL issued while analyzing the run-time environment

4. One or more tables of measurement data produced by the coNCePTuaL program

The first three items appear within comment lines. The measurement data is written in CSV format.

Extracting information from coNCePTuaL log files   It is common to want to extract information (especially measurement data) from log files. For simple formatting operations, a one-line awk or Perl script suffices. However, as the complexity of the formatting increases, the complexity of these scripts increases even more. That's where `ncptl-logextract` fits in. `ncptl-logextract` makes it easy to extract any of the four types of log data described above and format it in variety of ways. Although the number of options that `ncptl-logextract` supports may be somewhat daunting, it is well worth learning how to use `ncptl-logextract` to avoid reinventing the wheel every time a coNCePTuaL log file needs to be processed. `ncptl-logextract` takes care of all sorts of special cases that crop up when manipulating coNCePTuaL log files.

OPTIONS

`ncptl-logextract` accepts the following command-line options regardless of what data is extracted from the log file and what formatting occurs:

-h , **--help**
> Output the Synopsis section and the Options section then exit the program.

-m , **--man**
> Output a complete Unix man ("manual") page for `ncptl-logextract` then exit the program.

-e *info*, **--extract** =*info*
> Specify what sort of data should be extracted from the log file. Acceptable values for *info* are listed and described in the Additional Options section and include `data`, `params`, `env`, and `source`.

-f *format*, **--format** =*format*
> Specify how the extracted data should be formatted. Valid arguments depend upon the value passed to **--extract** and include such formats as `csv`, `html`, `latex`, `text`, and `bash`. See the Additional Options section for details, explanations, and descriptions of applicability.

-b *string*, **--before** =*string*
> Output an arbitrary string of text before any other output. *string* can contain escape characters such as \n for newline, \t for tab, and \\ for backslash.

-a *string*, **--after** =*string*
> Output an arbitrary string of text after all other output. *string* can contain escape characters such as \n for newline, \t for tab, and \\ for backslash.

-F [*number*], **--force-merge** [=*number*]
> Try extra hard to merge multiple log files, even if they seem to have been produced by different programs or in different execution environments. This generally implies padding empty rows and columns with blanks. However, if **--force-merge** is given a numeric argument, the value of that argument is used instead of blanks to pad empty locations. Note that **--force-merge** is different from **--force-merge=0** because data-merging functions (`mean`, `max`, etc.) ignore blanks but consider zeroes.

-p *string*, **--procs** =*string*
> When given a "merged" log file, unmerge only the data corresponding to the comma-separated processor ranges in *string*. For example, **--procs=0,16-20,25** unmerges the data for processors 0, 16, 17, 18, 19, 20, and 25. By default, `ncptl-logextract` uses all of the data from a merged log file.

-q , **--quiet**
> Suppress progress output. Normally, `ncptl-logextract` outputs status information regarding its operation. The **--quiet** option instruct `ncptl-logextract` to output only warning and error messages.

-v , *--verbose*

> Increase progress output. Normally, `ncptl-logextract` outputs basic status information regarding its operation. The *--verbose* option instruct `ncptl-logextract` to output more detailed information. Each time *--verbose* is specified, the program's verbosity increases (up to a maximum).

-o *filename*, *--output* =*filename*

> Redirect the output from `ncptl-logextract` to a file. By default, `ncptl-logextract` writes to the standard output device.

The above is merely a terse summary of the `ncptl-logextract` command-line options. The reader is directed to the Additional Options section for descriptions of the numerous ways that `ncptl-logextract` can format information. Note that *--extract* and *--format* are the two most common options as they specify what to extract and how to format it; most of the remaining options in the Additional Options section exist to provide precise control over formatting details.

ADDITIONAL OPTIONS

The `ncptl-logextract` command-line options follow a hierarchy. At the top level is *--extract*, which specifies which of the four types of data `ncptl-logextract` should extract. Next, *--format* specifies how the extracted data should be formatted. Valid values for *--format* differ based on the argument to *--extract*. Finally, there are various format-specific options that fine-tune the formatted output. Each output format accepts a different set of options. Many of the options appear at multiple places within the hierarchy, although usually with different default values.

The following hierarchical list describes all of the valid combinations of *--extract*, *--format*, and the various format-specific options:

--extract=data [default]

> Extract measurement data

>> *--format=csv [default]*

>>> Output each table in comma-separated-value format

>>> *--noheaders*

>>>> Do not output column headers

>>> *--colbegin=string*

>>>> Specify the text placed at the beginning of each data column [default: ""]

>>> *--colsep=string*

>>>> Specify the text used to separate data columns [default: ","]

>>> *--colend=string*

>>>> Specify the text placed at the end of each data column [default: ""]

>>> *--rowbegin=string*

>>>> Specify the text placed at the beginning of each data row [default: ""]

--rowsep=string
> Specify the text used to separate data rows [default: ""]

--rowend=string
> Specify the text placed at the end of each data row [default: "\\n"]

--hcolbegin=string
> Specify the text placed at the beginning of each header column [default: same as `colbegin`]

--hcolsep=string
> Specify the text used to separate header columns [default: same as `colsep`]

--hcolend=string
> Specify the text placed at the end of each header column [default: same as `colend`]

--hrowbegin=string
> Specify the text placed at the beginning of each header row [default: same as `rowbegin`]

--hrowsep=string
> Specify the text used to separate header rows [default: same as `rowsep`]

--hrowend=string
> Specify the text placed at the end of each header row [default: same as `rowend`]

--tablebegin=string
> Specify the text placed at the beginning of each table [default: ""]

--tablesep=string
> Specify the text used to separate tables [default: "\\n"]

--tableend=string
> Specify the text placed at the end of each table [default: ""]

--quote=string
> Specify the text used to begin quoted text [default: """]

--unquote=string
> Specify the text used to end quoted text [default: same as `quote`]

--excel   Output strings in a format readable by Microsoft Excel

--keep-columns=string
> Enumerate the columns that should be included in the output [default: all columns]

--merge=function

> Specify how to merge data from multiple files [default: "mean"]

--showfnames=option

> Add an extra header row showing the filename the data came from [default: "none"]

--format=tsv

> Output each table in tab-separated-value format

--noheaders

> Do not output column headers

--colbegin=string

> Specify the text placed at the beginning of each data column [default: ""]

--colsep=string

> Specify the text used to separate data columns [default: "\\t"]

--colend=string

> Specify the text placed at the end of each data column [default: ""]

--rowbegin=string

> Specify the text placed at the beginning of each data row [default: ""]

--rowsep=string

> Specify the text used to separate data rows [default: ""]

--rowend=string

> Specify the text placed at the end of each data row [default: "\\n"]

--hcolbegin=string

> Specify the text placed at the beginning of each header column [default: same as colbegin]

--hcolsep=string

> Specify the text used to separate header columns [default: same as colsep]

--hcolend=string

> Specify the text placed at the end of each header column [default: same as colend]

--hrowbegin=string

> Specify the text placed at the beginning of each header row [default: same as rowbegin]

--hrowsep=string
> Specify the text used to separate header rows [default: same as `rowsep`]

--hrowend=string
> Specify the text placed at the end of each header row [default: same as `rowend`]

--tablebegin=string
> Specify the text placed at the beginning of each table [default: ""]

--tablesep=string
> Specify the text used to separate tables [default: "`\\n`"]

--tableend=string
> Specify the text placed at the end of each table [default: ""]

--quote=string
> Specify the text used to begin quoted text [default: "`"`"]

--unquote=string
> Specify the text used to end quoted text [default: same as `quote`]

--excel   Output strings in a format readable by Microsoft Excel

--keep-columns=string
> Enumerate the columns that should be included in the output [default: all columns]

--merge=function
> Specify how to merge data from multiple files [default: "`mean`"]

--showfnames=option
> Add an extra header row showing the filename the data came from [default: "`none`"]

--format=html
> Output each table in HTML table format

--noheaders
> Do not output column headers

--colbegin=string
> Specify the text placed at the beginning of each data column [default: "`<td>`"]

--colsep=string
> Specify the text used to separate data columns [default: " "]

--colend=string
        Specify the text placed at the end of each data column
        [default: "`</td>`"]

--rowbegin=string
        Specify the text placed at the beginning of each data
        row [default: "`<tr>`"]

--rowsep=string
        Specify the text used to separate data rows [default:
        ""]

--rowend=string
        Specify the text placed at the end of each data row
        [default: "`</tr>\\n`"]

--hcolbegin=string
        Specify the text placed at the beginning of each header
        column [default: "`<th>`"]

--hcolsep=string
        Specify the text used to separate header columns [de-
        fault: same as `colsep`]

--hcolend=string
        Specify the text placed at the end of each header col-
        umn [default: "`</th>`"]

--hrowbegin=string
        Specify the text placed at the beginning of each header
        row [default: same as `rowbegin`]

--hrowsep=string
        Specify the text used to separate header rows [default:
        same as `rowsep`]

--hrowend=string
        Specify the text placed at the end of each header row
        [default: same as `rowend`]

--tablebegin=string
        Specify the text placed at the beginning of each table
        [default: "`<table>\\n`"]

--tablesep=string
        Specify the text used to separate tables [default: ""]

--tableend=string
        Specify the text placed at the end of each table [default:
        "`</table>\\n`"]

--quote=string
        Specify the text used to begin quoted text [default: ""]

--*unquote=string*
>   Specify the text used to end quoted text [default: same as `quote`]

--*keep-columns=string*
>   Enumerate the columns that should be included in the output [default: all columns]

--*merge=function*
>   Specify how to merge data from multiple files [default: "`mean`"]

--*showfnames=option*
>   Add an extra header row showing the filename the data came from [default: "`none`"]

--*format=gnuplot*
>   Output each table as a gnuplot data file

--*noheaders*
>   Do not output column headers

--*colbegin=string*
>   Specify the text placed at the beginning of each data column [default: ""]

--*colsep=string*
>   Specify the text used to separate data columns [default: " "]

--*colend=string*
>   Specify the text placed at the end of each data column [default: ""]

--*rowbegin=string*
>   Specify the text placed at the beginning of each data row [default: ""]

--*rowsep=string*
>   Specify the text used to separate data rows [default: ""]

--*rowend=string*
>   Specify the text placed at the end of each data row [default: "`\\n`"]

--*hcolbegin=string*
>   Specify the text placed at the beginning of each header column [default: same as `colbegin`]

--*hcolsep=string*
>   Specify the text used to separate header columns [default: same as `colsep`]

--hcolend=string
>>> Specify the text placed at the end of each header column [default: same as `colend`]

--hrowbegin=string
>>> Specify the text placed at the beginning of each header row [default: "`#` "

--hrowsep=string
>>> Specify the text used to separate header rows [default: same as `rowsep`]

--hrowend=string
>>> Specify the text placed at the end of each header row [default: same as `rowend`]

--tablebegin=string
>>> Specify the text placed at the beginning of each table [default: ""]

--tablesep=string
>>> Specify the text used to separate tables [default: "`\\n\\n`"]

--tableend=string
>>> Specify the text placed at the end of each table [default: ""]

--quote=string
>>> Specify the text used to begin quoted text [default: "`"`"]

--unquote=string
>>> Specify the text used to end quoted text [default: same as `quote`]

--keep-columns=string
>>> Enumerate the columns that should be included in the output [default: all columns]

--merge=function
>>> Specify how to merge data from multiple files [default: "`mean`"]

--showfnames=option
>>> Add an extra header row showing the filename the data came from [default: "`none`"]

--format=octave
>> Output each table as an Octave text-format data file

--noheaders
>>> Do not output column headers

`--colbegin=string`
> Specify the text placed at the beginning of each data column [default: ""]

`--colsep=string`
> Specify the text used to separate data columns [default: ""]

`--colend=string`
> Specify the text placed at the end of each data column [default: "\\n"]

`--rowbegin=string`
> Specify the text placed at the beginning of each data row [default: ""]

`--rowend=string`
> Specify the text placed at the end of each data row [default: ""]

`--hcolbegin=string`
> Specify the text placed at the beginning of each header column [default: ""]

`--hcolsep=string`
> Specify the text used to separate header columns [default: "_"]

`--hcolend=string`
> Specify the text placed at the end of each header column [default: ""]

`--hrowbegin=string`
> Specify the text placed at the beginning of each header row [default: "# "]

`--hrowsep=string`
> Specify the text used to separate header rows [default: ""]

`--hrowend=string`
> Specify the text placed at the end of each header row [default: "\\n"]

`--tablebegin=string`
> Specify the text placed at the beginning of each table [default: ""]

`--tablesep=string`
> Specify the text used to separate tables [default: "\\n"]

`--tableend=string`
> Specify the text placed at the end of each table [default: ""]

`--quote=string`
> Specify the text used to begin quoted text [default: ""]

`--unquote=string`
> Specify the text used to end quoted text [default: same as `quote`]

`--keep-columns=string`
> Enumerate the columns that should be included in the output [default: all columns]

`--merge=function`
> Specify how to merge data from multiple files [default: "`mean`"]

`--showfnames=option`
> Add an extra header row showing the filename the data came from [default: "`none`"]

`--format=custom`
> Output each table in a completely user-specified format

`--noheaders`
> Do not output column headers

`--colbegin=string`
> Specify the text placed at the beginning of each data column [default: ""]

`--colsep=string`
> Specify the text used to separate data columns [default: ""]

`--colend=string`
> Specify the text placed at the end of each data column [default: ""]

`--rowbegin=string`
> Specify the text placed at the beginning of each data row [default: ""]

`--rowsep=string`
> Specify the text used to separate data rows [default: ""]

`--rowend=string`
> Specify the text placed at the end of each data row [default: ""]

`--hcolbegin=string`
> Specify the text placed at the beginning of each header column [default: same as `colbegin`]

`--hcolsep=string`
> Specify the text used to separate header columns [default: same as `colsep`]

--hcolend=*string*

>
> Specify the text placed at the end of each header col-
> umn [default: same as `colend`]

--hrowbegin=*string*

>
> Specify the text placed at the beginning of each header
> row [default: same as `rowbegin`]

--hrowsep=*string*

>
> Specify the text used to separate header rows [default:
> same as `rowsep`]

--hrowend=*string*

>
> Specify the text placed at the end of each header row
> [default: same as `rowend`]

--tablebegin=*string*

>
> Specify the text placed at the beginning of each table
> [default: ""]

--tablesep=*string*

>
> Specify the text used to separate tables [default: ""]

--tableend=*string*

>
> Specify the text placed at the end of each table [default:
> ""]

--quote=*string*

>
> Specify the text used to begin quoted text [default: ""]

--unquote=*string*

>
> Specify the text used to end quoted text [default: same
> as `quote`]

--excel   Output strings in a format readable by Microsoft Excel

--keep-columns=*string*

>
> Enumerate the columns that should be included in the
> output [default: all columns]

--merge=*function*

>
> Specify how to merge data from multiple files [default:
> "`mean`"]

--showfnames=*option*

>
> Add an extra header row showing the filename the data
> came from [default: "`none`"]

--format=latex

>
> Output each table as a LaTeX tabular environment

--dcolumn

>
> Use the dcolumn package to align numbers on the dec-
> imal point

                    --booktabs
                         Use the booktabs package for a more professionally typeset look

                    --longtable
                         Use the longtable package to enable multi-page tables

                    --keep-columns=string
                         Enumerate the columns that should be included in the output [default: all columns]

                    --merge=function
                         Specify how to merge data from multiple files [default: "`mean`"]

                    --showfnames=option
                         Add an extra header row showing the filename the data came from [default: "`none`"]

--extract=params
          Extract the program's run-time parameters and environment variables

          --format=text [default]
                    Output the parameters in plain-text format

                    --include=filename
                       Read from a file the list of keys to output

                    --exclude=regexp
                       Ignore any keys whose name matches a regular expression

                    --sort   Sort the list of parameters alphabetically by key

                    --noenv   Exclude environment variables

                    --noparams
                     Exclude run-time parameters

                    --envformat=template
                     Format environment variable names using the given template [default: "`%s (environment variable)`"]

                    --columns=number
                     Output the parameters as a 1-, 2-, or 3-column table [default: 1]

                    --colsep=string
                     Specify the text used to separate data columns [default: "`: `"]

                    --rowbegin=string
                     Specify the text that's output at the start of each data row [default: ""]

`--rowend=string`
>   Specify the text that's output at the end of each data row [default: "`\\n`"]

`--format=dumpkeys`
>   Output a list of the keys only (i.e., no values)

`--include=filename`
>   Read the list of parameters to output from a given file

`--exclude=regexp`
>   Ignore any keys whose name matches a regular expression

`--envformat=template`
>   Format environment variable names using the given template [default: "`%s (environment variable)`"]

`--sort`   Sort the list of parameters alphabetically by key

`--noenv`   Exclude environment variables

`--noparams`
>   Exclude run-time parameters

`--format=latex`
>   Output the parameters as a LaTeX tabular environment

`--include=filename`
>   Read from a file the list of keys to output

`--exclude=regexp`
>   Ignore any keys whose name matches a regular expression

`--envformat=template`
>   Format environment variable names using the given template [default: "`%s (environment variable)`"]

`--sort`   Sort the list of parameters alphabetically by key

`--booktabs`
>   Use the booktabs package for a more professionally typeset look

`--tabularx`
>   Use the tabularx package to enable line wraps within the value column

`--longtable`
>   Use the longtable package to enable multi-page tables

`--noenv`   Exclude environment variables

`--noparams`
>   Exclude run-time parameters

`--extract=env`
> Extract the environment in which the program was run
>
> > `--format=sh [default]`
> > > Use Bourne shell syntax for setting environment variables
> > >
> > > > `--newlines`
> > > > > Separate commands with newlines instead of semicolons
> > > >
> > > > `--unset` Unset all other environment variables
> > > >
> > > > `--chdir` Switch to the program's original working directory
> >
> > `--format=bash`
> > > Use Bourne Again shell syntax for setting environment variables
> > >
> > > > `--newlines`
> > > > > Separate commands with newlines instead of semicolons
> > > >
> > > > `--unset` Unset all other environment variables
> > > >
> > > > `--chdir` Switch to the program's original working directory
> >
> > `--format=ksh`
> > > Use Korn shell syntax for setting environment variables
> > >
> > > > `--newlines`
> > > > > Separate commands with newlines instead of semicolons
> > > >
> > > > `--unset` Unset all other environment variables
> > > >
> > > > `--chdir` Switch to the program's original working directory
> >
> > `--format=csh`
> > > Use C shell syntax for setting environment variables
> > >
> > > > `--newlines`
> > > > > Separate commands with newlines instead of semicolons
> > > >
> > > > `--unset` Unset all other environment variables
> > > >
> > > > `--chdir` Switch to the program's original working directory
> >
> > `--format=zsh`
> > > Use Z shell syntax for setting environment variables
> > >
> > > > `--newlines`
> > > > > Separate commands with newlines instead of semicolons
> > > >
> > > > `--unset` Unset all other environment variables
> > > >
> > > > `--chdir` Switch to the program's original working directory

`--format=tcsh`
> Use tcsh syntax for setting environment variables
>
> `--newlines`
> > Separate commands with newlines instead of semicolons
>
> `--unset` Unset all other environment variables
>
> `--chdir` Switch to the program's original working directory

`--format=ash`
> Use ash syntax for setting environment variables
>
> `--newlines`
> > Separate commands with newlines instead of semicolons
>
> `--unset` Unset all other environment variables
>
> `--chdir` Switch to the program's original working directory

`--extract=source`
> Extract CONCEPTUAL source code
>
> `--format=text [default]`
> > Output the source code in plain-text format
> >
> > `--linebegin=string`
> > > Specify the text placed at the beginning of each line [default: ""]
> >
> > `--lineend=string`
> > > Specify the text placed at the end of each line [default: "\\n"]
> >
> > `--kwbegin=string`
> > > Specify the text placed before each keyword [default: ""]
> >
> > `--kwend=string`
> > > Specify the text placed after each keyword [default: ""]
> >
> > `--strbegin=string`
> > > Specify the text placed before each string [default: ""]
> >
> > `--strend=string`
> > > Specify the text placed after each string [default: ""]
> >
> > `--combegin=string`
> > > Specify the text placed before each comment [default: ""]
> >
> > `--comend=string`
> > > Specify the text placed after each comment [default: ""]

`--indent=number`

Indent each line by a given number of spaces

`--wrap=number`

Wrap the source code into a paragraph with a given character width

`--extract=warnings`

Extract a list of warnings the program issued during initialization

`--format=text [default]`

Output warnings in plain-text format

`--listbegin=string`

Specify text to appear at the beginning of the list [default: ""]

`--listend=string`

Specify text to appear at the end of the list [default: ""]

`--itembegin=string`

Specify text to appear before each warning [default: "* "]

`--itemend=string`

Specify text to appear after each warning [default: "\\n"]

`--format=html`

Output warnings as an HTML list

`--listbegin=string`

Specify text to appear at the beginning of the list [default: "<ul>\\n"]

`--listend=string`

Specify text to appear at the end of the list [default: "
\\n"]

`--itembegin=string`

Specify text to appear before each warning [default: " "]

`--itemend=string`

Specify text to appear after each warning [default: "
\\n"]

`--format=latex`

Output warnings as a LaTeX list

`--listbegin=string`

Specify text to appear at the beginning of the list [default: "\begin@{itemize@}\\n"]

> --*listend=string*
>> Specify text to appear at the end of the list [default: "\end@{itemize@}\\n"]
>
> --*itembegin=string*
>> Specify text to appear before each warning [default: " \item "]
>
> --*itemend=string*
>> Specify text to appear after each warning [default: "\\n"]

The following represent additional clarification for some of the above:

- If --*indent* is specified without an argument, the argument defaults to 2.

- If --*wrap* is specified without an argument, the argument defaults to 72.

- The following are examples of the different arguments to the --*columns* option:

  --*columns=1 (default)*

  ```
                  coNCePTuaL version: 1.0
                  coNCePTuaL backend: c_mpi
                  Average timer overhead [gettimeofday()]: <1 microsecond
                  Log creation time: Thu Mar 27 19:22:48 2003
                  Log completion time: Thu Mar 27 19:22:48 2003
  ```

  --*columns=2*

  ```
                  coNCePTuaL version:                      1.0
                  coNCePTuaL backend:                      c_mpi
                  Average timer overhead [gettimeofday()]: <1 microsecond
                  Log creation time:                       Thu Mar 27 19:22:48 2003
                  Log completion time:                     Thu Mar 27 19:22:48 2003
  ```

  --*columns=3*

  ```
                  coNCePTuaL version                     : 1.0
                  coNCePTuaL backend                     : c_mpi
                  Average timer overhead [gettimeofday()]: <1 microsecond
                  Log creation time                      : Thu Mar 27 19:22:48 2003
                  Log completion time                    : Thu Mar 27 19:22:48 2003
  ```

- --*dumpkeys* produces suitable input for the --*include* option.

- --*exclude* can be specified repeatedly on the command line.

- --*keep-columns* accepts a list of comma-separated column ranges. For example, --*keep-columns=1,3-6,8* tells ncptl-logextract to ignore all but the first, third through sixth, and eighth columns of data.

- --*merge* takes one of mean (arithmetic mean), hmean (harmonic mean), min (minimum), max (maximum), median (median), sum (sum), all (all values from each column), or concat (horizontal concatenation of all data), and applies the function to corresponding data values across all of the input files. --*merge* can also accept a comma-separated list of the above functions, one per data column. This enables a different merge operation to be used for each column. For example, --*merge=min,min,mean* will take the minimum value across all files of each element in the first and second columns and the arithmetic

mean across all files of each element in the third column. If the number of comma-separated values differs from the number of columns and `--force-merge` is specified, `ncptl-logextract` will cycle over the given values until all columns are accounted for. The `concat` merge type applies to all columns and therefore cannot be combined with any other merge type. The difference between `--merge=all` and `--merge=concat` is that the former merges three files each with columns A and B as $\{A, A, A, B, B, B\}$ while the latter merges the same files as $\{A, B, A, B, A, B\}$.

Note that `--merge` is applied *after* `--keep-columns`. Hence, if `--keep-columns` specifies that only three columns be kept, `--merge` should list exactly three operations (or a single operation that applies to all three columns).

- `--showfnames` prepends to each data table in the input file an extra header line indicating the log file the data was extracted from. This option makes sense only when data is being extracted and primarily when `--merge=all` is specified. `--showfnames` takes one of `none`, `all`, or `first`. The default is `none`, which doesn't add an extra header row. `all` repeats the filename in each column of the extra header row. `first` outputs the filename in only the first column, leaving the remaining columns with an empty string. The following examples show how a sample data table is formatted with `--showfnames` set in turn to each of `none`, `all`, and `first`:

  - Set to `none` (the default):

    ```
    "Size","Value"
    1,2
    2,4
    3,6
    ```

  - Set to `all` (filename repeated in each column of the first row):

    ```
    "mydata.log","mydata.log"
    "Size","Value"
    1,2
    2,4
    3,6
    ```

  - Set to `first` (filename shown only in the first column of the first row):

    ```
    "mydata.log",""
    "Size","Value"
    1,2
    2,4
    3,6
    ```

- If `--format=params` is used with both `--longtable` and `--tabularx`, the generated table will be formatted for use with the `ltxtable` LaTeX package. See the `ltxtable` documentation for more information.

NOTES

If no filenames are given, `ncptl-logextract` will read from the standard input device. If multiple log files are specified, CONCEPTUAL will merge the data values and take all other information from the first file specified. Note, however, that all of the log files must have been produced by the same CONCEPTUAL program and that that program must have been run in the same environment. In other words, only the data values may change across log

files; everything else must be invariant. See the description of `--merge` in the Additional Options section for more information about merging data values from multiple log files.

`ncptl-logextract` treats certain files specially:

- If `ncptl-logextract` is given a filename ending in `.gz`, `.bz2`, or `.Z` it automatically decompresses the file to a temporary location using `gunzip`, `bunzip2`, or `uncompress`, as appropriate, then recursively processes the decompressed file.

- If `ncptl-logextract` is given a filename ending in `.tar` or `.zip` it automatically extracts the file's contents to a temporary directory using `tar` or `unzip`, as appropriate, then recursively processes the temporary directory.

- If `ncptl-logextract` is given the name of a directory it processes all of the plain files found (recursively) beneath that directory.

- If an input file is a merged CONCEPTUAL log file (i.e., produced by `ncptl-logmerge`), `ncptl-logextract` automatically invokes `ncptl-logunmerge` to split the file into its constituent, ordinary log files then recursively processes those.

`ncptl-logmerge` treats filenames ending in `.tgz` as if they ended in `.tar.gz` and filenames ending in `.taz` as if they ended in `.tar.Z`.

If the argument provided to any `ncptl-logextract` option begins with an at sign ("@"), the value is treated as a filename and is replaced by the file's contents. To specify an non-filename argument that begins with an at sign, merely prepend an additional "@":

`--this=that`
> The option `this` is given the value "`that`".

`--this=@that`
> The option `this` is set to the contents of the file called `that`.

`--this=@@that`
> The option `this` is given the value "`@that`".

EXAMPLES

For the following examples, we assume that `results.log` is the name of a log file produced by a CONCEPTUAL program.

Extract the data in CSV format and write it to `results.csv`:

> `ncptl-logextract --extract=data results.log --output=results.csv`

Note that `--extract=data` is the default and therefore optional:

> `ncptl-logextract results.log --output=results.csv`

`ncptl-logextract` can combine data from multiple log files (using an arithmetic mean by default):

> `ncptl-logextract results-*.log --output=results.csv`

Put the data from all of the log files side-by-side and produce a CSV file that Microsoft Excel can read directly:

> `ncptl-logextract results-*.log --output=results.csv --merge=all \`
> `    --showfnames=first --excel`

Output the data from `result.log` in tab-separated-value format:

```
ncptl-logextract --format=tsv results.log
```
Output the data in space-separated-value format:
```
ncptl-logextract --colsep=" " results.log
```
Use **gnuplot** to draw a PostScript graph of the data:
```
ncptl-logextract results.log --format=gnuplot \
    --before=@params.gp | gnuplot > results.eps
```
In the above, the `params.gp` file might contain **gnuplot** commands such as the following:
```
set terminal postscript eps enhanced color "Times-Roman" 30
set output
set logscale xy
set data style linespoints
set pointsize 3
plot "-" title "Latency"
```
(There should be an extra blank line at the end of the file because `ncptl-logextract` strips off a trailing newline character whenever it reads a file using "`@`".)

Produce a complete HTML file of the data (noting that `--format=html` produces only tables, not complete documents):
```
ncptl-logextract --format=html
    --before='<html>\n<head>\n<title>Data</title>\n</head>\n<body>\n' \
    --after='</body>\n</html>\n' results.log
```
Output the data as a LaTeX `tabular`, relying on both the (standard) `dcolumn` and (non-standard) `booktabs` packages for more attractive formatting:
```
ncptl-logextract --format=latex --dcolumn --booktabs \
    --output=results.tex results.log
```
Output the run-time parameters in the form "*key* `--->` *value*" with all of the arrows aligned:
```
ncptl-logextract results.log --extract=params --columns=3 --colsep=" --> "
```
Output the run-time parameters as an HTML description list:
```
ncptl-logextract results.log --extract=params --before='<dl>' \
    --rowbegin='<dt>' --colsep='</dt><dd>' --rowend='</dd>\n' \
    --after='</dl>\n'
```
Restore the exact execution environment that was used to produce `results.log`, including the current working directory (assuming that `bash` is the current command shell):
```
eval `ncptl-logextract --extract=env --format=bash \
    --unset --chdir results.log`
```
Set all of the environment variables that were used to produce `results.log`, overwriting—but not removing—whatever environment variables are currently set (assuming that `tcsh` is the current command shell):
```
eval `ncptl-logextract --extract=env --format=tcsh results.log`
```
Extract the source code that produced `results.log`:
```
ncptl-logextract --extract=source results.log
```
Do the same, but indent the code by four spaces then re-wrap it into a 60-column paragraph:

```
ncptl-logextract --extract=source --indent=4 --wrap=60 results.log
```

Here are a variety of ways to express the same thing:

```
ncptl-logextract -e source --indent=4 --wrap=60 results.log

ncptl-logextract -e source --indent=4 results.log --wrap=60

cat results.log | ncptl-logextract --wrap=60 --indent=4 -e source
```

Output the source code wrapped to 72 columns, with no indentation, and formatted within an HTML preformatted-text block:

```
ncptl-logextract --extract=source --wrap --before="<PRE>\n" \
  after="</PRE>\n" results.log
```

List all of the warning messages which occur in `results.log`:

```
ncptl-logextract --extract=warnings results.log
```

SEE ALSO

ncptl-logmerge(1), ncptl-logunmerge(1), the CONCEPTUAL User's Guide

AUTHOR

Scott Pakin, pakin@lanl.gov

3.5.3 `ncptl-logmerge`

cONCePTuaL programs produce one log file per process. An unwieldy number of files can therefore be generated on large-scale computer systems. For the case in which only a single log file contains measurement data, `ncptl-logmerge` can merge a number of log files into a single file. Only lines that differ across log files are repeated, making the result fairly space-efficient. The primary advantage of `ncptl-logmerge` over an archiving program such as `tar` or `zip` is that the output of `ncptl-logmerge` is designed to be human-readable—in fact, easily readable.

`ncptl-logmerge` can also be used to highlight differences in log files. It is therefore an important diagnostic tool for unearthing subtle configuration discrepancies across nodes in a large-scale computer system.

Running *ncptl-logmerge --usage* causes `ncptl-logmerge` to list a synopsis of its core command-line options to the standard output device; running *ncptl-logmerge --help* produces basic usage information; and, running *ncptl-logmerge --man* outputs a complete manual page. See [ncptl-logmerge manual page], page 77, shows the `ncptl-logmerge` documentation as produced by *ncptl-logmerge --man*.

NAME

ncptl-logmerge - Merge CONCEPTUAL log files

SYNOPSIS

ncptl-logmerge *--usage* | *--help* | *--man*

ncptl-logmerge [*--output=filename*] [*--simplify*] *filename...*

DESCRIPTION

A CONCEPTUAL program produces one log file per process. For large numbers of processes the result can be unwieldy. `ncptl-logmerge` combines a large set of log files into a single, merged file which can later be expanded back into its constituent log files. There are a number of restrictions on the input to `ncptl-logmerge`; see the Restrictions section for details.

The merged output file does not modify lines which are identical in all of the input files. Lines which do differ across input files are prefixed with the processors and processor ranges in which they appeared.

As an example, the following text was extracted from a set of 186 CONCEPTUAL log files (from a 186-processor run):

```
# Microsecond timer type: PAPI_get_real_usec()
# Average microsecond timer overhead: <1 microsecond
#[0-4,6-12,14-16,18-52,54-78,80-94,96-101,103-121,123-140,142-169,
  171-185]# Microsecond timer increment: 1 +/- 0 microseconds
  (ideal: 1 +/- 0)
#[5]# Microsecond timer increment: 1.00229 +/- 0.15854
  microseconds (ideal: 1 +/- 0)
#[13]# Microsecond timer increment: 1.00228 +/- 0.158442
  microseconds (ideal: 1 +/- 0)
#[17,79]# Microsecond timer increment: 1.00228 +/- 0.158392
  microseconds (ideal: 1 +/- 0)
#[53]# Microsecond timer increment: 1.00228 +/- 0.158409
  microseconds (ideal: 1 +/- 0)
#[102]# Microsecond timer increment: 1.00228 +/- 0.158458
  microseconds (ideal: 1 +/- 0)
#[95,122]# Microsecond timer increment: 1.00228 +/- 0.158474
  microseconds (ideal: 1 +/- 0)
#[141]# Microsecond timer increment: 1.00228 +/- 0.158491
  microseconds (ideal: 1 +/- 0)
#[170]# Microsecond timer increment: 1.00228 +/- 0.158524
  microseconds (ideal: 1 +/- 0)
```

All of the input files contained the same `Microsecond timer type` and `Average microsecond timer overhead` lines. However, the measured `Microsecond timer increment` varied across input files. While many of the processors observed an increment of `1 +/- 0`, processor 5 was alone in observing `1.00229 +/- 0.15854`; processor 13 was

alone in observing `1.00228 +/- 0.158442`; and, both processor 17 and processor 79 observed `1.00228 +/- 0.158392` as the timer increment.

ncptl-logmerge can also be instructed to output only the lines which differ across files. Common lines are not output. This feature is useful for discovering misconfigured nodes in a large computer system. For example, on one computer system on which CONCEPTUAL was run, five processors were running at a higher clock rate than the remainder which naturally affected performance. `ncptl-logmerge` can be used to help identify such outliers.

OPTIONS

ncptl-logmerge accepts the following command-line options:

-u , *--usage*
>   Output the Synopsis section then exit the program.

-h , *--help*
>   Output the Synopsis section and the Options section then exit the program.

-m , *--man*
>   Output a complete Unix man ("manual") page for `ncptl-logmerge` then exit the program.

-o *filename*, *--output* =*filename*
>   ncptl-logmerge normally writes to the standard output device. The `--output` option redirects `ncptl-logmerge`'s output to a file.

-s , *--simplify*
>   Simplify the output by including only lines which differ across input files. No data is output, only prologue and epilogue comments. `--simplify` can be specified up to four times on the command line:
>
> | | |
> |---|---|
> | once | Omit all comments and all lines which are identical across all input files. |
> | twice | Lines which differ across *all* output files (e.g., `Processor (0<=P<tasks)`) are also omitted. |
> | three times | |
> | | The amount of output is further reduced by rounding to two significant digits all numbers appearing in all input files. Doing so makes `1.10644 +/- 0.593714` match `1.12511 +/- 0.58829`, for example. (Both are converted to `1.1 +/- 0.59`.) |
> | four times | Lists of processors are replaced by the list size. For example, `#[22,67,86,430]` becomes `#[4]`. |
>
> Note that `--simplify` is intended as a diagnostic tool; files output using `--simplify` cannot be un-merged to recover the original input files.

In addition to the preceding options `ncptl-logmerge` requires a list of log files to merge. If a directory is specified, all of the files immediately under that directory are used. (Note that `ncptl-logmerge` does not descend into subdirectories, however.) Files containing lists of filenames can be specified with a leading at sign ("@"). For example, `@filelist.txt` means to read a list of filenames from `@filelist.txt`. Filenames beginning with an at sign can be specified by doubling the at sign on the command line.

DIAGNOSTICS

filename `does not look like an unmerged coNCePTuaL log file`

> `ncptl-logmerge` accepts as input only log files produced directly by a CoN-CePTuaL program. It is not a general-purpose file combiner nor does it accept its own output as input. Unrecognized input files cause `ncptl-logmerge` to abort with the preceding error message.

`No process rank found in` *filename*

> `ncptl-logmerge` needs to map filenames to process ranks to indicate which ranks produced which lines of output. If an input file does not contain a `Rank (0<=P<tasks)` comment, `ncptl-logmerge` aborts with the preceding error message.

EXAMPLES

Merge a set of CoNCePTuaL log files:

> `ncptl-logmerge mybenchmark-[0-9]*.log > mybenchmark-all.log`

The following command is equivalent to the preceding one:

> `ncptl-logmerge mybenchmark-[0-9]*.log --output=mybenchmark-all.log`

Show only "interesting" differences among the input files:

> `ncptl-logmerge --simplify --simplify mybenchmark-[0-9]*.log`

For convenience, one can abbreviate *--simplify --simplify --simplify --simplify* to *-s -s -s -s* or even *-ssss*:

> `ncptl-logmerge -ssss mybenchmark-[0-9]*.log`

RESTRICTIONS

The log files passed to `ncptl-logmerge` are subject to the following restrictions:

- All files must be produced by the same run of the same CoNCePTuaL program.
- None of the files can have been previously merged by `ncptl-logmerge` (i.e., `ncptl-logmerge` can't read its own output).
- Only the first filename passed to `ncptl-logmerge` is allowed to contain data. Data from all other files is discarded with a warning message.

BUGS

`ncptl-logmerge` is not a particularly robust script. Specifically, it is confused when input files contain different numbers of comment lines. For example, if one input file includes more environment variables than another or issued a warning about a timer where another input file didn't, `ncptl-logmerge` will erroneously report all subsequent lines as being mismatched across input files.

SEE ALSO

ncptl-logunmerge(1), ncptl-logextract(1), the CoNCePTuaL User's Guide

AUTHOR

Scott Pakin, pakin@lanl.gov

3.5.4 `ncptl-logunmerge`

The primary capability of `ncptl-logmerge` (see Section 3.5.3 [ncptl-logmerge], page 76) is to merge multiple CONCEPTUAL log files into a more maintainable single file. `ncptl-logunmerge` performs the complementary operation of splitting that merged file back into the original set of CONCEPTUAL log files.

Running *ncptl-logunmerge --usage* causes `ncptl-logunmerge` to list a synopsis of its core command-line options to the standard output device; running *ncptl-logunmerge --help* produces basic usage information; and, running *ncptl-logunmerge --man* outputs a complete manual page. See [ncptl-logunmerge manual page], page 81, shows the `ncptl-logunmerge` documentation as produced by *ncptl-logunmerge --man*.

NAME

ncptl-logunmerge - Recover individual CONCEPTUAL log files from ncptl-logmerge output

SYNOPSIS

ncptl-logunmerge *--usage* | *--help* | *--man*

ncptl-logunmerge   [*--logfile=template*]   [*--procs=process_list*]   [*--quiet*]
[*--memcache=megabytes*] *filename*

DESCRIPTION

While `ncptl-logmerge` merges a set of CONCEPTUAL log files into a more convenient,
single file, `ncptl-logunmerge` performs the inverse operation, splitting a merged file into
separate CONCEPTUAL log files. Specifically, unadorned comment lines such as the fol-
lowing are written to all log files:

        # Executable name: /home/me/mybenchmark

Comment lines which specify processor ranges are written to the appropriate log files.
For example, the following line—with the leading `#[35,43,89]` stripped—is written only
to the log files corresponding to processes 35, 43, and 89:

        #[35,43,89]# Minimum sleep time: 9.68 +/- 0.556776
          microseconds (ideal: 1 +/- 0)

Non-comment lines (i.e., measurement data) such as the following are written only to
process 0's log file:

        "Contention factor","Msg. size (B)","1/2 RTT (us)","MB/s"
        "(all data)","(all data)","(all data)","(all data)"
        0,1048576,1368.0295,730.978389
        0,524288,691.9115,722.6357706

OPTIONS

`ncptl-logunmerge` accepts the following command-line options:

-u , *--usage*

         Output the Synopsis section then exit the program.

-h , *--help*

         Output the Synopsis section and the Options section then exit the program.

-m , *--man*

         Output a complete Unix man ("manual") page for `ncptl-logunmerge` then exit
         the program.

-L *template*, *--logfile* =*template*

         Specify a template for the names of the generated log files. The template must
         contain the literal string `%p` which will be replaced by the appropriate processor
         number. The template may contain the literal string `%r` (run number) which
         will be replaced by the smallest integer which produces a filename that does not
         already exist. In addition, `printf()`-style field widths can be used with `%p` and

%r. For example, `%04p` outputs the processor number as a four-digit number padded on the left with zeroes.

If *--logfile* is not specified, `ncptl-logunmerge` takes the default template from the merged log file's `Log-file template` line, discards the directory component of the filename, and uses the result as the log-file template.

-p *process_list*, --procs =*process_list*

Identify a subset of log files to extract from the merged log file. By default, `ncptl-logunmerge` extracts all of the constituent log files. *process_list* is a comma-separated list of process number or process ranges.

-q , --quiet

Suppress progress output. Normally, `ncptl-logunmerge` outputs status information regarding its operation. The *--quiet* option instruct `ncptl-logunmerge` to output only warning and error messages.

-M *megabytes*, --memcache =*megabytes*

Specify the size of the in-memory file cache. By default the program keeps up to 8 MB of extracted file data resident in memory to improve performance. The *--memcache* option enables more or less data to be cached in memory. For example, *--memcache=512* specifies that 512 MB of memory should be reserved for the file cache.

In addition to the preceding options `ncptl-logunmerge` requires the name of a merged CoNCePTuaL log file. If not provided, `ncptl-logunmerge` reads the contents of the merged CoNCePTuaL log file from standard input.

DIAGNOSTICS

`The input file does not look like a merged coNCePTuaL log file; assuming simple concatenation`

`ncptl-logunmerge` expects its input file to contain the output from `ncptl-logmerge`. However, `ncptl-logunmerge` can also accept an input file produced by concatenating a collection of CoNCePTuaL log files end-to-end (e.g., by using the `cat` command). The preceding warning message serves merely to alert the user in case the wrong input file was provided to `ncptl-logunmerge`.

`Unable to find a unique number of tasks in the input file`

`ncptl-logunmerge` determines the number of files to generate from the `Number of tasks` prologue comment. If that comment does not appear or takes on different values, `ncptl-logunmerge` rejects the input file.

EXAMPLES

Extract a set of CoNCePTuaL log files from a merged log file and name the extracted files `happy-0.log`, `happy-1.log`, `happy-2.log`, etc.:

    ncptl-logunmerge --logfile=happy-%p.log mybenchmark-all.log

Extract only `mybenchmark-0.log`, `mybenchmark-50.log`, `mybenchmark-51.log`, `mybenchmark-52.log`, and `mybenchmark-100.log` from `mybenchmark-all.log` (assuming that `mybenchmark-all.log` contains the line `# Log-file template: mybenchmark-%p.log`):

```
ncptl-logunmerge --procs=0,50-52,100 mybenchmark-all.log
```

SEE ALSO

ncptl-logmerge(1), ncptl-logextract(1), printf(3), the coNCePTuaL User's Guide

AUTHOR

Scott Pakin, pakin@lanl.gov

4 Grammar

The CONCEPTUAL language was designed to produce precise specifications of network correctness and performance tests yet read like an English-language document that contains a hint of mathematical notation. Unlike more traditional programming languages, CONCEPTUAL is more descriptive than imperative. There are no classes, functions, arrays, pointers, or even variable assignments (although expressions can be let-bound to identifiers).[1] The language operates primarily on integers, with support for string constants and floating-point numbers in only a few constructs. A CONCEPTUAL program merely describes a communication pattern and the CONCEPTUAL compiler generates code to implement that pattern.

As a domain-specific language, CONCEPTUAL contains primitives to send and receive messages. It is capable of measuring time, computing statistics, and logging results. It knows that it will be run in a shared-nothing SPMD[2] style with explicit message-passing. As a result of its special-purpose design CONCEPTUAL can express communication patterns in a clearer and terser style than is possible using a general-purpose programming language.

The CONCEPTUAL language is case-insensitive. `Hello` is the same as `HELLO` or `hello`. Furthermore, whitespace is insignificant; one space has the same meaning as multiple spaces. Comments are designated with a '`#`' character and extend to the end of the line.

We now describe the CONCEPTUAL grammar in a bottom-up manner, i.e., starting from primitives and working up to complete programs. Note that many of the sections in this chapter use the following syntax to formally describe language elements:

<*nonterminal*>

        a placeholder for a list of language primitives and additional placeholders

::=   "is defined as"

KEYWORD   a primitive with special meaning to the language

[...]   optional items

(...)   grouping of multiple items into one

*   zero or more occurrences of the preceding item

+   one or more occurrences of the preceding item

| either the item to the left or the item to the right but not both

---

4.1 Primitives

At the lowest level, CONCEPTUAL programs are composed of identifiers, strings, and integers (and a modicum of punctuation). Identifiers consist of a letter followed by zero or more alphanumerics or underscores. '`potato`', '`x`', and '`This_is_program_123`' are all examples of valid identifiers. Identifiers are used for two purposes: variables and keywords.

---

[1] CONCEPTUAL is not even Turing-complete. That is, it cannot perform arbitrary computations.
[2] Single Program, Multiple Data

Variables—referred to in the formal grammar as <*ident*>s—can be bound but not assigned. That is, once a variable is given a value it retains that value for the entire scope although it may be given a different value within a subordinate scope for the duration of that scope. All variables are of integer type. There are a number of variables that are predeclared and maintained automatically by CoNCePTuaL. These are listed and described in Section A.2 [Predeclared variables], page 207. Predeclared variables can be used by CoNCePTuaL programs but cannot be redeclared; an attempt to do so will result in a compile-time error message.

Keywords introduce actions. For example, SEND and RECEIVE are keywords. (A complete list of CoNCePTuaL keywords is presented in Section A.1 [Keywords], page 202.) Most keywords can appear in multiple forms. For example, OUTPUT and OUTPUTS are synonymous, as are COMPUTE and COMPUTES, A and AN, and TASK and TASKS. The intention is for programs to use whichever sounds better in an English-language sentence. Keywords may not be used as variable names; an attempt to do so will cause the compiler to output a parse error.

As a special case to increase program readability, a single '-' preceding a keyword is treated as a whitespace character. Hence, 'INTEGER-SIZED PAGE-ALIGNED MESSAGE' is equivalent to 'INTEGER SIZED PAGE ALIGNED MESSAGE' and '10 64-BYTE MESSAGES' is equivalent to '10 64 BYTE MESSAGES'. However, 'x-3' and '3-x' still represent subtraction operations.

Although identifiers are case insensitive—'SEND' is the same as 'send' is the same as 'sENd'—to increase clarity, this manual presents keywords in uppercase and variables in lowercase.

Strings consist of double-quoted text. Within a string—and only within a string—whitespace and case are significant. In particular, a literal newline is honored but can be suppressed by preceding it with a backslash as in the following example:

```
"This string\
contains some
newline characters."
```

⇒ This string contains some
newline characters.

Use '\"' for a double-quote character, '\\' for a backslash, '\t' for a tab character, '\r' for a carriage-return character, and '\n' for a newline character. All other escape sequences produce a warning message and are discarded. As examples of valid escape-sequence usage, the string "March 2019" represents the text "March 2019" and "I store \"stuff\" in C:\\MyStuff." represents the text "I store "stuff" in C:\MyStuff."

Integers consist of an optional '+' or '-'[3] followed by one or more digits followed by an optional multiplier. This multiplier is unique to CoNCePTuaL and consists of one of the following four letters:

'K' (kilo)   multiplies the integer by 1,024

'M' (mega)   multiplies the integer by 1,048,576

---

[3] From the lexer's perspective, integers are always unsigned and '+' or '-' are merely operators (see Section 4.2.1 [Arithmetic expressions], page 86). The parser, however, applies unary operators to integer literals. This alteration is evident in the output of the dot_ast backend (see Section 3.3.11 [The dot_ast backend], page 41).

'`G`' (giga)   multiplies the integer by 1,073,741,824

'`T`' (tera)   multiplies the integer by 1,099,511,627,776

A multiplier can also be '`E`' (exponent) followed by a positive integer. An '`E`' multiplier multiplies the base integer by 10 raised to the power of the positive integer following the '`E`'.

Integers can also be written as ordinals, with an '`ST`', '`ND`', '`RD`', or '`TH`' suffix. These suffixes are ignored by the lexer but can make the `PERCENTILE` aggregate function (see) look more English-like (as in, '`THE 95TH PERCENTILE`').

Some examples of valid integers include '`2016`', '`-42`', '`64K`' (= 65,536), '`8E3`' (= 8,000), and '`38TH`'.

4.2 Expressions

Expressions, as in any language, are a combination of primitives and other expressions in a semantically meaningful juxtaposition. CONCEPTUAL provides arithmetic expressions, which evaluate to a number, and relational expressions, which evaluate to either TRUE or FALSE. In addition, CONCEPTUAL provides the notion of an *aggregate expression*, which represents a function (e.g., statistical mean) applied to every value taken on by an arithmetic expression during the run of a program.

4.2.1 Arithmetic expressions

CONCEPTUAL supports a variety of arithmetic expressions. The following is the language's order of operations from highest to lowest precedence:

unary   '`+`', '`-`', `NOT`, '`<function>(<expr>, ...)`', '`REAL(<expr>)`'

power   '`**`'

multiplicative '`*`', '`/`', `MOD`, '`<<`', '`>>`', '`&`'

additive   '`+`', '`-`', '`|`', `XOR`

conditional   '`<expr> IF <rel_expr> OTHERWISE <expr>`'

In addition, as in most programming languages, parentheses can be used to group subexpressions.

The '`&`' ("and"), '`|`' ("or"), `XOR`, and `NOT` operators perform bitwise, not logical, operations. That is, they accept numerical arguments, not truth-value arguments. Hence, for example, '`3 | 5`' is equal to '`7`'.

'`<<`' and '`>>`' are bit-shift operators. That is, '`a << b`' is the CONCEPTUAL equivalent of the mathematical expression $a \times 2^b$ and '`a >> b`' is the CONCEPTUAL equivalent of the mathematical expression $a \div 2^b$. Consequently, negative values of b are valid and correspond to a shift in the opposite direction. (In contrast, C and Perl treat a negative shift amount as a—usually large—unsigned number and Python raises a `ValueError` exception on negative shifts.)

MOD is a modulo (i.e., remainder) operator: '10 MOD 3' returns '1'. 'MOD' is guaranteed to return a nonnegative remainder. Hence, '16 MOD 7' and '16 MOD -7' both return '2' even though '-5' is also mathematically valid. Similarly, '-16 MOD 7' and '-16 MOD -7' both return '5' even though '-2' is also mathematically valid.

The function calls allowed in '<function>(<expr>, ...)' are listed and described in Section 4.2.2 [Built-in functions], page 88. All functions take one or more arithmetic expressions as an argument. The operator '*' represents multiplication; '/' represents division; and '**' represents exponentiation (i.e., 'x ** y' $\equiv \lfloor x^y \rfloor$). Note that 0^y generates a run-time error for $y \leq 0$.

A conditional expression '<expr1> IF <rel_expr> OTHERWISE <expr2>' evaluates to <expr1> if the relational expression <rel_expr> evaluates to TRUE and <expr2> if <rel_expr> evaluates to FALSE.[4] Relational expressions are described in Section 4.2.5 [Relational expressions], page 101. As some examples of conditional expressions, '666 IF 2+2=5 OTHERWISE 777' returns '777' while '666 IF 2+2=4 OTHERWISE 777' returns '666'.

All operations proceed left-to-right except power and conditional expressions, which proceed right-to-left. That is, '4-3-2' means $(4 - 3) - 2$ but '4**3**2' means $4^{(3^2)}$. Similarly, '2 IF p=0 OTHERWISE 1 IF p=1 OTHERWISE 0' associates like '2 IF p=0 OTHERWISE (1 IF p=1 OTHERWISE 0)', not like '(2 IF p=0 OTHERWISE 1) IF p=1 OTHERWISE 0'.

Evaluation contexts

CONCEPTUAL normally evaluates arithmetic expressions in "integer context", meaning that each subexpression is truncated to the nearest integer after being evaluated. Hence, '24/5*5' is '20', not '24', because '24/5*5' $= \lfloor 24 \div 5 \rfloor \times 5 = 4 \times 5 = 20$. There are a few situations, however, in which CONCEPTUAL evaluates expressions in "floating-point context", meaning that no truncation occurs:

- within an OUTPUTS statement (see Section 4.5.2 [Writing to standard output], page 120)
- within a LOGS statement (see Section 4.5.3 [Writing to a log file], page 120)
- within a BACKEND EXECUTES statement (see Section 4.8.5 [Injecting arbitrary code], page 138)
- within a range in a FOR EACH statement (see [Range loops], page 127) when CONCEPTUAL is unable to find an arithmetic or geometric progression by evaluating the component <expr>s in integer context

Within any of the preceding statements, the expression '24/5*5' evaluates to 24. Furthermore, the expression '24/5' evaluates to 4.8, which is a number that can't be entered directly in a CONCEPTUAL program. (The language supports only integral constants, as mentioned in Section 4.1 [Primitives], page 84.)

The CONCEPTUAL language provides a special form called REAL, which resembles a single-argument function. When evaluated in floating-point context, REAL returns its argument evaluated normally, as if 'REAL' were absent. When evaluated in integer context, however, REAL evalutes its argument in floating-point context and then rounds the result to the nearest integer. As an example, '9/2 + 1/2' is '4' in integer context because '9/2 + 1/2' $= \lfloor 9/2 \rfloor + \lfloor 1/2 \rfloor = 4 + 0 = 4$. However, 'REAL(9/2 + 1/2)' is '5' in integer context because 'REAL(9/2 + 1/2)' $= \lfloor 9/2 + 1/2 + 0.5 \rfloor = \lfloor 5 + 0.5 \rfloor = 5$.

---

[4] It is therefore analogous to '<rel_expr> ? <expr1> : <expr2>' in the C programming language.

Formal grammar for arithmetic expressions

For completeness, the following productions formalize the process by which CoNCePTuaL parses arithmetic expressions:

<expr>	::=	*<cond_expr>*

<cond_expr>	::=	*<add_expr>* IF *<rel_expr>* OTHERWISE *<add_expr>*

<add_expr>	::=	*<mult_expr>*
	|	*<add_expr>* '+' *<mult_expr>*
	|	*<add_expr>* '-' *<mult_expr>*
	|	*<add_expr>* '|' *<mult_expr>*
	|	*<add_expr>* XOR *<mult_expr>*

<mult_expr>	::=	*<unary_expr>*
	|	*<mult_expr>* '*' *<unary_expr>*
	|	*<mult_expr>* '/' *<unary_expr>*
	|	*<mult_expr>* MOD *<unary_expr>*
	|	*<mult_expr>* '>>' *<unary_expr>*
	|	*<mult_expr>* '<<' *<unary_expr>*
	|	*<mult_expr>* '&' *<unary_expr>*

<power_expr>	::=	*<primary_expr>* ['**' *<unary_expr>*]

<unary_expr>	::=	*<power_expr>*
	|	*<unary_operator>* *<unary_expr>*

<unary_operator>	::=	'+' | '-' | NOT

<primary_expr>	::=	'(' *<expr>* ')'
	|	*<ident>*
	|	*<integer>*
	|	*<func_name>* '(' *<enumerated_exprs>* ')'
	|	REAL '(' *<expr>* ')'

<enumerated_exprs>	::=	*<expr>* [',' *<expr>*]*

4.2.2 Built-in functions

In addition to the operators described in Section 4.2.1 [Arithmetic expressions], page 86, CoNCePTuaL contains a number of built-in functions that perform a variety of arithmetic operations that are often found to be useful in network correctness and performance testing codes. These include simple functions that map one number to another as well as a set of topology-specific functions that help implement communication across various topologies, specifically *n*-ary trees, meshes, tori, and *k*-nomial trees. CoNCePTuaL currently supports the following functions:

<func_name>	::=	ABS | BITS | CBRT | FACTOR10 | LOG10 | MAX | MIN | ROOT | SQRT
	|	CEILING | FLOOR | ROUND
	|	TREE_PARENT | TREE_CHILD
	|	KNOMIAL_PARENT | KNOMIAL_CHILD | KNOMIAL_CHILDREN
	|	MESH_NEIGHBOR | MESH_COORDINATE | MESH_DISTANCE
	|	RANDOM_UNIFORM | RANDOM_GAUSSIAN
	|	RANDOM_POISSON | RANDOM_PARETO
	|	PROCESSOR_OF | TASK_OF

```
        |     FILE_DATA | STATIC_FILE_DATA
```

With a few exceptions, all of the above take as an argument one or more numeric values (which may be the result of an arithmetic expression). The following sections describe each function in turn.

Integer functions

`ABS` returns the absolute value of its argument. For example, '`ABS(99)`' and '`ABS(-99)`' are both '`99`'.

`BITS` returns the minimum number of bits needed to store its argument. For example, '`BITS(12345)`' is '`14`' because 2^{14} is 16,384, which is larger than 12,345, while 2^{13} is 8,192, which is too small. '`BITS(0)`' is defined to be '`0`'. Essentially, '`BITS($x$)`' represents $\lceil \log_2 x \rceil$, i.e., the ceiling of the base-2 logarithm of x. Negative numbers are treated as their two's-complement equivalent. For example, '`BITS(-1)`' returns '`32`' on a 32-bit system and '`64`' on a 64-bit system.

`CBRT` is an integer cube root function. It is essentially just syntactic sugar for the more general `ROOT` function: '`CBRT($x$)`' $\equiv$ '`ROOT(3, $x$)`'.

`FACTOR10` rounds its argument down (more precisely, towards zero) to the largest single-digit factor of an integral power of 10. '`FACTOR10(4975)`' is therefore '`4000`'. Similarly, '`FACTOR10(-4975)`' is '`-4000`'.

'`LOG10($x$)`' is $\lfloor \log x \rfloor$, i.e., the floor of the base-10 logarithm of x. For instance, '`LOG10(12345)`' is '`4`' because 10^4 is the largest integral power of 10 that does not exceed 12,345.

`MIN` and `MAX` return, respectively, the minimum and maximum value in a list of numbers. Unlike the other built-in functions, `MIN` and `MAX` accept an arbitrary number of arguments (but at least one). For example, '`MIN(8,6,7,5,3,0,9)`' is '`0`' and '`MAX(8,6,7,5,3,0,9)`' is '`9`'.

'`ROOT($n$, $x$)`' returns $\sqrt[n]{x}$, i.e., the nth root of x. More precisely, it returns the largest integer r such that $r^n \leq x$. `ROOT` is not currently defined on negative values of x. As an example of `ROOT` usage, '`ROOT(5, 245)`' is '`3`' because $3^5 = 243 \leq 245$ while $4^5 = 1024 > 245$. Similarly, '`ROOT(2, 16)`' = '`4`'; '`ROOT(3, 27)`' = '`3`'; '`ROOT(0, 0)`' and '`ROOT(4, -245)`' each return a run-time error; and, '`ROOT(-3, 8)`' = '`0`' (because '`ROOT(-3, 8)`' = '`1/ROOT(3, 8)`' = '`1/2`' = '`0`').

`SQRT` is an integer square root function. It is essentially just syntactic sugar for the more general `ROOT` function: '`SQRT($x$)`' $\equiv$ '`ROOT(2, $x$)`'.

Floating-point functions

As stated in Section 4.2.1 [Arithmetic expressions], page 86, there are certain constructs in which expressions are evaluated in floating-point context instead of integer context. In such constructs, all of CoNCePTuaL's built-in functions accept and return floating-point values. Furthermore, the `CBRT`, `LOG10`, `ROOT`, and `SQRT` functions (see [Integer functions], page 89) compute floating-point results, not integer results that are coerced into floating-point format.

The following functions are not meaningful in integer context but are in floating-point context:

- CEILING
- FLOOR
- ROUND

CEILING returns the smallest integer not less than its argument. For example, 'CEILING(-7777/10)' is '-777'. (-778 is less than -777.7 while -777 is not less than -777.7.)

FLOOR returns the largest integer not greater than its argument. For example, 'FLOOR(-7777/10)' is '-778'. (-778 is not greater than -777.7 while -777 is greater than -777.7.)

ROUND rounds its argument to the nearest integer. For example, 'ROUND(-7777/10)' is '-778'.

It is not an error to use CEILING, FLOOR, and ROUND in an integer context; each function merely return its argument unmodified.

n-ary tree functions

n-ary trees are used quite frequently in communication patterns because they require only logarithmic time (in the number of tasks) for a message to propagate from the root to a leaf. cONCePTuaL supports n-ary trees in the form of the TREE_PARENT and TREE_CHILD functions.

TREE_PARENT (*task_ID* [, *fan-out*])   [Function]
> TREE_PARENT takes a task number and an optional tree fan-out (n) and returns the task's parent in an n-ary tree. n defaults to '2', i.e., a binary tree. Taking the TREE_PARENT of any task less than 1 returns the value '-1'.

TREE_CHILD (*task_ID*, *child* [, *fan-out*])   [Function]
> TREE_CHILD takes a task number, a child number ($0 \leq i < N$), and an optional tree fan-out (n), which again defaults to '2'. It returns the task number corresponding to the given task's *child*th child.

The following illustrations show how tasks are numbered in, respectively, a 2-ary and a 3-ary tree:

As shown by the 2-ary tree, task 1's children are task 3 and task 4. Therefore, 'TREE_PARENT(3)' and 'TREE_PARENT(4)' are both '1'; 'TREE_CHILD(1, 0)' is '3'; and, 'TREE_CHILD(1, 1)' is '4'. In a 3-ary tree, each task has three children. Hence, the following expressions hold:

- 'TREE_PARENT(7, 3)' ⇒ '2'
- 'TREE_PARENT(8, 3)' ⇒ '2'
- 'TREE_PARENT(9, 3)' ⇒ '2'
- 'TREE_CHILD(2, 0, 3)' ⇒ '7'
- 'TREE_CHILD(2, 1, 3)' ⇒ '8'
- 'TREE_CHILD(2, 2, 3)' ⇒ '9'

k-nomial tree functions

k-nomial trees are an efficient way to implement collective-communication operations in software. Unlike in an *n*-ary tree, the number of children in a *k*-nomial tree decreases with increasing task depth (i.e., no task has more children than the root). The advantage is that the tasks that start communicating earlier perform more work, which reduces the total latency of the collective operation. In contrast, in an *n*-ary tree, the tasks that start communicating earlier finish earlier, at the expense of increased total latency. cONCEPTuaL supports *k*-nomial trees via the KNOMIAL_PARENT, KNOMIAL_CHILDREN, and KNOMIAL_CHILD functions, as described below.

KNOMIAL_PARENT (*task_ID* [, *fan_out* [, *num_tasks*]])   [Function]
    KNOMIAL_PARENT takes a task number, the tree fan-out factor (the "*k*" in "*k*-ary"), and the number of tasks in the tree. It returns the task ID of the given task's parent. *fan_out* defaults to '2' and the number of tasks defaults to num_tasks (see Section A.2 [Predeclared variables], page 207).

KNOMIAL_CHILDREN (*task_ID* [, *fan_out* [, *num_tasks*]])   [Function]
    KNOMIAL_CHILDREN takes the same arguments as KNOMIAL_PARENT but returns the number of immediate descendents the given task has.

KNOMIAL_CHILD (*task_ID*, *child* [, *fan_out* [, *num_tasks*]])   [Function]
    KNOMIAL_CHILD takes a task number, a child number ($0 \leq i <$ 'KNOMIAL_CHILDREN(...)'), the tree fan-out factor, and the number of tasks in the tree. It returns the task number corresponding to the given task's *i*th child. As in KNOMIAL_PARENT and KNOMIAL_CHILDREN, *fan_out* defaults to '2' and the number of tasks defaults to num_tasks (see Section A.2 [Predeclared variables], page 207).

The following figure shows how CONCEPTUAL numbers tasks in a k-nomial tree with $k = 2$ (a.k.a. a 2-nomial or binomial tree).

The figure is structured with time flowing downwards. That is, for a multicast operation expressed over a 2-nomial tree, task 0 sends a message to task 1 in the first time step. Then, task 0 sends to task 2 while task 1 sends to task 3. In the final step, task 0 sends to task 4, task 1 sends to task 5, task 2 sends to task 6, and task 3 sends to task 7—all concurrently. The following expressions also hold, assuming there are a total of eight tasks in the computation:

- 'KNOMIAL_PARENT(0)' ⇒ '-1'
- 'KNOMIAL_PARENT(1)' ⇒ '0'
- 'KNOMIAL_CHILDREN(1)' ⇒ '2'
- 'KNOMIAL_CHILD(1, 0)' ⇒ '3'
- 'KNOMIAL_CHILD(1, 1)' ⇒ '5'
- 'KNOMIAL_CHILDREN(7)' ⇒ '0'
- 'KNOMIAL_CHILD(7, 0)' ⇒ '-1'

k-nomial trees for $k > 2$ are much less common in practice than 2-nomial trees. However, they may perform well when a task has sufficient bandwidth to support multiple, simultaneous, outgoing messages. For example, a trinomial tree (i.e., a k-nomial tree with $k = 3$) should exhibit good performance if there is enough bandwidth to send two messages simultaneously. The following illustration shows how CONCEPTUAL constructs a 27-task trinomial tree:

As before, time flows downward (assuming a multicast operation) and tasks are expected to communicate with their children in order. The following are some CONCEPTUAL k-nomial tree expressions and their evaluations, assuming `num_tasks` is '27':

- ‘KNOMIAL_PARENT(0, 3)’ $\Rightarrow$ ‘-1’
- ‘KNOMIAL_PARENT(2, 3)’ $\Rightarrow$ ‘0’
- ‘KNOMIAL_CHILDREN(2, 3)’ $\Rightarrow$ ‘4’
- ‘KNOMIAL_CHILD(2, 0, 3)’ $\Rightarrow$ ‘5’
- ‘KNOMIAL_CHILD(2, 1, 3)’ $\Rightarrow$ ‘8’
- ‘KNOMIAL_CHILD(2, 2, 3)’ $\Rightarrow$ ‘11’
- ‘KNOMIAL_CHILD(2, 3, 3)’ $\Rightarrow$ ‘20’
- ‘KNOMIAL_CHILD(2, 4, 3)’ $\Rightarrow$ ‘-1’
- ‘KNOMIAL_CHILDREN(8, 3)’ $\Rightarrow$ ‘2’
- ‘KNOMIAL_CHILDREN(8, 3, 26)’ $\Rightarrow$ ‘1’
- ‘KNOMIAL_CHILDREN(8, 3, 10)’ $\Rightarrow$ ‘0’

Mesh/torus functions

cONCEPTuaL provides three functions—MESH_NEIGHBOR, MESH_COORDINATE, and MESH_DISTANCE—that help treat (linear) task IDs as positions on a multidimensional mesh or torus. Each of these functions takes a variable number of arguments, determined by the dimensionality of the mesh (1-D, 2-D, or 3-D).

> *Terminology note*: In the context of network topologies in general and MESH_NEIGHBOR, MESH_COORDINATE, and MESH_DISTANCE in particular, "torus" refers to a mesh topology, of any number of dimensions, that contains wraparound links. That is, nodes on the right side of the mesh are directly connected to nodes on the left side of the mesh, nodes on the top of the mesh are directly connected to nodes on the bottom of the mesh, and so forth. A "partial torus" refers to a mesh topology of two or more dimensions in which at least one dimension contains wraparound links and at least one dimension does not.

MESH_NEIGHBOR (‘(’ *width* ‘*’? [, *height* ‘*’? [, *depth* ‘*’?]] ‘)’,   [Function]
        task_ID, ‘(’ *x_offset* [, *y_offset* [, *z_offset*]] ‘)’)
MESH_NEIGHBOR returns a task's neighbor on a 1-D, 2-D, or 3-D mesh or torus. It takes exactly three arguments: a list of the mesh/torus's dimensions, a task number, and a list of the neighbor's offset in each dimension from the given task. The two list arguments must be parenthesized, even if they contain only a single element. Each dimension in the dimension list may be followed by an asterisk to indicate that the mesh wraps around in that dimension. If not specified, *height* and *depth* default to ‘1’, and *y_offset* and *z_offset* default to ‘0’. In the absence of wraparound links, offsets that move off the mesh cause MESH_NEIGHBOR to return the value ‘-1’.

MESH_COORDINATE (‘(’ *width* ‘*’? [, *height* ‘*’? [, *depth* ‘*’?]] ‘)’,   [Function]
        task_ID, *coordinate*)
MESH_COORDINATE returns a task's x, y, or z coordinate on a 1-D, 2-D, or 3-D mesh/torus. The first argument to MESH_COORDINATE is a list of the mesh/torus's dimensions. The second argument is a task number. The third argument should be ‘0’ to calculate an x coordinate, ‘1’ to calculate a y coordinate, or ‘2’ to calculate a z coordinate. Coordinates are zero-origined. Each dimension in the dimension list may be followed by an asterisk to indicate that the mesh wraps around in that dimension,

but this has no impact on the result. (Asterisks are allowed solely for consistency with the other mesh functions.)

MESH_DISTANCE ('(' *width* '*'? [, *height* '*'? [, *depth* '*'?]] ')',   [Function]
      task_ID_1, *task_ID_2*)
      MESH_DISTANCE returns the shortest Manhattan distance between two tasks on a 1-D, 2-D, or 3-D mesh or torus. It takes exactly three arguments: a list of the mesh/torus's dimensions and two task numbers. The list argument must be parenthesized, even if it contains only a single element. Each dimension in the dimension list may be followed by an asterisk to indicate that the mesh wraps around in that dimension. If not specified, *height* and *depth* default to '1'.

MESH_NEIGHBOR, MESH_COORDINATE, and MESH_DISTANCE number tasks following the right-hand rule: left-to-right, then top-to-bottom, and finally back-to-front, as shown in the following illustrations of a 4-element (1-D) mesh, a 4×3 (2-D) mesh, and a $4 \times 3 \times 2$ (3-D) mesh. Examples of MESH_NEIGHBOR, MESH_COORDINATE, and MESH_DISTANCE for 1-D, 2-D, and 3-D meshes follow the corresponding illustration.

$$0 \longrightarrow 1 \longrightarrow 2 \longrightarrow 3$$

The following examples show how to use MESH_NEIGHBOR and MESH_COORDINATE to calculate neighbors and coordinates on the 1-D mesh shown above:

- 'MESH_NEIGHBOR((4), 0, (-1))' $\Rightarrow$ '-1'
- 'MESH_NEIGHBOR((4), 0, (+1))' $\Rightarrow$ '1'
- 'MESH_NEIGHBOR((4), 1, (-1))' $\Rightarrow$ '0'
- 'MESH_NEIGHBOR((4), 1, (+1))' $\Rightarrow$ '2'
- 'MESH_NEIGHBOR((4), 2, (-1))' $\Rightarrow$ '1'
- 'MESH_NEIGHBOR((4), 2, (+1))' $\Rightarrow$ '3'
- 'MESH_NEIGHBOR((4), 3, (-1))' $\Rightarrow$ '2'
- 'MESH_NEIGHBOR((4), 3, (+1))' $\Rightarrow$ '-1'

- 'MESH_COORDINATE((4), -1, 0)' $\Rightarrow$ '-1'
- 'MESH_COORDINATE((4), 0, 0)' $\Rightarrow$ '0'
- 'MESH_COORDINATE((4), 1, 0)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4), 2, 0)' $\Rightarrow$ '2'
- 'MESH_COORDINATE((4), 3, 0)' $\Rightarrow$ '3'
- 'MESH_COORDINATE((4), 4, 0)' $\Rightarrow$ '-1'
- 'MESH_COORDINATE((4), 2, 1)' $\Rightarrow$ '0'
- 'MESH_COORDINATE((4), 2, 2)' $\Rightarrow$ '0'

- 'MESH_DISTANCE((4), -1, 0)' $\Rightarrow$ '-1'
- 'MESH_DISTANCE((4), 1, 2)' $\Rightarrow$ '1'

- 'MESH_DISTANCE((4), 0, 3)' $\Rightarrow$ '3'

We can treat the 1-D mesh as a 1-D torus (a ring) by putting a '*' after the length of the x dimension:

- 'MESH_NEIGHBOR((4*), 0, (-1))' $\Rightarrow$ '3'
- 'MESH_NEIGHBOR((4*), 3, (+1))' $\Rightarrow$ '0'

- 'MESH_DISTANCE((4*), -1, 0)' $\Rightarrow$ '-1'
- 'MESH_DISTANCE((4*), 1, 2)' $\Rightarrow$ '1'
- 'MESH_DISTANCE((4*), 0, 3)' $\Rightarrow$ '1'

The remaining MESH_NEIGHBOR examples return the same value as before. MESH_COORDINATE *always* returns the same value regardless of the presence or absence of wraparound links.

```
0 --------- 1 --------- 2 --------- 3
|           |           |           |
|           |           |           |
4 --------- 5 --------- 6 --------- 7
|           |           |           |
|           |           |           |
8 --------- 9 --------- 10 -------- 11
```

The next set of examples shows how to use MESH_NEIGHBOR, MESH_COORDINATE, and MESH_DISTANCE to calculate neighbors, coordinates, and distances on the 2-D mesh shown above:

- 'MESH_NEIGHBOR((4, 3), 5, (-1, -1))' $\Rightarrow$ '0'
- 'MESH_NEIGHBOR((4, 3), 5, (0, -1))' $\Rightarrow$ '1'
- 'MESH_NEIGHBOR((4, 3), 5, (+1, -1))' $\Rightarrow$ '2'
- 'MESH_NEIGHBOR((4, 3), 5, (-1, 0))' $\Rightarrow$ '4'
- 'MESH_NEIGHBOR((4, 3), 5, (0, 0))' $\Rightarrow$ '5'
- 'MESH_NEIGHBOR((4, 3), 5, (+1, 0))' $\Rightarrow$ '6'
- 'MESH_NEIGHBOR((4, 3), 5, (-1, +1))' $\Rightarrow$ '8'
- 'MESH_NEIGHBOR((4, 3), 5, (0, +1))' $\Rightarrow$ '9'
- 'MESH_NEIGHBOR((4, 3), 5, (+1, +1))' $\Rightarrow$ '10'

- 'MESH_COORDINATE((4, 3), 1, 0)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4, 3), 6, 0)' $\Rightarrow$ '2'
- 'MESH_COORDINATE((4, 3), 6, 1)' $\Rightarrow$ '1'

- 'MESH_COORDINATE((4, 3), 6, 2)' ⇒ '0'
- 'MESH_COORDINATE((4, 3), 8, 0)' ⇒ '0'
- 'MESH_COORDINATE((4, 3), 8, 1)' ⇒ '2'
- 'MESH_COORDINATE((4, 3), 12, 0)' ⇒ '-1'

- 'MESH_DISTANCE((4, 3), 1, 2)' ⇒ '1'
- 'MESH_DISTANCE((4, 3), 2, 6)' ⇒ '1'
- 'MESH_DISTANCE((4, 3), 0, 5)' ⇒ '2'
- 'MESH_DISTANCE((4, 3), 3, 9)' ⇒ '4'
- 'MESH_DISTANCE((4, 3), 11, 0)' ⇒ '5'

Wraparound links turn some out-of-bounds ('-1') neighbor values into in-bounds values and reduce some of the shortest-path distances between tasks:

- 'MESH_NEIGHBOR((4, 3), 5, (-2, 0))' ⇒ '-1'
- 'MESH_NEIGHBOR((4, 3), 5, (-2, -2))' ⇒ '-1'
- 'MESH_NEIGHBOR((4, 3*), 5, (-2, 0))' ⇒ '-1'
- 'MESH_NEIGHBOR((4, 3*), 5, (-2, -2))' ⇒ '-1'
- 'MESH_NEIGHBOR((4*, 3), 5, (-2, 0))' ⇒ '7'
- 'MESH_NEIGHBOR((4*, 3), 5, (-2, -2))' ⇒ '-1'
- 'MESH_NEIGHBOR((4*, 3*), 5, (-2, 0))' ⇒ '7'
- 'MESH_NEIGHBOR((4*, 3*), 5, (-2, -2))' ⇒ '11'

- 'MESH_DISTANCE((4*, 3*), 1, 2)' ⇒ '1'
- 'MESH_DISTANCE((4*, 3*), 2, 6)' ⇒ '1'
- 'MESH_DISTANCE((4*, 3*), 0, 5)' ⇒ '2'
- 'MESH_DISTANCE((4*, 3*), 3, 9)' ⇒ '3'
- 'MESH_DISTANCE((4*, 3*), 11, 0)' ⇒ '2'

The final set of examples in this section shows how to use `MESH_NEIGHBOR`, `MESH_COORDINATE`, and `MESH_DISTANCE` to calculate neighbors, coordinates, and distances on the 3-D mesh shown above:

- 'MESH_NEIGHBOR((4, 3, 2), 0, (0, 0, +1))' $\Rightarrow$ '12'
- 'MESH_NEIGHBOR((4, 3, 2), 0, (0, +1, 0))' $\Rightarrow$ '4'
- 'MESH_NEIGHBOR((4, 3, 2), 0, (+1, 0, 0))' $\Rightarrow$ '1'
- 'MESH_NEIGHBOR((4, 3, 2), 0, (+1, +1, +1))' $\Rightarrow$ '17'
- 'MESH_NEIGHBOR((4, 3, 2), 17, (+2, -1, -1))' $\Rightarrow$ '3'
- 'MESH_NEIGHBOR((4, 3, 2), 23, (+1, +1, +1))' $\Rightarrow$ '-1'
- 'MESH_NEIGHBOR((4, 3, 2), 7, (+100, -200, +300))' $\Rightarrow$ '-1'

- 'MESH_COORDINATE((4, 3, 2), -5, 0)' $\Rightarrow$ '-1'
- 'MESH_COORDINATE((4, 3, 2), 1, 0)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4, 3, 2), 6, 0)' $\Rightarrow$ '2'
- 'MESH_COORDINATE((4, 3, 2), 6, 1)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4, 3, 2), 6, 2)' $\Rightarrow$ '0'
- 'MESH_COORDINATE((4, 3, 2), 18, 0)' $\Rightarrow$ '2'
- 'MESH_COORDINATE((4, 3, 2), 18, 1)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4, 3, 2), 18, 2)' $\Rightarrow$ '1'
- 'MESH_COORDINATE((4, 3, 2), 18, 3)' $\Rightarrow$ $\boxed{\text{error}}$ Invalid coordinate

- 'MESH_DISTANCE((4, 3, 2), 23, 24)' $\Rightarrow$ '-1'
- 'MESH_DISTANCE((4, 3, 2), 5, 17)' $\Rightarrow$ '1'
- 'MESH_DISTANCE((4, 3, 2), 1, 16)' $\Rightarrow$ '3'
- 'MESH_DISTANCE((4, 3, 2), 14, 8)' $\Rightarrow$ '5'
- 'MESH_DISTANCE((4, 3, 2), 0, 23)' $\Rightarrow$ '6'

As before, wraparound links affect some of the return values:

- 'MESH_NEIGHBOR((4*, 3*, 2*), 23, (+1, +1, +1))' $\Rightarrow$ '0'
- 'MESH_NEIGHBOR((4*, 3*, 2*), 7, (+100, -200, +300))' $\Rightarrow$ '11'

- 'MESH_DISTANCE((4*, 3*, 2*), 23, 24)' $\Rightarrow$ '-1'
- 'MESH_DISTANCE((4*, 3*, 2*), 5, 17)' $\Rightarrow$ '1'
- 'MESH_DISTANCE((4*, 3*, 2*), 1, 16)' $\Rightarrow$ '3'
- 'MESH_DISTANCE((4*, 3*, 2*), 14, 8)' $\Rightarrow$ '4'
- 'MESH_DISTANCE((4*, 3*, 2*), 0, 23)' $\Rightarrow$ '3'

Random-number functions

CONCEPTUAL programs can utilize randomness in one of two ways. The functions described below are *unsynchronized* across tasks. That is, they can—and usually do—return a different value to each task on each invocation. One consequence is that these functions are not permitted within a task expression (see Section 4.3 [Task descriptions], page 106) because randomness would cause the tasks to disagree about who the sources and targets of an operation are. In contrast, the `A RANDOM TASK` construct described in Section 4.7.3 [Binding variables], page 130 returns a value guaranteed to be synchronized across tasks and thereby enables random-task selection.

RANDOM_UNIFORM (`lower_bound`, `upper_bound`)   [Function]
> Return a number selected at random from a uniform distribution over the range [*lower_bound*, *upper_bound*).

RANDOM_GAUSSIAN (`mean`, `stddev`)   [Function]
> Return a number selected at random from a Gaussian distribution with mean *mean* and standard deviation *stddev*.

RANDOM_POISSON (`mean`)   [Function]
> Return an integer selected at random from a Poisson distribution with mean *mean* and standard deviation $\sqrt{mean}$.

RANDOM_PARETO (`shape`, `scale`)   [Function]
RANDOM_PARETO (`shape`, `lower_bound`, `upper_bound`)   [Function]
> With the two-argument form, return an integer selected at random from a Pareto distribution with shape *shape* and scale *scale*. With the three-argument form, return an integer selected at random from a bounded Pareto distribution with shape *shape* and bounded by the range [*lower_bound*, *upper_bound*].

Processor/task mapping functions

CONCEPTUAL supports a level of indirection in task numbering. The language defines a "processor" as a value that does not change for the task's entire lifetime and a "task" as a value that can be remapped at will. By default, task numbers match processor numbers, but if task numbers are remapped (using the `IS ASSIGNED TO` construct described in Section 4.8.4 [Reordering task IDs], page 137), the following functions can be used to query the current mapping:

> `TASK_OF` returns the task number associated with a given processor number.

> `PROCESSOR_OF` returns the processor number associated with a given task number.

> These functions are especially useful when a task `IS ASSIGNED TO A RANDOM PROCESSOR`, because there is no other mechanism for determining the resulting task mapping.

File-reading functions

`FILE_DATA` interacts with the surrounding system by parsing a named file as a table and returning the contents of a specified row and column.

FILE_DATA ('(' *filename* [, *column* [, *row* [, *column-seps* [,   [Function]
        row-seps]]]] ')')
> Read the contents of file *filename* and parse the data read into rows and columns.
> Rows are separated by one or more characters appearing in the string *row-seps* (de-
> fault: '\n', i.e., a newline character). Columns are separated by one or more charac-
> ters appearing in the string *col-seps* (default: ' \t', i.e., a space and a tab). Once the
> data are parsed in this manner, the cell at column *column* (default: '1') and row *row*
> (default: '1') is converted to a number and returned.
>
> Columns and rows are each numbered starting from 1. Negative numbers represent
> offsets from the *last* column or row. Zero values cause the program to abort with an
> error message.

Consider a data file, `example.dat`, with the following contents:

```
Eir    231  58 490 703  19 644 967
Gunnr  901 875 141 372 374 618  77
Herja   17 979 363 930 977
Hrist  100 132 981 246  93 614
Svipul 688 536 965 746 209
```

Then, the following uses of `FILE_DATA` return the results shown:

- 'FILE_DATA("example.dat")' ⇒ '0'
  (The default cell is (1, 1), which contains 'Eir', and this is converted to 0.)
- 'FILE_DATA("example.dat", 3)' ⇒ '58'
- 'FILE_DATA("example.dat", 3, 2)' ⇒ '875'
- 'FILE_DATA("example.dat", 3, -4)' ⇒ '875'
- 'FILE_DATA("example.dat", 3, -6)' ⇒ ⬚error⬚ Too few rows
- 'FILE_DATA("example.dat", -3, -2)' ⇒ '246'
  (The function first finds the row, then the column within the row.)
- 'FILE_DATA("example.dat", 0)' ⇒ ⬚error⬚ Column number is zero

The use of `FILE_DATA` is discouraged in coNCePTuaL programs because of its reliance
on external files, whose contents are not automatically preserved in any coNCePTuaL log
files (see Section 3.5 [Interpreting coNCePTuaL log files], page 49) and which may exhibit
system-specific behavior (e.g., if the file is in fact a named pipe or a file-like interface
to a kernel data structure). However, `FILE_DATA` can be useful in certain circumstances
for logging changes in system state of which coNCePTuaL would otherwise be unaware.
For example, one might log values from various pseudo-files appearing in Linux's `/proc`
filesystem or from system log files (e.g., those within Linux's `/var/log` directory).

Because coNCePTuaL cannot determine a priori if multiple reads of the same row
and column from the same file will produce the same value it pessimistically assumes that
each read produces a different value. (Consider even the simple case of `FILE_DATA` reading
the last row of a system log file to which another process is asynchronously appending.)
Consequently, `FILE_DATA` cannot be used in any context that requires multiple tasks to
agree on a value, such as within message specifications (see Section 4.4.1 [Message specifica-
tions], page 108) or let bindings (see Section 4.7.3 [Binding variables], page 130). For cases

where the programmer can assert that reads from a particular row and column will always return the same value (on penalty of undefined behavior, including hangs and crashes), CONCEPTUAL provides the following function:

STATIC_FILE_DATA (‘(’ *filename* [, *column* [, *row* [, *column-seps* [,   [Function]
        row-seps]]]] ‘)’)
> Perform exactly the same operation as FILE_DATA but with the programmer's guarantee that the file contents will not change from invocation to invocation. STATIC_FILE_DATA can therefore be used in places where FILE_DATA is disallowed.

Given the example.dat file shown previously, the following let binding is legitimate:

```
LET msg_size BE STATIC_FILE_DATA("example.dat", 2, iter) AND rank_ofs
BE STATIC_FILE_DATA("example.dat", 3, iter) MOD num_tasks WHILE ...
```

However, the following, which uses FILE_DATA instead of STATIC_FILE_DATA, is not:

```
LET msg_size BE FILE_DATA("example.dat", 2, iter) AND rank_ofs BE
FILE_DATA("example.dat", 3, iter) MOD num_tasks WHILE ...
```

STATIC_FILE_DATA can be used as above to implement complex, irregular communication patterns that vary across runs of the program. That is, each iteration, a different value read from the file can be used to specify the message size and the distance between senders and receivers.

The use of STATIC_FILE_DATA is even more discouraged than the use of FILE_DATA because an incorrect claim that a file's contents are unvarying could lead to abnormal program behavior—and not necessarily easily detectable. It is best to avoid both FILE_DATA and STATIC_FILE_DATA if possible. If not, use FILE_DATA unless (a) the compiler rejects the program on the grounds that FILE_DATA is being used in an invalid location and (b) the data being read is known to be invariant throughout the lifetime of the program.

4.2.3 Aggregate expressions

Aggregate expressions (<*aggr_expr*>s) are currently used exclusively by the LOGS statement. They represent an expression with a given function applied to the aggregate of all (dynamic) instances of that expression. <*aggr_expr*>s take one of four forms:

<*aggr_expr*>   ::=   [EACH] <*expr*>
                | THE <*expr*>
                | THE <*aggr_func*> [AND THE <*aggr_func*>]* [OF [THE]] <*expr*>
                | A HISTOGRAM OF [THE] <*expr*>

(In the above, <*expr*> refers to an arithmetic expression defined in Section 4.2.1 [Arithmetic expressions], page 86 and <*aggr_func*> refers to one of the functions defined in Section 4.2.4 [Aggregate functions], page 101.)

The first form does not summarize <*expr*>; every individual instance of <*expr*> is utilized. The second form asserts that <*expr*> is a constant (i.e., all values are identical) and utilizes that constant.[5] The third form applies each <*aggr_func*> to the set of all values of <*expr*> and utilizes the result of that function. The fourth form produces a histogram of all values of <*expr*>, i.e., a list of {*unique value*, *tally*} pairs, sorted by *unique value*.

---

[5] The program aborts with a run-time error if <*expr*> is not a constant.

4.2.4 Aggregate functions

The following functions, referred to collectively as `<aggr_func>`s, may be used in an aggregate expression (see Section 4.2.3 [Aggregate expressions], page 100):

`<aggr_func>`   ::=   [ARITHMETIC] MEAN | HARMONIC MEAN | GEOMETRIC MEAN | MEDIAN |
                      STANDARD DEVIATION | VARIANCE | MEDIAN ABSOLUTE DEVIATION |
                      SUM | MINIMUM | MAXIMUM | FINAL | `<expr>` PERCENTILE

`MEAN` and `ARITHMETIC MEAN` are equivalent. `MEDIAN` is the value such that there are as many larger as smaller values. If there are an even number of values, `MEDIAN` is the arithmetic mean of the two medians. `FINAL` returns only the final value measured. `PERCENTILE` returns a given percentile, `<expr>`, of the values. `<expr>` must lie in the range [0, 100]. Percentile values are computed as $x_{\lfloor h \rfloor} + (h - \lfloor h \rfloor)(x_{\lfloor h \rfloor + 1} - x_{\lfloor h \rfloor})$, where x is the data in sorted order and h is $(|x| - 1) \cdot p/100 + 1$ when computing the value at the pth percentile. (See Wikipedia's entry on *quantile* for more information.) The interpretation of the remaining aggregate functions should be unambiguous.

4.2.5 Relational expressions

Relational expressions (`<rel_expr>`s) compare two arithmetic expressions (see Section 4.2.1 [Arithmetic expressions], page 86) or test an arithmetic expression for a property. A relational expression can be either TRUE or FALSE.

CONCEPTUAL supports a variety of relational expressions. The following is the language's order of operations from highest to lowest precedence:

unary/   IS EVEN, IS ODD
binary/   '=', '<', '>', '<=', '>=', '<>', DIVIDES, IS IN, IS NOT IN

conjunctive   '/\'

disjunctive   '\/'

In addition, as in most programming languages, parentheses can be used to group subexpressions.

The unary relation `IS EVEN` is TRUE if a given arithmetic expression represents an even number and the unary relation `IS ODD` is TRUE if a given arithmetic expression represents an odd number. For example, '456 IS EVEN' is TRUE and '64 MOD 6 IS ODD' is FALSE.

The CONCEPTUAL operators '=', '<', '>', '<=', '>=', and '<>' represent, respectively, the mathematical relations $=$, $<$, $>$, $\leq$, $\geq$, and $\neq$. These are all binary relations that operate on arithmetic expressions (see Section 4.2.1 [Arithmetic expressions], page 86). For example, '2+2 = 4' is TRUE and '2**3 > 2**4' is FALSE. The `DIVIDES` relation is TRUE if the first expression evenly divides the second, i.e., that $e_2 \equiv 0 \pmod{e_1}$. Hence, '2 DIVIDES 1234' (equivalent to '1234 MOD 2 = 0') is TRUE while '2 DIVIDES 4321' (equivalent to '4321 MOD 2 = 0') is FALSE.

The binary relation `IS IN` has the form

            `<expr>` IS IN `<range>` [, `<range>`]*

A `<range>` represents a range expression. Range expressions are described in Section 4.2.6 [Range expressions], page 102. In short, a range expression specifies a list of values by explicit enumeration, numeric progression, or predicated combinations of

other range expressions. As an example, the relational expression 'x IS IN {1, ..., 5}'
is TRUE if and only if x is one of 1, 2, 3, 4, or 5. As a more complex example, 'p*2 IS
IN {0}, {1, 2, 4, ..., num_tasks*2}' is TRUE if and only if twice p is either zero or a
power of two less than or equal to twice the number of tasks being used.

The complementary operation to IS IN is the binary relation IS NOT IN. Hence, '4 IS
NOT IN {3, ..., 5}' is FALSE while '6 IS NOT IN {3, ..., 5}' is TRUE.

Conjunction ($\wedge$) and disjunction ($\vee$) combine multiple relational expressions.
<rel_expr> '/\' *<rel_expr>* is TRUE if and only if both *<rel_expr>*s are TRUE, and
<rel_expr> '\/' *<rel_expr>* is TRUE if and only if either *<rel_expr>* is TRUE. For
example, '456 IS EVEN \/ 2**3 > 2**4' is TRUE and '456 IS EVEN /\ 2**3 > 2**4' is
FALSE. Conjunction and disjunction are both short-circuiting operations. Evaluation
proceeds left-to-right. Expressions such as 'x<>0 /\ 1/x=1' will therefore not result in a
divide-by-zero error.

CONCEPTUAL does not currently have a logical negation operator.

Formal grammar for relational expressions

For completeness, the following productions formalize the process by which CONCEPTUAL
parses relational expressions:

<rel_expr>   ::=   *<rel_disj_expr>*

<rel_disj_expr>   ::=   [*<rel_disj_expr>* '\/'] *<rel_conj_expr>*

<rel_conj_expr>   ::=   [*<rel_conj_expr>* '/\'] *<rel_primary_expr>*

<rel_primary_expr>   ::=   *<eq_expr>*
                        | '(' *<rel_expr>* ')'

<eq_expr>   ::=   *<expr>* '=' *<expr>*
                | *<expr>* '<' *<expr>*
                | *<expr>* '>' *<expr>*
                | *<expr>* '<=' *<expr>*
                | *<expr>* '>=' *<expr>*
                | *<expr>* '<>' *<expr>*
                | *<expr>* DIVIDES *<expr>*
                | *<expr>* IS EVEN
                | *<expr>* IS ODD
                | *<expr>* IS IN *<range>* [, *<range>*]*
                | *<expr>* IS NOT IN *<range>* [, *<range>*]*

4.2.6 Range expressions

Ranges are a powerful way of describing an ordered list of tasks. Ranges come in three
forms:

1. Enumerated lists

2. Arithmetic and geometric sequences

3. List comprehensions

These are described with the following syntax:

<range>   ::=   '{' *<expr>* [',' *<expr>*]* [', ... ,' *<expr>*] '}'

```
|     '{' <expr> [FOR EACH <ident> IN <range> [',' <range>]* ]+
          [WHERE <rel_expr>] '}'
```

Enumerated lists are the simplest type of range. They represent a list of explicitly specified values. For example, the range '{2**3, 2+4, 14/2, 8-3, 1+1+1, 0, 5*2-1}' represents each of the seven numbers 8, 6, 7, 5, 3, 0, and 9 in turn. Duplicate values are allowed in enumerated lists so '{9, 3, 9, 5, 5, 5, 0, 1, 1, 3}' represents ten numbers even though not all of them are unique.

Arithmetic and geometric progressions are specified by including the first few values of the progression, followed by an ellipsis, followed by the final value of the progression. A single initial value implies an arithmetic progression with increment of ± 1, depending on whether the final value is greater or less than the initial value. Two initial values imply an arithmetic progression with an increment of $second-first$. If three or more initial values are provided, CONCEPTUAL first looks for an arithmetic progression, then a geometric progression. If CONCEPTUAL finds neither an arithmetic nor a geometric progression, it re-evaluates all of the <expr>s in floating-point context (see Section 4.2.1 [Arithmetic expressions], page 86) and tries once again to find a geometric progression. If a pattern is still not found, a run-time error message is generated. As some examples of arithmetic and geometric range expressions, the range '{10, ..., 1E6}' represents the values 10, 11, 12, 13, 14, 15, and so forth by ones up to 1,000,000; the range '{10, 12, ..., 1E6}' represents the values 10, 12, 14, 16, and so forth by twos up to 1,000,000; and the range '{10, 100, 1000, ..., 1E6}' represents the values 10, 100, 1000, 10000, 100000, and 1000000.

Arithmetic and geometric progressions do not necessarily include the final value. For example, the range '{1, 4, ..., 15}' represents the values 1, 4, 7, 10, and 13 but not 15, even though 15 is specified as the final value of the range. Similarly, '{15, 12, ..., 1}' represents the values 15, 12, 9, 6, and 3 but not 1, even though 1 is specified as the final value of the range. Progressions stop at or before the final value, never after. If the number following the ellipsis is less than (respectively, greater than) the first number in an increasing (respectively, decreasing) range (as in '{15, 25, ..., 5}'), then the range represents an empty list of values.

List comprehensions provide a way to combine and filter ranges to describe complex sequences of values with comparatively little code.[6] A simple example is '{n FOR EACH n IN {1, ..., 10}}', but this is not particularly useful, as the range '{1, ..., 10}' represents the same list of values. However, replacing the expression 'n' with 'n/2' to get '{n/2 FOR EACH n IN {1, ..., 10}}' returns the values 0, 1, 1, 2, 2, 3, 3, 4, 4, and 5, which would be more difficult to express without a list comprehension. We can even utilize only a subset of those values by applying a filter. For example, '{n/2 FOR EACH n IN {1, ..., 10} WHERE n/2 IS ODD}' represents the values 1, 1, 3, 3, and 5. List comprehensions can be multivariate, as in the following example:

```
{diag+ofs
   FOR EACH diag IN {0, 11, ..., 99}
     FOR EACH ofs IN {-1, 0, 1}
       WHERE diag+ofs IS IN {0, ..., 99}}
```

---

[6] Readers unfamilier with the concept of a list comprehension may want to consult the Wikipedia article on list comprehensions for more information.

That single range expression represents all cells on the tridiagonal of a 10×10 matrix, with cell 0 in the upper left and cell 99 in the lower right. The first FOR EACH generator binds 'diag' to each cell on the diagonal: 0, 11, 22, 33, and so forth. The second FOR EACH generator binds 'ofs' to each of -1, 0, and 1. The 'WHERE' predicate, which itself uses a range expression, selects only those values of 'diag' and 'ofs' which which $0 \leq$ diag+ofs ≤ 99. The value 'diag+ofs' is returned for all values for which the predicate is true. The following figure illustrates the values described by the preceding range expression:

0	1	2	3	4	5	6	7	8	9
10	11	12	13	14	15	16	17	18	19
20	21	22	23	24	25	26	27	28	29
30	31	32	33	34	35	36	37	38	39
40	41	42	43	44	45	46	47	48	49
50	51	52	53	54	55	56	57	58	59
60	61	62	63	64	65	66	67	68	69
70	71	72	73	74	75	76	77	78	79
80	81	82	83	84	85	86	87	88	89
90	91	92	93	94	95	96	97	98	99

Here are a number of additional examples of range expressions and the values they represent:

'{2, 2, 2, 2, 2, 2}'
$\Rightarrow$ 2, 2, 2, 2, 2, 2

'{2, 2, ..., 2}'
$\Rightarrow$ 2

'{2, 2, ..., 1000}'
$\Rightarrow$ 2

'{20, 30, 40, ..., 55}'
$\Rightarrow$ 20, 30, 40, 50

'{20, 30, 40, ..., 30}'
$\Rightarrow$ 20, 30

'{20, 30, 40, ..., 20}'
$\Rightarrow$ 20

'{20, 30, 40, ..., 10}'
$\Rightarrow \emptyset$

'{2, ..., 50}'
$\Rightarrow$ 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50

'`{2, 4, ..., 50}`'
> ⇒ 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 42, 44, 46, 48, 50

'`{2, 4, 4, ..., 50}`'
> ⇒ ⌐error⌐ Unable to find either an arithmetic or geometric pattern

'`{2, 4, 6, ..., 50}`'
> ⇒ 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26, 28, 30, 32, 34, 36, 38, 40, 42, 44, 46, 48, 50

'`{2, 4, 8, ..., 50}`'
> ⇒ 2, 4, 8, 16, 32

'`{1, 3, ..., 25}`'
> ⇒ 1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25

'`{1, 3, 5, 7, 9, 11, 13, ..., 25}`'
> ⇒ 1, 3, 5, 7, 9, 11, 13, 15, 17, 19, 21, 23, 25

'`{6561, 2187, 729, ..., 1}`'
> ⇒ 6561, 2187, 729, 243, 81, 27, 9, 3, 1

'`{100, 150, 225, ..., 10000}`'
> ⇒ 100, 150, 225, 337, 506, 759, 1139, 1708, 2562, 3844, 5766, 8649

'`{1, 4, 9, ..., 81}`'
> ⇒ ⌐error⌐ Unable to find either an arithmetic or geometric pattern

'`{sqrt(s)**2 FOR EACH s IN {1, ..., 16}}`'
> ⇒ 1, 1, 1, 4, 4, 4, 4, 4, 9, 9, 9, 9, 9, 9, 9, 16

'`{y*10+x`
  `FOR EACH y IN {0, ..., 6}`
    `FOR EACH x IN {0, ..., 6}`
      `WHERE (x IS EVEN /\ y IS ODD) \/ (y IS EVEN /\ x IS ODD)}`'
> ⇒ 1, 3, 5, 10, 12, 14, 16, 21, 23, 25, 30, 32, 34, 36, 41, 43, 45, 50, 52, 54, 56, 61, 63, 65

The preceding examples used exclusively constant expressions. However, as indicated by the definition of <*range*> above, any arbitrary arithmetic expression (see Section 4.2.1 [Arithmetic expressions], page 86) is valid. Assuming that 'x' is currently bound to 123, the following expansion holds:

'`{x*10, x*9, x*8, ..., x*3}`'
> ⇒ 1230, 1107, 984, 861, 738, 615, 492

4.3 Task descriptions

Task descriptions are a powerful way of tersely describing the sources and targets of CON-
CEPTUAL operations. Task IDs range from 0 to '`num_tasks-1`' (see
). Operations involving out-of-bound task IDs are silently
ignored.

As a side effect, a task description can declare a variable that can be used in subsequent
expressions. (See .) There are two types of task descrip-
tions: one for "source" tasks and one for "target" tasks. The two are syntactically similar
but semantically different. Specifically, the scope of a variable declared in a <*target_tasks*>
specification is more limited than one declared in a <*source_task*> specification.

Before introducing <*source_task*> and <*target_tasks*> specifications we first introduce
the notion of a <*restricted_ident*>, which is a variable declaration that can be used to define
a set of tasks. We then present CONCEPTUAL's complete set of mechanisms for describing
sets of source and target tasks.

4.3.1 Restricted identifiers

A *restricted identifier* declares a variable, restricting it to the set of tasks that satisfy
a given relational expression (see). The
syntax, shown below, represents the mathematical notion of "$\forall \langle ident \rangle | (\langle rel\_expr \rangle \wedge (0 \leq \langle ident \rangle < \langle \#tasks \rangle)))$". That is, "for all <*ident*> such that <*rel_expr*> is TRUE and <*ident*>
is between zero and the number of tasks. . .".

<*restricted_ident*>   ::=   <*ident*> SUCH THAT <*rel_expr*>

As an example, '`evno SUCH THAT evno IS EVEN`' describes all even-numbered tasks. On
each such task, the variable '`evno`' takes on that task's ID. Similarly, '`thr SUCH THAT 3
DIVIDES thr-1`' describes tasks 1, 4, 7, 10, 13. . . . On each of those tasks, '`thr`' will be
bound to the task ID. On all other tasks, '`thr`' will be undefined. When order matters (as
in the cases described in and
), <*ident*> takes on task IDs in increasing order.

4.3.2 Source tasks

A <*source_task*> specification takes one of six forms:

<*source_task*>   ::=   ALL TASKS
               | ALL TASKS <*ident*>
               | TASK <*expr*>
               | TASKS <*restricted_ident*>
               | TASK GROUP <*ident*>
               | TASKS <*range*> [, <*range*>]*

ALL TASKS specifies that each task will perform a given operation. If followed by a
variable name (<*ident*>), each task will individually bind <*ident*> to its task ID—a number
from zero to one less than the total number of tasks. That is, '`ALL TASKS me`' will bind '`me`'
to '`0`' on task 0, '`1`' on task 1, and so forth.

'`TASK <expr>`' specifies that only the task described by arithmetic expression <*expr*> will
perform the given operation. For example, '`TASK 2*3+1`' says that only task 7 will act; the
other tasks will do nothing.

'TASKS `<restricted_ident>`' describes a set of tasks that will perform a given operation. For instance, 'TASKS x SUCH THAT x>0 /\ x<num_tasks-1'—read as "tasks x such that x is greater than zero and x is less than *num_tasks* minus one"—expresses that a given operation should be performed on all tasks except the first and last in the computation. On each task that satisfies the relational expression, 'x' will be bound to the task ID as in ALL TASKS above. Hence, 'x' will be undefined on task 0, '1' on task 1, '2' on task 2, and so forth up to task 'num_tasks-1', on which 'x' will again be undefined.

'TASK GROUP `<ident>`' specifies that a previously defined group of tasks will perform a given operation. This form of *<source_task>* is described further in Section 4.7.3 [Binding variables], page 130.

'TASKS `<range>`' is syntactic sugar for the *<restricted_ident>* case with an anonymous *<ident>* and an IS IN relation. For example, 'TASKS {3, 5, 8}' is equivalent to 'TASKS var SUCH THAT var IS IN {3, 5, 8}' except that it does not bind variable 'var'.

As per the definitions in Section 4.1 [Primitives], page 84 and Section 4.3.1 [Restricted identifiers], page 106, respectively, *<ident>*s and *<restricted_ident>*s do not accept parentheses. Hence, 'TASKS (bad SUCH THAT bad IS EVEN)' and 'ALL TASKS (no_good)' result in parse errors while 'TASKS fine SUCH THAT fine IS EVEN' and 'ALL TASKS dandy' are acceptable constructs. As an analogy, 'x = 3' is valid in many general-purpose programming languages while '(x) = 3' is not.

Variables declared in a 'source_task' specification are limited in scope to the surrounding statement.

4.3.3 Target tasks

A *<target_tasks>* specification takes one of six forms:

<target_tasks>	::=	ALL OTHER TASKS
		| TASK *<expr>*
		| TASKS *<restricted_ident>*
		| TASK GROUP *<ident>*
		| TASKS *<range>* [, *<range>*]*

ALL OTHER TASKS is just like ALL TASKS in a *<source_task>* specification (see Section 4.3.2 [Source tasks], page 106) but applies to all tasks *except* the source task. Also, unlike ALL TASKS, ALL OTHER TASKS does not accept an *<ident>* term.

The remaining *<target_tasks>* variants behave like their *<source_task>* counterparts (see Section 4.3.2 [Source tasks], page 106), but their variables are evaluated in the scope of the corresponding *<source_task>*. For example, in the *<source_task>* 'TASKS src SUCH THAT 3 DIVIDES src', the variable 'src' can be used in a *<target_tasks>* such as 'TASKS dst SUCH THAT dst>src'.

4.4 Communication statements

Communication statements are the core of any CONCEPTUAL program. The CONCEPTUAL language makes it easy to express a variety of communication features:

- synchronous or asynchronous communication
- unaligned, aligned (to arbitrary byte boundaries), or misaligned (from a page boundary) message buffers

- ignored, touched, or verified message contents
- unique or recycled message buffers
- point-to-point or collective operations

Communication statements are performed by an arbitrary *<source_task>* (see Section 4.3.2 [Source tasks], page 106) and may involve arbitrary *<target_tasks>* (see Section 4.3.3 [Target tasks], page 107). After explaining how to describe a message to cONCEPTuAL (see Section 4.4.1 [Message specifications], page 108) this section presents each communication statement in turn and explains its purpose, syntax, and semantics.

4.4.1 Message specifications

A *message specification* describes a set of messages. The following is a formal definition:

```
<message_spec>        ::=     <item_count>
                              [NONUNIQUE | UNIQUE]
                              <item_size>
                              [UNALIGNED |
                               <message_alignment> ALIGNED |
                               <message_alignment> MISALIGNED]
                              MESSAGES
                              [WITH VERIFICATION | WITH DATA TOUCHING |
                               WITHOUT VERIFICATION | WITHOUT DATA TOUCHING]
                              [USING TAG
                               <expr> | <string>]
                              [FROM [<expr> <data_multiplier> INTO]
                               BUFFER <expr> | THE DEFAULT BUFFER]
```

Within a `RECEIVE` statement (see Section 4.4.3 [Receiving], page 116), a *<message_spec>*'s `FROM` keyword must be replaced with `INTO`.

A `SEND` statement's `WHO RECEIVES IT` clause (see Section 4.4.2 [Sending], page 113) utilizes a slightly different message specification, which is referred to here as a *<recv_message_spec>*:

```
<recv_message_spec>   ::=     [SYNCHRONOUSLY | ASYNCHRONOUSLY]
                              [AS [A|AN]
                               [NONUNIQUE | UNIQUE]
                               [UNALIGNED |
                                <message_alignment> ALIGNED |
                                <message_alignment> MISALIGNED]
                               MESSAGES]
                              [WITH VERIFICATION | WITH DATA TOUCHING |
                               WITHOUT VERIFICATION | WITHOUT DATA TOUCHING]
                              [USING TAG
                               <expr> | <string>]
                              [INTO [<expr> <data_multiplier> INTO]
                               BUFFER <expr> | THE DEFAULT BUFFER]
```

Although not indicated by the preceding grammatical rule, a *<recv_message_spec>* is not allowed to be empty. That is, at least one of the optional clauses must be specified.

A `REDUCE` statement (see Section 4.4.6 [Reducing], page 118) utilizes the following variations of *<message_spec>* and *<recv_message_spec>*, respectively:

```
<reduce_message_spec>    ::=   <item_count>
                               [NONUNIQUE | UNIQUE]
                               [UNALIGNED |
                                <message_alignment> ALIGNED |
                                <message_alignment> MISALIGNED]
                               INTEGERS | DOUBLEWORDS
                               [WITH DATA TOUCHING | WITHOUT DATA TOUCHING]
                               [USING TAG
                                <expr> | <string>]
                               [FROM [<expr> <data_multiplier> INTO]
                               BUFFER <expr> | THE DEFAULT BUFFER]

<reduce_target_message_spec>  ::=   [AS <item_count>
                                     [NONUNIQUE | UNIQUE]
                                     [<message_alignment> ALIGNED |
                                      <message_alignment> MISALIGNED]
                                     INTEGERS | DOUBLEWORDS]
                                     [WITH DATA TOUCHING | WITHOUT DATA TOUCHING]
                                     [USING TAG
                                      <expr> | <string>]
                                     [INTO [<expr> <data_multiplier> INTO]
                                     BUFFER <expr> | THE DEFAULT BUFFER]
```

We now describe in turn each component of a *<message_spec>*, *<recv_message_spec>*, *<reduce_message_spec>*, and *<reduce_target_message_spec>*.

Item count

The *<item_count>* says how many messages the *<message_spec>* represents:

```
<item_count>   ::=   A | AN | <expr>
```

`A` and `AN` are synonyms for the value '`1`'.

Unique messages

Normally, the CONCEPTUAL backends recycle message memory to reduce the program's memory requirements and improve performance. By adding the keyword `UNIQUE`, every message buffer will reside in a unique memory region. `NONUNIQUE` explicitly specifies the default, buffer-recycling behavior.

Item size

The message size is represented by the *<item_size>* nonterminal. It can be empty or expressed in one of two other ways:

```
<item_size>   ::=   ε
              |     <expr> <data_multiplier>
              |     <data_type> SIZED
```

A <*data_multiplier*> is a scaling factor that converts a unitless number into a number of bytes. The following are the valid possibilities for <*data_multiplier*> and the number of bytes by which they multiply <*expr*>:

<*data_multiplier*>   ::=   BIT | BYTE | HALFWORD | WORD | INTEGER | DOUBLEWORD |
                           QUADWORD | PAGE | KILOBYTE | MEGABYTE | GIGABYTE

BIT   1/8 bytes, rounded up to the nearest integral number of bytes

BYTE   1 byte

HALFWORD   2 bytes

WORD   4 bytes

INTEGER   the number of bytes in the backend's fundamental integer type

DOUBLEWORD
           8 bytes

QUADWORD   16 bytes

PAGE   the number of bytes in an operating-system page

KILOBYTE   1,024 bytes

MEGABYTE   1,048,576 bytes

GIGABYTE   1,073,741,824 bytes

A <*data_type*> is an "atomic" unit of data. It can be any of the following:

<*data_type*>   ::=   BYTE | HALFWORD | WORD | INTEGER | DOUBLEWORD |
                     QUADWORD | PAGE

BYTE   1 byte

HALFWORD   2 bytes

WORD   4 bytes

INTEGER   the number of bytes in the backend's fundamental integer type

DOUBLEWORD
           8 bytes

QUADWORD   16 bytes

PAGE   the number of bytes in an operating-system page

Hence, valid <*item_size*>s include, for example, '16 MEGABYTE' or 'PAGE SIZED'. Empty <*item_size*>s (shown in the grammar as ε) are equivalent to '0 BYTE'. Note that INTEGER varies in size based on the backend, backend compiler, and CPU architecture but is commonly either 4 or 8 bytes; PAGE varies in size from operating system to operating system; each of the other <*data_type*>s has a fixed size, as indicated above.

Message alignment

Messages are normally allocated with arbitrary alignment in memory. However, CoNCeP-TuaL can force a specific alignment relative to the operating-system page size (commonly 4 KB or 8 KB, but significantly larger sizes are gaining popularity). A *<message_alignment>* is represented as follows:

<message_alignment>   ::=   *<data_type>*
                          | *<expr>* *<data_multiplier>*

'64 BYTE', '3 MEGABYTE', and 'QUADWORD' are therefore all valid examples of *<message_alignment>*s. Bit counts are rounded up to the nearest byte count, so '27 BITS' is in fact equivalent to '4 BYTES'.

The ALIGNED keyword forces CoNCePTuaL to align messages on *exactly* the specified alignment. Hence, a 'HALFWORD ALIGNED' message can begin at memory locations 0, 2, 4, 6, 8, ..., $2k$ ($k \in \mathbb{Z}^+$). In contrast, the MISALIGNED keyword forces CoNCePTuaL to align messages the given number of bytes (positive or negative) past a page boundary. For example, if pages are 8192 bytes in size then a message described as 'HALFWORD MISALIGNED' can begin at memory locations 2, 8194, 16386, 24578, ..., $8192k + 2$ ($k \in \mathbb{Z}^+$). Unlike ALIGNED, MISALIGNED supports negative alignments. If the page size is 4096 bytes, then '-10 BYTE MISALIGNED' enables a message to begin at memory locations 4086, 8182, 12278, etc. The MISALIGNED alignment is taken modulo the page size. Therefore, with a 4096-byte page size, '10000 BYTE MISALIGNED' is the same as '1808 BYTE MISALIGNED'.

The UNALIGNED keyword explicitly specifies the default behavior, with messages aligned on arbitrary boundaries.

Data touching

A *<message_spec>* described as being WITH DATA TOUCHING will force every word in a message to be both read and written ("touched"). When *<message_spec>* describes an outgoing message, the data will be touched before transmission. When *<message_spec>* describes an incoming message, the data will be touched after reception. In a sense, WITH DATA TOUCHING presents a more realistic assessment of network performance, as real applications almost always access the data they send or receive. It also distinguishes between messaging layers that implicitly touch data and those that can transmit data without having to touch it. One would expect the latter to perform better when the data is not touched, as the former may be paying a penalty for touching the data. However, either could perform better when messages are sent WITH DATA TOUCHING, because the latter now has to pay the penalty that the former has already paid.

Another form of data-touching supported by CoNCePTuaL is WITH VERIFICATION. This causes the source task to write known, but randomly generated, data into the message before transmission and the target task to verify that every bit was correctly received. When a message is received WITH VERIFICATION, the bit_errors variable (see Section A.2 [Predeclared variables], page 207) is updated appropriately.

WITHOUT DATA TOUCHING and WITHOUT VERIFICATION are synonymous. Both explicitly specify the default behavior of neither touching nor verifying message contents.

Tag matching

Messages sent from task A to task B are normally received in order. However, the `USING TAG` clause enables a program to enforce a particular order on message reception. `USING TAG` takes a tag as an argument. This tag can be either an *<expr>* or a *<string>*. In the latter case, CONCEPTUAL will automatically hash the string to an *<expr>*. A message sent with a given tag *tag* is received only by a `RECEIVE` statement that also specifies tag *tag*. If the `USING TAG` clause is omitted, *tag* defaults to '0'.

As an example of the difference `USING TAG` can make, consider the following code that does not use tags:

```
TASK 0 SENDS A MESSAGE TO UNSUSPECTING TASK 1 THEN
TASK 0 SLEEPS FOR 2 SECONDS THEN
TASK 0 SENDS A MESSAGE TO UNSUSPECTING TASK 1 THEN
TASK 1 RECEIVES A MESSAGE FROM TASK 0 THEN
TASK 1 SLEEPS FOR 2 SECONDS THEN
TASK 1 RECEIVES A MESSAGE FROM TASK 0
```

In that example, the first send matches the first receive, then both tasks simultaneously sleep for two seconds, and finally the second send matches the second receive. The total execution time should therefore be just over two seconds. Now consider the following variation that uniquely tags each of the messages:

```
TASK 0 SENDS A MESSAGE USING TAG 123 TO UNSUSPECTING TASK 1 THEN
TASK 0 SLEEPS FOR 2 SECONDS THEN
TASK 0 SENDS A MESSAGE USING TAG "stuff" TO UNSUSPECTING TASK 1 THEN
TASK 1 RECEIVES A MESSAGE USING TAG "stuff" FROM TASK 0 THEN
TASK 1 SLEEPS FOR 2 SECONDS THEN
TASK 1 RECEIVES A MESSAGE USING TAG 123 FROM TASK 0
```

The first message is sent with tag '123' while the second message is sent with tag '"stuff"'. Task 1 receives the messages in the reverse order: first the '"stuff"' message and then the '123' message. Consequently, task 1's first `RECEIVE` cannot complete until task 0 has sent its first message, slept for two seconds, and sent its second message. Task 1 can then receive a message, sleep for two seconds, and receive its other message. The total execution time should therefore be just over four seconds, as the new message ordering forces a serialization on the two `SLEEPS`.

The following are a few things to keep in mind when using message tags:

1. Using strings for tag values can benefit program readability and is recommended for programs that utilize a small, fixed set of tags, while using numerical expressions for tag values supports the implementation of a variety of sophisticated communication patterns.

2. *<string>* tags are hashed, not coerced, to numerical values. Hence, a message sent 'USING TAG 42' will not be matched by a message received 'USING TAG "42"'.

3. Different backends place different restrictions on the range of accceptable tag values. For example, the `c_mpi` backend is limited by MPI's tag values, which range from 0 to at least 32,767—but usually more. (The range $[0, 2^{31} - 1]$ is probably typical.) CONCEPTUAL will automatically map any *<expr>* or *<string>* used as a tag value into whatever range the backend supports. One implication is that different tags in

a CONCEPTUAL program can map to the same tag at run time. For example, if a backend limits tag values to the range [0, 32,767], then 'USING TAG 131068' is equivalent to 'USING TAG 0'—or even not using any tag at all. Also, in that case, two arbitrary strings have a 0.003% chance of mapping to the same tag (or 0.00000005% in the more typical case for c_mpi), which is unlikely but not impossible.

Buffer control

The CONCEPTUAL run-time library allocates a unique message buffer for each message sent/received with the UNIQUE keyword (see [Unique messages], page 109). The message buffers for NONUNIQUE messages are recycled subject to the constraint that no two concurrent transmissions will reference the same buffer. For example, if a task performs a synchronous send followed by a synchronous receive, those operations must be executed serially and will therefore share a message buffer. If, instead, a task performs an asynchronous send followed by an asynchronous receive, those operations may overlap, so CONCEPTUAL will use different message buffers for the two operations.

Message specifications enable the programmer to override the default buffer-allocation behavior. If a message is sent FROM BUFFER <expr> or received INTO BUFFER <expr>, the message is guaranteed to be sent/received using the specified buffer number. For example, FROM BUFFER and INTO BUFFER can be used to force a synchronous send and synchronous receive to use different buffers or an asynchronous send and asynchronous receive to use the same buffer. If <expr> is negative, the behavior is the same as if FROM BUFFER/INTO BUFFER was not specified. FROM THE DEFAULT BUFFER and INTO THE DEFAULT BUFFER also explicitly specify the default buffer-allocation behavior.

It is also possible to specify byte offsets into a buffer. For example, a message sent 'FROM 4 WORDS INTO BUFFER 3' will be sent from 4 words (12 bytes) past buffer 3's normal starting memory address.

Blocking semantics

By default—or if SYNCHRONOUSLY is specified—messages are sent synchronously. That is, a sender blocks (i.e., waits) until the message buffer is safe to reuse before it continues and a receiver blocks until it actually receives the message. The ASYNCHRONOUSLY keyword specifies that messages should be sent and received asynchronously. That is, the program merely posts the message (i.e., declares that it should eventually be sent and/or received) and immediately continues executing. Asynchronous messages must be *completed* as described in Section 4.4.4 [Awaiting completion], page 117.

4.4.2 Sending

The SEND statement is fundamental to CONCEPTUAL. It is used to send a multiple messages from multiple source tasks to multiple target tasks. The syntax is formally specified as follows:

<send_stmt>   ::=   <source_task>
                      [ASYNCHRONOUSLY] SENDS
                      <message_spec>
                      TO [UNSUSPECTING] <target_tasks>
               | <source_task>
                      [ASYNCHRONOUSLY] SENDS

<*message_spec*>
TO <*target_tasks*>
WHO RECEIVE IT
<*recv_message_spec*>

<*source_task*> is described in Section 4.3.2 [Source tasks], page 106; <*message_spec*> and <*recv_message_spec*> are described in Section 4.4.1 [Message specifications], page 108; and, <*target_tasks*> is described in Section 4.3.3 [Target tasks], page 107.

The SEND statement's simplest form, "<*source_task*> SENDS <*message_spec*> TO <*target_tasks*>", is fairly straightforward. The following is a example:

    TASK 0 SENDS A 0 BYTE MESSAGE TO TASK 1

The only subtlety in the preceding statement is that it implicitly causes task 1 to perform a corresponding receive. This receive can be suppressed by adding the keyword UNSUSPECTING before the <*target_tasks*> description:

    TASK 0 SENDS A 0 BYTE MESSAGE TO UNSUSPECTING TASK 1

Here are some further examples of valid <*send_stmt*>s:

- ALL TASKS SEND A 64 KILOBYTE MESSAGE TO TASK 0

- TASK num_tasks-1 SENDS 5 53 BYTE PAGE ALIGNED MESSAGES TO ALL OTHER TASKS

- TASKS upper SUCH THAT upper>=num_tasks/2 ASYNCHRONOUSLY SEND A 0 BYTE MESSAGE TO TASK upper/2

- TASKS nonzero SUCH THAT nonzero>0 SEND nonzero 1E3 BYTE MESSAGES TO UNSUSPECTING TASK 0

There are a number of attributes associated with every message transmission:

- synchronous vs. asynchronous operation
- unique vs. recycled message buffers
- unaligned vs. aligned vs. misaligned message buffers
- no data touching vs. data touching vs. data verification
- first-in, first-out vs. selective message ordering based on message tags
- implicit vs. explicit message-buffer selection
- explicit byte offsets into message buffers

A few of those appear in the preceding examples.

When UNSUSPECTING is omitted, the implicit RECEIVE statement normally inherits all of the attributes of the corresponding SEND. However, the second form of a <*send_stmt*>, which contains a WHO RECEIVES IT (or WHO RECEIVES THEM) clause, enables the receiver's attributes to be overridden on a per-attribute basis. For instance, consider the following SEND statement:

    TASK 0 SENDS A 1 MEGABYTE MESSAGE TO TASK 1 WHO RECEIVES IT
    ASYNCHRONOUSLY

The alternative sequence of statements that does not use WHO RECEIVES IT is less straightforward to read because it splits a single message transmission into two statements:

1. TASK 1 ASYNCHRONOUSLY RECEIVES A 1 MEGABYTE MESSAGE FROM TASK 0

2. TASK 0 SENDS A 1 MEGABYTE MESSAGE TO UNSUSPECTING TASK 1

Some further examples of `WHO RECEIVES IT` follow:

```
TASKS left SUCH THAT left IS EVEN SEND 5 2 KILOBYTE 64 BYTE ALIGNED
MESSAGES TO TASKS left+1 WHO RECEIVE THEM AS UNALIGNED MESSAGES WITH
DATA TOUCHING

TASK num_tasks-1 ASYNCHRONOUSLY SENDS A 1E5 BYTE MESSAGE WITH
VERIFICATION TO TASK 0 WHO RECEIVES IT SYNCHRONOUSLY

TASK leaf SUCH THAT KNOMIAL_CHILDREN(leaf,2)=0 SENDS A UNIQUE 1536
BYTE MESSAGE WITH DATA TOUCHING TO TASK KNOMIAL_PARENT(leaf,2) WHO
RECEIVES IT ASYNCHRONOUSLY AS A NONUNIQUE QUADWORD ALIGNED MESSAGE
WITHOUT DATA TOUCHING INTO BUFFER KNOMIAL_PARENT(leaf,2)
```

One subtlety of the `SEND` statement when used without `UNSUSPECTING` involves the orderings of the sends and receives. The rule is that receives are posted before sends. Furthermore, *<restricted_ident>*s (see Section 4.3.1 [Restricted identifiers], page 106) are evaluated in order from 0 to *num_tasks* -1. The implication is that a statement such as '`TASKS ev SUCH THAT ev IS EVEN /\ ev<6 SEND A 4 WORD MESSAGE TO TASK ev+2`' is exactly equivalent to the following ordered sequence of statements (assuming *num_tasks* ≥ 5):

1. `TASK 2 RECEIVES A 4 WORD MESSAGE FROM TASK 0`
2. `TASK 4 RECEIVES A 4 WORD MESSAGE FROM TASK 2`
3. `TASK 6 RECEIVES A 4 WORD MESSAGE FROM TASK 4`
4. `TASK 0 SENDS A 4 WORD MESSAGE TO UNSUSPECTING TASK 2`
5. `TASK 2 SENDS A 4 WORD MESSAGE TO UNSUSPECTING TASK 4`
6. `TASK 4 SENDS A 4 WORD MESSAGE TO UNSUSPECTING TASK 6`

(The `RECEIVE` statement is described in Section 4.4.3 [Receiving], page 116.)

If the above sequence were executed, tasks 2, 4, and 6 would immediately block on their receives (steps 1–3). Task 0 would awaken task 2 by sending it a message (step 4). Then, task 2 would be able to continue to step 5 at which point it would send a message to task 4. Task 4 would then finally be able to send a message to task 6 (step 6). Hence, even though the original CONCEPTUAL statement encapsulates multiple communication operations, the component communications proceed sequentially because of data dependences and because the operations are blocking.

As another example of send/receive ordering, consider the statement, '`TASKS x SUCH THAT x IS IN {1,4,5} SEND A 768 BYTE MESSAGE FROM 4 DOUBLEWORDS INTO THE DEFAULT BUFFER TO TASKS y SUCH THAT y IS IN {6,2,3}`'. This statement causes nine messages to be sent and received and in the following order: $1 \Rightarrow 2$, $1 \Rightarrow 3$, $1 \Rightarrow 6$, $4 \Rightarrow 2$, $4 \Rightarrow 3$, $4 \Rightarrow 6$, $5 \Rightarrow 2$, $5 \Rightarrow 3$, $5 \Rightarrow 6$. The reason that task 6 receives from each sender *after* tasks 2 and 3 relates to the comment above that *<restricted_ident>*s are evaluated in order from 0 to *num_tasks* -1. One can therefore think of the preceding CONCEPTUAL statement as being implemented with the following pseudocode:

for s := 0 **to** num_tasks-1 **do**
  if s **is in** $\{1, 4, 5\}$ **then**
    for r := 0 **to** num_tasks-1 **do**
      if r **is in** $\{6, 2, 3\}$ **then**

r **receives from** s
for s := 0 **to** num_tasks-1 **do**
  if s **is in** {1, 4, 5} **then**
    for r := 0 **to** num_tasks-1 **do**
      if r **is in** {6, 2, 3} **then**
        s **sends to** r

4.4.3 Receiving

Section 4.4.2 [Sending], page 113, mentioned the `<send_stmt>`'s `UNSUSPECTING` keyword, which specifies that the targets should not implicitly perform a receive operation. Because every send must have a matching receive, CONCEPTUAL offers a `RECEIVE` statement that explicitly receives a set of messages. A `<receive_stmt>` is much like a `<send_stmt>` (see Section 4.4.2 [Sending], page 113) with the `<source_task>` and `<target_tasks>` in the reverse order:

`<receive_stmt>`   ::=   `<target_tasks>`
                          [`ASYNCHRONOUSLY`] `RECEIVE`
                          `<message_spec>`
                          `FROM` `<source_task>`

`<target_tasks>` is described in Section 4.3.3 [Target tasks], page 107; `<message_spec>` is described in Section 4.4.1 [Message specifications], page 108; and, `<source_task>` is described in Section 4.3.2 [Source tasks], page 106.

Like other statements that contain both a `<source_task>` and a `<target_tasks>`, the `RECEIVE` statement propagates all variables declared by the `<source_task>` to the `<target_tasks>`, as in

```
TASKS src-1 RECEIVE A 1 MEGABYTE MESSAGE FROM ALL TASKS src.
```

However, the `RECEIVE` statement is unique among CONCEPTUAL statements in that it can alternatively propagate all variables declared by the `<target_tasks>` to the `<source_task>`, as in

```
ALL TASKS dst RECEIVE A 1 MEGABYTE MESSAGE FROM TASKS dst+1.
```

The `RECEIVE` statement automatically determines the direction in which variables should be propagated. If the direction cannot unambiguously be determined, the CONCEPTUAL compiler aborts with an error message. For example, the `RECEIVE` statement in the following code is ambiguous because the 'bbb' in the expression 'aaa < bbb' can refer to either the one in the `LET` statement or the one in the 'TASKS bbb SUCH THAT bbb > 3' expression:

```
LET bbb BE 5 WHILE {
  TASKS aaa SUCH THAT aaa < bbb RECEIVE A MESSAGE FROM TASKS bbb SUCH
  THAT bbb > 3 THEN
  TASKS ccc SUCH THAT ccc > 3 SEND A MESSAGE TO UNSUSPECTING TASKS
  ddd SUCH THAT ddd < ccc
}
```

For each message sent via a `SEND...TO UNSUSPECTING` statement there must be a `RECEIVE` statement that receives a message of the same size. The `<target_tasks>`'s `<message_spec>` can, however, specify different values for message uniqueness, message alignment, and data touching. In addition, the source and target do not need to agree on the use of the `ASYNCHRONOUSLY` keyword. The only restriction is that `WITH VERIFICATION`

will return spurious results if used by the target but not by the source. Hence, the following <*send_stmt*> and <*receive_stmt*> correctly match each other:

```
TASK 0 SENDS 3 4 KILOBYTE MESSAGES TO UNSUSPECTING TASK 1


TASK 1 ASYNCHRONOUSLY RECEIVES 3 UNIQUE 4 KILOBYTE 48 BYTE ALIGNED
MESSAGES WITH DATA TOUCHING FROM TASK 0.
```

In general, it is better to use a single `SEND` statement with a `WHO RECEIVES IT` clause (see Section 4.4.2 [Sending], page 113) than a `RECEIVE` plus a matching `SEND...TO UNSUSPECTING`; the former is less error-prone than the latter. However, the latter is useful for programs in which a set of receives is posted, then the tasks perform various communication, computation, and synchronization operations, and—towards the end of the program—the matching sends are posted. That sort of split-phase structure requires separate `SEND` and `RECEIVE` statements.

4.4.4 Awaiting completion

When a message is sent or received asynchronously it must eventually be *completed*. In some messaging layers, asynchronous messages are not even sent or received until completion time. CONCEPTUAL provides the following statement for completing messages that were send/received asynchronously:

<*wait_stmt*>   ::=   <*source_task*>
                     `AWAITS COMPLETION`

That is, a <*wait_stmt*> simply specifies the set of tasks that should block until all of their pending communications complete. <*source_task*> is as defined in Section 4.3.2 [Source tasks], page 106. Note that a <*wait_stmt*> blocks until *all* pending communications complete. CONCEPTUAL does not provide finer-grained control over completions. It is safe, however, for a task to `AWAIT COMPLETION` even if it has no asynchronous messages pending.

4.4.5 Multicasting

Although a single <*send_stmt*> (see Section 4.4.2 [Sending], page 113) can specify multiple messages at once, these messages are sent one at a time. *Multicasting* is a form of collective communication in which a set of tasks collaborates to deliver a message from a source to multiple targets. With many messaging layers, multicasting a message to N tasks is more efficient than sending a sequence of N individual messages. CONCEPTUAL supports multicasting as follows:

<*mcast_stmt*>   ::=   <*source_task*>
                      [`ASYNCHRONOUSLY`] `MULTICASTS`
                      <*message_spec*>
                      `TO` <*target_tasks*>

Unlike <*send_stmt*>s, <*mcast_stmt*>s do not support the `UNSUSPECTING` keyword. This is because `MULTICASTS` is a collective operation: all parties are active participants in delivering messages to the <*target_tasks*>.

<*source_task*> (see Section 4.3.2 [Source tasks], page 106) and <*target_tasks*> (see Section 4.3.3 [Target tasks], page 107) can be either disjoint or overlapping sets. That is, either of the following is legal:

```
      TASK 0 MULTICASTS A 16 BYTE MESSAGE TO TASKS recip SUCH THAT recip<4
      TASK 0 MULTICASTS A 16 BYTE MESSAGE TO TASKS recip SUCH THAT recip>=4
```

Note that in the first <mcast_stmt>, task 0 both sends and receives a message, while in the second <mcast_stmt>, task 0 sends but does not receive.

The MULTICASTS statement supports not only traditional one-to-many multicasting as shown in the preceding example but also many-to-many multicasting—and the less interesting but still valid many-to-one and one-to-one multicasting. The constituent pairs of communicating tasks are identical to those induced by the SENDS statement (see Section 4.4.2 [Sending], page 113). For example, the statement

```
      TASKS s SUCH THAT s<3 MULTICAST A 1 KILOBYTE MESSAGE TO TASKS s SUCH
      THAT s>6
```

performs the collective analogue of

```
      TASKS s SUCH THAT s<3 SEND A 1 KILOBYTE MESSAGE TO TASKS s SUCH
      THAT s>6
```

Specifically, task 0 will send a message to each of tasks 7, 8, 9, . . .; task 1 will send a message to each of tasks 7, 8, 9, . . .; and task 2 will send a message to each of tasks 7, 8, 9, . . .

4.4.6 Reducing

A reduction operation is, in a sense, a complementary operation to a multicast (see Section 4.4.5 [Multicasting], page 117). While a multicast delivers a message from one source to multiple targets, a reduction combines messages from multiple sources (by applying a commutative/associative operator to corresponding elements) to a single target. Reduction is a collective operation: all parties collaborate to calculate the reduced value(s). As an example, if tasks 0, 1, and 2 collectively reduce the messages $\{5, 1\}$, $\{2, 7\}$, and $\{3, 4\}$ to task 2 using the "+" operator, then task 2 will receive the message $\{10, 12\}$. In fact, CONCEPTUAL's implementation of reductions also supports reductions to multiple targets with each target receiving a copy of the reduced value.

The following grammatical rules define CONCEPTUAL's many-to-many and many-to-one reduction facilities:

<reduce_stmt>   ::=   <source_task>
                       REDUCES
                       <reduce_message_spec>
                       TO <source_task>
                       [WHO RECEIVES THE RESULT <reduce_target_message_spec>]
                 | <source_task>
                       REDUCES
                       <reduce_message_spec>
                       [TO <reduce_message_spec>]

<reduce_message_spec> is defined in Section 4.4.1 [Message specifications], page 108. Both the data providers and data receivers are specified as <source_task> nonterminals (see Section 4.3.2 [Source tasks], page 106). This design enables any set of tasks to provide the data to reduce and any disjoint or overlapping set of tasks to receive the reduced data. As with all communication in CONCEPTUAL, message contents are opaque. Furthermore, the

grammar does not currently enable the programmer to specify the commutative/associative operator to use.

A simple many-to-one reduction can be expressed in CONCEPTUAL with 'ALL TASKS REDUCE 5 DOUBLEWORDS TO TASK 0'. Note that the definition of <reduce_message_spec> and <reduce_target_message_spec> (see Section 4.4.1 [Message specifications], page 108) supports reductions only of INTEGERs and DOUBLEWORDs, not arbitrary <data_type> values. Omitting the optional 'TO <reduce_message_spec>', as in 'TASKS rt SUCH THAT 3 DIVIDES rt REDUCE AN INTEGER', specifies that all tasks performing the reduction will receive a copy of the reduced value. The sources and targets can also be designated explicitly as in 'TASKS xyz SUCH THAT xyz<num_tasks/2 REDUCE 100 DOUBLEWORDS TO TASKS xyz+num_tasks/4'. When that code is run with 8 tasks, tasks 0–3 reduce 100 doublewords (800 bytes) apiece and tasks 2–5 each receive identical copies of the 100 doublewords of reduced data.

Message data used with REDUCES can be transferred WITH DATA TOUCHING (but not WITH VERIFICATION); data alignment can be specified; and, messages buffers can be named explicitly. The following example represents fairly complex many-to-many usage of REDUCES:

```
TASKS rsrc SUCH THAT rsrc IS EVEN REDUCE 32 64-BYTE-ALIGNED INTEGERS
WITH DATA TOUCHING FROM BUFFER 2 TO TASKS rtarg SUCH THAT
rtarg<num_tasks/4 \/ rtarg>(3*num_tasks)/4 WHO RECEIVE THE RESULT AS
32 UNIQUE PAGE-ALIGNED INTEGERS WITHOUT DATA TOUCHING.
```

4.4.7 Synchronizing

CONCEPTUAL enables sets of tasks to perform *barrier synchronization*. The semantics are that no task can finish synchronizing until all tasks have started synchronizing. The syntax is as follows:

<sync_stmt>   ::=   <source_task>
                     SYNCHRONIZES

A <sync_stmt> can be used to ensure that one set of statements has completed before beginning another set. For example, a CONCEPTUAL program might have a set of tasks post a series of asynchronous receives (see Section 4.4.3 [Receiving], page 116), then make 'ALL TASKS SYNCHRONIZE' before having another set of tasks perform the corresponding UNSUSPECTING sends (see Section 4.4.2 [Sending], page 113). This procedure ensures that all of the target tasks are ready to receive before the source tasks start sending to them.

4.5 I/O statements

CONCEPTUAL provides two statements for presenting information. One statement writes simple messages to the standard output device and is intended to be used for providing status information during the run of a program. The other statement provides a powerful mechanism for storing performance and correctness data to a log file.

4.5.1 Utilizing log-file comments

<output_stmt>s and <log_stmt>s have limited access to the ⟨key:value⟩ pairs that are written as comments at the top of every log file as shown in Section 3.5.1 [Log-file format], page 49. Given a key, *key*, the string expression 'THE VALUE OF *key*' represents the value associated with that key or the empty string if *key* does not appear in the log-file comments:

<string_or_log_comment>   ::=   <string>

| THE VALUE OF <*string*>

That is, '"`CPU frequency`"' means the literal string "CPU frequency" while '`THE VALUE OF` "`CPU frequency`"' translates to a string like "1300000000 Hz (1.3 GHz)". Environment variables are also considered keys and are therefore acceptable input to a `THE VALUE OF` construct.

4.5.2 Writing to standard output

CONCEPTUAL's `OUTPUT` keyword is used to write a message from one or more source tasks (see Section 4.3.2 [Source tasks], page 106) to the standard output device. This is useful for providing progress reports during the execution of long-running CONCEPTUAL programs. An <*output_stmt*> looks like this:

<*output_stmt*>   ::=   <*source_task*>
                         `OUTPUTS`
                         <*expr*> | <*string_or_log_comment*>
                         [`AND`   <*expr*> | <*string_or_log_comment*>]*

The following are some sample <*output_stmt*>s:

```
TASK 0 OUTPUTS "Hello, world!"

TASKS nr SUCH THAT nr>0 OUTPUT nr AND "'s parent is " AND nr>>1 AND
" and its children are " AND nr<<1 AND " and " AND nr<<1+1

ALL TASKS me OUTPUT "Task " AND me AND " is running on host " AND THE
VALUE OF "Host name" AND " and plans to send to task " AND (me+1) MOD
num_tasks
```

`OUTPUT` does not implicitly output spaces between terms. Hence, '`OUTPUT "Yes" AND "No"`' will output "YesNo", not "Yes No". Although it is unlikely that a program would ever need to output two arithmetic expressions with no intervening text, an empty string can be used for this purpose: '`OUTPUT 6 AND "" AND 3`'.

An <*output_stmt*> implicitly outputs a newline character at the end. Additional newline characters can be output by embedding '`\n`' in a string. (see Section 4.1 [Primitives], page 84.) CONCEPTUAL does not provide a means for suppressing the newline, however.

4.5.3 Writing to a log file

After performing a network correctness or performance test it is almost always desirable to store the results in a file. CONCEPTUAL has language support for writing tabular data to a log file. The <*log_stmt*> command does the bulk of the work:

<*log_stmt*>   ::=   <*source_task*>
                      `LOGS`
                      <*aggr_expr*> `AS` <*string_or_log_comment*>
                      [`AND` <*aggr_expr*> `AS` <*string_or_log_comment*>]*

The idea behind a <*log_stmt*> is that a set of source tasks (see Section 4.3.2 [Source tasks], page 106) log an aggregate expression (see Section 4.2.3 [Aggregate expressions], page 100) to a log file under the column heading <*string_or_log_comment*>. Each task individually maintains a separately named log file so there is no ambiguity over which task wrote which entries.

Each (static) `LOGS` statement in a CoNCePTuaL program specifies one or more columns of the log file. Every dynamic execution of a `LOGS` statement writes a single row to the log file. A single `LOGS` statement should suffice for most CoNCePTuaL programs.

The following are some examples of <*log_stmt*>s:

```
ALL TASKS LOG bit_errors AS "Bit errors"

TASK 0 LOGS THE msgsize AS "Bytes" AND
            THE MEDIAN OF (1E6*bytes_sent)/(1M*elapsed_usecs) AS "MB/s"
```

The first example produces a log file like the following:

```
"Bit errors"
"(all data)"
3
```

The second example produces a log file like this:

```
"Bytes","MB/s"
"(only value)","(median)"
65536,179.9416266
```

Note that in each log file, the CoNCePTuaL run-time system writes two rows of column headers for each column. The first row contains <*string_or_log_comment*> as is. The second row describes the <*aggr_func*> (see Section 4.2.4 [Aggregate functions], page 101) used to aggregate the data. One or more rows of data follow.

Assume that the second <*log_stmt*> presented above appears within a loop (see Section 4.7.2 [Iterating], page 126). It is therefore important to include the `THE` keyword before 'msgsize' to assert that the expression 'msgsize' is constant across invocations of the <*log_stmt*> and that, consequently, only a single row of data should be written to the log file. Using 'msgsize' without the `THE` would produce a column of data with one row per <*log_stmt*> invocation:

```
"Bytes","MB/s"
"(all data)","(median)"
65536,179.9416266
65536,
65536,
65536,
65536,
     .
     .
     .
```

The rules that determine how `LOGS` statements produce rows and columns of a log file are presented below:

1. Each *static* `LOGS` statement (and `AND` clause within a `LOGS` statement) in a program produces a unique column.

2. Each *dynamic* execution of a `LOGS` statement appends a row to the column(s) it describes.

3. Each top-level complex statement (see Section 4.10 [Complete programs], page 142) produces a new table in the log file.

Note that the choice of column name is inconsequential for determining what columns are written to the log file:

```
TASK 0 LOGS 314/100 AS "Pi" AND 22/7 AS "Pi"
```

```
"Pi","Pi"
"(all data)","(all data)"
3.14,3.142857143
```

As mentioned in Section 4.2.3 [Aggregate expressions], page 100, multiple *<aggr_func>* values can be specified per *<aggr_expr>*. This is a convenient feature when a CONCEPTUAL program needs to compute multiple statistics over the same values while reusing the same description. For example,

```
ALL TASKS LOG THE MINIMUM OF bytes_received/elapsed_usecs AS
"Bandwidth (B/us)" AND THE MAXIMUM OF bytes_received/elapsed_usecs AS
"Bandwidth (B/us)" AND THE MEAN OF bytes_received/elapsed_usecs AS
"Bandwidth (B/us)" AND THE MEDIAN OF bytes_received/elapsed_usecs AS
"Bandwidth (B/us)"
```

can be shortened to

```
ALL TASKS LOG THE MINIMUM AND THE MAXIMUM AND THE MEAN AND THE MEDIAN
OF bytes_received/elapsed_usecs AS "Bandwidth (B/us)"
```

Computing aggregates

What if '`msgsize`' takes on a number of values throughout the execution of the program and for each value a number of runs is performed? How would one log the median of each set of data? Using '`THE msgsize`' won't work because the message size is not constant. Using '`msgsize`' alone won't work either because CONCEPTUAL would then take the median of the times gathered across *all* message sizes, which is undesirable. The solution is for the program to specify explicitly when aggregate functions (`MEDIAN` and all of the other functions listed in Section 4.2.4 [Aggregate functions], page 101) compute a value:

<flush_stmt>   ::=   *<source_task>*
                      COMPUTES AGGREGATES

The intention is that an inner loop might `LOG` data after every iteration and an outer loop would '`COMPUTE AGGREGATES`' after each iteration.

4.6 Counter and timer statements

Critical to any performance or correctness test is the ability to specify which operations represent the test itself and should be measured and which are setup or other uninteresting operations and should not. CONCEPTUAL automatically maintains a number of "counters"—variables that represent message counts, byte counts, bit-error counts, and elapsed time. The complete list is presented in Section A.2 [Predeclared variables], page 207.

Normally, a CONCEPTUAL program performs some setup operations, RESETS ITS COUNTERS to zero, executes a communication pattern, and logs some function of the resulting changes in counter values (see Section 4.5.3 [Writing to a log file], page 120). If additional setup work needs to be performed during an experiment, a program can STORE ITS COUNTERS, perform any arbitrarily costly operations, RESTORE ITS COUNTERS, and continue the experiment as if those operations never happened.

Some CONCEPTUAL statements implicitly store and restore counters. For example, the LOGS statement see Section 4.5.3 [Writing to a log file], page 120) takes up no time from the program's perspective, and counted loops (see [Counted loops], page 126) bracket any warmup repetitions and post-warmup synchronizations between a counter store and restore so no delays, bit errors, or messaging operations contribute to the totals measured by the experiment.

4.6.1 Resetting counters

At the start of an experiment, after all setup processing has completed, all tasks that will eventually log measurement results should zero out their counters:

<reset_stmt>   ::=   <source_task>
                         RESETS ITS COUNTERS

Hence, writing 'ALL TASKS RESET THEIR COUNTERS' causes each task to reset all of the variables listed in Section A.2 [Predeclared variables], page 207—with the exception of num_tasks—to zero. Note that ITS and THEIR, like RESET and RESETS, are considered synonyms (see Section 4.1 [Primitives], page 84).

4.6.2 Storing counter values

A program can store the current values of all of the variables listed in Section A.2 [Predeclared variables], page 207 as follows:

<store_stmt>   ::=   <source_task>
                         STORES ITS COUNTERS

For example, writing 'TASK 0 STORES ITS COUNTERS' causes task 0 to store the current values of elapsed_usecs, total_msgs, bit_errors, etc. The values are not modified. Note that ITS and THEIR, like STORE and STORES, are considered synonyms (see Section 4.1 [Primitives], page 84).

4.6.3 Restoring counter values

Counters can be restored to their most recently saved values with a <restore_stmt>:

<restore_stmt>   ::=   <source_task>
                           RESTORES ITS COUNTERS

For example, writing 'TASKS t SUCH THAT 3 DIVIDES t RESTORE THEIR COUNTERS' causes every third task to replace `elapsed_usecs`, `total_msgs`, `bit_errors`, etc. with the values stored by a corresponding <*store_stmt*> (see Section 4.6.2 [Storing counter values], page 123). Note that ITS and THEIR, like RESTORE and RESTORES, are considered synonyms (see Section 4.1 [Primitives], page 84).

An important feature of STORES and RESTORES is that they can be nested. That is, each STORE pushes a set of counter values on a stack, and each RESTORE pops a set of counter values from the stack. Consequently, one can write code like the following:

```
ALL TASKS RESET THEIR COUNTERS THEN                                    #  1
ALL TASKS COMPUTE FOR 2 SECONDS THEN                                   #  2
ALL TASKS STORE THEIR COUNTERS THEN                                    #  3
  ALL TASKS COMPUTE FOR 5 SECONDS THEN                                 #  4
  ALL TASKS STORE THEIR COUNTERS THEN                                  #  5
    ALL TASKS COMPUTE FOR 9 SECONDS THEN                              #  6
    ALL TASKS LOG ROUND(elapsed_usecs/1E6) AS "Should be 16" THEN     #  7
  ALL TASKS RESTORE THEIR COUNTERS THEN                                #  8
  ALL TASKS LOG ROUND(elapsed_usecs/1E6) AS "Should be 7" THEN        #  9
  ALL TASKS STORE THEIR COUNTERS THEN                                  # 10
    ALL TASKS COMPUTE FOR 1 SECOND THEN                               # 11
    ALL TASKS LOG ROUND(elapsed_usecs/1E6) AS "Should be 8" THEN      # 12
  ALL TASKS RESTORE THEIR COUNTERS THEN                                # 13
  ALL TASKS LOG ROUND(elapsed_usecs/1E6) AS "Should be 7" THEN        # 14
  ALL TASKS RESTORE THEIR COUNTERS THEN                                # 15
ALL TASKS LOG ROUND(elapsed_usecs/1E6) AS "Should be 2".               # 16
```

Indentation is used in the above to clarify which RESTORE operations match which STORE operations. The first STORE statement (line 3) occurs after 2 seconds have elapsed. The second STORE statement (line 5) occurs after $2 + 5 = 7$ seconds have elapsed. When the LOG statement in line 7 is executed, it reports that $2 + 5 + 9 = 16$ seconds have elapsed. The RESTORE statement in line 8 then "winds back the clock" to the previous STORE statement, the one in line 5. The next LOG statement (line 9) executes as if lines 5–8 never ran and therefore reports that only $2 + 5 = 7$ seconds have elapsed. The LOG statement in line 12 sees an additional second of elapsed time due to line 11's COMPUTE statement, for a total of $2 + 5 + 1 = 8$ seconds. However, the RESTORE in line 13 makes it as if that COMPUTE never happened. Consequently, the LOG statement in line 14 reports that only $2 + 5 = 7$ seconds have elapsed. Finally, the RESTORE in line 15 sets the timer to the value it had all the way back at line 3. The final LOG statement (line 16) therefore reports that only 2 seconds have elapsed because line 2 contains the only COMPUTE statement whose execution time has not been discarded.

Because the `interpret` backend and those derived from it use logical time instead of physical time, the code listed above will report all zeroes. Replacing 'ROUND(elapsed_usecs/1E6)' with just 'elapsed_usecs' will log the logical times $\{6, 5, 8, 7, 3\}$. That is, the LOG statement in line 7 sees six events after the initial RESET;[7]

---

[7] Each statement in the example corresponds to a single event and therefore counts as one unit of time.

the `LOG` statement in line 9 sees only five events after the `RESET` (corresponding to lines 2, 3, 4, 5, and 9); and so forth up to the final `LOG` statement, which sees only three events: those produced by lines 2, 3, and 16.

A program that calls `RESTORE` more times than it calls `STORE` will abort with a fatal run-time error.

4.7 Complex statements

The CONCEPTUAL statements presented in Section 4.4 [Communication statements], page 107, Section 4.5 [I/O statements], page 119, and Section 4.8 [Other statements], page 133 are all known as *simple statements*. This section expands upon the statements already introduced by presenting *complex statements*. In its most basic form, a *<complex_stmt>* is just a *<simple_stmt>*. However, the primary purpose of a *<complex_stmt>* is to juxtapose simple statements and other complex statements into more expressive forms.

Complex statements take the following form:

<complex_stmt>   ::=   *<simple_stmt>* [`THEN` *<complex_stmt>*]

The constituent simple statements include `FOR` loops, `LET` bindings, `IF` conditionals, grouping constructs, and all of the statements introduced in Section 4.4 [Communication statements], page 107, Section 4.5 [I/O statements], page 119, and Section 4.8 [Other statements], page 133:

<simple_stmt>   ::=   `FOR` *<expr>* `REPETITIONS` [`PLUS` *<expr>* `WARMUP REPETITIONS`
                        [`AND A SYNCHRONIZATION`]] *<simple_stmt>*
                    | `FOR EACH` *<ident>* `IN` *<range>* [`,` *<range>*]* *<simple_stmt>*
                    | `FOR` *<expr>* *<time_unit>* [`PLUS` *<expr>* `WARMUP` *<time_unit>* [`AND`
                        `A SYNCHRONIZATION`]] *<simple_stmt>*
                    | `LET` *<let_binding>* [`AND` *<let_binding>*]* `WHILE` *<simple_stmt>*
                    | `IF` *<rel_expr>* `THEN` *<simple_stmt>* [`OTHERWISE` *<simple_stmt>*]
                    | `{` [*<complex_stmt>*] `}`
                    | *<send_stmt>*
                    | *<receive_stmt>*
                    | *<wait_stmt>*
                    | *<mcast_stmt>*
                    | *<reduce_stmt>*
                    | *<sync_stmt>*
                    | *<output_stmt>*
                    | *<log_stmt>*
                    | *<flush_stmt>*
                    | *<reset_stmt>*
                    | *<store_stmt>*
                    | *<restore_stmt>*
                    | *<assert_stmt>*
                    | *<delay_stmt>*
                    | *<touch_stmt>*
                    | *<touch_buffer_stmt>*
                    | *<processor_stmt>*
                    | *<backend_stmt>*

The remainder of this section describes in turn the `THEN` construct and each of the just-introduced *<simple_stmt>* types.

4.7.1 Combining statements

The `THEN` keyword separates statements that are to be performed sequentially. For example, a simple ping-pong communication can be expressed as follows:

```
ALL TASKS RESET ALL COUNTERS THEN
TASK 0 SENDS A 0 BYTE MESSAGE TO TASK 1 THEN
TASK 1 SENDS A 0 BYTE MESSAGE TO TASK 0 THEN
TASK 0 LOGS elapsed_usecs/2 AS "One-way latency"
```

There is no implicit intertask synchronization across `THEN` statements. Consequently, the two communications specified in the following statement will be performed concurrently:

```
TASK 0 ASYNCHRONOUSLY SENDS AN 8 KILOBYTE MESSAGE TO TASK 1 THEN
TASK 1 ASYNCHRONOUSLY SENDS AN 8 KILOBYTE MESSAGE TO TASK 0 THEN
ALL TASKS AWAIT COMPLETION
```

4.7.2 Iterating

CONCEPTUAL provides a variety of looping constructs designed to repeatedly execute a *<simple_stmt>*.

Counted loops

The simplest form of iteration in CONCEPTUAL repeats a *<simple_stmt>* a given number of times. In its most basic form the syntax is simply "`FOR <expr> REPETITIONS <simple_stmt>`". As could be expected, the *<expr>* term (see Section 4.2.1 [Arithmetic expressions], page 86) specifies the number of repetitions to perform. Hence, the following *<simple_stmt>* outputs the phrase "I will not talk in class" 100 times:

```
FOR 100 REPETITIONS ALL TASKS OUTPUT "I will not talk in class."
```

`FOR...REPETITIONS` can optionally specify a number of "warmup" repetitions to perform in addition to the base number of repetitions. The syntax is "`FOR <expr> REPETITIONS PLUS <expr> WARMUP REPETITIONS <simple_stmt>`". During warmup repetitions, the `OUTPUTS` statement (see Section 4.5.2 [Writing to standard output], page 120), the `LOGS` statement (see Section 4.5.3 [Writing to a log file], page 120), and the `COMPUTES AGGREGATES` statement (see [Computing aggregates], page 122) are all suppressed (i.e., they have no effect) and none of the special variables predeclared by CONCEPTUAL (see Section A.2 [Predeclared variables], page 207) are updated. Many benchmarks synchronize all tasks after performing a set of warmup repetitions. This behavior can be expressed conveniently as part of a CONCEPTUAL `FOR` loop by appending the `AND A SYNCHRONIZATION` clause:

```
FOR 1000 REPETITIONS PLUS 3 WARMUP REPETITIONS AND A SYNCHRONIZATION
TASK 0 MULTICASTS A 1 MEGABYTE MESSAGE TO ALL OTHER TASKS
```

CONCEPTUAL also provides a separate `SYNCHRONIZES` statement. This is described in Section 4.4.7 [Synchronizing], page 119.

The importance of performing warmup repetitions is that many communication layers give atypically poor performance on the first few transmissions. This may be because the messages miss in the cache; because the communication layer needs to establish connections between pairs of communicating tasks; or, because the operating system needs to "register"

message buffers with the network interface. Regardless of the reason, specifying warmup
repetitions helps make performance measurements less variable.

Range loops

Range loops are cONCEPTuAL's most powerful looping construct. Unlike the counted-
loop construct presented in [Counted loops], page 126, a range loop binds a variable to a
different value on each iteration. Range loops have the following syntax:

> FOR EACH <*ident*>
> IN <*range*> [, <*range*>]*
> <*simple_stmt*>

A <*range*> represents a range expression. Range expressions are described in
Section 4.2.6 [Range expressions], page 102. In short, a range expression specifies a list of
values by explicit enumeration, numeric progression, or predicated combinations of other
range expressions. For a range loop, cONCEPTuAL successively binds a specified variable
to each <*expr*> in each <*range*> and evaluates the given <*simple_stmt*>.

The following are some examples of the FOR EACH statement from simplest to most elab-
orate. Each example uses 'i' as the loop variable and 'TASK 0 OUTPUTS i' as the loop body.
The output from each example is shown with a "⊣" symbol preceding each line.

```
FOR EACH i IN {8, 7, 5, 4, 5} TASK 0 OUTPUTS i
 ⊣ 8
 ⊣ 7
 ⊣ 5
 ⊣ 4
 ⊣ 5


FOR EACH i IN {1, ..., 5} TASK 0 OUTPUTS i
 ⊣ 1
 ⊣ 2
 ⊣ 3
 ⊣ 4
 ⊣ 5


FOR EACH i IN {5, ..., 1} TASK 0 OUTPUTS i
 ⊣ 5
 ⊣ 4
 ⊣ 3
 ⊣ 2
 ⊣ 1


FOR EACH i IN {1, ..., 5}, {8, 7, 5, 4, 5} TASK 0 OUTPUTS i
 ⊣ 1
 ⊣ 2
 ⊣ 3
 ⊣ 4
 ⊣ 5
 ⊣ 8
```

```
⊣ 7
⊣ 5
⊣ 4
⊣ 5

FOR EACH i IN {1, 4, 7, ..., 30} TASK 0 OUTPUTS i
⊣ 1
⊣ 4
⊣ 7
⊣ 10
⊣ 13
⊣ 16
⊣ 19
⊣ 22
⊣ 25
⊣ 28

FOR EACH i IN {3**1, 3**2, 3**3, ..., 3**7} TASK 0 OUTPUTS i
⊣ 3
⊣ 9
⊣ 27
⊣ 81
⊣ 243
⊣ 729
⊣ 2187

FOR EACH i IN {0}, {1, 2, 4, ..., 256} TASK 0 OUTPUTS i
⊣ 0
⊣ 1
⊣ 2
⊣ 4
⊣ 8
⊣ 16
⊣ 32
⊣ 64
⊣ 128
⊣ 256

FOR EACH i IN {neigh FOR EACH neigh IN {0,...,100} WHERE
MESH_DISTANCE((10,10), 55, neigh) IS IN {1, 2}} TASK 0 OUTPUTS i.
⊣ 35
⊣ 44
⊣ 45
⊣ 46
⊣ 53
⊣ 54
⊣ 56
```

      ⊣ 57
      ⊣ 64
      ⊣ 65
      ⊣ 66
      ⊣ 75

That final example loops over all values that are a Manhattan distance of either 1 or 2 from the number 55 in a 10×10 layout of the numbers 0–99.

FOR EACH loops with constant progressions are executed exactly once. Note that if the *<range>* does not contain an ellipsis then all values are used, regardless of order or constancy:

```
FOR EACH i IN {4, 4, 4, ..., 4} TASK 0 OUTPUTS i
⊣ 4

FOR EACH i IN {4, 4, 4, 4, 4} TASK 0 OUTPUTS i
⊣ 4
⊣ 4
⊣ 4
⊣ 4
⊣ 4
```

Timed loops

A *timed loop* is similar to a counted loop (see [Counted loops], page 126) but instead of running for a given number of iterations it runs for a given length of time. Timed loops are absent from all general-purpose programming languages but can be quite useful in the context of network correctness and performance testing. The syntax of cONCEPTuaL's timed-loop construct is "FOR *<expr>* *<time_unit>* *<simple_stmt>*". *<time_unit>* is unit of time as listed in Section 4.8.2 [Delaying execution], page 134 and *<expr>* specifies the number of *<time_unit>*s for which to execute.

The following example shows how to spend three seconds sending messages from task 0 to task 1:

```
FOR 3 SECONDS TASK 0 SENDS A 1 MEGABYTE MESSAGE TO TASK 1
```

Although cONCEPTuaL tries its best to run for exactly the specified length of time there will invariably be some error in the process. Always use elapsed_usecs (see Section A.2 [Predeclared variables], page 207) as the indicator of actual time instead of the time requested in the loop.

Timed loops can optionally run for a length of "warmup" time in addition to the base execution time. The syntax is "FOR *<expr>* *<time_unit>* PLUS *<expr>* WARMUP *<time_unit>* *<simple_stmt>*". During warmup repetitions, the OUTPUTS statement (see Section 4.5.2 [Writing to standard output], page 120), the LOGS statement (see Section 4.5.3 [Writing to a log file], page 120), and the COMPUTES AGGREGATES statement (see [Computing aggregates], page 122) are all suppressed (i.e., they have no effect) and none of the special variables predeclared by cONCEPTuaL (see Section A.2 [Predeclared variables], page 207) are updated. Many benchmarks synchronize all tasks after performing a set of warmup repetitions. This behavior can be expressed conveniently as part of a timed loop by appending the AND A SYNCHRONIZATION clause:

```
    FOR 3 SECOND PLUS 1 WARMUP MILLISECOND AND A SYNCHRONIZATION TASK 0
    SENDS A 1 MEGABYTE MESSAGE TO TASK 1
```

CONCEPTUAL also provides a separate SYNCHRONIZES statement. This is described in Section 4.4.7 [Synchronizing], page 119.

The importance of performing warmup repetitions is that many communication layers give atypically poor performance on the first few transmissions. This may be because the messages miss in the cache; because the communication layer needs to establish connections between pairs of communicating tasks; or, because the operating system needs to "register" message buffers with the network interface. Regardless of the reason, specifying warmup repetitions helps make performance measurements less variable.

4.7.3 Binding variables

There are four ways to bind a value to a variable:

1. as part a source or target task description (see Section 4.3 [Task descriptions], page 106)

2. as part of a range loop (see [Range loops], page 127)

3. via a command-line argument (see Section 4.9.2 [Command-line arguments], page 139)

4. explicitly using the LET keyword (this section)

The LET statement has the following form:

```
        LET <let_binding>
        [AND <let_binding>]*
        WHILE <simple_stmt>
```

where <*let_binding*> is defined as follows:

<*let_binding*>   ::=   <*ident*> BE   expr
                  | <*source_task*>
                  | A RANDOM TASK [<*random_task_constraints*>]

The optional <*random_task_constraints*> limits the set of tasks from which A RANDOM TASK selects:

<*random_task_constraints*>   ::=   OTHER THAN <*expr*>
                        | LESS THAN <*expr*> [BUT NOT <*expr*>]
                        | GREATER THAN <*expr*> [BUT NOT <*expr*>]
                        | IN '[' <*expr*> ',' <*expr*> ']' [BUT NOT <*expr*>]

An <*ident*> bound using the form 'LET <*ident*> BE <*source_task*>' represents a group of tasks and must therefore be accessed using the TASK GROUP construct (see Section 4.3.2 [Source tasks], page 106). For example, in the CONCEPTUAL program

```
    LET left BE TASKS t SUCH THAT t < num_tasks/2 WHILE TASK GROUP left
    SENDS A 1 KILOBYTE MESSAGE TO TASK num_tasks - left - 1
```

the 'TASK GROUP left' expression defines 'left' as the group of tasks described by 'TASKS t SUCH THAT t < num_tasks/2'. The 'left' in 'num_tasks - left - 1' refers back to *that* definition of 'left', not the let-bound 'left'. The 't' in the let-binding is used only to specify the group of tasks that will be bound to 'left'. Its scope is limited to the expression 't < num_tasks/2'. In other words, the preceding example is equivalent to

```
TASKS left SUCH THAT left < num_tasks/2 SEND A 1 KILOBYTE MESSAGE TO
TASK num_tasks - left - 1
```

The advantage of using LET to define a task group is that the same—possible complex—description of a set of tasks can be reused in multiple statements without having to type it repeatedly.

Here are some further examples of LET:

```
LET reps BE 3 WHILE FOR reps REPETITIONS TASK 0 OUTPUTS "Laissez les
bons temps rouler."

LET src BE num_tasks-1 AND dest BE num_tasks/2 WHILE TASK src SENDS A
55E6 BIT MESSAGE TO TASK dst

LET hotspot BE A RANDOM TASK WHILE TASKS other SUCH THAT
other<>hotspot SEND 1000 1 MEGABYTE MESSAGES TO TASK hotspot

LET target BE A RANDOM TASK OTHER THAN 3 WHILE TASK 3 SENDS A 24 BYTE
MESSAGE TO TASK target

LET x BE A RANDOM TASK AND y be a RANDOM TASK GREATER THAN x WHILE
TASK 0 OUTPUTS "Did you know that " AND x AND " is less than " AND y
AND "?"

LET nonzero BE A RANDOM TASK LESS THAN 10 BUT NOT 0 WHILE TASK nonzero
SLEEPS FOR 2 SECONDS

LET middles BE A RANDOM TASK IN [1, num_tasks-2] WHILE TASK middles
ASYNCHRONOUSLY SENDS A 10 BYTE MESSAGE TO TASK ends SUCH THAT ends IS
NOT IN [1, num_tasks-2]

LET num BE 1000 AND num BE num*2 WHILE TASK 0 OUTPUTS num
⊣ 2000
```

That last example demonstrates that LET can bind a variable to a function of its previous value. It is important to remember, though, that variables in CONCEPTUAL cannot be assigned, only bound to a value for the duration of the current scope. They can, however, be bound to a different value for the duration of a child scope. The following example is an attempt to clarify the distinction between binding and assignment:

```
LET var BE 123 WHILE FOR 5 REPETITIONS LET var BE var+1 WHILE TASK 0
OUTPUTS var
⊣ 124
⊣ 124
⊣ 124
⊣ 124
⊣ 124
```

In that example, 'var' is bound to '123' for the scope containing the FOR statement. Then, within the FOR statement, a new scope begins with 'var' being given one plus the value it

had in the outer scope, resulting in '124'. If CONCEPTUAL supported assignment instead of variable-binding, the program would have output '124', '125', '126', '127', and '128'. Note that if A RANDOM TASK were used in the example instead of 'var+1', 'var' would get a different value in each iteration.

When a variable is LET-bound to A RANDOM TASK, all tasks agree on the random number. Otherwise, task A might send a message to task B but task B might be expecting to receive the message from task C, thereby leading to a variety of problems. If there are no valid random tasks, as in the following example, A RANDOM TASK will return '-1':

```
LET invalid_task BE A RANDOM TASK GREATER THAN num_tasks WHILE TASK 0
OUTPUTS invalid_task
⊣ -1
```

Furthermore, the <expr> passed to GREATER THAN is bounded from below by '0' and the <expr> passed to LESS THAN is bounded from above by 'num_tasks-1'. Hence, the following CONCEPTUAL statement will always output values less than or equal to 'num_tasks-1' (unless num_tasks is greater than 1×10^6, of course, in which case it will always output values less than 1×10^6):

```
LET valid_task BE A RANDOM TASK LESS THAN 1E6 WHILE TASK valid_task
OUTPUTS "Hello from " AND valid_task
```

4.7.4 Conditional execution

Like most programming languages, CONCEPTUAL supports conditional code execution:

<if_stmt>   ::=   IF <rel_expr>
                   THEN <simple_stmt>
                   [OTHERWISE <simple_stmt>]

The semantics of an <if_stmt> are that if <rel_expr> (see Section 4.2.5 [Relational expressions], page 101) is TRUE then the first <simple_stmt> is executed. If <rel_expr> is FALSE then the second <simple_stmt> is executed. One restriction is that <rel_expr> must return the same truth value to every task. Consequently, functions that involve task-specific random numbers (see [Random-number functions], page 98) are forbidden within <rel_expr>.

The following is an example of an <if_stmt>:

```
IF this>that THEN TASK 0 SENDS A 3 KILOBYTE MESSAGE TO TASK this
OTHERWISE TASK num_tasks-1 SENDS A 4 KILOBYTE MESSAGE TO TASK that
```

4.7.5 Grouping

FOR loops, LET bindings, and IF statements operate on a single <simple_stmt> (or two <simple_stmt>s in the case of IF...OTHERWISE). Operating on multiple <simple_stmt>s—or, more precisely, operating on a single <complex_stmt> that may consist of multiple <simple_stmt>s—is a simple matter of placing the <simple_stmt>s within curly braces. Contrast the following:

```
FOR 3 REPETITIONS TASK 0 OUTPUTS "She loves me." THEN TASK 0 OUTPUTS
"She loves me not."
⊣ She loves me.
⊣ She loves me.
⊣ She loves me.
⊣ She loves me not.
```

```
FOR 3 REPETITIONS {TASK 0 OUTPUTS "She loves me." THEN TASK 0 OUTPUTS
"She loves me not."}
⊣ She loves me.
⊣ She loves me not.
⊣ She loves me.
⊣ She loves me not.
⊣ She loves me.
⊣ She loves me not.
```

In other words, everything between '{' and '}' is treated as if it were a single statement. Hence, the `FOR` loop applies only to the "She loves me" output in the first statement above, while the `FOR` loop applies to both "She loves me" and "She loves me not" in the second statement.

Variable scoping is limited to the *<simple_stmt>* in the body of a `LET`:

```
LET year BE 1984 WHILE LET year BE 2084 WHILE TASK 0 OUTPUTS year THEN
TASK 0 OUTPUTS year
```
<kbd>error</kbd>   The second '`year`' is outside the scope of both '`LET`' statements.

```
LET year BE 1984 WHILE {{LET year BE 2084 WHILE TASK 0 OUTPUTS year}
THEN TASK 0 OUTPUTS year}
⊣ 2084
⊣ 1984
```

As indicated by the grammatical rules presented at the beginning of Section 4.7 [Complex statements], page 125, CONCEPTUAL does support empty pairs of curly braces, which represent a statement that does nothing and takes no time to execute. While never strictly needed, '`{}`' may be a convenient mechanism for mechanically produced CONCEPTUAL programs.

4.8 Other statements

CONCEPTUAL contains a few more statements than those described in Section 4.4 [Communication statements], page 107 and Section 4.5 [I/O statements], page 119. As there is no category that clearly describes the remaining statements, they are listed here in this "catch-all" section.

4.8.1 Asserting conditions

CONCEPTUAL programs can encode the run-time conditions that must hold in order for the test to run properly. This is achieved through *assertions*, which are expressed as follows:

<assert_stmt>   ::=   `ASSERT THAT` *<string>*
                         `WITH` *<rel_expr>*

<string> is a message to be reported to the user if the assertion fails. *<rel_expr>* is a relational expression (as described in Section 4.2.5 [Relational expressions], page 101) that must evaluate to TRUE for the program to continue running. Assertion failures are considered fatal errors. They cause the CONCEPTUAL program to abort immediately.

Here are some sample *<assert_stmt>*s:

```
ASSERT THAT "the bandwidth test requires at least two tasks" WITH
num_tasks >= 2

ASSERT THAT "pairwise ping-pongs require an even number of task"
WITH num_tasks IS EVEN

ASSERT THAT "this program requires a square number of tasks" WITH
SQRT(num_tasks)**2 = num_tasks
```

(For the last example, recall that CONCEPTUAL expressions are of integer type. Hence, the example's <rel_expr> is mathematically equivalent to $\lfloor\sqrt{N}\rfloor^2 = N$, which is TRUE if and only if N is a square.)

4.8.2 Delaying execution

It is sometimes interesting to measure the progress of a communication pattern when delays are inserted at various times on various tasks. CONCEPTUAL provides two mechanisms for inserting delays: one that relinquishes the CPU while delaying (SLEEP) and one that hoards it (COMPUTE).

<delay_stmt>   ::=   <source_task>
                      SLEEPS | COMPUTES
                      FOR <expr> <time_unit>

   <source_task> (see Section 4.3.2 [Source tasks], page 106) specifies the set of tasks that will stall. <expr> (see Section 4.2.1 [Arithmetic expressions], page 86) specifies the number of <time_unit>s for which to delay and <time_unit> represents any of the following measures of time:

<time_unit>   ::=   MICROSECONDS | MILLISECONDS | SECONDS | MINUTES | HOURS | DAYS

   Delay times are only approximate. SLEEP's accuracy depends upon the operating-system's clock resolution or length of time quantum (commonly measured in milliseconds or tens of milliseconds). COMPUTE, which is implemented by repeatedly reading a variable until the desired amount of time elapses, is calibrated during the CONCEPTUAL run-time system's initialization phase and can be adversely affected by intermittant system load. Both forms of <delay_stmt> attempt to measure wall-clock time ("real time"), not just the time the program is running ("virtual time"). Because the delay times are approximate, it is strongly recommended that the elapsed_usecs variable (see Section A.2 [Predeclared variables], page 207) be employed to determine the actual elapsed time.

4.8.3 Touching memory

"Touching" memory means reading and writing it. The CONCEPTUAL TOUCHES statement enables a program to touch memory for one of two purposes: either to simulate computation in an application by thrashing some or all of the memory hierarchy or to preload message buffers into the upper levels of the memory hierarchy in order to better separate communication costs from memory costs.

Simulating computation

While the statements described in Section 4.8.2 [Delaying execution], page 134 delay for a specified length of time, it is also possible to delay for the duration of a specified amount of "work". "Work" is expressed in terms of memory accesses. That is, a CONCEPTUAL program can touch (i.e., read plus write) data with a given stride from a memory region of a given size. By varying these parameters, a program can emulate an application's computation by hoarding the CPU or any level of the memory hierarchy.

\<touch_stmt\>	::=	*\<source_task\>*
		`TOUCHES`
		[*\<expr\>* *\<data_type\>* `OF`]
		`AN` *\<item_size\>* `MEMORY REGION`
		[*\<expr\>* `TIMES`]
		[`WITH STRIDE` *\<expr\>* *\<data_type\>* `|` `WITH RANDOM STRIDE`]

\<item_size\> and *\<data_type\>* are described in [Item size], page 109 and *\<expr\>* is described in Section 4.2.1 [Arithmetic expressions], page 86.

As shown by the formal definition of *\<touch_stmt\>* the required components are a *\<source_task\>* and the size of the memory region to touch. By default, every `WORD` (see [Item size], page 109) of memory in the region is touched exactly once. The type of data that is touched can be varied with an '`<expr> <data_type> OF`' clause. For instance, '`100 BYTES OF`' of a memory region will touch individual bytes. An optional repeat count enables the memory region (or subset thereof) to be touched multiple times. Hence, if '`TASK 0 TOUCHES A 6 MEGABYTE MEMORY REGION 5 TIMES`', then the touch will be performed as if '`TASK 0`' were told to '`TOUCH 5*6M BYTES OF A 6 MEGABYTE MEMORY REGION 1 TIME`' or simply to '`TOUCH 5*6M BYTES OF A 6 MEGABYTE MEMORY REGION`'.

By default, every *\<data_type\>* of data is touched. However, a *\<touch_stmt\>* provides for touching only a subset of the *\<data_type\>*s in the memory region. By writing '`WITH STRIDE <expr> <data_type>`', only the first *\<data_type\>* out of every *\<expr\>* will be touched. Instead of specifying an exact stride, the memory region can be accessed in random order using the `WITH RANDOM STRIDE` clause.

Unless the number of touches and data type are specified explicitly, the number of `WORD`s that are touched is equal to the size of the memory region divided by the stride length then multiplied by the repeat count. Therefore, if '`TASK 0 TOUCHES AN 8 MEGABYTE MEMORY REGION 2 TIMES WITH STRIDE 8 WORDS`', then a total of $(2^{23}/(4 \times 8)) \times 2 = 524288$ touches will be performed. For the purpose of the preceding calculation, '`WITH RANDOM STRIDE`' should be treated as if it were '`WITH STRIDE 1 WORD`' (again, unless the number of touches and data type are specified explicitly).

To save memory, all `TOUCH` statements in a CONCEPTUAL program access subsets of the same region of memory, whose size is determined by the maximum needed. However, each dynamic execution of a *\<touch_stmt\>* starts touching from where the previous execution left off. For example, consider the following statement:

```
TASK 0 TOUCHES 100 WORDS OF A 200 WORD MEMORY REGION
```

The first time that that statement is executed within a loop (see Section 4.7.2 [Iterating], page 126), the first 200 words are touched. The second time, the second 200 words are touched. The third time, the index into the region wraps around and the first 200 words are touched again.

Each static <touch_stmt> maintains its own index into the memory region. Therefore, the first of the following two statements will terminate successfully (assuming it's not executed in the body of a loop) while the second will result in a run-time error because the final byte of the final word does not fit within the given memory region.

```
TASK 0 TOUCHES 100 WORDS OF A 799 BYTE MEMORY REGION THEN
TASK 0 TOUCHES 100 WORDS OF A 799 BYTE MEMORY REGION
```

```
FOR 2 REPETITIONS TASK 0 TOUCHES 100 WORDS OF A 799 BYTE MEMORY REGION
```

(THEN is described in Section 4.7.1 [Combining statements], page 126, and FOR...REPETITIONS is described in [Counted loops], page 126.) The first statement shown above touches the same 100 words (400 bytes) in each of the two <touch_stmt>s. The second statement touches the first 100 words the first time the <touch_stmt> is executed and fails when trying to touch the (only partially extant) second 100 words.

Priming message buffers

[Data touching], page 111, describes how a message sent or received WITH DATA TOUCHING will have all of its data touched before or after transmission. Sometimes, however, a program may want to touch message data without actually transmitting a message. For example, a task could touch message data to load it into the cache, then RESET ITS COUNTERS (see Section 4.6.1 [Resetting counters], page 123), and finally send or receive a message without further data touching. The TOUCHES statement has an alternate form that touches message buffers instead of isolated memory regions:

<touch_buffer_stmt>   ::=   <source_task>
                               TOUCHES
                                   ALL MESSAGE BUFFERS
                               | MESSAGE BUFFER <expr>
                               | THE CURRENT MESSAGE BUFFER

The first form, 'ALL MESSAGE BUFFERS', touches all message buffers available to the program, even those not yet used at the time TOUCHES is invoked. In the following statement, for example, TOUCHES "knows" that 10 message buffers will be used and touches the data in all 10 of them:

```
ALL TASKS TOUCH ALL MESSAGE BUFFERS THEN
FOR EACH SZ IN {0, ..., 9}
  TASK 0 ASYNCHRONOUSLY SENDS AN SZ MEGABYTE MESSAGE TO TASK 1 THEN
ALL TASKS AWAIT COMPLETION
```

([Buffer control], page 113, explains why the preceding statement requires 10 buffers.) One caveat is that in CONCEPTUAL version 1.5.1b, messages sent or received with the UNIQUE keyword (see [Unique messages], page 109) are not touched. This limitation may be lifted in a later release of CONCEPTUAL.

The second form of the <touch_buffer_stmt> statement, 'MESSAGE BUFFER <expr>', touches a specific message buffer. It is expected to be used in programs that send 'FROM BUFFER <expr>' or receive 'INTO BUFFER <expr>'.

'THE CURRENT MESSAGE BUFFER', the third and final form of the <touch_buffer_stmt> statement, touches the first message buffer that is not currently in use (i.e., is not the source or destination of an asynchronous operation). Usually, this is whichever message buffer will

next be sent from or received into. (The exception is when a task `AWAITS COMPLETION` after touching the current message buffer but before sending or receiving a message—probably a somewhat contrived situation.)

4.8.4 Reordering task IDs

CONCEPTUAL distinguishes between "task IDs", which are used in task descriptions (see Section 4.3 [Task descriptions], page 106) and "processor IDs", which are assigned by the underlying communication layer. As stated in Section 3.4 [Running coNCePTuaL programs], page 46, a CONCEPTUAL program has no control over how processor IDs map to physical processors. It therefore has no way to specify, for instance, that a set of tasks must run on the same multiprocessor node (or on different nodes, for that matter). Initially, every task's task ID is set equal to its processor ID. However, while processor IDs are immutable, task IDs can be changed dynamically during the execution of a program. Altering task IDs can simplify CONCEPTUAL programs that might otherwise need to evaluate complex expressions to determine peer tasks. CONCEPTUAL enables either a specific or a randomly selected task to be assigned to a given processor:

<*processor_stmt*>   ::=   <*source_task*>
                            IS ASSIGNED TO
                            PROCESSOR <*expr*> | A RANDOM PROCESSOR

In addition to performing the specified processor assignment, CONCEPTUAL will perform an additional, implicit processor assignment in order to maintain a bijection between task IDs and processor IDs (i.e., every task ID corresponds to exactly one processor ID and every processor ID corresponds to exactly one task ID). Consider the following statement:

```
TASK n SUCH THAT n<(num_tasks+1)/2 IS ASSIGNED TO PROCESSOR n*2
```

If `num_tasks` is '8' the preceding statement will cause '`TASK 0`' to refer to processor 0, '`TASK 1`' to refer to processor processor 2, '`TASK 2`' to refer to processor 4, and '`TASK 3`' to refer to processor 6. What may be unintuitive is that the remaining tasks will not map to their original processors, as doing so would violate the bijection invariant. To clarify CONCEPTUAL's implicit processor assignments the following timeline illustrates the execution of '`TASK n SUCH THAT n<(num_tasks+1)/2 IS ASSIGNED TO PROCESSOR n*2`' when `num_tasks` is '8':

'n'	'TASK 0'	'TASK 1'	'TASK 2'	'TASK 3'	'TASK 4'	'TASK 5'	'TASK 6'	'TASK 7'
—	0	1	2	3	4	5	6	7
0	0	1	2	3	4	5	6	7
1	0	**2**	**1**	3	4	5	6	7
2	0	2	**4**	3	**1**	5	6	7
3	0	2	4	**6**	1	5	**3**	7

Initially, task and processor IDs are equal. When `n` takes on the value 0, CONCEPTUAL performs the equivalent of '`TASK 0 IS ASSIGNED TO PROCESSOR 0`', which does not change the task ID to processor ID mapping. When `n` is 1, CONCEPTUAL performs the equivalent of '`TASK 1 IS ASSIGNED TO PROCESSOR 2`', which sets task 1's processor to 2. However, because task 2 also has processor 2, CONCEPTUAL implicitly performs the equivalent of '`TASK 2 IS ASSIGNED TO PROCESSOR 1`' in order to preserve the unique task ID to processor ID mapping. When `n` is 2, CONCEPTUAL performs the equivalent of '`TASK 2 IS`

ASSIGNED TO PROCESSOR 4' and, because task 4 also has processor 4, the equivalent of 'TASK 4 IS ASSIGNED TO PROCESSOR 1'. Finally, when n is 3, CONCEPTUAL performs the equivalent of 'TASK 3 IS ASSIGNED TO PROCESSOR 6' and, because task 6 also has processor 6, the equivalent of 'TASK 6 IS ASSIGNED TO PROCESSOR 3'. Thus, CONCEPTUAL maintains the invariant that after any processor assignment every task corresponds to a unique processor and every processor corresponds to a unique task.

4.8.5 Injecting arbitrary code

There are some features that are outside the scope of the CONCEPTUAL language. CONCEPTUAL does, however, provide a mechanism for inserting backend-specific statements into the control flow of a CONCEPTUAL program. This feature is intended for users with specific needs that can't be satisfied through the conventional CONCEPTUAL statements.

<backend_stmt>   ::=   <source_task>
                         BACKEND EXECUTES
                         <expr> | <string>
                         [AND <expr> | <string>]*

The following example assumes a C-based backend:

```
ALL TASKS taskID BACKEND EXECUTE "my_c_function(" AND taskID AND ");"
```

The my_c_function() function needs be defined in some object file and linked with the CONCEPTUAL-generated code.

Most users will never need to BACKEND EXECUTE code. In fact, most users should *not* use <backend_stmt>s as they produce nonportable code. One of CONCEPTUAL's goals is for programs to be understandable by people unfamiliar with the language, and <backend_stmt>s thwart that goal. However, <backend_stmt>s do help ensure that all of the target language/library's features are available to CONCEPTUAL.

All <expr>s are passed to the backend in floating-point context (see [Evaluation contexts], page 87). Consequently, all backend code that takes a CONCEPTUAL <expr> needs to expect a floating-point value (which the backend code can of course cast explicitly to an integer if necessary).

There is an important special case defined for the argument to the BACKEND EXECUTES statement: once all of the <string> and <expr> arguments are concatenated into a single string, all occurrences of the substring '[MESSAGE BUFFER *expr*]' are replaced by a pointer to the *expr*th message buffer created using the FROM BUFFER or INTO BUFFER keywords (presented in [Buffer control], page 113). This special case makes it easy for a CONCEPTUAL program to invoke communication functions provided by the underlying communications library. Note that the notion of a default buffer is not meaningful in the context of BACKEND EXECUTES, which is not a communication statement.

See Section 4.9.3 [Backend-specific declarations], page 140, for a description of BACKEND DECLARES, a companion statement to BACKEND EXECUTES that enables CONCEPTUAL programs to directly declare variables and functions in the target language.

4.9 Header declarations

CoNCePTuaL programs may contain a header section that precedes the first statement in a CoNCePTuaL program. The header section contains three types of declarations that affect the remainder of the program: language versioning declarations, declarations of command-line options, and backend-specific variable and function declarations.

4.9.1 Language versioning

Because the CoNCePTuaL language is still under development, the programmer is forewarned that major changes are likely. Changes may prevent old code from compiling or, even worse, may cause old code to produce incorrect results (e.g., if scoping or block structuring are altered). To mitigate future language changes CoNCePTuaL enables programs to specify which version of the language they were written to. The syntax is straightforward:

<version_decl>   ::=   REQUIRE LANGUAGE VERSION <string>

The parser issues a warning message if <string> does not exactly match the language version supported by the compiler. If the program successfully compiles after a version-mismatch warning, the programmer should check the output very carefully to ensure that the program behaved as expected.

The current version of the CoNCePTuaL language is '1.5'. Note that the language version does not necessarily correspond to the version of the CoNCePTuaL toolset (see Chapter 3 [Usage], page 14) as a whole.

4.9.2 Command-line arguments

CoNCePTuaL makes it easy to declare command-line parameters, although the syntax is a bit verbose:

<param_decl>   ::=   <ident>
                     IS <string>
                     AND COMES FROM <string> OR <string>
                     WITH DEFAULT <expr>

<ident> is the CoNCePTuaL variable being declared. The first <string> is a descriptive string that is provided when the user runs the program with --help or -?. The '<string> OR <string>' terms list the long name for the command-line option, preceded by '--', and the short (single-character) name, preceded by '-'. Finally, <expr> specifies the value that will be assigned to <ident> if the command-line option is not used. <expr> must be a constant expression and may not utilize any of the random-number functions listed in [Random-number functions], page 98. Note that short names (also long names) must be unique.

For instance, the declaration 'nummsgs IS "Number of messages to send" AND COMES FROM "--messages" OR "-m" WITH DEFAULT 25*4' declares a new CoNCePTuaL variable called 'nummsgs'. 'nummsgs' is given the value '100' ('25*4') by default. However, if the user running the program specifies, for example, --messages=55 (or, equivalently, -m 55), then 'nummsgs' will be given the value '55'. The following is an example of the output that might be produced if the program is run with --help or -?:

```
Usage: a.out [OPTION...]
  -m, --messages=<number>     Number of messages to send [default: 100]
  -C, --comment=<string>      Additional commentary to write to the log
                              file, @FILE to import commentary from FILE,
                              or !COMMAND to import commentary from COMMAND
                              (may be specified repeatedly)
  -L, --logfile=<string>      Log-file template [default: "a.out-%p.log"]
  -N, --no-trap=<string>      List of signals that should not be trapped
                              [default: ""]


Help options:
  -?, --help                  Show this help message
  --usage                     Display brief usage message
```

The above is only an example. Depending on what libraries were available when the
CONCEPTUAL run-time system was configured, the output could be somewhat different.
Also, long options may not be supported if a suitable argument-processing library was not
available at configuration time. The above example does indicate one way that help strings
could be formatted. It also shows that the CONCEPTUAL run-time system reserves some
command-line options for its own purposes. Currently, these all use uppercase letters for
their short forms so it should be safe for programs to use any lowercase letter.

4.9.3 Backend-specific declarations

The BACKEND EXECUTES statement (see Section 4.8.5 [Injecting arbitrary code], page 138)
provides support for executing non-CONCEPTUAL code from a CONCEPTUAL program. A
related construct, BACKEND DECLARES, provides support for embedding non-CONCEPTUAL
variable and function declarations in a CONCEPTUAL program:

<*backend_decl*>   ::=   THE BACKEND DECLARES <*string*>

Like BACKEND EXECUTES, BACKEND DECLARES produces nonportable code. Its use is there-
fore strongly discouraged. However, BACKEND DECLARES and BACKEND EXECUTES together
help ensure that all of the target language/library's features are available to CONCEP-
TUAL.

The following example uses BACKEND DECLARES to declare a C global variable and two C
functions that access that variable:

```
THE BACKEND DECLARES "
int tally = 0;

void increment_tally (void)
{
  tally++;
}

void show_tally (char *msg)
{
  printf(\"%s%d.\\n\", msg, tally);
}
".

FOR 3 REPETITIONS PLUS 2 WARMUP REPETITIONS {
  ALL TASKS src SEND A 512-BYTE MESSAGE TO TASK src+1 THEN
  ALL TASKS BACKEND EXECUTE "increment_tally();"
} THEN
TASK 0 BACKEND EXECUTES "show_tally(\"Tally is\\n==> \");".
```

The preceding code works only when using a C-based backend such as `c_mpi` or `c_udgram`. Eliciting the same behavior from a Python-based backend such as `interpret` or `latex_vis` requires a complete rewrite of the coNCePTuaL code:

```
THE BACKEND DECLARES "
global tally
tally = 0

def increment_tally():
    global tally
    tally = tally + 1

def show_tally(msg):
    global tally
    print \"%s%d.\\n\" % (msg, tally)
"

FOR 3 REPETITIONS PLUS 2 WARMUP REPETITIONS {
  ALL TASKS src SEND A 512-BYTE MESSAGE TO TASK src+1 THEN
  ALL TASKS BACKEND EXECUTE "increment_tally()"
} THEN
TASK 0 BACKEND EXECUTES "show_tally(\"Tally is\\n==> \")".
```

It is because of this need to rewrite programs for each set of backends that `BACKEND DECLARES` and `BACKEND EXECUTES` should be used only when absolutely necessary.

4.10 Complete programs

A complete CONCEPTUAL program consists of zero or more header declarations (see
Section 4.9 [Header declarations], page 139), each terminated with a '.', followed by one
or more complex statements (see Section 4.7 [Complex statements], page 125), each also
terminated with a '.'. More formally, CONCEPTUAL's top-level nonterminal is the *<program>*:

<program>   ::=   (*<version_decl>* | *<param_decl>* | *<backend_decl>* ['.'])*
                         (*<top_level_complex_stmt>* ['.'])+

<top_level_complex_stmt>   ::=   *<complex_stmt>*

    Because a *<complex_stmt>* can reduce to a *<simple_stmt>*, the most basic, complete
CONCEPTUAL program would be a *<simple_stmt>* with a terminating period:

```
ALL TASKS self OUTPUT "Hello from task " AND self AND "!".
```

A fuller example might contain multiple header declarations and multiple *<complex_stmt>*s:

```
# A complete coNCePTuaL program
# By Scott Pakin <pakin@lanl.gov>

REQUIRE LANGUAGE VERSION "1.5".

maxval IS "Maximum value to loop to" AND COMES FROM "--maxval" OR
  "-v" WITH DEFAULT 100.

step IS "Increment after each iteration" AND COMES FROM "--step" OR
  "-s" WITH DEFAULT 1.

TASK 0 OUTPUTS "Looping from 0 to " AND maxval AND " by " AND step
  AND "...".

FOR EACH loopvar IN {0, step, ..., maxval}
  TASK 0 OUTPUTS "    " AND loopvar.

TASK 0 OUTPUTS "Wasn't that fun?".
```

    Technically, the '.' is optional; the language is unambiguous without it. However, for
aesthetic purposes it is recommended that you terminate sentences with a period, just like
in a natural language. An exception would be when a *<complex_stmt>* ends with a curly
brace. The '}.' syntax is unappealing so a simple '}' should be used instead. See Chapter 5
[Examples], page 150, for further examples.

Top-level statements and log files

The reason that CONCEPTUAL distinguishes between *<top_level_complex_stmt>*s and
*<complex_stmt>*s is that *<top_level_complex_stmt>*s begin a new table in the log file (see
Section 4.5.3 [Writing to a log file], page 120) while *<complex_stmt>*s add columns to the
current table. Consider the following piece of code:

```
TASK 0 LOGS 111 AS "First" AND
            222 AS "Second".
```

Because 'First' and 'Second' are logged within the same *<simple_stmt>* they appear in the log file within the same table but as separate columns:

```
"First","Second"
"(all data)","(all data)"
111,222
```

The same rule holds when LOGS is used repeatedly across *<simple_stmt>*s but within the same *<complex_stmt>*:

```
TASK 0 LOGS 111 AS "First" THEN
TASK 0 LOGS 222 AS "Second".
```

However, if 'First' and 'Second' are logged from separate *<top_level_complex_stmt>*s, the CONCEPTUAL run-time library stores them in separate tables:

```
TASK 0 LOGS 111 AS "First".
TASK 0 LOGS 222 AS "Second".
```

```
"First"
"(all data)"
111

"Second"
"(all data)"
222
```

4.11 Summary of the grammar

The following is the complete grammar for the CONCEPTUAL language. The EBNF productions appear here in the order that they were presented in the rest of the chapter.

<expr>   ::=   *<cond_expr>*

<cond_expr>   ::=   *<add_expr>* IF *<rel_expr>* OTHERWISE *<add_expr>*

<add_expr>   ::=   *<mult_expr>*
    | *<add_expr>* '+' *<mult_expr>*
    | *<add_expr>* '-' *<mult_expr>*
    | *<add_expr>* '|' *<mult_expr>*
    | *<add_expr>* XOR *<mult_expr>*

<mult_expr>   ::=   *<unary_expr>*
    | *<mult_expr>* '*' *<unary_expr>*
    | *<mult_expr>* '/' *<unary_expr>*
    | *<mult_expr>* MOD *<unary_expr>*
    | *<mult_expr>* '>>' *<unary_expr>*
    | *<mult_expr>* '<<' *<unary_expr>*
    | *<mult_expr>* '&' *<unary_expr>*

<power_expr>   ::=   *<primary_expr>* ['**' *<unary_expr>*]

<unary_expr>   ::=   *<power_expr>*

```
                  |        <unary_operator> <unary_expr>
<unary_operator>   ::=    '+' | '−' | NOT
<primary_expr>    ::=    '(' <expr> ')'
                  |        <ident>
                  |        <integer>
                  |        <func_name> '(' <enumerated_exprs> ')'
                  |        REAL '(' <expr> ')'
<enumerated_exprs>   ::=    <expr> [',' <expr>]*
<func_name>   ::=    ABS | BITS | CBRT | FACTOR10 | LOG10 | MAX | MIN | ROOT | SQRT
                  |        CEILING | FLOOR | ROUND
                  |        TREE_PARENT | TREE_CHILD
                  |        KNOMIAL_PARENT | KNOMIAL_CHILD | KNOMIAL_CHILDREN
                  |        MESH_NEIGHBOR | MESH_COORDINATE | MESH_DISTANCE
                  |        RANDOM_UNIFORM | RANDOM_GAUSSIAN
                  |        RANDOM_POISSON | RANDOM_PARETO
                  |        PROCESSOR_OF | TASK_OF
                  |        FILE_DATA | STATIC_FILE_DATA
<aggr_expr>   ::=    [EACH] <expr>
                  |        THE <expr>
                  |        THE <aggr_func> [AND THE <aggr_func>]* [OF [THE]] <expr>
                  |        A HISTOGRAM OF [THE] <expr>
<aggr_func>   ::=    [ARITHMETIC] MEAN | HARMONIC MEAN | GEOMETRIC MEAN | MEDIAN |
                       STANDARD DEVIATION | VARIANCE | MEDIAN ABSOLUTE DEVIATION |
                       SUM | MINIMUM | MAXIMUM | FINAL | <expr> PERCENTILE
<rel_expr>   ::=    <rel_disj_expr>
<rel_disj_expr>   ::=    [<rel_disj_expr> '\/'] <rel_conj_expr>
<rel_conj_expr>   ::=    [<rel_conj_expr> '/\'] <rel_primary_expr>
<rel_primary_expr>   ::=    <eq_expr>
                          |        '(' <rel_expr> ')'
<eq_expr>   ::=    <expr> '=' <expr>
                  |        <expr> '<' <expr>
                  |        <expr> '>' <expr>
                  |        <expr> '<=' <expr>
                  |        <expr> '>=' <expr>
                  |        <expr> '<>' <expr>
                  |        <expr> DIVIDES <expr>
                  |        <expr> IS EVEN
                  |        <expr> IS ODD
                  |        <expr> IS IN <range> [, <range>]*
                  |        <expr> IS NOT IN <range> [, <range>]*
<range>   ::=    '{' <expr> [',' <expr>]* [', ... ,' <expr>] '}'
           |        '{' <expr> [FOR EACH <ident> IN <range> [',' <range>]* ]+
                       [WHERE <rel_expr>] '}'
<restricted_ident>   ::=    <ident> SUCH THAT <rel_expr>
```

```
<source_task>     ::=   ALL TASKS
                  |     ALL TASKS <ident>
                  |     TASK <expr>
                  |     TASKS <restricted_ident>
                  |     TASK GROUP <ident>
                  |     TASKS <range> [, <range>]*

<target_tasks>    ::=   ALL OTHER TASKS
                  |     TASK <expr>
                  |     TASKS <restricted_ident>
                  |     TASK GROUP <ident>
                  |     TASKS <range> [, <range>]*

<message_spec>         ::=   <item_count>
                            [NONUNIQUE | UNIQUE]
                            <item_size>
                            [UNALIGNED |
                             <message_alignment> ALIGNED |
                             <message_alignment> MISALIGNED]
                            MESSAGES
                            [WITH VERIFICATION | WITH DATA TOUCHING |
                             WITHOUT VERIFICATION | WITHOUT DATA TOUCHING]
                            [USING TAG
                             <expr> | <string>]
                            [FROM [<expr> <data_multiplier> INTO]
                             BUFFER <expr> | THE DEFAULT BUFFER]

<recv_message_spec>   ::=   [SYNCHRONOUSLY | ASYNCHRONOUSLY]
                            [AS [A|AN]
                             [NONUNIQUE | UNIQUE]
                             [UNALIGNED |
                              <message_alignment> ALIGNED |
                              <message_alignment> MISALIGNED]
                             MESSAGES]
                            [WITH VERIFICATION | WITH DATA TOUCHING |
                             WITHOUT VERIFICATION | WITHOUT DATA TOUCHING]
                            [USING TAG
                             <expr> | <string>]
                            [INTO [<expr> <data_multiplier> INTO]
                             BUFFER <expr> | THE DEFAULT BUFFER]

<reduce_message_spec>  ::=   <item_count>
                             [NONUNIQUE | UNIQUE]
                             [UNALIGNED |
                              <message_alignment> ALIGNED |
                              <message_alignment> MISALIGNED]
                             INTEGERS | DOUBLEWORDS
                             [WITH DATA TOUCHING | WITHOUT DATA TOUCHING]
                             [USING TAG
                              <expr> | <string>]
```

```
                                          [FROM [<expr> <data_multiplier> INTO]
                                          BUFFER <expr> | THE DEFAULT BUFFER]
<reduce_target_message_spec>    ::=    [AS <item_count>
                                          [NONUNIQUE | UNIQUE]
                                          [<message_alignment> ALIGNED |
                                           <message_alignment> MISALIGNED]
                                          INTEGERS | DOUBLEWORDS]
                                         [WITH DATA TOUCHING | WITHOUT DATA TOUCHING]
                                         [USING TAG
                                          <expr> | <string>]
                                         [INTO [<expr> <data_multiplier> INTO]
                                          BUFFER <expr> | THE DEFAULT BUFFER]
```

```
<item_count>   ::=   A | AN | <expr>
```

```
<item_size>   ::=    ε
            |       <expr> <data_multiplier>
            |       <data_type> SIZED
```

```
<data_multiplier>   ::=    BIT | BYTE | HALFWORD | WORD | INTEGER | DOUBLEWORD |
                           QUADWORD | PAGE | KILOBYTE | MEGABYTE | GIGABYTE
```

```
<data_type>   ::=   BYTE | HALFWORD | WORD | INTEGER | DOUBLEWORD |
                    QUADWORD | PAGE
```

```
<message_alignment>   ::=    <data_type>
                        |       <expr> <data_multiplier>
```

```
<send_stmt>   ::=    <source_task>
                    [ASYNCHRONOUSLY] SENDS
                    <message_spec>
                    TO [UNSUSPECTING] <target_tasks>
              |     <source_task>
                    [ASYNCHRONOUSLY] SENDS
                    <message_spec>
                    TO <target_tasks>
                    WHO RECEIVE IT
                    <recv_message_spec>
```

```
<receive_stmt>   ::=    <target_tasks>
                       [ASYNCHRONOUSLY] RECEIVE
                       <message_spec>
                       FROM <source_task>
```

```
<wait_stmt>   ::=    <source_task>
                    AWAITS COMPLETION
```

```
<mcast_stmt>   ::=    <source_task>
                     [ASYNCHRONOUSLY] MULTICASTS
                     <message_spec>
                     TO <target_tasks>
```

```
<reduce_stmt>   ::=    <source_task>
                      REDUCES
```

                                                <reduce_message_spec>
                                                TO <source_task>
                                                [WHO RECEIVES THE RESULT <reduce_target_message_spec>]
                           | <source_task>
                                                REDUCES
                                                <reduce_message_spec>
                                                [TO <reduce_message_spec>]

<sync_stmt>   ::=   <source_task>
                               SYNCHRONIZES

<string_or_log_comment>   ::=   <string>
                                       | THE VALUE OF <string>

<output_stmt>   ::=   <source_task>
                               OUTPUTS
                               <expr> | <string_or_log_comment>
                               [AND   <expr> | <string_or_log_comment>]*

<log_stmt>   ::=   <source_task>
                               LOGS
                               <aggr_expr> AS <string_or_log_comment>
                               [AND <aggr_expr> AS <string_or_log_comment>]*

<flush_stmt>   ::=   <source_task>
                               COMPUTES AGGREGATES

<reset_stmt>   ::=   <source_task>
                               RESETS ITS COUNTERS

<store_stmt>   ::=   <source_task>
                               STORES ITS COUNTERS

<restore_stmt>   ::=   <source_task>
                               RESTORES ITS COUNTERS

<complex_stmt>   ::=   <simple_stmt> [THEN <complex_stmt>]

<simple_stmt>   ::=   FOR <expr> REPETITIONS [PLUS <expr> WARMUP REPETITIONS
                               [AND A SYNCHRONIZATION]] <simple_stmt>
                           | FOR EACH <ident> IN <range> [',' <range>]* <simple_stmt>
                           | FOR <expr> <time_unit> [PLUS <expr> WARMUP <time_unit> [AND
                               A SYNCHRONIZATION]] <simple_stmt>
                           | LET <let_binding> [AND <let_binding>]* WHILE <simple_stmt>
                           | IF <rel_expr> THEN <simple_stmt> [OTHERWISE <simple_stmt>]
                           | '{' [<complex_stmt>] '}'
                           | <send_stmt>
                           | <receive_stmt>
                           | <wait_stmt>
                           | <mcast_stmt>
                           | <reduce_stmt>
                           | <sync_stmt>
                           | <output_stmt>
                           | <log_stmt>

```
                    |       <flush_stmt>
                    |       <reset_stmt>
                    |       <store_stmt>
                    |       <restore_stmt>
                    |       <assert_stmt>
                    |       <delay_stmt>
                    |       <touch_stmt>
                    |       <touch_buffer_stmt>
                    |       <processor_stmt>
                    |       <backend_stmt>
```

<let_binding>   ::=   <ident> BE   expr
                                   | <source_task>
                                   | A RANDOM TASK [<random_task_constraints>]

<random_task_constraints>   ::=   OTHER THAN <expr>
                                   | LESS THAN <expr> [BUT NOT <expr>]
                                   | GREATER THAN <expr> [BUT NOT <expr>]
                                   | IN '[' <expr> ',' <expr> ']' [BUT NOT <expr>]

<if_stmt>   ::=   IF <rel_expr>
                   THEN <simple_stmt>
                   [OTHERWISE <simple_stmt>]

<assert_stmt>   ::=   ASSERT THAT <string>
                        WITH <rel_expr>

<delay_stmt>   ::=   <source_task>
                       SLEEPS | COMPUTES
                       FOR <expr> <time_unit>

<time_unit>   ::=   MICROSECONDS | MILLISECONDS | SECONDS | MINUTES | HOURS | DAYS

<touch_stmt>   ::=   <source_task>
                       TOUCHES
                       [<expr> <data_type> OF]
                       AN <item_size> MEMORY REGION
                       [<expr> TIMES]
                       [WITH STRIDE <expr> <data_type> | WITH RANDOM STRIDE]

<touch_buffer_stmt>   ::=   <source_task>
                              TOUCHES
                                 ALL MESSAGE BUFFERS
                               | MESSAGE BUFFER <expr>
                               | THE CURRENT MESSAGE BUFFER

<processor_stmt>   ::=   <source_task>
                           IS ASSIGNED TO
                           PROCESSOR <expr> | A RANDOM PROCESSOR

<backend_stmt>   ::=   <source_task>
                         BACKEND EXECUTES
                         <expr> | <string>
                         [AND <expr> | <string>]*

<*version_decl*>   ::=   `REQUIRE LANGUAGE VERSION` <*string*>

<*param_decl*>   ::=   <*ident*>
                              `IS` <*string*>
                              `AND COMES FROM` <*string*> `OR` <*string*>
                              `WITH DEFAULT` <*expr*>

<*backend_decl*>   ::=   `THE BACKEND DECLARES` <*string*>

<*program*>   ::=   (<*version_decl*> | <*param_decl*> | <*backend_decl*> ['.'])*
                       (<*top_level_complex_stmt*> ['.'])+

<*top_level_complex_stmt*>   ::=   <*complex_stmt*>

The primitives <*ident*>, <*string*>, and <*integer*> are described in

5 Examples

This chapter presents a variety of examples of complete CoNCePTuaL programs. The purpose is to put in context the grammatical elements described in Chapter 4 [Grammar], page 84 and also to illustrate CoNCePTuaL's power and expressiveness.

5.1 Latency

One of the most common network performance benchmarks is a ping-pong latency test. Not surprisingly, such a test is straightforward to implement in CoNCePTuaL:

```
# A ping-pong latency test written in coNCePTuaL

Require language version "1.5".

# Parse the command line.
reps is "Number of repetitions of each message size" and comes from
 "--reps" or "-r" with default 1000.
maxbytes is "Maximum number of bytes to transmit" and comes from
 "--maxbytes" or "-m" with default 1M.

# Ensure the we have a peer with whom to communicate.
Assert that "the latency test requires at least two tasks" with
  num_tasks>=2.

# Perform the benchmark.
For each msgsize in {0}, {1, 2, 4, ..., maxbytes} {
  for reps repetitions {
    task 0 resets its counters then
    task 0 sends a msgsize byte message to task 1 then
    task 1 sends a msgsize byte message to task 0 then
    task 0 logs the msgsize as "Bytes" and
               the median of elapsed_usecs/2 as "1/2 RTT (usecs)"
  } then
  task 0 computes aggregates
}
```

Note that the outer FOR loop specifies two <range>s (see Section 4.2.6 [Range expressions], page 102). This is because '{0, 1, 2, 4, ..., maxbytes}' is not a geometric progression. Hence, that incorrect <range> is split into the singleton '{0}' and the geometric progression '{1, 2, 4, ..., maxbytes}'.

5.2 Hot potato

One way to measure performance variance on a parallel system is with a "hot potato" test. The idea is that the tasks send a message in a ring pattern, then the first task logs the minimum, mean, and variance of the per-hop latency. Ideally, the minimum should

equal the mean and these should both maintain a constant value as the number of tasks increases. Also, the variance should be small and constant. The following CoNCePTuaL code implements a hot-potato test.

```
# Virtual ring "hot potato" test

Require language version "1.5".

trials is "Number of trials to perform" and comes from "--trials" or
  "-t" with default 100000.

Assert that "the hot-potato test requires at least two tasks" with
  num_tasks>=2.

Let len be 0 while {
  for each task_count in {2, ..., num_tasks} {
    task 0 outputs "Performing " and trials and " " and
                   task_count and "-task runs...." then
    for trials repetitions plus 5 warmup repetitions {
      task 0 resets its counters then
      task 0 sends a len byte message to unsuspecting task 1 then
      task (n+1) mod task_count receives a len byte message from task
        n such that n<task_count then
      task n such that n>0 /\ n<task_count sends a len byte message
        to unsuspecting task (n+1) mod task_count then
      task 0 logs the task_count as "# of tasks" and
                  the minimum of elapsed_usecs/task_count as
                    "Latency (usecs)" and
                  the mean of elapsed_usecs/task_count as
                    "Latency (usecs)" and
                  the variance of elapsed_usecs/task_count as
                    "Latency (usecs)"
    } then
    task 0 computes aggregates
  }
}
```

All tasks receive from their left neighbor and send to their right neighbor. However, in order to avoid a deadlock situation, task 0 sends then receives while all of the other tasks receive then send.

5.3 Hot spot

Different systems react differently to resource contention. A hot-spot test attempts to measure the performance degradation that occurs when a task is flooded with data. That is, all tasks except 0 concurrently send a batch of messages to task 0. Task 0 reports the incoming bandwidth, i.e., the number of bytes it received divided by the time it took to

receive that many bytes. The two independent variables are the message size and the number of tasks.

```
# Hot-spot bandwidth

Require language version "1.5".

maxbytes is "Maximum message size in bytes" and comes from
  "--maxbytes" or "-x" with default 1024.
numtrials is "Number of bursts of each size" and comes from "--trials"
  or "-t" with default 100.
burst is "Number of messages in each burst" and comes from
  "--burstsize" or "-b" with default 1000.

Assert that "the hot-spot test requires at least two tasks" with
  num_tasks>=2.

For each maxtask in {2, ..., num_tasks}
  for each msgsize in {1, 2, 4, ..., maxbytes} {
    task 0 outputs "Performing " and numtrials and " " and
                   maxtask and "-task trials with " and
                   msgsize and "-byte messages" then
    for numtrials repetitions plus 3 warmup repetitions {
      task 0 resets its counters then
      task sender such that sender>0 /\ sender<maxtask asynchronously
        sends burst msgsize byte messages to task 0 then
      all tasks await completion then
      task 0 logs the maxtask as "Tasks" and
                   the msgsize as "Message size (B)" and
                   the mean of (1E6*bytes_received)/(1M*elapsed_usecs)
                     as "Incoming BW (MB/s)"
    } then
    task 0 computes aggregates
  }
```

5.4 Multicast trees

It may be worth comparing the performance of a native multicast operation to the performance achieved by multicasting over a *k*-nomial tree to gauge how well the underlying communication layer implements multicasts. The following code records a wealth of data, varying the tree arity (i.e., *k*), the number of tasks receiving the multicast, and the message size. It provides a good demonstration of how to use the KNOMIAL_CHILDREN and KNOMIAL_CHILD functions.

```
# Test the performance of multicasting over various k-nomial trees
# By Scott Pakin <pakin@lanl.gov>

Require language version "1.5".

# Parse the command line.
minsize is "Min. message size (bytes)" and comes from "--minbytes" or "-n"
  with default 1.
maxsize is "Max. message size (bytes)" and comes from "--maxbytes" or "-x"
  with default 1M.
reps is "Repetitions to perform" and comes from "--reps" or "-r" with
  default 100.
maxarity is "Max. arity of the tree" and comes from "--maxarity" or "-a"
  with default 2.

Assert that "this program requires at least two processors" with
  num_tasks>=2.

# Send messages from task 0 to 1, 2, 3, ... other tasks in a k-nomial tree.
For each arity in {2, ..., maxarity} {
  for each num_targets in {1, ..., num_tasks-1} {
    for each msgsize in {minsize, minsize*2, minsize*4, ..., maxsize} {
      task 0 outputs "Multicasting a " and msgsize and "-byte message to "
        and num_targets and " target(s) over a " and arity and
        "-nomial tree ..." then
      for reps repetitions {
        task 0 resets its counters then
        for each src in {0, ..., num_tasks}
          for each dstnum in {0, ..., knomial_children(src, arity,
              num_targets+1)}
            task src sends a msgsize byte message to task
              knomial_child(src, dstnum, arity) then
        all tasks synchronize then
        task 0 logs the arity as "k-nomial arity" and
                    the num_targets as "# of recipients" and
                    the msgsize as "Message size (bytes)" and
                    the median of (1E6/1M)*(msgsize/elapsed_usecs) as
                        "Incoming bandwidth (MB/s)" and
                    the median of (num_targets*msgsize/elapsed_usecs)*
                        (1E6/1M) as "Outgoing bandwidth (MB/s)"
      } then
      task 0 computes aggregates
    }
  }
}
```

5.5 Calling MPI functions

The CONCEPTUAL language is designed to be highly portable. Any CONCEPTUAL program can be compiled using any of the backends listed in Section 3.3 [Supplied backends], page 20. A consequence of this portability is that CONCEPTUAL does not include primitives that are specific to any particular target language or communication library.

The BACKEND EXECUTES and BACKEND DECLARES statements (see Section 4.8.5 [Injecting arbitrary code], page 138, and Section 4.9.3 [Backend-specific declarations], page 140) give a programmer the ability to sacrifice portability for the ability to measure the performance of features provided by a specific target language or communication library. Hence, it is possible to write the core parts of a benchmark in a lower-level language while letting CONCEPTUAL handle the setup, measurement, logging, and other mundane operations.

The following program uses BACKEND EXECUTES to measure the performance of the MPI_Allgather() function provided by an MPI library. Because it utilizes C code to call an MPI function, the code builds only with the c_mpi backend.

```
# Measure the performance of MPI_Allgather()
# By Scott Pakin <pakin@lanl.gov>
#
# N.B. Requires the c_mpi backend.

Require language version "1.5".

# Parse the command line.
numwords is "Message size (words)" and comes from "--msgsize" or "-s" with
default 1.

# Allocate a send buffer and a receive buffer.
Task 0 multicasts a numwords-word message from buffer 0 to all other tasks.
Task 0 multicasts a numwords*num_tasks word message from buffer 1 to all
other tasks.

# Measure the performance of MPI_Allgather().
Task 0 resets its counters then
for 100 repetitions plus 3 warmup repetitions
  all tasks backend execute "
    MPI_Allgather([MESSAGE BUFFER 0], (int)" and numwords and ", MPI_INT,
                  [MESSAGE BUFFER 1], (int)" and numwords and ", MPI_INT,
                  MPI_COMM_WORLD);
  " then
task 0 logs elapsed_usecs/100 as "Gather time (us)".
```

The preceding code demonstrates a few useful techniques:

- A pair of MULTICASTS statements (see Section 4.4.5 [Multicasting], page 117) are used to allocate the two message buffers.

- The `BACKEND EXECUTE` statement uses a value provided on the command line (`numwords`, via `--msgsize`) in the call to `MPI_Allgather()`.
- Pointers to the two message buffers are passed to `MPI_Allgather()` using the '`[MESSAGE BUFFER expr]`' substitution described in Section 4.8.5 [Injecting arbitrary code], page 138.

6 Implementation

CONCEPTUAL could have been implemented as a benchmarking library instead of as a special-purpose language. In addition to improved readability and the practicality of including entire source programs in every log file, one advantage of the language approach is that the same CONCEPTUAL source code can be used to compare the performance of multiple communication libraries. A compiler command-line option selects a particular backend module to use to generate code. Each backend outputs code for a different combination of low-level language and communication library.

The CONCEPTUAL compiler is structured into a pipeline of modules. Thus, the backend can be replaced without altering the front end, lexer, or parser modules. CONCEPTUAL ensures consistency across backends by providing a run-time library that generated code can link to. The run-time library encapsulates many of the mundane tasks a network correctness or performance test needs to perform.

6.1 Overview

Compiler

The CONCEPTUAL compiler is written in Python and is based on the PLY (Python Lex–Yacc) compiler framework. Compiler execution follows a basic pipeline structure. Compilation starts with the top-level file (`ncptl.py`), which processes the command line then transfers control to the lexer (`ncptl_lexer.py`). The lexer inputs CONCEPTUAL source code and outputs a stream of tokens (`ncptl_token.py`). Next, the parser (`ncptl_parser.py`) finds structure in those tokens based on CONCEPTUAL's grammatical rules and outputs an abstract syntax tree (`ncptl_ast.py`). Finally, the code generator (`codegen_language_library.py`) that was designated on the command line walks the abstract syntax tree, converting it to code in the target language and for the target communication library.

Run-time library

CONCEPTUAL makes a large run-time library (`runtimelib.c`) available to generated programs. The CONCEPTUAL run-time library, which is written in C, provides consistent functionality across target languages and communication layers as well as across hardware architectures and operating systems. The library also simplifies code generation by implementing functions for such tasks as memory allocation, queue management, and data logging. The functions in this library are described in Section 6.3 [Run-time library functions], page 165.

Build process

CONCEPTUAL is built using the GNU Autotools (Autoconf, Automake, and Libtool). Consequently, changes should be made to original files, not generated files. Specifically, `configure.ac` and `acinclude.m4` should be edited in place of `configure`; `ncptl.h.in` should be edited in place of `ncptl.h`; and, the various `Makefile.am` files should be edited in place of the corresponding `Makefile`s. See the *Autoconf* documentation, the *Automake* documentation, and the *Libtool* documentation for information about how these various tools operate.

If `configure` is given the `--enable-maintainer-mode` option, `make` will automatically re-run `aclocal`, `autoheader`, `automake`, `autoconf`, and/or `./configure` as needed. Developers who plan to modify any of the "maintainer" files (`acinclude.m4`, `configure.ac`, and the various `Makefile.am` files) are strongly encouraged to configure CoNCePTuaL with `--enable-maintainer-mode` in order to ensure that the build process is kept current with any changes.

6.2 Backend creation

The CoNCePTuaL compiler's backend generates code from an abstract syntax tree (AST). The compiler was designed to support a variety of code generators, each targeting a particular programming language and communication library. There are two ways to create a new CoNCePTuaL backend. Either a `codegen_language_library.py` backend supporting an arbitrary language and communication library can be written from scratch or a C-based `codegen_c_library.py` backend can be derived from `codegen_c_generic.py`.

In the former case, the backend must define an `NCPTL_CodeGen` class. `NCPTL_CodeGen` class must contain a `generate` method with the following signature:

```
def generate(self, ast, filesource='<stdin>', filetarget='-',
             sourcecode=None):
```

That is, `generate` takes as arguments a class object, the root of an abstract-syntax tree (as defined in `ncptl_ast.py`), the name of the input file containing CoNCePTuaL code (to be used for outputting error messages), the name of the output file to produce, and the complete CoNCePTuaL source code (which is both stored in prologue comments and passed to the run-time library). `generate` should invoke `self.postorder_traversal` to traverse the AST in a post-order fashion, calling various code-generating methods as it proceeds. The `NCPTL_CodeGen` must implement all of the methods listed in Section B.1 [Method calls], page 209, each of which corresponds to some component of the CoNCePTuaL grammar. Each method takes a "self" class object an a node of the AST (of type `AST`).

The compiler front-end, `ncptl`, invokes the following two methods, which must be defined by the backend's `NCPTL_CodeGen` class:

```
def compile_only(self, progfilename, codelines, outfilename,
                 verbose, keepints):
```

```
def compile_and_link(self, progfilename, codelines, outfilename,
                     verbose, keepints):
```

The `compile_only` method compiles the backend-specific code into an object file. The `compile_and_link` method compiles the backend-specific code into an object file and links it into an executable file. For some backends, the notions of "compile" and "link" may not be appropriate. In that situation, the backend should perform the closest meaningful operations. For example, the `dot_ast` backend (see Section 3.3.11 [The dot_ast backend], page 41) compiles to a `.dot` file and links into the target graphics format (`.ps` by default).

For both the `compile_only` and `compile_and_link` methods, `progfilename` is the name of the CoNCePTuaL input file specified on the `ncptl` command line or the string '`<command line>`' if a program was specified literally with `--program`. `codelines` is the output from the `generate` method, i.e., a list of lines of backend-specific code. `outfilename` is the name of the target file specified on the `ncptl` command line with `--output` or the string '`-`' if `--output` was not used. If `verbose` is '1', the method should write each operation it plans to perform to the standard-error device. For consistency, comment lines should begin with '`#`'; shell commands should be output verbatim. If `verbose` is '0', corresponding to the `ncptl` `--quiet` option, the method should output nothing but error messages. Finally, `keepints` corresponds to the `--keep-ints` option to `ncptl`. If equal to '0', all intermediate files should be deleted before returning; if equal to '1', intermediate files should be preserved. See Section 3.2 [Compiling coNCePTuaL programs], page 18, for a description of the various command-line options to `ncptl`.

As long as `NCPTL_CodeGen` implements all of the required functions it is free to generate code in any way that it sees fit. However, Section B.1 [Method calls], page 209, lists a large number of methods, many of which will be identical across multiple code generators for the same language but different communication libraries. To simplify a common case, C plus some messaging library, CoNCePTuaL provides `codegen_c_generic.py`, to which the remainder of the Implementation chapter is devoted to explaining.

6.2.1 Hook methods

Multiple code generators for the same language but different communication libraries are apt to contain much code in common. Because C is a popular language, CoNCePTuaL provides a `codegen_c_generic.py` module that implements a virtual `NCPTL_CodeGen` base class. This base class implements all of the methods listed in Section B.1 [Method calls], page 209. However, rather than support a particular communication library, the `codegen_c_generic.py` implementation of `NCPTL_CodeGen` merely provides a number of calls to "hook" methods—placeholders that implement library-specific functionality. See Section B.2 [C hooks], page 211, for a list of all of the hooks that `codegen_c_generic.py` defines. For clarity, hooks are named after the method from which they're called but with an all-uppercase tag appended. Hook methods take a single parameter, a read-only dictionary (the result of invoking Python's `locals()` function) of all of the local variables in the caller's scope. They return C code in the form of a list with one line of C per element. A hook method is invoked only if it exists, which gives the backend developer some flexibility in selecting places at which to insert code. Of course, for coarser-grained control, a backend developer can override complete methods in `codegen_c_generic.py` if desired. Generally, this will not be necessary as hook invocations are scattered liberally throughout the file.

An example

`codegen_c_generic.py` defines a method named `code_specify_include_files`. (`codegen_c_generic.py` names all of its code-generating helper methods '`code_something`'.) `code_specify_include_files` pushes a sequence of `#include` directives onto a queue of lines of C code. The method is shown below in its entirety:

```
def code_specify_include_files(self, node):
    "Load all of the C header files the generated code may need."

    # Output a section comment.
    self.pushmany([
        "/*****************",
        " * Include files *",
        " *****************/",
        ""])

    # Enable hooks both before and after the common includes.
    self.pushmany(self.invoke_hook("code_specify_include_files_PRE",
                                   locals(),
                                   before=[
        "/* Header files specific to the %s backend */" %
        self.backend_name],
                                   after=[""]))
    self.pushmany([
        "/* Header files needed by all C-based backends */",
        "#include <stdio.h>",
        "#include <string.h>",
        "#include <ncptl/ncptl.h>"])
    self.pushmany(self.invoke_hook("code_specify_include_files_POST",
                                   locals(),
                                   before=[
        "",
        "/* Header files specific to the %s backend */" %
        self.backend_name]))
```

code_specify_include_files uses the pushmany method (see Section 6.2.4 [Internals], page 163) to push each element in a list of lines of C code onto the output queue. It starts by pushing a section comment—codegen_c_generic.py outputs fully commented C code. Next, it invokes the code_specify_include_files_PRE hook if it exists and pushes that method's return value onto the queue. Then, it pushes all of the #includes needed by the generated C code. Finally, it invokes the code_specify_include_files_POST hook if it exists and pushes that method's return value onto the queue.

A backend that requires additional header files from those included by code_specify_include_files need only define code_specify_include_files_PRE to add extra header files before the standard ones or code_specify_include_files_POST to add extra header files after them. The following is a sample (hypothetical) hook definition:

```
def code_specify_include_files_POST(self, localvars):
    "Specify extra header files needed by the c_pthreads backend."
    return [
        "#include <errno.h>",
        "#include <pthread.h>"]
```

Although the top-level structure of `codegen_c_generic.py` is described in Section 6.2.4 [Internals], page 163, a backend developer will normally need to study the `codegen_c_generic.py` source code to discern the purpose of each hook method and its relation to the surrounding code.

6.2.2 A minimal C-based backend

A backend derived from `codegen_c_generic.py` starts by defining an `NCPTL_CodeGen` child class that inherits much of its functionality from the parent `NCPTL_CodeGen` class. There are only two items that a C-based backend *must* define: `backend_name`, the name of the backend in the form '`c_library`'; and, `backend_desc`, a brief phrase describing the backend. (These are used for error messages and file comments.) Also, a backend's `__init__` method must accept an `options` parameter, which is given a list of command-line parameters not recognized by `ncptl.py`. After `NCPTL_CodeGen`'s `__init__` method processes the entries in `options` that it recognizes, it should pass the remaining options to its parent class's `__init__` method for further processing. (For proper initialization, the parent class's `__init__` method must be called, even if there are no remaining options to process.)

The following is the complete source code to a minimal coNCePTuaL backend. This backend, `codegen_c_seq.py`, supports only sequential coNCePTuaL programs (e.g., '`TASK 0 OUTPUTS "Hello, world!"`'); any attempt to use communication statements (see Section 4.4 [Communication statements], page 107) will result in a compile-time error.

```
#! /usr/bin/env python

############################################################
# Code generation module for the coNCePTuaL language: #
# Minimal C-based backend -- all communication         #
# operations result in a compiler error                #
#                                                       #
# By Scott Pakin <pakin@lanl.gov>                       #
############################################################


import codegen_c_generic

class NCPTL_CodeGen(codegen_c_generic.NCPTL_CodeGen):
    def __init__(self, options):
        "Initialize the sequential C code generation module."
        self.backend_name = "c_seq"
        self.backend_desc = "C, sequential code only"
        codegen_c_generic.NCPTL_CodeGen.__init__(self, options)

        # We don't have our own command-line options but we handle
        # --help, nevertheless.
        for arg in range(0, len(options)):
            if options[arg] == "--help":
                # Output a help message.
                self.show_help()
                raise SystemExit, 0
```

The `c_seq` backend can be used like any other:

```
ncptl --backend=c_seq \
  --program='For each i in {10, ..., 1} task 0 outputs i.' | \
  indent > myprogram.c
```

(`codegen_c_generic.py` outputs unindented code, deferring attractive formatting to the
Unix `indent` utility.)

One sequential construct the `c_seq` backend does not support is randomness, as needed
by A RANDOM PROCESSOR (see Section 4.8.4 [Reordering task IDs], page 137) and A RANDOM
TASK (see Section 4.7.3 [Binding variables], page 130). `codegen_c_generic.py` cannot
support randomness itself because doing so requires broadcasting the seed for the random-
number generator to all tasks. Broadcasting requires messaging-layer support, which a de-
rived backend provides through the `code_def_init_reseed_BCAST` hook (see Section 6.2.1
[Hook methods], page 158). For the sequential backend presented above, a broadcast can
be implemented as a no-op:

```
def code_def_init_reseed_BCAST(self, localvars):
    '"Broadcast" a random-number seed to all tasks.'
    return []
```

In fact, that same do-nothing hook method is used by the `c_udgram` backend. `c_udgram` seeds the random-number generator before calling `fork()`, thereby ensuring that all tasks have the same seed without requiring an explicit broadcast.

6.2.3 Generated code

`codegen_c_generic.py` generates thoroughly commented C code. However, the overall structure of the generated code may be somewhat unintuitive, as it does not resemble the code that a human would write to accomplish a similar task. The basic idea behind the generated C code is that it expands the entire program into a list of "events", then starts the clock, then executes all of the events in a single loop. Regardless of the CONCEPTUAL program being compiled, the body of the generated C code will look like this:

```
for (i=0; i<numevents; i++) {
  CONC_EVENT *thisev = &eventlist[i];

  switch (thisev->type) {
    case event_1:



                                      ⋮



    case event_2:



                                      ⋮



  }
}
```

Programs generated by `codegen_c_generic.py` define a number of event types that are summarized in Section B.3 [Event types], page 215. The `EV_CODE` event is used, for example, by the `BACKEND EXECUTES` (see Section 4.8.5 [Injecting arbitrary code], page 138), `LOGS` (see Section 4.5.3 [Writing to a log file], page 120), and `OUTPUTS` (see Section 4.5.2 [Writing to standard output], page 120) statements. Note that there are no loop events—in fact, there are no complex statements (see Section 4.7 [Complex statements], page 125) whatsoever. Complex statements are expanded into multiple simple statements at initialization time.

The advantage of completely expanding a CONCEPTUAL program during the initialization phase—essentially, "pre-executing" the entire program—is that that enables all of the expensive, non-communication-related setup to be hoisted out of the timing loop, which is how a human would normally express a network benchmark. Pre-execution is possible because the CONCEPTUAL language is not a Turing machine; infinite loops are not expressible by the language and message contents and timings cannot affect program behavior, for instance. During its initialization phase, the generated C code allocates memory for message buffers, evaluates numerical expressions, verifies program assertions, unrolls loops, and does

everything else that's not directly relevant to communication performance. For instance, the CONCEPTUAL program 'TASK tx SUCH THAT tx>4 SENDS 10 1 MEGABYTE MESSAGES TO TASK tx/2' would cause each task to perform the following steps during initialization:

- determine if its task ID is greater than 4, making the task a sender
- determine if its task ID is equal to 'tx/4' (rounded down to the nearest integer) for some task 'tx' in the program, making the task a receiver
- allocate 1 MB for a message buffer
- allocate and initialize a repeat event, specifying that the subsequent event should repeat 10 times
- allocate a send or receive event

The final two of those steps repeat as necessary. For example, task 3 receives 10 messages from each of task 6 and task 7.

Note that each task's receive events (if any) are allocated before its send events (if any), as described Section 4.4.2 [Sending], page 113. Also, note that only a single message buffer is allocated because the CONCEPTUAL source did not specify the UNIQUE keyword (see [Unique messages], page 109).

An event is implemented as a C struct that contains all of the state needed to perform a particular operation. For example, an event corresponding to a synchronous or asynchronous send operation (CONC_SEND_EVENT) stores the destination task ID, the number of bytes to send, the message alignment, the number of outstanding asynchronous sends and receives, a flag indicating whether the data is to be touched, and a flag indicating that the message should be filled with data the receiver can verify. In addition, the code_declare_datatypes_SEND_STATE hook (see Section 6.2.1 [Hook methods], page 158) enables a backend to include additional, backend-specific state in the (CONC_SEND_EVENT) data structure.

6.2.4 Internals

codegen_c_generic.py is a fairly substantial piece of code. It is divided into ten sections:

1. methods exported to the compiler front end
2. utility functions that do not generate code
3. utility functions that do generate code
4. methods for outputting language atoms (see Section 4.1 [Primitives], page 84)
5. methods for outputting miscellaneous language constructs (e.g., restricted identifiers; see Section 4.3.1 [Restricted identifiers], page 106)
6. methods for outputting expressions (see Section 4.2 [Expressions], page 86)
7. methods for outputting complete programs (see Section 4.10 [Complete programs], page 142)
8. methods for outputting complex statements (see Section 4.7 [Complex statements], page 125)
9. methods for outputting simple statements (e.g., communication statements; see Section 4.4 [Communication statements], page 107)
10. methods for outputting nodes with non-textual names (e.g., '...' and various operators)

The `NCPTL_CodeGen` class defined in `codegen_c_generic.py` generates code as follows. The `generate` method, which is invoked by `ncptl.py`, calls upon PLY to process the abstract-syntax tree (AST) in postorder fashion. `NCPTL_CodeGen` maintains a stack (`codestack`) on which code fragments are pushed and popped but that ends up containing a complete line of C code in each element. For example, in the CONCEPTUAL program 'TASK 0 OUTPUTS 1+2*3', the `n_outputs` method will pop '[('expr', '(1)+((2)*(3))')]' (a list containing the single expression '1+2*3') and '('task_expr', '0')' (a tuple designating a task by the expression '0') and push multiple lines of code that prepare task 0 to evaluate and output the given expression.

The utility functions are the most useful for backend developers to understand, as they are frequently called from hook methods (see Section 6.2.1 [Hook methods], page 158). The following should be of particular importance:

`push`
`pushmany`   Push a single value (typically a string of C code) or each value in a list of values onto a stack.

`error_fatal`
`error_internal`
         Output a generic error message or an "internal error" error message and abort the program.

`code_declare_var`
         Push (using the `push` method) a line of C code that declares a variable with an optionally specified type, name, initial value, and comment. Return the variable name actually used.

See the definitions in `codegen_c_generic.py` of each of the above to determine required and optional parameters. The following, adapted from `codegen_c_udgram.py` demonstrates some of the preceding methods:

```
def n_for_count_SYNC_ALL(self, localvars):
    "Synchronize all of the tasks in the job."
    synccode = []
    self.push("{", synccode)
    loopvar = self.code_declare_var(suffix="task",
      comment="Loop variable that iterates over all (physical) ranks",
      stack=synccode)
    self.pushmany([
      "thisev_sync->s.sync.peerqueue = ncptl_queue_init (sizeof(int));",
      "for (%s=0; %s<var_num_tasks; %s++)" %
      (loopvar, loopvar, loopvar),
      "*(int *)ncptl_queue_allocate(thisev_sync->s.sync.peerqueue) = %s;" %
      loopvar,
      "thisev_sync->s.sync.syncrank = physrank;",
      "}"],
                  stack=synccode)
    return synccode
```

That definition of the `n_for_count_SYNC_ALL` hook method defines a new stack (`synccode`) and pushes a '{' onto it. It then declares a loop variable, letting `code_declare_var` select a name but dictating that it end in '`_task`'. The hook method then pushes some additional C code onto the `synccode` stack and finally returns the stack (which is really just a list of lines of C code).

Some useful variables defined by `NCPTL_CodeGen` include the following:

`base_global_parameters`

> a list of 6-ary tuples defining extra command-line parameters to parse (format: {*type*, *variable*, *long_name*, *short_name*, *description*, *default_value*})

`events_used`

> a dictionary containing the names of events actually used by the program being compiled

Some methods in `codegen_c_generic.py` that are worth understanding but are unlikely to be used directly in a derived backend include the following:

`pop`   Pop a value from a stack.

`push_marker`

> Push a specially designated "marker" value onto a stack.

`combine_to_marker`

> Pop all items off a stack up to the first marker value found; discard the marker; then, push the popped items as a single list of items. This is used, for example, by a complex statement (see Section 4.7 [Complex statements], page 125) that applies to a list of statements, which can be popped as a unit using `combine_to_marker`.

`invoke_hook`

> Call a hook method, specifying code to be pushed before/after the hook-produced code and alternative text (or Python code) to be pushed (or executed) in the case that a hook method is not provided.

6.3 Run-time library functions

To simplify the backend developer's task and to provide consistent functionality across backends, cONCEPTuAL provides a run-time library that encapsulates many of the common operations needed for network-correctness and performance-testing programs. This section describes all of the functions that the library exports (plus a few important types and variables). The library is written in C, so all of the type/variable/function prototypes are expressed with C syntax. The library includes, among others, functions that manage heap-allocated memory, accurately read the time, write results to log files, control queues of arbitrary data, and implement various arithmetic operations. All of these functions should be considered "slow" and should therefore generally not be invoked while execution is being timed.[1]

---

[1] Some notable exceptions are the functions described in Section 6.3.4 [Message-buffer manipulation functions], page 169, which implement cONCEPTuAL's `WITH DATA TOUCHING` and `WITH VERIFICATION` constructs.

6.3.1 Constants, variables, and data types

The following constants, variables, and data types are used by various run-time library functions and directly by backends.

`ncptl_int`   [Data type]

    The internal data type of the CONCEPTUAL run-time library is `ncptl_int`. This is normally a 64-bit signed integer type selected automatically by `configure` (see Section 2.1 [configure], page 4) but can be overridden with the `--with-datatype` option to `configure`. `ncptl.h` defines a string macro called `NICS` that can be used to output an `ncptl_int` regardless of how the `ncptl_int` type is declared:

```
ncptl_fatal ("My variable contains a negative value (%" NICS ")",
             my_ncptl_int_var);
```

    `ncptl_int` constants declared by backends derived from `codegen_c_generic.py` are given an explicit suffix that defaults to 'LL' but can be overridden at configuration time using the `--with-const-suffix` option.

`NCPTL_INT_MIN`   [Constant]

    `NCPTL_INT_MIN` is a C preprocessor macro that represents the smallest (i.e., most negative) number that can be assigned to an `ncptl_int`. For example, if `ncptl_int` is a 64-bit signed integer type, then `NCPTL_INT_MIN` will be the value '-9223372036854775808'.

`NCPTL_CMDLINE`   [Data type]

    The `NCPTL_CMDLINE` structure describes an acceptable command-line option. It contains a type, which is either `NCPTL_TYPE_INT` for an `ncptl_int` or `NCPTL_TYPE_STRING` for a `char *`, a pointer to a variable that will receive the value specified on the command line, the long name of the argument (without the '--'), the one-letter short name of the argument (without the '-'), a textual description of what the argument represents, and a default value to use if the option is not specified on the command line.

`NCPTL_QUEUE`   [Data type]

    An `NCPTL_QUEUE` is an opaque data type that represents a dynamically growing queue that can be flattened to an array for more convenient access. `NCPTL_QUEUE`s have proved to be quite useful when implementing CONCEPTUAL backends.

`NCPTL_LOG_FILE_STATE`   [Data type]

    Every CONCEPTUAL log file is backed by a unique `NCPTL_LOG_FILE_STATE` opaque data type. An `NCPTL_LOG_FILE_STATE` data type represents all of the state needed to maintain that file, such as file descriptors, prologue comments, and data that has not yet been aggregated.

`int ncptl_pagesize`   [Variable]

    This variable is initialized by `ncptl_init()` to the number of bytes in an operating-system memory page. `ncptl_pagesize` can be used by backends to implement CONCEPTUAL's `PAGE SIZED` (see [Item size], page 109) and `PAGE ALIGNED` (see [Message alignment], page 111) keywords.

`int ncptl_fast_init`   [Variable]

> The `ncptl_init()` function (see Section 6.3.2 [Initialization functions], page 167) can take many seconds to complete. Much of this time is spent calibrating and measuring the quality of the various timers the run-time library uses. For backends such as `picl` (see Section 3.3.8 [The picl backend], page 34) that do not measure real time there is little need to have an accurate timer. Setting `ncptl_fast_init` to '1' before invoking `ncptl_init()` skips the timer calibration and measurement steps, thereby leading to faster initialization times. A user can also override the setting of `ncptl_fast_init` at run time by setting the `NCPTL_FAST_INIT` environment variable to either '0' or '1', as appropriate.

6.3.2 Initialization functions

The following functions are intended to be called fairly early in the generated code.

`void ncptl_init` (*int* `version`, *char* `*program_name`)   [Function]

> Initialize the CoNCePTuaL run-time library. *version* is used to verify that `runtimelib.c` corresponds to the version of `ncptl.h` used by the generated code. The caller must pass in `NCPTL_RUN_TIME_VERSION` for *version*. *program_name* is the name of the executable program and is used for outputting error messages. The caller should pass in `argv[0]` for *program_name*. `ncptl_init()` must be the first library function called by the generated code (with a few exceptions, as indicated below).

`void ncptl_permit_signal` (*int* `signalnum`)   [Function]

> Indicate that the backend relies on signal *signalnum* for correct operation. Because signal handling has performance implications, the CoNCePTuaL run-time library normally terminates the program upon receiving a signal. Hence, the user can be assured that if a program runs to completion then no signals have affected its performance. See Section 3.4 [Running coNCePTuaL programs], page 46, for a description of the `--no-trap` command-line option, which enables a user to permit additional signals to be delivered to the program (e.g., when linking with a particular implementation of a communication library that relies on signals). `ncptl_permit_signal()` must be invoked before `ncptl_parse_command_line()` to have any effect.

`void ncptl_parse_command_line` (*int* `argc`, *char* `*argv[]`,   [Function]
>   *NCPTL_CMDLINE* `*arglist`, *int* `numargs`)

> Parse the command line. *argc* and *argv* should be the argument count and argument vector passed to the generated code by the operating system. *arglist* is a list of descriptions of acceptable command-line arguments and *numargs* is the length of that list.

Because `ncptl_init()` takes many seconds to run, it is common for generated code to scan the command line for `--help` or `-?` and, if found, skip `ncptl_init()` and immediately invoke `ncptl_parse_command_line()`. Doing so gives the user immediate feedback when requesting program usage information. Skipping `ncptl_init()` is safe in this context because `ncptl_parse_command_line()` terminates the program after displaying usage information; it does not require any information discovered by `ncptl_init()`.

Most generated programs have a --*seed*/-*S* option that enables the user to specify explicitly a seed for the random-number generator with --*help*/-*?* showing the default seed. `ncptl_seed_random_task()` must therefore be called before `ncptl_parse_command_line()` which, as stated in the previous paragraph, can be invoked without a prior invocation of `ncptl_init()`. Consequently, it can be considered safe also to invoke `ncptl_seed_random_task()` before `ncptl_init()`.

A generated program's initialization routine will generally exhibit a structure based on the following pseudocode:

> **if** "--`help`" **or** "-`?`" **in** command-line options **then**
>   *only_help* := TRUE
> **else**
>   *only_help* := FALSE
>   `ncptl_init(...)`
> **end if**
> *random_seed* := `ncptl_seed_random_task(0)`
> `ncptl_parse_command_line(...)`
> **if** *only_help* = TRUE **then**
>   `ncptl_error("Internal error; should have exited")`
> **end if**
> `ncptl_seed_random_task(`*random_seed*`)`

6.3.3 Memory-allocation functions

The CONCEPTUAL run-time library provides its own wrappers for `malloc()`, `free()`, `realloc()`, and `strdup()` as well as a specialized `malloc()` designed specifically for allocating message buffers. The wrappers' "value added" is that they support the explicit data alignments needed by `ALIGNED` messages (see [Message alignment], page 111) and that they automatically call `ncptl_fatal()` on failure, so the return value does not need to be checked for `NULL`.

void * ncptl_malloc (*ncptl_int* `numbytes`, *ncptl_int* `alignment`)   [Function]
> Allocate *numbytes* bytes of memory aligned to an *alignment*-byte boundary. If *alignment* is '0', `ncptl_malloc()` will use whatever alignment is "natural" for the underlying architecture. `ncptl_malloc()` will automatically call `ncptl_fatal()` if memory allocation fails. Therefore, unlike `malloc()`, there is no need to check the return value for `NULL`.

void * ncptl_malloc_misaligned (*ncptl_int* `numbytes`, *ncptl_int*   [Function]
    `misalignment`)
> Allocate *numbytes* bytes of memory from the heap aligned *misalignment* bytes past a page boundary. If *alignment* is '0', `ncptl_malloc_misaligned()` will return page-aligned memory. `ncptl_malloc_misaligned()` will automatically call `ncptl_fatal()` if memory allocation fails. Therefore, unlike `malloc()`, there is no need to check the return value for `NULL`.

void ncptl_free (*void* `*pointer`)   [Function]
> Free memory previously allocated by `ncptl_malloc()`. It is an error to pass `ncptl_free()` memory not allocated by `ncptl_malloc()`.

void * ncptl_realloc (*void *pointer*, *ncptl_int* **numbytes**, *ncptl_int*   [Function]
        alignment)

Given a pointer returned by `ncptl_malloc()`, change its size to *numbytes* and byte-alignment to *alignment* without altering the contents (except for truncation in the case of a smaller target size). If *alignment* is '0', `ncptl_realloc()` will use whatever alignment is "natural" for the underlying architecture. `ncptl_realloc()` will automatically call `ncptl_fatal()` if memory allocation fails. Therefore, unlike `realloc()`, there is no need to check the return value for `NULL`.

char * ncptl_strdup (*const char *string*)   [Function]

`ncptl_strdup()` copies a string as does the standard C `strdup()` function. However, `ncptl_strdup()` uses `ncptl_malloc()` instead of `malloc()` to allocate memory for the copy, which must therefore be deallocated using `ncptl_free()`.

void * ncptl_malloc_message (*ncptl_int* **numbytes**, *ncptl_int*   [Function]
        alignment, *ncptl_int* **outstanding**, *int* **misaligned**)

Allocate *numbytes* bytes of memory from the heap either aligned on an *alignment*-byte boundary (if *misaligned* is '0') or *alignment* bytes past a page boundary (if *misaligned* is '1'). All calls with the same value of *outstanding* will share a buffer. `ncptl_malloc_message()` is intended to be used in two passes. The first time the function is called on a set of messages it merely determines how much memory to allocate. The second time, it returns valid memory buffers. Note that the returned pointer can be neither `free()`d nor `ncptl_free()`d.

void * ncptl_get_message_buffer (*ncptl_int* **buffernum**)   [Function]

Return a pointer to a message buffer previously allocated (and finalized) by `ncptl_malloc_message()`. The *buffernum* argument to `ncptl_get_message_buffer()`, which corresponds to the *outstanding* argument to `ncptl_malloc_message()`, specifies the number of the buffer to return. `ncptl_get_message_buffer()` returns `NULL` if buffer *buffernum* is either unallocated or uninitialized.

6.3.4 Message-buffer manipulation functions

The CONCEPTUAL language facilitates verifying message contents and touching every word in a message (see [Data touching], page 111). The following functions implement those features.

void ncptl_fill_buffer (*void *buffer*, *ncptl_int* **numbytes**, *int*   [Function]
        validity)

Fill a region of memory with known values. If *validity* is '+1', `ncptl_fill_buffer()` will fill the first *numbytes* bytes of *buffer* with a verifiable sequence of integers (see [Data touching], page 111). If *validity* is '-1', `ncptl_fill_buffer()` will pollute the first *numbytes* bytes of *buffer*. Receive buffers should be polluted before reception to avoid false negatives caused, for example, by an inadvertently dropped message destined for a previously validated buffer.

ncptl_int ncptl_verify (*void *buffer*, *ncptl_int* **numbytes**)   [Function]

Verify the contents of memory filled by `ncptl_fill_buffer()`. The function returns the number of erroneous bits. `ncptl_verify()` is used to implement CONCEPTUAL's `WITH VERIFICATION` construct (see [Data touching], page 111).

void ncptl_touch_data (*void *buffer*, *ncptl_int **numbytes***)   [Function]
     Touch every byte in a given buffer. `ncptl_touch_data()` is used to implement the
     `WITH DATA TOUCHING` construct described in [Data touching], page 111.

The following function, `ncptl_touch_memory()`, does not actually manipulate message
buffers. It is included in this section because of its similarity to `ncptl_touch_data()`.
`ncptl_touch_data()` touches message buffers to implement the `WITH DATA TOUCHING` con-
struct; `ncptl_touch_memory()` touches a hidden memory region to implement the `TOUCHES`
statement.

void ncptl_touch_memory (*void *buffer*, *ncptl_int **bufferbytes***,   [Function]
          *ncptl_int **wordbytes***, *ncptl_int **firstbyte***, *ncptl_int **numaccesses***, *ncptl_int*
          bytestride)
     Walk a memory region *buffer* of size *bufferbytes* bytes. Each "word" to touch con-
     tains *wordbytes* bytes. *firstbyte* indicates the byte index into *buffer* from which
     to start touching. The function will read and write *numaccesses* "words" with
     stride *bytestride* bytes. For example, '`ncptl_touch_memory (mybuffer, 1048576,`
     `64, 192, 10000, 4096)`' will read and write (but otherwise do nothing with) the
     64 bytes that lie 192 bytes into a 1 MB memory region, then the 64 bytes starting at
     offset 4288, then the 64 bytes starting at $192 + 4096 \times 2$, then at $192 + 4096 \times 3$, then
     at $192 + 4096 \times 4$, and so forth up to $192 + 4096 \times 10000$, wrapping around the 1 MB
     region as necessary.

          A *bytestride* of `NCPTL_INT_MIN` implies a random stride.

          As an important special case, if *firstbyte* is '`-1`', then `ncptl_touch_memory()` will
     touch one or more message buffers (cf. `ncptl_malloc_message()` in Section 6.3.3
     [Memory-allocation functions], page 168) instead of the given *buffer*. In that case,
     bufferbytes stores the buffer number (a nonnegative number). If *bufferbytes* is '`-1`',
     however, then *all* message buffers are touched. Note that when *firstbyte* is '`-1`', all
     parameters to `ncptl_touch_memory()` other than *bufferbytes* are ignored.

          `ncptl_touch_memory()` is intended to to be used to implement the `TOUCHES` state-
     ment (see Section 4.8.3 [Touching memory], page 134).

6.3.5 Time-related functions

An essential component of any benchmarking system is an accurate timer. CoNCePTuaL's
`ncptl_time()` function selects from a variety of timers at configuration time, first favoring
lower-overhead cycle-accurate timers, then higher-overhead cycle-accurate, and finally non-
cycle-accurate timers. `ncptl_init()` measures the actual timer overhead and resolution
and `ncptl_log_write_prologue()` writes this information to the log file. Furthermore,
the `validatetimer` program (see [Validating the CoNCePTuaL timer], page 9) can be used
to verify that the timer used by `ncptl_init()` truly does correspond to wall-clock time.

The CoNCePTuaL language provides a few time-related functions. These are supported
by the functions described below.

uint64_t ncptl_time (*void*)   [Function]
     Return the time in microseconds. The timer ticks even when the program
     is not currently scheduled. No assumptions can be made about the relation
     of the value returned to the time of day; `ncptl_time()` is intended to be

used strictly for computing elapsed time. The timer's resolution and accuracy are logged to the log file by `ncptl_log_write_prologue()` (more precisely, by the internal `log_write_prologue_timer()` function, which is called by `ncptl_log_write_prologue()`). Note that `ncptl_time()` always returns a 64-bit unsigned value, regardless of how `ncptl_int` is declared.

The CONCEPTUAL `configure` script (see Section 2.1 [configure], page 4) searches for a number of high-resolution timers and selects the best timer mechanism from among the ones available. The selection criteria is as follows:

1. If `./configure` was passed `--with-gettimeofday` (see Section 2.1 [configure], page 4) then `ncptl_time()` uses `gettimeofday()` as its timer mechanism.

2. If `./configure` was passed `--with-mpi-wtime` (see Section 2.1 [configure], page 4) then `ncptl_time()` uses `MPI_Wtime()` as its timer mechanism.

3. If `./configure` recognizes the CPU architecture, knows how to instruct the C compiler to insert inline assembly code, and can determine the number of clock cycles per second, then `ncptl_time()` reads the timer using inline assembly code. If the cycle counter is likely to wrap around during a moderately long benchmark (i.e., because the cycle counter is a 32-bit register), `ncptl_time()` augments the inline assembly code with calls to `gettimeofday()` in an attempt to produce accurate timings that don't suffer from clock wraparound.

4. If the Linux `get_cycles()` function is available and `./configure` can determine the number of clock cycles per second then `ncptl_time()` uses `get_cycles()` to measure execution time.

5. If the PAPI library is available, `ncptl_time()` becomes a call to PAPI's `PAPI_get_real_usec()` function.

6. If the System V `clock_gettime()` function is available and the `CLOCK_SGI_CYCLE` macro is defined, `ncptl_time()` invokes `clock_gettime()` with the `CLOCK_SGI_CYCLE` argument. If `CLOCK_SGI_CYCLE` is not defined but `CLOCK_REALTIME` is, then `ncptl_time()` invokes `clock_gettime()` with the `CLOCK_REALTIME` argument.

7. Intel's (now obsolete) supercomputers provide a `dclock()` function for reading the time. `ncptl_time()` makes use of `dclock()` if it's available.

8. Microsoft Windows provides functions for reading a high-resolution timer (`QueryPerformanceCounter()`) and for determining the number of ticks per second that the timer measures (`QueryPerformanceFrequency()`). If those functions are available, `ncptl_time()` uses them.

9. As a last resort, `ncptl_time()` uses `gettimeofday()` to measure execution time.

Furthermore, CONCEPTUAL makes use of the High-Precision Event Timers (HPET) device if and only if all of the following conditions hold at run time:

- neither `--with-gettimeofday`, `--with-mpi-wtime`, nor `--disable-hpet` was specified to `./configure`

- `/dev/hpet`—or an alternative device specified at configuration time with `--enable-hpet`—exists, is readable, and can be memory-mapped into user space

- the HPET device's main counter is a 64-bit value, not a 32-bit value

- the period of the HPET counter is within the ranges required by the HPET specification, i.e., between 0 and 100 nanoseconds

Failing any of those conditions, CONCEPTUAL falls back to the timer selected at configuration time. See the HPET specification for more information on HPET.

void ncptl_set_flag_after_usecs (*volatile int *flag*, *uint64_t*   [Function]
        delay)
ncptl_set_flag_after_usecs() uses the operating system's interval timer to asynchronously set a variable to '1' after a given number of microseconds. This function is intended to be used to support the 'FOR *time*' construct (see [Timed loops], page 129). Note that *delay* is a 64-bit unsigned value, regardless of how ncptl_int is declared.

ncptl_set_flag_after_usecs() is implemented in terms of the setitimer() function and issues a run-time error if the setitimer() function is not available.

void ncptl_udelay (*uint64_t delay*, *int spin0block1*)   [Function]
If *spin0block1* is '0', ncptl_udelay() spins for *delay* microseconds (i.e., using the CPU). If *spin0block1* is '1', ncptl_udelay() sleeps for *delay* microseconds (i.e., relinquishing the CPU). Note that *delay* is a 64-bit unsigned value, regardless of how ncptl_int is declared. ncptl_udelay() is intended to be used to support the CONCEPTUAL language's SLEEPS and COMPUTES statements (see Section 4.8.2 [Delaying execution], page 134).

When *spin0block1* is '0', ncptl_udelay() uses ncptl_time() to determine when *delay* microseconds have elapsed. Unless ncptl_time() is known to utilize an extremely low-overhead timer, ncptl_udelay() intersperses calls to ncptl_time() with writes to a dummy variable. When *spin0block1* is '0', ncptl_udelay() invokes nanosleep() to introduce delays. ncptl_udelay() issues a run-time error if the nanosleep() function is not available.

6.3.6 Log-file functions

Benchmarking has limited value without a proper record of the performance results. The CONCEPTUAL run-time library provides functions for writing data to log files. It takes care of much of the work needed to calculate statistics on data columns and to log a thorough experimental setup to every log file.

The library treats a log file as a collection of tables of data. Each table contains a number of rows, one per dynamic invocation of the LOGS statement (see Section 4.5.3 [Writing to a log file], page 120). Each row contains a number of columns, one per aggregate expression (see Section 4.2.3 [Aggregate expressions], page 100) expressed statically in a CONCEPTUAL program.[2] Log-file functions should be called only if the CONCEPTUAL source code accesses a log file (see Section 4.5.3 [Writing to a log file], page 120).

void ncptl_log_add_comment (*const char *key*, *const char *value*)   [Function]
ncptl_log_add_comment() makes it possible for a backend to add backend-specific ⟨*key*:*value*⟩ pairs to the set of prologue or epilogue comments that get written to a log file (see Section 3.5.1 [Log-file format], page 49). ncptl_log_add_comment() can be called repeatedly. All calls that precede ncptl_log_open() are included in the

---

[2] Writing A HISTOGRAM OF THE <*expr*> produces two columns, one for values and one for tallies.

log-file prologue. All calls that follow `ncptl_log_open()` are included in the log-file epilogue. Note that `ncptl_log_add_comment()` makes a copy of *key* and *value*, so these need not be heap-allocated.

NCPTL_LOG_FILE_STATE * ncptl_log_open (*char *template*,   [Function]
     ncptl_int **processor**)

Given a filename template containing a '`%d`' placeholder and a processor number (i.e., the process's physical rank in the computation), `ncptl_log_open()` creates and opens a log file named by the template with '`%d`' replaced by *processor*. For example, if *template* is `/home/me/myprog-%d.log` and *processor* is '3', the resulting filename will be `/home/me/myprog-3.log`. `ncptl_log_open()` must be called before any of the other `ncptl_log_`*something*`()` functions—except for `ncptl_log_add_comment()`, which should be called before `ncptl_log_open()`. `ncptl_log_open()` returns a pointer to an opaque `NCPTL_LOG_FILE_STATE` value; the backend will need to pass this pointer to nearly all of the other log-file functions described in this section.

There are three special cases for *template*. First, if *template* points to an empty string, all log-file output is sent to the null device (i.e., `/dev/null` on Unix and Unix-like operating systems). Second, if *template* is a single dash ('`-`'), all log-file output is sent to the standard output device. Third, if *template* is a single dollar sign ('`$`'), all log-file output is buffered in a library-internal string. The string can be retrieved using `ncptl_log_get_contents()`.

char * ncptl_log_generate_uuid (*void*)   [Function]

Return a random string of hexadecimal digits formatted as "*xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx*" (36 bytes plus a `NULL` byte) to pass to `ncptl_log_write_prologue()`. The caller should `ncptl_free()` the string when it is no longer needed (generally, as soon as `ncptl_log_write_prologue()` returns).

void ncptl_log_write_prologue (*NCPTL_LOG_FILE_STATE*   [Function]
     *logstate, *char *progname*, *char *uuid*, *char *backend_name*, *char*
     *backend_desc, *ncptl_int* **numtasks**, *NCPTL_CMDLINE *arglist*, *int*
     numargs, *char **sourcecode*)

`ncptl_log_write_prologue()` standardizes the prologue with which all log files begin. *progname* is the name of the program executable (`argv[0]` in C). *uuid* is a value returned by `ncptl_log_generate_uuid()`. Note that every process in a program must pass the same value of *uuid* to `ncptl_log_write_prologue()`. *backend_name* is the name of the backend in '`language_library`' format (e.g., '`java_rmi`'). *backend_desc* is a brief description of the backend (e.g., '`Java + RMI`'). *numtasks* is the total number of tasks in the program. *arglist* is the list of arguments passed to `ncptl_parse_command_line()` and *numargs* is the number of entries in that list. *sourcecode* is the complete CONCEPTUAL source code stored as a `NULL`-terminated list of `NULL`-terminated strings.

char * ncptl_log_lookup_string (*NCPTL_LOG_FILE_STATE*   [Function]
     *logstate, *char *key*)

`ncptl_log_write_prologue()` stores every ⟨*key*:*value*⟩ comment it writes into an in-memory database. `ncptl_log_lookup_string()` searches the comment database for a key and returns the corresponding value. The function returns the empty string

if the key is not found in the database. In either case, the caller should not deallocate the result. `ncptl_log_lookup_string()` is intended to be used to implement CON- CEPTUAL's `THE VALUE OF` construct (see Section 4.5.1 [Utilizing log-file comments], page 119).

void ncptl_log_write (*NCPTL_LOG_FILE_STATE *`logstate`, int*   [Function]
         *`logcolumn`, char *`description`, LOG_AGGREGATE `aggregate`, double*
         `aggregate_param`, double `value`)

Push value *value* onto column *logcolumn* of the current table. Gaps between columns are automatically elided. *description* is used as the column header for column *logcolumn*. *aggregate_param* is a parameter needed by certain aggregates. Currently, it used only to represent the target percentile for `NCPTL_FUNC_PERCENTILE`. Acceptable values for *aggregate* are defined in Section B.4 [Representing aggregate functions], page 216.

void ncptl_log_compute_aggregates (*NCPTL_LOG_FILE_STATE*   [Function]
         `*logstate`)

`ncptl_log_compute_aggregates()` implements the `COMPUTES AGGREGATES` construct described in [Computing aggregates], page 122. When `ncptl_log_compute_aggregates()` is invoked, the CONCEPTUAL run- time library uses the aggregate function specified by `ncptl_log_write()` to aggregate all of the data that accumulated in each column since the last invocation of `ncptl_log_compute_aggregates()`. Note that `ncptl_log_compute_aggregates()` is called implicitly by `ncptl_log_commit_data()`.

void ncptl_log_commit_data (*NCPTL_LOG_FILE_STATE*   [Function]
         `*logstate`)

The CONCEPTUAL run-time library keeps the current data table in memory and doesn't write anything to the log file until `ncptl_log_commit_data()` is called, at which point the run-time library writes all accumulated data to the log file and be- gins a new data table. Note that `ncptl_log_commit_data()` is called implicitly by `ncptl_log_close()`. Furthermore, a backend should call `ncptl_log_commit_data()` when beginning execution of a new statement in a CONCEPTUAL program. For instance, the `codegen_c_generic.py` backend invokes `ncptl_log_commit_data()` from `code_def_main_newstmt`.

void ncptl_log_write_epilogue (*NCPTL_LOG_FILE_STATE*   [Function]
         `*logstate`)

Write a stock set of comments as an epilogue to the log file.

const char * ncptl_log_get_contents   [Function]
         (*NCPTL_LOG_FILE_STATE *`logstate`*)

Return the current contents of the log file as a string. The caller must not use the string after the next run-time library call that modifies the log file, modify the string, or free the string. The caller should copy the string (e.g., with `ncptl_strdup()`) if any of those actions are necessary. If the the log file does not support random access (because it goes to the standard-output device or the null device), `NULL` is returned.

void ncptl_log_close (*NCPTL_LOG_FILE_STATE *logstate*)   [Function]
>   Close the log file. No ncptl_log_*something*() function should be called after
>   ncptl_log_close() is invoked.

6.3.7 Random-task functions

Randomness appears in various forms in the CONCEPTUAL language, such as when assigning a task to A RANDOM PROCESSOR (see Section 4.8.4 [Reordering task IDs], page 137) or when let-binding A RANDOM TASK or A RANDOM TASK OTHER THAN a given task ID to a variable (see Section 4.7.3 [Binding variables], page 130). The following functions are used to select tasks at random. CONCEPTUAL currently uses the Mersenne Twister as its random-number generator. Hence, given the same seed, a CONCEPTUAL program will see the same random-number sequence on every platform.

int ncptl_seed_random_task (*int seed, ncptl_int procID*)   [Function]
>   Initialize the random-number generator needed by ncptl_random_task().
>   If *seed* is zero, ncptl_seed_random_task() selects an arbitrary seed value.
>   ncptl_seed_random_task() returns the seed that was used. *procID* specifies the
>   (physical) processor ID of the calling task and is needed to seed the task-local
>   random-number generators used by some of the functions in Section 6.3.11
>   [Language-visible functions], page 178.

ncptl_int ncptl_random_task (*ncptl_int lower_bound, ncptl_int*   [Function]
>   *upper_bound, ncptl_int excluded*)
>   Return a randomly selected task number from *lower_bound* to *upper_bound* (both
>   inclusive). If *excluded* is nonnegative then that task number will never be selected,
>   even if it's within range.

6.3.8 Task-mapping functions

NCPTL_VIRT_PHYS_MAP * ncptl_allocate_task_map (*ncptl_int*   [Function]
>   *numtasks*)
>   Allocate and initialize an opaque NCPTL_VIRT_PHYS_MAP object and return a pointer
>   to it.

NCPTL_VIRT_PHYS_MAP * ncptl_point_to_task_map   [Function]
>   (*NCPTL_VIRT_PHYS_MAP *oldmap*)
>   Store a pointer to a task map, and mark the task map as "live".

NCPTL_VIRT_PHYS_MAP * ncptl_conditionally_copy_task_map   [Function]
>   (*NCPTL_VIRT_PHYS_MAP *oldmap*)
>   Replicate an existing task map if it is "live" (i.e., something points to it) or return
>   the input parameter if not.

ncptl_int ncptl_virtual_to_physical (*NCPTL_VIRT_PHYS_MAP*   [Function]
>   **procmap, ncptl_int virtID*)
>   Given a process map allocated by ncptl_allocate_task_map() and a (virtual) task
>   ID, return the corresponding (physical) process ID.

`ncptl_int ncptl_physical_to_virtual` (*NCPTL_VIRT_PHYS_MAP* [Function]
   *`*procmap`, *ncptl_int* `physID`)
  Given a process map allocated by `ncptl_allocate_task_map()` and a (physical)
  process ID, return the corresponding (virtual) task ID.

`ncptl_int ncptl_assign_processor` (*ncptl_int* `virtID`, *ncptl_int*   [Function]
   `physID`, *NCPTL_VIRT_PHYS_MAP* `*procmap`, *ncptl_int* `physrank`)
  Assign a (physical) processor ID, *physID*, to a (virtual) task ID, *virtID* given a virtual-
  to-physical mapping table, *procmap*. Return a new task ID for the caller's proces-
  sor given its processor ID, *physrank*. `ncptl_assign_processor()` is intended to
  implement the IS ASSIGNED TO construct (see Section 4.8.4 [Reordering task IDs],
  page 137).

6.3.9 Queue functions

Because queues are a widely applicable construct, the run-time library provides support for
queues of arbitrary datatypes. In the current implementation, these can more precisely be
termed "dynamically growing lists" than "queues". However, they may be extended in a
future version of the library to support more queue-like functionality.

`NCPTL_QUEUE * ncptl_queue_init` (*ncptl_int* `eltbytes`)   [Function]
  `ncptl_queue_init()` creates and initializes a dynamically growing queue in which
  each element occupies *eltbytes* bytes of memory.

`void * ncptl_queue_allocate` (*NCPTL_QUEUE* `*queue`)   [Function]
  Allocate a new data element at the end of queue *queue*. The queue passed
  to `ncptl_queue_allocate()` must be one returned by `ncptl_queue_init()`.
  `ncptl_queue_allocate()` returns a pointer to the data element allocated.

`void * ncptl_queue_push` (*NCPTL_QUEUE* `*queue`, *void* `*element`) [Function]
  Push (via a memory copy) the element pointed to by *element* onto the
  end of queue *queue* and return a pointer to the copy in the queue. The
  queue passed to `ncptl_queue_allocate()` must be one returned by
  `ncptl_queue_init()`. (`ncptl_queue_push()` is actually implemented in terms of
  `ncptl_queue_allocate()`.)

`void ncptl_queue_push_all` (*NCPTL_QUEUE* `*target_queue`,   [Function]
   NCPTL_QUEUE `*source_queue`)
  Push (via a memory copy) all of the elements in *source_queue* onto the end of *tar-
  get_queue*. A fatal error will occur if the two queues were not initialized with the
  same *eltbytes* argument to `ncptl_queue_init()`. `ncptl_queue_push_all()` does
  not modify *source_queue*.

`void * ncptl_queue_pop` (*NCPTL_QUEUE* `*queue`)   [Function]
  Pop a pointer to the element at the head of queue *queue*. If *queue* is empty, return
  NULL. The pointer returned by `ncptl_queue_pop()` is guaranteed to be valid until
  the next invocation of `ncptl_queue_empty()`.

void * ncptl_queue_pop_tail (*NCPTL_QUEUE* `*queue`)   [Function]
> Pop a pointer to the element at the tail of queue *queue*. If *queue* is empty, return
> `NULL`. In essence, this lets the caller treat the queue as a stack. The pointer returned
> by `ncptl_queue_pop_tail()` is guaranteed to be valid until the next invocation of
> `ncptl_queue_empty()`, `ncptl_queue_allocate()` or `ncptl_queue_push()`.

void * ncptl_queue_contents (*NCPTL_QUEUE* `*queue`, *int*   [Function]
> `copyelts`)
> Return queue *queue* as an array of elements. If `ncptl_queue_contents()` is passed
> '1' for *copyelts*, a new array is allocated using `ncptl_malloc()`; the queue's internal
> array is copied to the newly allocated array; and, this new array is returned to the
> caller. It is the caller's responsibility to pass the result to `ncptl_free()` when the
> array is no longer needed. If `ncptl_queue_contents()` is passed '0' for *copyelts*, a
> pointer to the queue's internal array is returned without first copying it. This pointer
> should not be passed to `ncptl_free()` as it is still needed by *queue*.

ncptl_int ncptl_queue_length (*NCPTL_QUEUE* `*queue`)   [Function]
> Return the number of elements in queue *queue*.

void ncptl_queue_empty (*NCPTL_QUEUE* `*queue`)   [Function]
> Empty a queue, freeing the memory that had been allocated for its elements. Queue
> contents returned by `ncptl_queue_contents()` with *copyelts* set to '0' are also in-
> validated. The queue itself can continue to be used and should be deallocated with
> `ncptl_free()` (see Section 6.3.3 [Memory-allocation functions], page 168) when no
> longer needed.

6.3.10 Unordered-set functions

Because unordered collections of data are a widely applicable construct, the run-time library
provides support for sets. A set contains zero or more keys, each of which must be unique
within the set. Furthermore, each key can be associated with a data value. Sets are currently
implemented in the CONCEPTUAL run-time library as hash tables.

NCPTL_SET * ncptl_set_init (*ncptl_int* `numelts`, *ncptl_int*   [Function]
> `keybytes`, *ncptl_int* `valuebytes`)
> Allocate and initialize a set object and return a pointer to it. Each element in the set
> maps a *keybytes*-byte key to a *valuebytes*-byte value. The *numelts* parameter is an
> estimate of the maximum number of elements in the set. `ncptl_set_init()` returns
> a pointer to the set.

void * ncptl_set_find (*NCPTL_SET* `*set`, *void* `*key`)   [Function]
> Given a set and a pointer to a key, return a pointer to the associated value or `NULL`
> if the key is not found in the set.

void ncptl_set_insert (*NCPTL_SET* `*set`, *void* `*key`, *void* `*value`) [Function]
> Insert a key into a set and associate a value with it. `ncptl_set_insert()` copies
> both *key* and *value* so stack-allocated keys and values are acceptable inputs. The
> run-time library aborts with an error message if the key is already in the set.

void **ncptl_set_walk** (*NCPTL_SET* ***set**, *void* (***userfunc**)(*void **,   [Function]
        void *))

> Execute function *userfunc* for every ⟨*key*:*value*⟩ pair in a set. *userfunc* must take two
> **void *** values as input: a pointer to a key and a pointer to a value. The order in
> which keys and values are passed to the function is unspecified.

void **ncptl_set_remove** (*NCPTL_SET* ***set**, *void* ***key**)   [Function]

> Remove a key and its associated value from a set. The run-time library aborts with
> an error message if the key is not found in the set.

void **ncptl_set_empty** (*NCPTL_SET* ***set**)   [Function]

> Empty a set, freeing the memory that had been allocated for its contents. The set
> itself can continue to be used and should be deallocated with **ncptl_free()** (see
> Section 6.3.3 [Memory-allocation functions], page 168) when no longer needed.

ncptl_int **ncptl_set_length** (*NCPTL_SET* ***set**)   [Function]

> Return the number of ⟨*key*:*value*⟩ pairs in set *set*.

6.3.11 Language-visible functions

The CONCEPTUAL language contains a number of built-in functions that perform various
operations on floating-point numbers (used when writing to a log file or the standard output
device) and integers (used at all other times) and that determine the IDs of neighboring
tasks on a variety of topologies. Each function occurs in two forms: **ncptl_func_*function***,
which maps **ncptl_int**s to **ncptl_int**s, and **ncptl_dfunc_*function***, which maps **doubles**
to **doubles**. See Section 4.2.2 [Built-in functions], page 88, for additional details about each
function's semantics.

   Although some of the functions described in this section are fairly simple, including them
in the run-time library ensures that each function returns the same value across different
backends and across different platforms.

Integer functions

ncptl_int **ncptl_func_sqrt** (*ncptl_int* **num**)   [Function]
double **ncptl_dfunc_sqrt** (*double* **num**)   [Function]

> **ncptl_func_sqrt()** returns the unique integer x such that $x^2 \leq num \;\wedge\; (x+1)^2 >$
> *num*. **ncptl_dfunc_sqrt()** returns $\sqrt{num}$ in double-precision arithmetic.

ncptl_int **ncptl_func_cbrt** (*ncptl_int* **num**)   [Function]
double **ncptl_dfunc_cbrt** (*double* **num**)   [Function]

> **ncptl_func_cbrt()** returns the unique integer x such that $x^3 \leq num \;\wedge\; (x+1)^3 >$
> *num*. **ncptl_dfunc_cbrt()** returns $\sqrt[3]{num}$ in double-precision arithmetic.

ncptl_int **ncptl_func_root** (*ncptl_int* **root**, *ncptl_int* **num**)   [Function]
double **ncptl_dfunc_root** (*double* **root**, *double* **num**)   [Function]

> Return the *root*th root of a number *num*. More precisely, **ncptl_dfunc_root()**
> returns $\sqrt[root]{num}$ while **ncptl_func_root()** returns the largest-in-magnitude integer
> x with the same sign as *num* such that $|x^{root}| \leq |num|$. Currently, *num* must be
> nonnegative but this may change in a future release of CONCEPTUAL.

`ncptl_int ncptl_func_bits` (*ncptl_int* `num`)   [Function]
`double ncptl_dfunc_bits` (*double* `num`)   [Function]
>   Return the minimum number of bits needed to represent a given integer. (*num* is
>   rounded up to the nearest integer in the case of `ncptl_dfunc_bits()`.)

`ncptl_int ncptl_func_shift_left` (*ncptl_int* `num`, *ncptl_int* `bits`)   [Function]
`double ncptl_dfunc_shift_left` (*double* `num`, *double* `bits`)   [Function]
>   Shift a number to the left by *bits* bits. This corresponds to multiplying *num* by
> 2^{bits}. In the case of `ncptl_dfunc_shift_left()`, *num* and *bits* are first converted
>   to `ncptl_int` values. There are no corresponding `ncptl_func_shift_right()` or
>   `ncptl_dfunc_shift_right()` functions because shifting right by x is defined to be
>   equivalent to shifting left by $-x$.

`ncptl_int ncptl_func_log10` (*ncptl_int* `num`)   [Function]
`double ncptl_dfunc_log10` (*double* `num`)   [Function]
>   Return $\log_{10}(num)$ In the case of `ncptl_func_log10()`, this value is rounded down
>   to the nearest integer.

`ncptl_int ncptl_func_factor10` (*ncptl_int* `num`)   [Function]
`double ncptl_dfunc_factor10` (*double* `num`)   [Function]
>   Return *num* rounded down to the nearest single-digit factor of a power of 10.

`ncptl_int ncptl_func_abs` (*ncptl_int* `num`)   [Function]
`double ncptl_dfunc_abs` (*double* `num`)   [Function]
>   Return $|num|$. In the case of `ncptl_func_log10()`, this value is rounded down to the
>   nearest integer.

`ncptl_int ncptl_func_power` (*ncptl_int* `base`, *ncptl_int* `exponent`)   [Function]
`double ncptl_dfunc_power` (*double* `base`, *double* `exponent`)   [Function]
>   Return *base* raised to the power of *exponent*.

`ncptl_int ncptl_func_modulo` (*ncptl_int* `numerator`, *ncptl_int*   [Function]
>   `denominator`)
`double ncptl_dfunc_modulo` (*double* `numerator`, *double*   [Function]
>   `denominator`)
>   Return the remainder of dividing *numerator* by *denominator*. The result is guaranteed
>   to be a nonnegative integer. `ncptl_dfunc_modulo()` rounds each of *numerator* and
>   *denominator* down to the nearest integer before dividing and taking the remainder.

`ncptl_int ncptl_func_min` (*ncptl_int* `count`, . . .)   [Function]
`double ncptl_dfunc_min` (*double* `count`, . . .)   [Function]
>   Return the minimum of a list of numbers. The first argument specifies the number
>   of remaining arguments and must be a positive integer.

`ncptl_int ncptl_func_max` (*ncptl_int* `count`, . . .)   [Function]
`double ncptl_dfunc_max` (*double* `count`, . . .)   [Function]
>   Return the maximum of a list of numbers. The first argument specifies the number
>   of remaining arguments and must be a positive integer.

Floating-point functions

ncptl_int ncptl_func_floor (*ncptl_int* **num**)   [Function]
double ncptl_dfunc_floor (*double* **num**)   [Function]
    Return $\lfloor num \rfloor$. (This is the identity function in the case of `ncptl_func_floor()`.)

ncptl_int ncptl_func_ceiling (*ncptl_int* **num**)   [Function]
double ncptl_dfunc_ceiling (*double* **num**)   [Function]
    Return $\lceil num \rceil$. (This is the identity function in the case of `ncptl_func_ceiling()`.)

ncptl_int ncptl_func_round (*ncptl_int* **num**)   [Function]
double ncptl_dfunc_round (*double* **num**)   [Function]
    Return *num* rounded to the nearest integer. (This is the identity function in the case
    of `ncptl_func_round()`.)

Topology functions

In the following functions, the 'dfunc' versions merely cast their arguments to `ncptl_int`s
and call the corresponding 'func' versions.

ncptl_int ncptl_func_tree_parent (*ncptl_int* **task**, *ncptl_int*   [Function]
      arity)
double ncptl_dfunc_tree_parent (*double* **task**, *double* **arity**)   [Function]
    Return task *task*'s parent in an *arity*-ary tree.

ncptl_int ncptl_func_tree_child (*ncptl_int* **task**, *ncptl_int* **child**,   [Function]
      ncptl_int **arity**)
double ncptl_dfunc_tree_child (*double* **task**, *double* **child**, *double*   [Function]
      arity)
    Return child *child* of task *task* in an *arity*-ary tree.

ncptl_int ncptl_func_mesh_coord (*ncptl_int* **vartask**, *ncptl_int*   [Function]
      coord, *ncptl_int* **width**, *ncptl_int* **height**, *ncptl_int* **depth**)
double ncptl_dfunc_mesh_coord (*double* **vartask**, *double* **coord**,   [Function]
      double **width**, *double* **height**, *double* **depth**)
    Return task *task*'s x coordinate ($coord = 0$), y coordinate ($coord = 1$), or z coordinate
    ($coord = 2$) on a *width* $\times$ *height* $\times$ *depth* mesh (or torus).

ncptl_int ncptl_func_mesh_neighbor (*ncptl_int* **width**, *ncptl_int*   [Function]
      height, *ncptl_int* **depth**, *ncptl_int* **xtorus**, *ncptl_int* **ytorus**, *ncptl_int*
      ztorus, *ncptl_int* **task**, *ncptl_int* **xdelta**, *ncptl_int* **ydelta**, *ncptl_int*
      zdelta)
double ncptl_dfunc_mesh_neighbor (*double* **width**, *double* **height**,   [Function]
      double **depth**, *double* **xtorus**, *double* **ytorus**, *double* **ztorus**, *double* **task**,
      double **xdelta**, *double* **ydelta**, *double* **zdelta**)
    Return one of task *task*'s neighbors—not necessarily an immediate neighbor—on a
    3-D mesh or torus. For the following explanation, assume that task *task* lies at
    coordinates (x, y, z) on a *width* $\times$ *height* $\times$ *depth* mesh or torus. In the mesh case
    (*xtorus* = *ytorus* = *ztorus* = 0), the value returned is the task ID corresponding
    to coordinates $(x + xdelta, y + ydelta, z + zdelta)$. In the full-torus case (*xtorus* =
    ytorus = *ztorus* = 1), the value returned is the task ID corresponding to coordinates
    $((x + xdelta)$ mod *width*, $(y + ydelta)$ mod *height*, $(z + zdelta)$ mod *depth*).

Note that there are no 1-D or 2-D grid functions. Instead, the appropriate 3-D function should be used with *depth* and—in the 1-D case—*height* set to '1'.

`ncptl_int ncptl_func_knomial_parent` (*ncptl_int* `task`, *ncptl_int*   [Function]
        `arity`, *ncptl_int* `numtasks`)
`double ncptl_dfunc_knomial_parent` (*double* `task`, *double* `arity`,   [Function]
        double `numtasks`)
> Return task *task*'s parent in an *arity*-nomial tree of *numtasks* tasks.

`ncptl_int ncptl_func_knomial_child` (*ncptl_int* `task`, *ncptl_int*   [Function]
        `child`, *ncptl_int* `arity`, *ncptl_int* `numtasks`, *ncptl_int* `count_only`)
`double ncptl_dfunc_knomial_child` (*double* `task`, *double* `child`,   [Function]
        double `arity`, *double* `numtasks`, *double* `count_only`)
> If *count_only* is '0', return task *task*'s *child*th child in an *arity*-nomial tree of *numtasks* tasks. If *count_only* is '1', return the number of children task *task* has in an *arity*-nomial tree of *numtasks* tasks.

Random-number functions

`ncptl_int ncptl_func_random_uniform` (*ncptl_int* `lower_bound`,   [Function]
        ncptl_int `upper_bound`)
`double ncptl_dfunc_random_uniform` (*double* `lower_bound`, *double*   [Function]
        `upper_bound`)
> Return a number in the interval [*lower_bound*, *upper_bound*) selected at random with a uniform distribution.

`ncptl_int ncptl_func_random_gaussian` (*ncptl_int* `mean`, *ncptl_int*   [Function]
        `stddev`)
`double ncptl_dfunc_random_gaussian` (*double* `mean`, *double* `stddev`)   [Function]
> Return a number selected at random from a Gaussian distribution with mean *mean* and standard deviation *stddev*.

`ncptl_int ncptl_func_random_poisson` (*ncptl_int* `mean`)   [Function]
`double ncptl_dfunc_random_poisson` (*double* `mean`)   [Function]
> Return an integer selected at random from a Poisson distribution with mean *mean* and standard deviation $\sqrt{mean}$.

`ncptl_int ncptl_func_random_pareto` (*ncptl_int* `shape`, *ncptl_int*   [Function]
        `lower_bound`, *ncptl_int* `upper_bound`)
`double ncptl_dfunc_random_pareto` (*double* `shape`, *double*   [Function]
        `lower_bound`, *double* `upper_bound`)
> If *lower_bound*=*upper_bound*, return a value selected at random from a Pareto distribution with shape *shape* and scale *lower_bound*. If *lower_bound*≠*upper_bound*, return a value selected at random from a bounded Pareto distribution with shape *shape* and bounded by the range [*lower_bound*, *upper_bound*].

6.3.12 Finalization functions

The following function should be called towards the end of the generated code's execution.

void ncptl_finalize (*void*)   [Function]
> Shut down the CONCEPTUAL run-time library. No run-time library functions should be invoked after `ncptl_finalize()`.

void ncptl_fatal (*char *format*, . . .)   [Function]
> Output an error message and abort the program. `ncptl_fatal()` takes the same types of arguments as C's `printf()` routine.

7 Tips and Tricks

The following sections present some ways to make better use of cONCEPTuaL in terms of producing simpler, more efficient programs or being able to run on complex computer systems.

7.1 Using out-of-bound task IDs to simplify code

See Section 4.3 [Task descriptions], page 106, mentions a language feature that can substantially simplify cONCEPTuaL programs: Operations involving out-of-bound task IDs are silently ignored. The beauty of this feature is that it reduces the need for special cases at network boundaries. Consider, for example, a simple pipeline pattern in which each task in turn sends a message to the subsequent task:

```
ALL TASKS t SEND A 64 DOUBLEWORD MESSAGE TO TASK t+1.
```

Because implicit receives are posted before the corresponding sends (see Section 4.4.2 [Sending], page 113), all tasks except task 0 start by posting a blocking receive. (No task is sending to task 0.) Task 0 is therefore free to send a message to task 1. Receipt of that message unblocks task 1, who then sends a message to task 2, thereby unblocking task 3, and so forth. Without needing an explicit special case in the program, task 'num_tasks-1' receives a message from task 'num_tasks-2' but does not attempt to send a message to nonexistent task 'num_tasks', thanks to the rule that communication with nonexistent tasks turns into a no-op (i.e., is elided from the program).

As a more complex variation of the same program, consider a wavefront communication pattern that progresses from the upper-left corner of a mesh to the lower-right corner. Such a pattern can be expressed in just four lines of cONCEPTuaL (receive left, receive up, send right, send down) by relying on the property that communication with a nonexistent task is simply not executed:

```
TASK MESH_NEIGHBOR(src, xsize, +1, ysize,  0) RECEIVES A
    64 DOUBLEWORD MESSAGE FROM ALL TASKS src THEN
TASK MESH_NEIGHBOR(src, xsize,  0, ysize, +1) RECEIVES A
    64 DOUBLEWORD MESSAGE FROM ALL TASKS src THEN
ALL TASKS src SEND A 64 DOUBLEWORD MESSAGE TO
    UNSUSPECTING TASK MESH_NEIGHBOR(src, xsize, +1, ysize,  0) THEN
ALL TASKS src SEND A 64 DOUBLEWORD MESSAGE TO
    UNSUSPECTING TASK MESH_NEIGHBOR(src, xsize,  0, ysize, +1).
```

To understand the preceding program recall that MESH_NEIGHBOR returns '-1' for nonexistent neighbors. Because '-1' is outside of the range [0, num_tasks) communication with a nonexistent neighbor is ignored. To help the reader understand the preceding program, we present a trace of the events it posts as it runs with a 2×2 arrangement of tasks:

```
[TRACE] phys: 0 | virt: 0 | action: SEND | event: 1 / 2 | lines: 3 - 3
[TRACE] phys: 1 | virt: 1 | action: RECV | event: 1 / 2 | lines: 1 - 1
[TRACE] phys: 2 | virt: 2 | action: RECV | event: 1 / 2 | lines: 2 - 2
[TRACE] phys: 3 | virt: 3 | action: RECV | event: 1 / 2 | lines: 1 - 1

[TRACE] phys: 0 | virt: 0 | action: SEND | event: 2 / 2 | lines: 4 - 4
[TRACE] phys: 1 | virt: 1 | action: SEND | event: 2 / 2 | lines: 4 - 4
```

```
    [TRACE] phys: 2 | virt: 2 | action: SEND | event: 2 / 2 | lines: 3 - 3

    [TRACE] phys: 3 | virt: 3 | action: RECV | event: 2 / 2 | lines: 2 - 2
```

The `c_trace` backend (see Section 3.3.4 [The c_trace backend], page 24) was used to produce that trace. To increase clarity, we manually added blank lines to group concurrent events (i.e., there is no significance to the order of the TRACE lines within each group). The important thing to notice is that there are exactly four receives and exactly four sends:

- Although all tasks are instructed to receive a message from the left, only tasks 1 and 3 actually do so;
- although all tasks are instructed to receive a message from above, only tasks 2 and 3 actually do so;
- although all tasks are instructed to send a message to the right, only tasks 0 and 2 actually do so; and,
- although all tasks are instructed to send a message downwards, only tasks 0 and 1 actually do so.

Because communication with nonexistent tasks is elided at program initialization time there is no run-time cost for such operations—as evidenced by the `c_trace` output presented above. Furthermore, there is no reliance on the backend to drop messages from nonexistent senders or to nonexistent receivers; it is perfectly safe to utilize no-op'ed communcation in any CONCEPTUAL program and when using any backend.

7.2 Proper use of conditionals

CONCEPTUAL supports two forms of conditional execution: conditional expressions (see Section 4.2.1 [Arithmetic expressions], page 86) and conditional statements (see Section 4.7.4 [Conditional execution], page 132). From the perspective of code readability and "thinking in CONCEPTUAL" it is generally preferable to use restricted identifiers (see Section 4.3.1 [Restricted identifiers], page 106) to select groups of tasks rather than a loop with a conditional as would be typical in other programming languages. For example, consider the following code in which certain even-numbered tasks each send a message to the right:

```
FOR EACH evtask IN {0, ..., num_tasks-1}
  IF evtask IS EVEN /\ evtask MOD 3 <> 2 THEN
    TASK evtask SENDS A 64 BYTE MESSAGE TO TASK evtask+1
```

While the preceding control flow is representative of that in other programming languages, CONCEPTUAL can express the same communication pattern without needing either a loop or an explicit conditional statement:

```
TASK evtask SUCH THAT evtask IS EVEN /\ evtask MOD 3 <> 2 SENDS A 64
BYTE MESSAGE TO TASK evtask+1
```

One situation in which conditional statements do not have a convenient analogue is when a program selects among multiple disparate subprograms based on a command-line parameter:

```
func IS "Operation to perform (1=op1; 2=op2; 3=op3)" AND COMES FROM
"--function" OR "-f" WITH DEFAULT 1.

ASSERT THAT "the function must be 1, 2, or 3" WITH func>=1 /\ func<=3.

IF func = 1 THEN {
  Perform operation op1.
}
OTHERWISE IF func = 2 THEN {
  Perform operation op2.
}
OTHERWISE IF func = 3 THEN {
  Perform operation op3.
}
```

7.3 Memory efficiency

As described in Section 6.2.3 [Generated code], page 162, the `c_generic` backend (and therefore all derived backends) generates programs that run by executing a sequence of events in an event list. While this form of program execution makes it possible to hoist a significant amount of computation out of the timing loop, it does imply that a program's memory requirements are proportional to the number of statements that the program executes.

coNCePTuaL's memory usage can be reduced by taking advantage of repeat counts within statements that support such a construct. The language's <send_stmt> (see Section 4.4.2 [Sending], page 113), <receive_stmt> (see Section 4.4.3 [Receiving], page 116), and <touch_stmt> (see Section 4.8.3 [Touching memory], page 134) are all examples of statements that accept repeat counts. For other statements and for groups of statements that repeat, the FOR...REPETITIONS statement produces a single EV_REPEAT event followed by a single instance of the events in the loop body. This technique is valid because coNCePTuaL knows *a priori* that every iteration is identical to every other iteration. In contrast, the more general FOR EACH statement can induce different behavior each iteration based on the value of the loop variable so programs must conservatively instantiate the events in the loop body for every iteration. Consider the following examples:

Least efficient:
> 'FOR EACH i IN {1, ..., 1000} TASK 0 TOUCHES A 1 WORD MEMORY REGION'
> (1000 EV_TOUCH events on task 0)

More efficient:
> 'FOR 1000 REPETITIONS TASK 0 TOUCHES A 1 WORD MEMORY REGION'
> (an EV_REPEAT event and an EV_TOUCH event on task 0)

Most efficient:
> 'TASK 0 TOUCHES A 1 WORD MEMORY REGION 1000 TIMES'
> (one EV_TOUCH event on task 0)

Least efficient:

> 'FOR EACH i IN {1, ..., 1000} TASK 0 SENDS A 32 KILOBYTE MESSAGE TO
> TASK 1'
> (1000 `EV_SEND` events on task 0 and 1000 `EV_RECV` events on task 1)

More efficient:

> 'FOR 1000 REPETITIONS TASK 0 SENDS A 32 KILOBYTE MESSAGE TO TASK 1'
> (an `EV_REPEAT` event and an `EV_SEND` event on task 0 plus an `EV_REPEAT` event
> and an `EV_RECV` event on task 1)

Most efficient:

> 'TASK 0 SENDS 1000 32 KILOBYTE MESSAGES TO TASK 1'
> (currently the same as the above although a future release of CoNCePTuaL
> may reduce this to a single `EV_SEND` event on task 0 and a single `EV_RECV` event
> on task 1)

7.4 Cross-compilation

Some experimental architectures lack native compilers and therefore require programs to be compiled on an architecturally distinct front-end machine. CoNCePTuaL can be configured to support such an arrangement. As an example, we follow how CoNCePTuaL was recently configured to run on a Cray/SNL Red Storm prototype. The Red Storm prototype requires programs to be cross-compiled from a compile server then launched from an execution server to run on the compute nodes. The CPUs in all three node types use the x86-64 architecture but the compile server and execution server run Linux while the compute nodes run the lightweight Catamount kernel. Cross-compilation is necessary to prevent the *./configure* script that runs on the compile server from attempting to execute small programs to test run-time features—these would be guaranteed to fail.

   ./configure (see Section 2.1 [configure], page 4) enters cross-compilation mode if the build system—specified with *--build*—has a different CPU, comes from a different vendor, or runs a different operating system from the execution system—specified with *--host*. The `config.guess` script outputs the build-system type in the form *CPU-vendor-operating system*:

```
% ./config.guess
x86_64-unknown-linux-gnu
```

(That is, the CPU architecture is '`x86_64`'; the computer vendor is '`unknown`'; and, the operating system is '`linux-gnu`'.)

   As a special case of cross-compilation, if the build system and execution system utilize the same CPU and operating system[1] (but presumably differ in terms of other feature), *./configure* provides a *--with-cross-compilation* option to force a cross-compile. The alternative is to modify one of the build-system name components; best is to modify the vendor component as that's used solely for informational purposes. For our Red Storm

---

[1] As long as the host and build operating systems have moderately similar interfaces (e.g., both are Unix-like) and the same CPU architecture (e.g., both are x86-64 variants), *--with-cross-compilation* should be applicable.

configuration (which predated the `--with-cross-compilation` option) we renamed *vendor* from '`unknown`' to '`cray`'.

Because `configure` assumes it cannot execute small test programs on the execution system, it is unable to determine valid memory-buffer alignments. Consequently, the `--with-alignment` option must also be passed to `configure` to specify explicitly the minimum number of bytes at which data must be aligned. (If not specified, the minimum alignment defaults to '8'.) We know *a priori* that the x86-64 architecture can support byte-aligned data. Hence, we specify `--with-alignment=1`.

A command line like the following was used to configure CONCEPTUAL for the Red Storm prototype:

```
% ./configure --build=x86_64-unknown-linux-gnu --host=x86_64-cray-linux-gnu
    --with-alignment=1 CC=mpicc
```

Once CONCEPTUAL is configured to cross-compile, there is nothing special about performing the compilation itself. The `make` command runs unmodified from its description in Section 2.2 [make], page 7.

Running `make check` can be tricky because it involves both compilation and execution. As stated previously, these cannot be performed on the same servers in the Red Storm prototype. The solution is first to run `make check` on the compile server. This compiles all of the CONCEPTUAL regression tests—and unsuccessfully attempts to run them. Then, when `make check` is run on the execution server it does not need to compile any of the tests (which it can't do successfully, anyway) but can simply run each of them.

One catch is that the Red Storm execution server cannot directly run compute-node programs. Rather, it needs to spawn a job launcher (`yod`) for each test program. The mechanism for doing this is the `TESTS_ENVIRONMENT` variable, which `make check` prepends to every command it executes. We were therefore able to regression-test the CONCEPTUAL run-time library on the Red Storm prototype with the following command:

```
% make TESTS_ENVIRONMENT="yod -list 4" check
```

If the system for which you're cross-compiling is unable to run Python or unable to build the interpreter-based backends, you might consider building CONCEPTUAL twice—cross-compiled for the target system and compiled regularly for the front-end system. The `--prefix` option to `configure` (see Section 2.1 [configure], page 4) specifies the top-level directory for the CONCEPTUAL installation.

7.5 Implicit dynamic-library search paths

By default, CONCEPTUAL installs both a static and dynamic library with the dynamic library taking precedence (at least on most Unix and Unix-like operating systems). While dynamic libraries offer a number of benefits—such as not requiring applications to be relinked every time a library is upgraded—one inconvenience is the need to set the `LD_LIBRARY_PATH` environment variable to point to the CONCEPTUAL library directory if the CONCEPTUAL run-time library is not installed in a standard location (e.g., `/usr/lib`).

An alternative to setting *LD_LIBRARY_PATH* is to specify the target library directory at configuration time via the *LDFLAGS* variable and let the linker embed that directory into the CONCEPTUAL run-time library's dynamic search path. Many linkers accept a *-rpath* option for exactly that purpose. With the GNU C compiler and linker the appropriate option is *LDFLAGS="-Wl,-rpath,directory"* (e.g., 'LDFLAGS="-Wl,-rpath,/usr/local/ncptl/lib"'). Other compilers may have analogous mechanisms for passing flags directly to the linker.

7.6 Running without installing

You can compile CONCEPTUAL programs without having first to do a *make install* by using the same mechanism as was discussed in Section 7.5 [Implicit dynamic-library search paths], page 187.

First, Libtool builds the CONCEPTUAL run-time library in the .libs subdirectory. Hence, you should add 'LDFLAGS="-L`pwd`/.libs -Wl,-rpath,`pwd`/.libs"' to the *./configure* command line (see Section 2.1 [configure], page 4) to point both the static and dynamic linkers to the CONCEPTUAL build directory.

Second, CONCEPTUAL-generated C code specifies '#include <ncptl/ncptl.h>'. To ensure that the generated code can find ncptl.h you should add 'CPPFLAGS="-I`pwd`"' to the *./configure* command line to point the C compiler to the CONCEPTUAL build directory. Then, create a symbolic link from the CONCEPTUAL build directory to an ncptl directory by running *ln -s `pwd` ncptl* from the build directory.

After running *./configure* and *make*, you can compile CONCEPTUAL programs by invoking the compiler as *./ncptl.py* instead of the usual *ncptl*.

7.7 Reporting configuration information

On platforms that support it, the CONCEPTUAL run-time library is also an executable program that can be run from the command line. Executing the library outputs to the standard output device a complete log-file prologue and epilogue but no data.[2] This feature makes it quite convenient to determine all of the configuration options, compiler features, etc. that were used to build the run-time library.

---

[2] This is the same information produced by *make empty.log* (see Section 2.2 [make], page 7).

8 Troubleshooting

In any complex system, things are bound to go wrong. The following sections present solutions to various problems that have been encountered when building CoNCePTuaL and running CoNCePTuaL programs.

8.1 Problems with configure

For typical CoNCePTuaL usage, the most important function of the `configure` script is to prepare the system to build the CoNCePTuaL run-time library. CoNCePTuaL's functionality is severely restricted without that library. This section provides information to help ensure that configuration succeeds.

8.1.1 Interpreting configure warnings

The `configure` script performs a large number of tests to ensure that CoNCePTuaL will compile properly and function as expected. In particular, any missing or improperly functioning feature upon which the C run-time library relies causes *./configure* to issue a '`not building the C run-time library`' warning. Without its run-time library, CoNCePTuaL's functionality is severely limited so it's worth every effort to get *./configure* to build that.

Like all Autoconf scripts, `configure` logs detailed information to a `config.log` file. As a general diagnostic technique one should search for puzzling output in `config.log` and examine the surrounding context. For instance, on one particular system, *./configure* output '`no`' following '`checking if we can run a trivial program linked with "-lrt -lm -lpopt "`' and then refused to build the run-time library. The following relevant lines appeared in `config.log`:

```
configure:12845: checking if we can run a trivial program linked with "-lrt
-lm -lpopt "
configure:12862: /usr/local/bin/gcc  -o conftest -g -O2    conftest.c -lrt
-lm -lpopt  >&5
configure:12865: $? = 0
configure:12867: ./conftest
ld.so.1: ./conftest: fatal: libpopt.so.0: open failed: No such file or
directory
./configure: line 1:  5264 Killed                       ./conftest$ac_exeext
configure:12870: $? = 137
configure: program exited with status 137
configure: failed program was:
#line 12851 "configure"
#include "confdefs.h"

int
main (int argc, char *argv[])
{
  return 0;
}
configure:12879: result: no
```

Note the error message from `ld.so.1` about `libpopt.so.0` not being found. Further investigation revealed that although `/usr/local/bin/gcc` knew to look in `/usr/local/lib/` for shared libraries, that directory was not in the search path utilized by `ld.so.1`. Consequently, it couldn't find `/usr/local/lib/libpopt.so.0`. The solution in this case was to add `/usr/local/lib/` to the *LD_LIBRARY_PATH* environment variable before running *./configure*.

In general, `config.log` should be the first place to look when trying to interpret warnings issued by *./configure*. Furthermore, note that certain command-line options to *./configure* (see Section 2.1 [configure], page 4) may help bypass problematic operations that the script stumbles over.

8.1.2 'PRId64 is not a valid printf conversion specifier'

The following configuration warning was encountered on various BSD Unix systems (OpenBSD 3.4 and FreeBSD 4.10-BETA, both IA-32) and inhibited the building of the cONCePTuaL run-time library:

```
PRId64 is not a valid printf conversion specifier for values of type int64_t
```

The `config.log` file indicated the source of the problem was a 'syntax error before 'PRId64'' that was reported when compiling a sample program. A brief search revealed that the `PRId64` macro is not defined in any of the standard C header files on the systems in question. The solution turned out to be to configure with *./configure --with-printf-format='"lld"'* to instruct the C compiler to use `printf()` format strings such as '%lld' when outputting 64-bit signed integers. The extra pair of double quotes is required because the conversion specifier is used in constructs like the following:

```
    printf ("The number is %10" conversion-specifier "!\n", num);
```

Most—but apparently not all—C compilers define the `PRId64` macro ("PRInt signed decimal number of length 64 bits") in one the standard header files. Typical values of `PRId64` include the strings `"lld"` and `"ld"`.

8.1.3 Header is 'present but cannot be compiled'

On some platforms `./configure` may output one or more warnings of the following form:

```
WARNING: filename.h: present but cannot be compiled
WARNING: filename.h:     check for missing prerequisite headers?
WARNING: filename.h: see the Autoconf documentation
WARNING: filename.h:     section "Present But Cannot Be Compiled"
WARNING: filename.h: proceeding with the preprocessor's result
WARNING: filename.h: in the future, the compiler will take precedence
WARNING:     ## --------------------------- ##
WARNING:     ## Report this to pakin@lanl.gov ##
WARNING:     ## --------------------------- ##
```

Typically, such warnings indicate that the header file in question was written for one compiler but a different compiler is being used to build CONCEPTUAL. Unless the header file in question poses problems during the run of *make*, the 'present but cannot be compiled' warnings can be ignored and there is no need to report the compiler-preprocessor mismatch to the e-mail address shown.

8.1.4 Slow 'checking the maximum length of command line arguments...'

Recent versions of Libtool need to know the maximum supported command-line length. Normally, `configure` determines this value by invoking a test script with successively longer command lines until such an invocation fails. On most platforms, the maximum command-line length is determined almost instantaneously. However, on one test system, a 600 MHz $x86$ running the original—not GNU—Bourne shell under Solaris, `./configure` was stuck 'checking the maximum length of command line arguments...' for approximately 45 *minutes*.

Inspecting the `configure` script revealed that the result of executing the length-checking code in `configure` is an assignment to the `lt_cv_sys_max_cmd_len` shell variable. Consequently, running `./configure lt_cv_sys_max_cmd_len=8192` (or some other conservative estimate of the maximum command-line length) caused `./configure` to skip the unduly slow length-checking test and use the given value instead.

8.1.5 'configure' is slow

On a few systems, `./configure` has been observed to take an extremely long time to run. A common source of the problem is a slow filesystem. During the course of its execution the `configure` script creates and compiles a large number of small files. If the current directory resides in a filesystem that exhibits poor small-file performance, then this may explain why `./configure` runs slowly. Try to configure CONCEPTUAL from a local filesystem (e.g., `/tmp`) and see if it runs faster.

A second source of poor CONCEPTUAL configuration speed regards poor implementations of the Bourne shell. One test system, a 1.9 GHz PowerPC running the original—not

GNU—Bourne shell under AIX, took several *hours* to run `./configure` to completion. Fortunately, the system adminstrator had installed the GNU Bourne Again shell (`bash`), which does not exhibit the same poor performance as the default Bourne shell. `./configure` can be instructed to use `bash`—or any Bourne-compatible shell—by setting the `CONFIG_SHELL` environment variable:

```
env CONFIG_SHELL=/usr/local/bin/bash ./configure
```

The preceding command reduced configuration time from over two hours to under two minutes.

8.1.6 Problems with 'C compiler used for Python extension modules'

The following message from `configure` is not uncommon:

```
* Not building the Python interface to the coNCePTuaL run-time
  library because the C compiler used for Python extension modules
  (gcc -pthread) can't link against the output of gcc
```

Without the Python interface, none of the Python-based backends can be built, either:

```
* Not installing the coNCePTuaL interpreter because it depends upon
  the Python interface to the coNCePTuaL run-time library
* Not installing the statistics backend because it depends upon the
  coNCePTuaL interpreter
* Not installing the PICL backend because it depends upon the
  coNCePTuaL interpreter
* Not installing the Paraver backend because it depends upon the
  coNCePTuaL interpreter
* Not installing the LaTeX visualization backend because it depends
  upon the coNCePTuaL interpreter
```

Note that coNCePTuaL is still quite usable without the Python-based backends; the C+MPI backend (see Section 3.3.2 [The c_mpi backend], page 22), for example, does not rely on Python.

A typical cause of the '`can't link`' message is that the `configure` script is unable to locate the Python development files that are needed to build a Python module. Most Python installations exclude the development files by default, requiring that they be installed separately. In some Linux distributions these files are provided by a package with a name like `python-dev`. A key file to look for is `Python.h`. If `./configure` cannot find `Python.h`, the '`can't link`' message will almost certainly appear.

8.1.7 Manual configuration

One of the most important by-products of running `./configure` is a `config.h` file that is used to build the coNCePTuaL run-time library. If the approach outlined in Section 8.1.1 [Interpreting configure warnings], page 189 proves unable to convince coNCePTuaL to build the library or other component, it may be possible to address the failed tests manually by editing any incorrect definitions in `config.h`. In addition, running `./configure` with the `--enable-broken-components` option will force `make` to attempt to build and install *everything*, no matter how unlikely the prospects of success are. You may need to `make`

individual rules from the makefile in order to skip over unsalvageable parts of the build process.

The `--with-header-code` option to `configure` lets you specify a single line of C code to insert into every test file that `configure` generates and also into the header files used to build the run-time library. This extra line of code can be used to make up for missing functionality or to load nonstandard header files. If you need to inject more than one line of code use `./configure --with-header=code='#include "myheaders.h"'` and create an appropriate `myheaders.h` file.

Like all Autoconf scripts, `configure` runs a large number of small test programs to characterize the user's system. It keeps track of the results of these tests in a set of *cache variables*. If a cache variable is assigned a value on the command line, `configure` will skip the associated test and simply assume that the cache variable's value represents the test result. Thus, cache variables provide a fine level of control over the behavior of the `configure` script and are a useful mechanism for working around the idiosyncrasies of any given system. See Appendix D [Cache Variables], page 219, for a list of many of the cache variables recognized by the coNCePTuaL `configure` script.

It is highly unlikely that `./configure` will fail completely. If it does, a working coN-CePTuaL system may still be possible. Manually edit `config.h` for your system. (If `config.h` could not be created, copy `config.h.in` to `config.h`.) Do likewise for `ncptl.h`. Finally, copy `Makefile.simple.in` to `Makefile.simple` if necessary and replace all text bracketed by at signs with appropriate values or blanks. In particular, '`@DEFS@`' should be replaced by the string '`-DHAVE_CONFIG_H`' as this instructs the various C files to include `config.h`. See Section 8.2.9 [Building on problematic platforms], page 196, for more information about `Makefile.simple`. Although `Makefile.in` can be copied to `Makefile` and edited, doing so requires many more string replacements. The advantage is that the result will support all of the options described in Section 2.2 [make], page 7.

8.2 Problems with make

Once coNCePTuaL is configured, the next step is to build the various components. This section explains what to do if a compile fails or `make` is otherwise unable to perform all of its operations.

8.2.1 Failure following `missing` check

The `Makefile` produced by `./configure` sometimes gets confused into thinking that it needs to recreate one or more of the GNU Autotools build files such as `aclocal.m4`, `config.h.in`, `configure`, or `Makefile.in`. A symptom of this confusion is an error message following an invocation of the `missing` script:

```
$ make
cd . && /bin/sh ./config.status ncptl.h
config.status: creating ncptl.h
config.status: ncptl.h is unchanged
cd . && /bin/sh /home/me/conceptual-1.5.1b/missing --run autoheader
                        [Error message appears here.]
```

The source of the problem is usually an incongruous timestamp on some file relative to that of some other file. Unfortunately, identifying the files in question can be difficult. The quick fix is to edit the `missing` script, putting 'exit 0' at the beginning of the file, right after the '#! /bin/sh' line. This tells the `Makefile` that none of the GNU Autotools build files need to be recreated.

8.2.2 'Too many columns in multitable item'

Very old versions of *makeinfo* are unable to process `conceptual.texi` into Emacs 'info'-format documentation; instead, they fail with a large number of 'Too many columns in multitable item (max 1)' errors. Simply re-running *make* should bypass the failing documentation-building step.

8.2.3 Can't find `compiler_version.h`

The CoNCePTuaL build process currently requires *make* to be executed before *make install*. Skipping the *make* step results in a `compiler_version.h` error resembling the following:

```
logfilefuncs.c:18:30: compiler_version.h: No such file or directory
```

Ergo, make sure always to run *make* before running *make install*.

8.2.4 'could not read symbols'

An attempt to link the CoNCePTuaL run-time library on one x86-64 Linux platform aborted with the following error:

```
/usr/lib/libpopt.so: could not read symbols: File in wrong format
```

An attempt on another x86-64 Linux platform aborted with a similar error:

```
/usr/lib/libpopt.so: could not read symbols: Invalid operation
```

The problem on both platforms turned out to be that the `/usr/lib/libpopt.so` is a 32-bit binary and could not be linked with a 64-bit library. The solution was to configure with *./configure LDFLAGS=-L/usr/lib64* to indicate that libraries should be read from `/usr/lib64` instead of the default of `/usr/lib64`.

8.2.5 Incorrect tools/flags are utilized

You might find that although you specified a particular tool or flag at configuration time (e.g., with 'MPICC=mympicompiler'; see Section 2.1 [configure], page 4), the `ncptl` compiler seems to ignore it. A likely culprit is that an environment variable with the same name as the configuration parameter is set (e.g., the *MPICC* environment variable may be set to 'othermpicompiler') and therefore overrides all prior settings. This situation sometimes arises when a compiler or MPI library is made available using the Environment Modules system (e.g., with the *module load* command).

The solution is simply to undefine or properly redefine the offending environment variable. In most Unix environments the `env` command can be used to redefine an environment variable for the duration of a single command:

```
env MPICC=mympicompiler ncptl --backend=c_mpi myprogram.ncptl
```

8.2.6 Compaq compilers on Alpha CPUs

Although CONCEPTUAL builds fine on Alpha-based systems when using a `gcc` compiler, Compaq's C compilers are sometimes problematic. For instance, the `libncptl_wrap.c` source file fails to compile on a system with the following versions of the operating system, C compiler, and Python interpreter:

```
% uname -a
OSF1 qsc14 V5.1 2650 alpha
% cc -V
Compaq C V6.5-011 on Compaq Tru64 UNIX V5.1B (Rev. 2650)
Compiler Driver V6.5-003 (sys) cc Driver
% python -V
Python 2.3
```

On the system that was tested, `cc` aborts with a set of '`Missing type specifier or type qualifier`' messages. The problem appears to be that some of Compaq's standard C header files fail to `#include` various header files they depend upon. A workaround is to insert the following C preprocessor directives in `libncptl_wrap.c` before the line reading '`#include "Python.h"`':

```
#include <sys/types.h>
#include <sys/time.h>
#include <sys/stat.h>
```

A second problem with Compaq compilers on Alpha-based systems occurs under Linux when using Compaq's `ccc` compiler:

```
% uname -a
Linux wi 2.4.21-3.7qsnet #2 SMP Fri Oct 17 14:08:00 MDT 2003 alpha unknown
% ccc -V
Compaq C T6.5-002 on Linux 2.4.21-3.7qsnet alpha
Compiler Driver T6.5-001 (Linux) cc Driver
Installed as ccc-6.5.6.002-1
Red Hat Linux release 7.2 (Enigma)
Using /usr/lib/gcc-lib/alpha-redhat-linux/2.96 (4).
```

When linking files into a shared object, `ccc` aborts with a '`[. . .]/.libs: file not recognized: Is a directory`' error message. The problem appears to be that `libtool` is confused about the arguments it's supposed to pass to the linker; `libtool` uses '`--rpath`'

(two hyphens) where the Compaq linker expects '`-rpath`' (one hyphen). As a workaround, you can edit the `libtool` file after running *`./configure`* but before running *`make`*. Simply replace '`--rpath`' with '`-rpath`' in the following `libtool` line and the problem should go away:

```
hardcode_libdir_flag_spec="\${wl}--rpath \${wl}\$libdir"
```

8.2.7 '`undefined type, found 'DEFINE_____'`'

The GCC 2.95.2 compiler on at least one MacOS 10.1 system complains repeatedly about an '`undefined type, found 'DEFINE_____'`' while trying to preprocess a data file generated by `makehelper.py` and based upon the contents of `config.h substitutions.dat`. Fortunately, the complaint is only a warning and can be ignored. The problem is that the C preprocessor performs "smart preprocessing" if given a filename on the command line (as is the case here) but "basic preprocessing" when reading from the standard input device. Because the input does not represent syntactically correct C code—the C preprocessor is used only as a convenient device for performing macro substitutions—the syntax-aware smart preprocessing fails. However, the C preprocessor then reprocesses the file with basic preprocessing (as indicated by the message '`cpp-precomp: warning: errors during smart preprocessing, retrying in basic mode`') and succeeds.

8.2.8 *`makehelper.py config`* fails

When `makehelper.py` is run with the '`config`' option, it generates a temporary `.c` file that it runs through the C preprocessor. Problems (e.g., a '`gcc.exe: no input files`' error message) may occur when running under Cygwin (a Unix-like user environment for Microsoft Windows) with a Cygwin Python interpreter but a non-Cygwin C preprocessor (e.g., MinGW's) because `makehelper.py` may try to pass a Unix-style filename to the C preprocessor, which expects to receive a Windows-style filename.

A workaround is to specify explicitly a temporary directory for `makehelper.py` to use. Use forward slashes and filenames without spaces (e.g., DOS-style 8.3 filenames) as in the following example:

```
make TEMP=C:/DOCUME~1/user/LOCALS~1/Temp
```

8.2.9 Building on problematic platforms

Some experimental systems require rather specialized build procedures that thwart CoN-CePTuaL's standard makefile. Unfortunately, `Makefile` is complex and difficult to edit by hand. Users comfortable with Automake should edit `Makefile.am`—which is used to generate `Makefile`—and re-run *`automake`*, *`autoconf`*, and *`./configure`* as described in .

CoNCePTuaL includes an alternate Makefile called `Makefile.simple` (generated at configuration time from `Makefile.simple.in`). `Makefile.simple` is a stripped-down version of `Makefile` that is designed to be easy to edit by hand. *`make -f Makefile.simple`* builds a static version of the CoNCePTuaL run-time library in the current directory. *`make -f Makefile.simple clean`* deletes the run-time library and all of the object files used to build it. `Makefile.simple` supports no other features. The intention is to provide the bare minimum needed to get backends such as `c_mpi` to produce executable programs even when running in unusual environments.

8.3 Problems running

After CONCEPTUAL is configured, compiled, built, and installed, there is still the chance that CONCEPTUAL-generated executables fail to run. This section addresses some common problems and presents their solutions.

As a quick tip, a very conservative way to run a CONCEPTUAL program is with the `NCPTL_NOFORK` environment variable set to '1', the `NCPTL_CHECKPOINT` environment variable set to '0', and with `--no-trap=1-63` on the command line. Such usage should work around some of the most common—and some of the hardest to diagnose—problems that may impact a CONCEPTUAL program. Read the corresponding sections (Section 8.3.2 [Miscellaneous mysterious hangs or crashes], page 198; Section 8.3.3 [Extremely noisy measurements], page 198; and, Section 8.3.4 [Keeping programs from dying on a signal], page 199) for information about the drawbacks of each of the preceding settings.

8.3.1 'cannot open shared object file'

A common problem on many workstation clusters is that the head node (on which `./configure` is run) has more libraries installed than do the compute nodes (on which CONCEPTUAL programs themselves are run). Consequently, CONCEPTUAL programs that compile and link properly on the head node will fail with a 'cannot open shared object file' error if they try to dynamically link a shared object that is absent on the compute nodes.

Ideally, the missing shared objects should be installed on each of the compute nodes. If they cannot be installed in their standard locations (e.g., `/usr/lib`), they can be installed elsewhere (e.g., in the user's home directory) and that location can be pointed to using `-rpath` or `LD_LIBRARY_PATH`, as described below. Alternatively, `configure` can be instructed to disregard certain libraries—even if they exist on the head node—by passing the `--with-ignored-libs` option to `configure` (see Section 2.1 [configure], page 4). For example, if `libpopt.so` is not installed on the compute nodes one can run `./configure --with-ignored-libs=popt` to prevent `libpopt.so` from being used.

The `--with-ignored-libs` option is of little use if the CONCEPTUAL library itself cannot be found at run time. By default, `make` will build and `make install` will install both a static and a dynamic version of the CONCEPTUAL run-time library. Most linkers give precedence to the dynamic library over the static library unless the static library is requested explicitly. As a consequence, the dynamic version of the CONCEPTUAL run-time library needs to be available at program-load time in order to avoid error messages like the following:

```
a.out: error while loading shared libraries: libncptl.so.0: cannot open
shared object file: No such file or directory
```

There are a few alternatives for pointing the dynamic loader to the CONCEPTUAL run-time library. On systems that support it, the best option is the `-rpath` approach described in Section 7.5 [Implicit dynamic-library search paths], page 187. The second-best option is to add the directory in which `libncptl.so` was installed (by default, `/usr/local/lib`) to your `LD_LIBRARY_PATH` environment variable. Finally, as a last resort, you can use

the `--disable-shared` configuration option (see Section 2.1 [configure], page 4) to prevent CoNCePTuaL from building the dynamic version of the run-time library altogether, thereby forcing the linker to use the static version:

```
make uninstall
make clean
./configure --disable-shared ...
make
make install
```

As mentioned in Section 2.1 [configure], page 4, however, `libncptlmodule.so` can't be built when `--disable-shared` is in effect.

8.3.2 Miscellaneous mysterious hangs or crashes

In some implementations of the software stack for InfiniBand—and possibly for some other networks as well—invocations of the `fork()` system call made while the network device is open can corrupt process memory and hang or crash the corresponding process. (On Unix and Unix-like operating systems this is typically via a segmentation fault.)

The workaround is to build CoNCePTuaL using the `--without-fork` option (see Section 2.1 [configure], page 4). Alternatively, you can set the *NCPTL_NOFORK* environment variable when running a CoNCePTuaL program. Either option suppresses the CoNCePTuaL run-time library's use of all process-spawning functions. See Appendix C [Environment Variables], page 217, for a description of *NCPTL_NOFORK*'s side effects. The same description also applies to `--without-fork`.

The C library that comes bundled with a lightweight run-time kernel or other custom operating system may provide broken versions of some of the functions on which the CoNCePTuaL run-time library relies. If a debugger or crash-analysis tool indicates that a particular function is problematic it may be possible to disable that function at configuration time. For example, the current (at the time of this writing) release of Red Storm's software stack includes a faulty `getrusage()` function that crashes consistently. Observing that `configure` checks for `getrusage()` we searched the generated `config.log` file and encountered the corresponding shell variable, '`ac_cv_func_getrusage`', which was set to '`yes`'. Rerunning *./configure* with '`ac_cv_func_getrusage=no`' on the command line disabled the CoNCePTuaL run-time library's use of `getrusage()` and thereby resulted in crash-free execution of CoNCePTuaL programs.

8.3.3 Extremely noisy measurements

Sometimes a CoNCePTuaL program runs to completion but the data written to the log files exhibit high levels of variability across runs or even across trials within a single run. A possible source of this variability—especially for long-running programs running on a large number of processors—is the run-time library's log-file checkpointing mechanism. Because each process in a CoNCePTuaL programs writes its own log file, poorly scalable shared filesystems, limited spare network bandwidth, and asynchronous operating-system buffer flushes may each impact program performance in an unpredictable manner and at unpredictable times.

Setting the `NCPTL_CHECKPOINT` environment variable to '0' disables log-file checkpointing and may thereby reduce some of the data variability. The caveat is that a program that aborts abnormally will leave behind *no* data in its log files. See Appendix C [Environment Variables], page 217, for more information about `NCPTL_CHECKPOINT`.

8.3.4 Keeping programs from dying on a signal

By default, CoNCePTuaL programs terminate upon receiving *any* unexpected signal. The error message list the signal number and, if available, a human-readable signal name:

```
myprogram: Received signal 28 (Window changed); specify --no-trap=28 to
ignore
```

The motivation behind this decision to abort on unexpected signals is that signal-handling adversely affects a program's performance. Hence, by running to completion, a program indicates that it did not receive any unexpected signals. However, some messaging layers use signals internally (most commonly `SIGUSR1` and `SIGUSR2`) to coordinate helper processes. To permit a program to deliver such signals to the messaging layer a user should use the program's `--no-trap` command-line option as described in Section 3.4 [Running coNCePTuaL programs], page 46.

8.3.5 'Unaligned access' warnings

On some platforms you may encounter messages like the following written to the console and/or various system log files (e.g., `/var/log/messages`):

```
myprog(25044): unaligned access to 0x6000000000001022,
ip=0x40000000000009e1
```

Alternatively:

```
Unaligned access pid=7890104 <myprog> va=0x140004221 pc=0x1200012b4
ra=0x1200012a4 inst=0xb449fff8
```

What's happening is that some CPUs require n-byte-wide data to be aligned on an n-byte boundary. For example, a 64-bit datatype can be accessed properly only from memory locations whose address is a multiple of 64 bits (8 bytes). On some platforms, misaligned accesses abnormally terminate the program, typically with a `SIGBUS` signal. On other platforms, misaligned accesses interrupt the operating system. The operating system fixes the access by splitting it into multiple aligned accesses plus some bit masking and shifting and then notifying the user and/or system administrator that a fixup occurred.

CoNCePTuaL's `configure` script automatically determines what data alignments are allowed by the architecture but it has no way to determine if fixups occurred as these are transparent to programs. The result is annoying "unaligned access" messages such as those quoted above. One solution is to use the `--with-alignment` option to `configure` to specify explicitly the minimum data alignment that CoNCePTuaL should be permitted to use. Alternatively, some operating systems provide a mechanism to cause misaligned accesses to result in a `SIGBUS` signal instead of a fixup and notification message. On Linux/IA-64 this is achieved with the command `prctl --unaligned=signal`. On OSF1/Alpha the equivalent command is `uac p sigbus`. Be sure to rerun `configure` after issuing those commands to make it reexamine the set of valid data alignments.

8.3.6 'Unable to determine the OS page size'

To implement `PAGE ALIGNED` messages (see [Message alignment], page 111) a program needs to be able to query the operating system's page size. The CoNCePTuaL run-time library performs this query using one of the `getpagesize()` or `sysconf()` operating-system calls. If neither call is available or functional the run-time library aborts with an 'Unable to determine the OS page size' error.

As a workaround, `configure` provides a `--with-page-size` option that enables the user to manually specify the page size. Because any manually specified value is prone to error, CoNCePTuaL log files include a 'Page size was specified manually at configuration time' warning if the run-time library was configured using `--with-page-size`.

8.3.7 Invalid timing measurements

Although the CoNCePTuaL `configure` script is usually good at selecting a mechanism for measuring elapsed time, there are a few systems that confuse the script. For example, different Intel processors use different mechanisms for mapping cycle counts to time (cf. http://en.wikipedia.org/wiki/RDTSC); the correct mapping cannot always be determined at configuration time and may require administrator privileges to calculate. Consequently, if `configure` determines that the CoNCePTuaL timer should read the cycle counter directly, incorrect times may be reported.

The best way to test the timer quality is to run the `validatetimer` program (see [Validating the coNCePTuaL timer], page 9). If the difference between wall-clock time and CoNCePTuaL-reported time is great, the `--with-gettimeofday` configuration option (see Section 2.1 [configure], page 4) is usually a safe bet for improving accuracy, albeit at a slight cost in measurement overhead.

8.3.8 'TeX capacity exceeded'

TeX—and by consequence, LaTeX—does not use dynamically allocated memory. Therefore, attempting to produce a very large diagram with the `latex_vis` backend (see Section 3.3.10 [The latex_vis backend], page 37) will likely exceed TeX's hardwired memory capacity:

```
! TeX capacity exceeded, sorry [main memory size=350001].
\psm@endnode@i ... \endgroup \psm@endmath \egroup
                                                  \use@par \@psttrue
l.489 \task
           {0} & \task{1} \\
No pages of output.
```

TeX/LaTeX's memory capacity can be increased but the mechanism for doing so varies from one TeX distribution to another and is rarely straightforward. See http://www.tex.ac.uk/cgi-bin/texfaq2html?label=enlarge for a few terse pointers. In general, it is best to try to minimize the number of loop repetitions when running programs through the `latex_vis` backend.

An alternative is to use a prebuilt large-memory `latex`. Some TeX distributions come with a `hugelatex` executable, which is just like `latex` but compiled with larger memory limits. `ncptl` can be told to use `hugelatex` by setting the *LATEX* environment variable

(e.g., with 'env LATEX=hugelatex ncptl ...'). This may be the most convenient way to produce complex diagrams with `latex_vis` when simplifying the run is not an option.

8.3.9 Bad bounding boxes from `latex_vis`

Very tall program visualizations are susceptible to Ghostscript bug #202735, "bbox device doesn't allow min coords < 0". `latex_vis` attempts to work around that bug by defining a large PostScript `ImagingBBox` as described in `http://bugs.ghostscript.com/show_bug.cgi?id=202735`. While no problems with the `latex_vis` workaround have yet been reported, if problems do occur it should be sufficient to set the *GS* environment variable to ':' or the name of a program known not to exist. `latex_vis` will then issue a warning message and generate an EPS file with an acceptable but slightly loose bounding box, as mentioned in Section 3.3.10 [The latex_vis backend], page 37.

8.4 When all else fails

The coNCePTuaL project pages on SourceForge.net provide a variety of mechanisms for providing feedback to the coNCePTuaL developers:

Mailing lists
>   Read and post messages on the coNCePTuaL mailing lists at `http://sourceforge.net/mail/?group_id=117615`.

Trackers   Search for old or post new bug reports, feature requests, or code patches at `http://sourceforge.net/tracker/?group_id=117615`.

Appendix A Reserved Words

As mentioned in Section 4.1 [Primitives], page 84, not all identifiers can be used as variables. The following sections provide a complete list of identifiers that are forbidden as variable names. These identifiers fall into two categories: keywords, which are never allowed as variable names, and predeclared variables, which are "read-only" variables; they can be utilized just like any other variables but cannot be redeclared.

A.1 Keywords

The following is a list of all currently defined keywords in the CONCEPTUAL language. It is an error to try to use any of these as identifiers.

- A
- ABS
- ABSOLUTE
- AGGREGATES
- ALIGNED
- ALL
- AN
- AND
- ARE
- ARITHMETIC
- AS
- ASSERT
- ASSIGNED
- ASYNCHRONOUSLY
- AWAIT
- AWAITS
- BACKEND
- BE
- BIT
- BITS
- BUFFER
- BUFFERS
- BUT
- BYTE
- BYTES
- CBRT
- CEILING
- COMES
- COMPLETION

- COMPLETIONS
- COMPUTE
- COMPUTES
- COUNTERS
- CURRENT
- DATA
- DAY
- DAYS
- DECLARES
- DEFAULT
- DEVIATION
- DIVIDES
- DOUBLEWORD
- DOUBLEWORDS
- EACH
- EVEN
- EXECUTE
- EXECUTES
- FACTOR10
- FILE_DATA
- FINAL
- FLOOR
- FOR
- FROM
- GEOMETRIC
- GIGABYTE
- GREATER
- GROUP
- HALFWORD
- HALFWORDS
- HARMONIC
- HISTOGRAM
- HOUR
- HOURS
- IF
- IN
- INTEGER
- INTEGERS

- INTO
- IS
- IT
- ITS
- KILOBYTE
- KNOMIAL_CHILD
- KNOMIAL_CHILDREN
- KNOMIAL_PARENT
- LANGUAGE
- LESS
- LET
- LOG
- LOG10
- LOGS
- MAX
- MAXIMUM
- MEAN
- MEDIAN
- MEGABYTE
- MEMORY
- MESH_COORDINATE
- MESH_DISTANCE
- MESH_NEIGHBOR
- MESSAGE
- MESSAGES
- MICROSECOND
- MICROSECONDS
- MILLISECOND
- MILLISECONDS
- MIN
- MINIMUM
- MINUTE
- MINUTES
- MISALIGNED
- MOD
- MULTICAST
- MULTICASTS
- MY

- NONUNIQUE
- NOT
- ODD
- OF
- OR
- OTHER
- OTHERWISE
- OUTPUT
- OUTPUTS
- PAGE
- PAGES
- PERCENTILE
- PLUS
- PROCESSOR
- PROCESSOR_OF
- PROCESSORS
- QUADWORD
- QUADWORDS
- RANDOM
- RANDOM_GAUSSIAN
- RANDOM_PARETO
- RANDOM_POISSON
- RANDOM_UNIFORM
- REAL
- RECEIVE
- RECEIVES
- REDUCE
- REDUCES
- REGION
- REPETITION
- REPETITIONS
- REQUIRE
- RESET
- RESETS
- RESTORE
- RESTORES
- RESULT
- RESULTS

- ROOT
- ROUND
- SECOND
- SECONDS
- SEND
- SENDS
- SIZED
- SLEEP
- SLEEPS
- SQRT
- STANDARD
- STATIC_FILE_DATA
- STORE
- STORES
- STRIDE
- SUCH
- SUM
- SYNCHRONIZATION
- SYNCHRONIZE
- SYNCHRONIZES
- SYNCHRONOUSLY
- TAG
- TASK
- TASK_OF
- TASKS
- THAN
- THAT
- THE
- THEIR
- THEM
- THEN
- TIME
- TIMES
- TO
- TOUCH
- TOUCHES
- TOUCHING
- TREE_CHILD

- `TREE_PARENT`
- `UNALIGNED`
- `UNIQUE`
- `UNSUSPECTING`
- `USING`
- `VALUE`
- `VARIANCE`
- `VERIFICATION`
- `VERSION`
- `WARMUP`
- `WHERE`
- `WHILE`
- `WHO`
- `WITH`
- `WITHOUT`
- `WORD`
- `WORDS`
- `XOR`

A.2 Predeclared variables

CONCEPTUAL predeclares a set of variables that programs can use but not redeclare. These variables and their descriptions are listed below.

`bit_errors`
>   Total number of bit errors observed

`bytes_received`
>   Total number of bytes received

`bytes_sent`
>   Total number of bytes sent

`elapsed_usecs`
>   Elapsed time in microseconds

`msgs_received`
>   Total number of messages received

`msgs_sent`
>   Total number of messages sent

`num_tasks`
>   Number of tasks running the program

`total_bytes`
>   Sum of bytes sent and bytes received

`total_msgs`

> Sum of messages sent and messages received

As should be evident from their descriptions, cONCEpTuAL's predeclared variables are updated dynamically. Each access can potentially return a different result. Furthermore, unlike user-declared variables, predeclared variables in can have a different value on each task.

The following notes clarify the semantics that relate to the updating of some of the preceding variables:

`ASYNCHRONOUSLY` (see [Blocking semantics], page 113)

> A message that is *sent* `ASYNCHRONOUSLY` immediately increments each of `msgs_sent` and `total_msgs` by 1 and each of `bytes_sent` and `total_bytes` by the message size. A message that is *received* `ASYNCHRONOUSLY` increments none of the predeclared variables. However, after the receiving task `AWAITS COMPLETION` (see Section 4.4.4 [Awaiting completion], page 117) it increments each of `msgs_received` and `total_msgs` by 1 and each of `bytes_received` and `total_bytes` by the message size.

`MULTICASTS` (see Section 4.4.5 [Multicasting], page 117)

> None of the byte or message variables are updated as part of a multicast operation, regardless of how the underlying messaging layer implements multicasts.

`SYNCHRONIZES` (see Section 4.4.7 [Synchronizing], page 119)

> None of the byte or message variables are updated as part of a barrier synchronization, regardless of how the underlying messaging layer implements barriers.

Appendix B Backend Developer's Reference

Programmers wishing to develop their own CONCEPTUAL backends can refer to the information presented in this appendix as a complement to the more tutorial-in-nature Chapter 6 [Implementation], page 156.

B.1 Method calls

The following method calls must be defined when writing a CONCEPTUAL backend from scratch. They are invoked indirectly as part of PLY's AST traversal. See Section 6.2 [Backend creation], page 157, for more information.

- `n_add_expr`
- `n_aggregate_expr`
- `n_aggregate_func`
- `n_an`
- `n_assert_stmt`
- `n_awaits_completion`
- `n_backend_decl`
- `n_backend_stmt`
- `n_buffer_number`
- `n_buffer_offset`
- `n_computes_for`
- `n_data_multiplier`
- `n_data_type`
- `n_dimension`
- `n_dimension_list`
- `n_empty_stmt`
- `n_eq_expr`
- `n_expr_list`
- `n_for_count`
- `n_for_each`
- `n_for_each_expr`
- `n_for_time`
- `n_func_call`
- `n_ident`
- `n_if_stmt`
- `n_ifelse_expr`
- `n_integer`
- `n_item_size`
- `n_let_binding`
- `n_let_binding_list`

- `n_let_stmt`
- `n_log_expr_list`
- `n_log_expr_list_elt`
- `n_log_flush_stmt`
- `n_log_stmt`
- `n_mcast_stmt`
- `n_message_alignment`
- `n_message_spec`
- `n_mult_expr`
- `n_my_task`
- `n_no_touching`
- `n_output_stmt`
- `n_param_decl`
- `n_power_expr`
- `n_processor_stmt`
- `n_program`
- `n_range`
- `n_range_list`
- `n_real`
- `n_receive_attrs`
- `n_receive_stmt`
- `n_recv_buffer_number`
- `n_reduce_message_spec`
- `n_reduce_stmt`
- `n_rel_conj_expr`
- `n_rel_disj_expr`
- `n_reset_stmt`
- `n_restore_stmt`
- `n_send_attrs`
- `n_send_stmt`
- `n_simple_stmt_list`
- `n_sleeps_for`
- `n_store_stmt`
- `n_stride`
- `n_string`
- `n_string_or_expr_list`
- `n_string_or_log_comment`
- `n_sync_stmt`

- `n_tag`
- `n_task_expr`
- `n_time_unit`
- `n_top_level_stmt`
- `n_touch_buffer_stmt`
- `n_touch_repeat_count`
- `n_touch_stmt`
- `n_touching`
- `n_trivial_node`
- `n_unary_expr`
- `n_unique`
- `n_verification`
- `n_version_decl`
- `n_where_expr`

B.2 C hooks

To save the backend developer from having to implement CONCEPTUAL backends entirely from scratch, CONCEPTUAL provides a `codegen_c_generic.py` module that defines a base class for code generators that output C code. The base class handles the features that are specific to C but independent of any messaging library. Derived classes need only define those "hook" functions that are needed to implement library-specific functionality.

Hooks are named after the method from which they're called but with an all-uppercase tag appended. The following list shows each hook-calling method in `codegen_c_generic.py` and the set of hooks it calls. See Section 6.2.1 [Hook methods], page 158, for more information.

- `code_declare_datatypes`
    - `code_declare_datatypes_EXTRA_EVENTS`
    - `code_declare_datatypes_EXTRA_EVENT_STATE`
    - `code_declare_datatypes_EXTRA_EVS`
    - `code_declare_datatypes_MCAST_STATE`
    - `code_declare_datatypes_POST`
    - `code_declare_datatypes_PRE`
    - `code_declare_datatypes_RECV_STATE`
    - `code_declare_datatypes_REDUCE_STATE`
    - `code_declare_datatypes_SEND_STATE`
    - `code_declare_datatypes_SYNC_STATE`
    - `code_declare_datatypes_WAIT_STATE`
- `code_declare_globals`
    - `code_declare_globals_DUMMY_VAR`
    - `code_declare_globals_EXTRA`

- `code_def_alloc_event`
  - `code_def_alloc_event_DECLS`
  - `code_def_alloc_event_POST`
  - `code_def_alloc_event_PRE`
- `code_def_exit_handler`
  - `code_def_exit_handler_BODY`
- `code_def_finalize`
  - `code_def_finalize_DECL`
  - `code_def_finalize_POST`
  - `code_def_finalize_PRE`
- `code_def_init_check_pending`
  - `code_def_init_check_pending_POST`
  - `code_def_init_check_pending_PRE`
- `code_def_init_cmd_line`
  - `code_def_init_cmd_line_POST_ARGS`
  - `code_def_init_cmd_line_POST_PARSE`
  - `code_def_init_cmd_line_PRE_ARGS`
  - `code_def_init_cmd_line_PRE_PARSE`
- `code_def_init_decls`
  - `code_def_init_decls_POST`
  - `code_def_init_decls_PRE`
- `code_def_init_init`
  - `code_def_init_init_POST`
  - `code_def_init_init_PRE`
- `code_def_init_misc`
  - `code_def_init_misc_EXTRA`
  - `code_def_init_misc_PRE_LOG_OPEN`
- `code_def_init_msg_mem`
  - `code_def_init_msg_mem_EACH_TAG`
  - `code_def_init_msg_mem_POST`
  - `code_def_init_msg_mem_PRE`
- `code_def_init_reseed`
  - `code_def_init_reseed_BCAST`
- `code_def_init_seed`
  - `code_def_init_seed_POST`
  - `code_def_init_seed_PRE`
- `code_def_init_uuid`
  - `code_def_init_uuid_BCAST`

- `code_def_mark_used`
  - `code_def_mark_used_POST`
  - `code_def_mark_used_PRE`
- `code_def_procev`
  - `code_def_procev_DECL`
  - `code_def_procev_EVENTS_DECL`
  - `code_def_procev_EXTRA_EVENTS`
  - `code_def_procev_POST`
  - `code_def_procev_POST_SWITCH`
  - `code_def_procev_PRE`
  - `code_def_procev_PRE_SWITCH`
- `code_def_procev_arecv`
  - `code_def_procev_arecv_BODY`
- `code_def_procev_asend`
  - `code_def_procev_asend_BODY`
- `code_def_procev_etime`
  - `code_def_procev_etime_REDUCE_MIN`
- `code_def_procev_mcast`
  - `code_def_procev_mcast_BODY`
- `code_def_procev_newstmt`
  - `code_def_procev_newstmt_BODY`
- `code_def_procev_recv`
  - `code_def_procev_recv_BODY`
- `code_def_procev_reduce`
  - `code_def_procev_reduce_BODY`
- `code_def_procev_repeat`
  - `code_def_procev_repeat_BODY`
- `code_def_procev_send`
  - `code_def_procev_send_BODY`
- `code_def_procev_sync`
  - `code_def_procev_sync_BODY`
- `code_def_procev_wait`
  - `code_def_procev_wait_BODY_RECVS`
  - `code_def_procev_wait_BODY_SENDS`
- `code_def_small_funcs`
  - `code_def_small_funcs_POST`
  - `code_def_small_funcs_PRE`
- `code_define_functions`

- – code_define_functions_INIT_COMM_1
  - – code_define_functions_INIT_COMM_2
  - – code_define_functions_INIT_COMM_3
  - – code_define_functions_POST
  - – code_define_functions_PRE
- code_define_macros
  - – code_define_macros_POST
  - – code_define_macros_PRE
- code_define_main
  - – code_define_main_DECL
  - – code_define_main_POST_EVENTS
  - – code_define_main_POST_INIT
  - – code_define_main_PRE_EVENTS
  - – code_define_main_PRE_INIT
- code_output_header_comments
  - – code_output_header_comments_EXTRA
- code_specify_include_files
  - – code_specify_include_files_POST
  - – code_specify_include_files_PRE
- code_synchronize_all
  - – code_synchronize_all_BODY
- n_for_count
  - – n_for_count_SYNC_ALL
- n_mcast_stmt
  - – n_mcast_stmt_DECL
  - – n_mcast_stmt_INIT
  - – n_mcast_stmt_MANY_MANY
  - – n_mcast_stmt_POST
  - – n_mcast_stmt_PRE
- n_recv_stmt
  - – n_recv_stmt_BODY
- n_reduce_stmt
  - – n_reduce_stmt_DECL
  - – n_reduce_stmt_HAVE_PEERS
  - – n_reduce_stmt_INIT
  - – n_reduce_stmt_INIT2
  - – n_reduce_stmt_POST
  - – n_reduce_stmt_PRE

- `n_send_stmt`
  - `n_send_stmt_BODY`
- `n_sync_stmt`
  - `n_sync_stmt_DECL`
  - `n_sync_stmt_INIT`
  - `n_sync_stmt_POST`
  - `n_sync_stmt_PRE`

B.3 Event types

Programs generated by `codegen_c_generic.py` define the following event types:

`EV_ARECV`   Asynchronous receive

`EV_ASEND`   Asynchronous send

`EV_BTIME`   Beginning of a timed loop

`EV_DELAY`   Spin or sleep

`EV_ETIME`   Ending of a timed loop

`EV_FLUSH`   Compute aggregate functions for log-file columns

`EV_MCAST`   Synchronous multicast

`EV_NEWSTMT`
            Beginning of a new top-level statement

`EV_RECV`   Synchronous receive

`EV_REDUCE`
            Reduction with or without a subsequent multicast

`EV_REPEAT`
            Repeatedly process the next N events

`EV_RESET`   Reset counters

`EV_RESTORE`
            Restore the previously pushed counter values

`EV_SEND`   Synchronous send

`EV_STORE`   Store all counters' current values

`EV_SUPPRESS`
            Suppress writing to the log and standard output

`EV_SYNC`   Barrier synchronization

`EV_TOUCH`   Touch a region of memory

`EV_WAIT`   Wait for all asynchronous sends/receives to complete

EV_CODE
            None of the above

motivates the use of event-based execution for coNCePTuaL programs.

B.4 Representing aggregate functions

The `LOG_AGGREGATE` enumerated type, defined in `ncptl.h`, accepts the following values:

`NCPTL_FUNC_NO_AGGREGATE`
> Log all data points.

`NCPTL_FUNC_MEAN`
> Log only the arithmetic mean.

`NCPTL_FUNC_HARMONIC_MEAN`
> Log only the harmonic mean.

`NCPTL_FUNC_GEOMETRIC_MEAN`
> Log only the geometric mean.

`NCPTL_FUNC_MEDIAN`
> Log only the median.

`NCPTL_FUNC_MAD`
> Log only the median absolute deviation.

`NCPTL_FUNC_STDEV`
> Log only the standard deviation.

`NCPTL_FUNC_VARIANCE`
> Log only the variance.

`NCPTL_FUNC_SUM`
> Log only the sum.

`NCPTL_FUNC_MINIMUM`
> Log only the minimum.

`NCPTL_FUNC_MAXIMUM`
> Log only the maximum.

`NCPTL_FUNC_FINAL`
> Log only the final measurement.

`NCPTL_FUNC_PERCENTILE`
> Log only the nth percentile.

`NCPTL_FUNC_ONLY`
> Log any data point, aborting if they're not all identical.

`NCPTL_FUNC_HISTOGRAM`
> Log a histogram of the data points

Appendix C Environment Variables

The CoNCePTuaL compiler (`ncptl`) honors the following environment variables:

NCPTL_BACKEND

> Name a default backend for the compiler to use. For example, setting *NCPTL_BACKEND* to 'c_udgram' tells `ncptl` to use the `c_udgram` backend unless the `--backend` compiler option designates a different backend. See Section 3.2 [Compiling coNCePTuaL programs], page 18, for more information.

NCPTL_PATH

> Specify a colon-separated list of directories in which to search for compiler backends. See Section 3.2 [Compiling coNCePTuaL programs], page 18, for more information.

The following environment variables are honored when running a CoNCePTuaL program (any backend):

NCPTL_CHECKPOINT

> Specify the minimum number of seconds between log-file checkpoints (default: 60). The CoNCePTuaL run-time library buffers logged data in memory because CoNCePTuaL programs are not restricted to writing data in a top-to-bottom, left-to-right format. The following program, for example, writes row 9, column 2 before writing row 5, column 1:
>
> ```
> FOR EACH i IN {1, ..., 10} {IF i IS EVEN THEN TASK 0 LOGS i
> AS "Even numbers" THEN TASK 0 LOGS i AS "All numbers"}
> ```
>
> Buffering data in memory enables that sort of "two-dimensional" logging. However, it is not robust to computer crashes or uncatchable signals (e.g., `SIGKILL`). Consequently, every `LOGS` statement (see Section 4.5.3 [Writing to a log file], page 120) that is executed at least *NCPTL_CHECKPOINT* seconds after the previous one forces the CoNCePTuaL run-time library to write its partial data to the log file. Specifically, the library flushes its partial-data buffers then rewinds the write pointer to the beginning of that partial data. Hence, the partial data is overwritten by a later checkpoint or by the complete data set.
>
> Smaller values of *NCPTL_CHECKPOINT* provide more robustness to crashes and uncatchable signals. Larger values put less stress on the filesystem. As a special case, if *NCPTL_CHECKPOINT* is set to '0' then log-file checkpointing is disabled altogether.

NCPTL_FAST_INIT

> If set to '1', more quickly initialize the run-time library by skipping the timer calibration and measurement steps. As a consequence, all timing measurements will be meaningless. *NCPTL_FAST_INIT* may be useful during the development of a CoNCePTuaL program or compiler backend to enable shorter turnaround times. If set to '0', *NCPTL_FAST_INIT* forces a thorough initialization even for backends that do not rely on timing measurements, for instance the `picl` backend (see Section 3.3.8 [The picl backend], page 34). See Section 6.3.1 [Variables and data types], page 166, for more information.

`NCPTL_LOG_DELAY`

> Artificially delay each log-file open and flush operation by a random number of milliseconds in the range 0...`NCPTL_LOG_DELAY` (default: '0'). For example, if 1000 CPUs share a single filesystem, an `NCPTL_LOG_DELAY` of '2000' (i.e., 2 seconds) will probabilistically ensure that a log-file creation request will be issued only once every $2000 \div 1000 = 2$ milliseconds instead of all at once. `NCPTL_LOG_DELAY` is intended to help CONCEPTUAL programs run atop broken filesystems that are unable to handle large numbers of concurrent accesses—an all-too-common problem on large-scale workstation clusters and parallel computers in which hundreds or thousands of diskless compute nodes compete for access to the same filesystem.

`NCPTL_LOG_ONLY`

> Limit the set of processes that produce log files. `NCPTL_LOG_ONLY` accepts a comma-separated list of dash-separated process ranges such as '0-3,12-16,24,25,32-48'. Only processes included in the list produce log files. See Section 3.4 [Running coNCePTuaL programs], page 46, for more information.

`NCPTL_NOFORK`

> If set, inhibit the use of `fork()`, `system()`, `popen()`, and other process-spawning functions. The result is that some information will be omitted from the log-file prologue. `NCPTL_NOFORK` is intended to be used on systems in which such functions corrupt messaging-layer state, hang or crash processes, or wreak other such havoc.

Appendix D `configure` Cache Variables

The `configure` script that is used to build CoNCePTuaL internally uses a number of Autoconf cache variables to store configuration information that it discovers while it runs. If these variables are set on the *./configure* command line, the given value will be used instead of an automatically determined value. This feature gives the user precise control over many aspects of CoNCePTuaL configuration. For example, to prevent the CoNCePTuaL run-time library from using the `gethostbyname()` function you can specify '`ac_cv_func_gethostbyname=no`' on the *./configure* command line to convince `configure` that `gethostbyname()` does not exist.

The variables and descriptions in the following list were extracted mechanically from `configure.ac` and `config.log`. Consequently, some variables may be missing, and some descriptions may be inaccurate or hard to understand.

`ac_cv_func__strtoi64`
> the `_strtoi64()` function is available ('yes' or 'no')

`ac_cv_func_cbrt`
> the `cbrt()` function is available ('yes' or 'no')

`ac_cv_func_floor`
> the `floor()` function is available ('yes' or 'no')

`ac_cv_func_fmod`
> the `fmod()` function is available ('yes' or 'no')

`ac_cv_func_fork`
> the `fork()` function is available ('yes' or 'no')

`ac_cv_func_fork_works`
> the `fork_works()` function is available ('yes' or 'no')

`ac_cv_func_getaddrinfo`
> the `getaddrinfo()` function is available ('yes' or 'no')

`ac_cv_func_getcwd`
> the `getcwd()` function is available ('yes' or 'no')

`ac_cv_func_geteuid`
> the `geteuid()` function is available ('yes' or 'no')

`ac_cv_func_gethostbyname`
> the `gethostbyname()` function is available ('yes' or 'no')

`ac_cv_func_getinvent`
> the `getinvent()` function is available ('yes' or 'no')

`ac_cv_func_getpagesize`
> the `getpagesize()` function is available ('yes' or 'no')

`ac_cv_func_getpwuid`
> the `getpwuid()` function is available ('yes' or 'no')

`ac_cv_func_getrusage`
> the `getrusage()` function is available ('yes' or 'no')

`ac_cv_func_getuid`
> the `getuid()` function is available ('yes' or 'no')

`ac_cv_func_glob`
> the `glob()` function is available ('yes' or 'no')

`ac_cv_func_ibv_get_device_list`
> the `ibv_get_device_list()` function is available ('yes' or 'no')

`ac_cv_func_ibv_get_devices`
> the `ibv_get_devices()` function is available ('yes' or 'no')

`ac_cv_func_kill`
> the `kill()` function is available ('yes' or 'no')

`ac_cv_func_llabs`
> the `llabs()` function is available ('yes' or 'no')

`ac_cv_func_nanosleep`
> the `nanosleep()` function is available ('yes' or 'no')

`ac_cv_func_nearbyint`
> the `nearbyint()` function is available ('yes' or 'no')

`ac_cv_func_popen`
> the `popen()` function is available ('yes' or 'no')

`ac_cv_func_realpath`
> the `realpath()` function is available ('yes' or 'no')

`ac_cv_func_rint`
> the `rint()` function is available ('yes' or 'no')

`ac_cv_func_round`
> the `round()` function is available ('yes' or 'no')

`ac_cv_func_sched_getaffinity`
> the `sched_getaffinity()` function is available ('yes' or 'no')

`ac_cv_func_setitimer`
> the `setitimer()` function is available ('yes' or 'no')

`ac_cv_func_sigaction`
> the `sigaction()` function is available ('yes' or 'no')

`ac_cv_func_signal`
> the `signal()` function is available ('yes' or 'no')

`ac_cv_func_sleep`
> the `sleep()` function is available ('yes' or 'no')

`ac_cv_func_strcasestr`
> the `strcasestr()` function is available ('yes' or 'no')

`ac_cv_func_strerror`
> the `strerror()` function is available ('yes' or 'no')

`ac_cv_func_strsignal`
   the `strsignal()` function is available ('yes' or 'no')

`ac_cv_func_strtok`
   the `strtok()` function is available ('yes' or 'no')

`ac_cv_func_strtoll`
   the `strtoll()` function is available ('yes' or 'no')

`ac_cv_func_strtoq`
   the `strtoq()` function is available ('yes' or 'no')

`ac_cv_func_sysconf`
   the `sysconf()` function is available ('yes' or 'no')

`ac_cv_func_sysctl`
   the `sysctl()` function is available ('yes' or 'no')

`ac_cv_func_sysmp`
   the `sysmp()` function is available ('yes' or 'no')

`ac_cv_func_syssgi`
   the `syssgi()` function is available ('yes' or 'no')

`ac_cv_func_time`
   the `time()` function is available ('yes' or 'no')

`ac_cv_func_trunc`
   the `trunc()` function is available ('yes' or 'no')

`ac_cv_func_uname`
   the `uname()` function is available ('yes' or 'no')

`ac_cv_func_vfork`
   the `vfork()` function is available ('yes' or 'no')

`ac_cv_func_vfork_works`
   the `vfork_works()` function is available ('yes' or 'no')

`ac_cv_func_vfprintf`
   the `vfprintf()` function is available ('yes' or 'no')

`ac_cv_func_vsnprintf`
   the `vsnprintf()` function is available ('yes' or 'no')

`ac_cv_func_vsprintf`
   the `vsprintf()` function is available ('yes' or 'no')

`ac_cv_header_bglpersonality_h`
   the `bglpersonality.h` header file is available ('yes' or 'no')

`ac_cv_header_cf_h`
   the `cf.h` header file is available ('yes' or 'no')

`ac_cv_header_common_bgp_personality_h`
   the `common/bgp_personality.h` header file is available ('yes' or 'no')

`ac_cv_header_dlfcn_h`
>   the `dlfcn.h` header file is available ('yes' or 'no')

`ac_cv_header_errno_h`
>   the `errno.h` header file is available ('yes' or 'no')

`ac_cv_header_glob_h`
>   the `glob.h` header file is available ('yes' or 'no')

`ac_cv_header_ieeefp_h`
>   the `ieeefp.h` header file is available ('yes' or 'no')

`ac_cv_header_infiniband_verbs_h`
>   the `infiniband/verbs.h` header file is available ('yes' or 'no')

`ac_cv_header_inttypes_h`
>   the `inttypes.h` header file is available ('yes' or 'no')

`ac_cv_header_invent_h`
>   the `invent.h` header file is available ('yes' or 'no')

`ac_cv_header_libhal_h`
>   the `libhal.h` header file is available ('yes' or 'no')

`ac_cv_header_machine_hal_sysinfo_h`
>   the `machine/hal_sysinfo.h` header file is available ('yes' or 'no')

`ac_cv_header_memory_h`
>   the `memory.h` header file is available ('yes' or 'no')

`ac_cv_header_mpi_h`
>   the `mpi.h` header file is available ('yes' or 'no')

`ac_cv_header_netdb_h`
>   the `netdb.h` header file is available ('yes' or 'no')

`ac_cv_header_odmi_h`
>   the `odmi.h` header file is available ('yes' or 'no')

`ac_cv_header_papi_h`
>   the `papi.h` header file is available ('yes' or 'no')

`ac_cv_header_pci_pci_h`
>   the `pci/pci.h` header file is available ('yes' or 'no')

`ac_cv_header_popt_h`
>   the `popt.h` header file is available ('yes' or 'no')

`ac_cv_header_pwd_h`
>   the `pwd.h` header file is available ('yes' or 'no')

`ac_cv_header_rms_rmscall_h`
>   the `rms/rmscall.h` header file is available ('yes' or 'no')

`ac_cv_header_sched_h`
>   the `sched.h` header file is available ('yes' or 'no')

`ac_cv_header_signal_h`
>   the `signal.h` header file is available ('yes' or 'no')

`ac_cv_header_stdint_h`
>   the `stdint.h` header file is available ('yes' or 'no')

`ac_cv_header_stdlib_h`
>   the `stdlib.h` header file is available ('yes' or 'no')

`ac_cv_header_string_h`
>   the `string.h` header file is available ('yes' or 'no')

`ac_cv_header_strings_h`
>   the `strings.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_cfgodm_h`
>   the `sys/cfgodm.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_mman_h`
>   the `sys/mman.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_param_h`
>   the `sys/param.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_resource_h`
>   the `sys/resource.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_select_h`
>   the `sys/select.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_stat_h`
>   the `sys/stat.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_sysctl_h`
>   the `sys/sysctl.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_sysinfo_h`
>   the `sys/sysinfo.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_sysmp_h`
>   the `sys/sysmp.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_syssgi_h`
>   the `sys/syssgi.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_time_h`
>   the `sys/time.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_types_h`
>   the `sys/types.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_utsname_h`
>   the `sys/utsname.h` header file is available ('yes' or 'no')

`ac_cv_header_sys_wait_h`
>   the `sys/wait.h` header file is available ('yes' or 'no')

`ac_cv_header_time_h`
> the `time.h` header file is available ('yes' or 'no')

`ac_cv_header_unistd_h`
> the `unistd.h` header file is available ('yes' or 'no')

`ac_cv_header_uuid_uuid_h`
> the `uuid/uuid.h` header file is available ('yes' or 'no')

`ac_cv_header_vfork_h`
> the `vfork.h` header file is available ('yes' or 'no')

`ac_cv_header_windows_h`
> the `windows.h` header file is available ('yes' or 'no')

`ac_cv_lib_ibverbs_ibv_query_device`
> `libibverbs.a` contains a function called `ibv_query_device()` ('yes' or 'no')

`ac_cv_lib_lipr_IprStart`
> `liblipr.a` contains a function called `IprStart()` ('yes' or 'no')

`ac_cv_lib_m_fmod`
> `libm.a` contains a function called `fmod()` ('yes' or 'no')

`ac_cv_lib_papi_PAPI_get_real_usec`
> `libpapi.a` contains a function called `PAPI_get_real_usec()` ('yes' or 'no')

`ac_cv_lib_perfctr_main`
> `libperfctr.a` contains a function called `main()` ('yes' or 'no')

`ac_cv_lib_popt_poptGetContext`
> `libpopt.a` contains a function called `poptGetContext()` ('yes' or 'no')

`ac_cv_lib_sysfs_sysfs_open_class`
> `libsysfs.a` contains a function called `sysfs_open_class()` ('yes' or 'no')

`ac_cv_lib_uuid_uuid_unparse`
> `libuuid.a` contains a function called `uuid_unparse()` ('yes' or 'no')

`ac_cv_member_struct_CuAt_name`
> the 'CuAt' struct contains a field called 'name' ('yes' or 'no')

`ac_cv_member_struct_addrinfo_ai_canonname`
> the 'addrinfo' struct contains a field called 'ai_canonname' ('yes' or 'no')

`ac_cv_member_struct_hostent_h_name`
> the 'hostent' struct contains a field called 'h_name' ('yes' or 'no')

`ac_cv_path_ASY`
> the fully qualified filename for `asy`

`ac_cv_path_DOT`
> the fully qualified filename for `dot`

`ac_cv_path_EGREP`
> the fully qualified filename for `egrep`

`ac_cv_path_FGREP`

        the fully qualified filename for `fgrep`

`ac_cv_path_GIMP`

        the fully qualified filename for `gimp`

`ac_cv_path_GPERF`

        the fully qualified filename for `gperf`

`ac_cv_path_GREP`

        the fully qualified filename for `grep`

`ac_cv_path_GS`

        the fully qualified filename for `gs`

`ac_cv_path_HTMLTIDY`

        the fully qualified filename for `htmltidy`

`ac_cv_path_JYTHONC`

        the fully qualified filename for `jythonc`

`ac_cv_path_LDD`

        the fully qualified filename for `ldd`

`ac_cv_path_PERL`

        the fully qualified filename for `perl`

`ac_cv_path_PYTHON`

        the fully qualified filename for `python`

`ac_cv_path_SED`

        the fully qualified filename for `sed`

`ac_cv_path_SWIG`

        the fully qualified filename for `swig`

`ac_cv_prog_AWK`

        command line for running `awk`

`ac_cv_prog_CPP`

        command line for running `cpp`

`ac_cv_prog_EMACS`

        command line for running `emacs`

`ac_cv_prog_INDENT`

        command line for running `indent`

`ac_cv_prog_JAR`

        command line for running `jar`

`ac_cv_prog_JAVAC`

        command line for running `javac`

`ac_cv_prog_LATEX`

        command line for running `latex`

`ac_cv_prog_MAKE`
> command line for running `make`

`ac_cv_prog_MPICC`
> command line for running `mpicc`

`ac_cv_prog_MSGDIGEST`
> command line for running `msgdigest`

`ac_cv_prog_PDFTEX`
> command line for running `pdftex`

`ac_cv_prog_PKGCONFIG`
> command line for running `pkgconfig`

`ac_cv_prog_TEX`
> command line for running `tex`

`ac_cv_search_MPI_Init`
> linker flags/libraries needed to use the `MPI_Init()` function or '`none required`'

`ac_cv_search_clock_gettime`
> linker flags/libraries needed to use the `clock_gettime()` function or '`none required`'

`ac_cv_search_kstat_data_lookup`
> linker flags/libraries needed to use the `kstat_data_lookup()` function or '`none required`'

`ac_cv_search_nanosleep`
> linker flags/libraries needed to use the `nanosleep()` function or '`none required`'

`ac_cv_sizeof_u_int128_t`
> size in bytes of the '`u_int128_t`' datatype

`ac_cv_sizeof_unsigned___int128`
> size in bytes of the '`unsigned __int128`' datatype

`ac_cv_sizeof_unsigned_long_long_long`
> size in bytes of the '`unsigned long long long`' datatype

`ac_cv_type_int64_t`
> the '`int64_t`' datatype exists ('`yes`' or '`no`')

`ac_cv_type_pid_t`
> the '`pid_t`' datatype exists ('`yes`' or '`no`')

`ac_cv_type_signal`
> the '`signal`' datatype exists ('`yes`' or '`no`')

`ac_cv_type_ssize_t`
> the '`ssize_t`' datatype exists ('`yes`' or '`no`')

`ac_cv_type_uid_t`
> the '`uid_t`' datatype exists ('`yes`' or '`no`')

`ac_cv_type_uint128_t`
> the 'uint128_t' datatype exists ('yes' or 'no')

`ac_cv_type_uint16_t`
> the 'uint16_t' datatype exists ('yes' or 'no')

`ac_cv_type_uint32_t`
> the 'uint32_t' datatype exists ('yes' or 'no')

`ac_cv_type_uint64_t`
> the 'uint64_t' datatype exists ('yes' or 'no')

`ac_cv_type_uint8_t`
> the 'uint8_t' datatype exists ('yes' or 'no')

`ac_cv_type_uintptr_t`
> the 'uintptr_t' datatype exists ('yes' or 'no')

`ax_cv_const_nics_valid`
> '$NICS' is a valid printf conversion specifier for 'ncptl_int'

`ax_cv_decl_double_smallest`
> a definition of the largest double value

`ax_cv_decl_suffix_valid`
> '$NCPTL_INT_SUFFIX' is a valid suffix for an 'ncptl_int' constant

`ax_cv_file_cray_xc_cname`
> node coordinates appear in a file

`ax_cv_file_cray_xt_nid`
> '_my_pnid' is definable from a file

`ax_cv_func_mpicc_requires_libm`
> '$MPICC' wants the math library to be linked explicitly

`ax_cv_func_round_works`
> 'round' is likely to work

`ax_cv_prog_cc_ansic`
> '$CC' supports some basic ANSI C features

`ax_cv_prog_eval_break_works`
> 'eval break' successfully breaks out of a loop

`ax_cv_prog_gperf_recent`
> gperf is sufficiently recent

`ax_cv_prog_jar_e`
> '$JAR' supports '-e'

`ax_cv_prog_ld_interp`
> we can store symbols in '.interp'

`ax_cv_prog_ld_so`
> the name of the dynamic linker

`ax_cv_prog_make_RM`
> '`$am_make`' predefines `\$(RM)`

`ax_cv_prog_make_RM_qeq`
> '`$am_make`' supports '`?=`'

`ax_cv_prog_perl_logextract`
> '`$PERL`' can run `ncptl-logextract`

`ax_cv_prog_swig_version`
> `swig` version number

`ax_cv_prog_swig_works`
> `swig` works

`ax_cv_sys_null_device`
> the null device is called

`ax_cv_type_MT_STATE_size`
> number of words of random-number-generator state

`ax_cv_type_long_big_enough`
> a long can hold both an '`int64_t`' and an '`ncptl_int`' type

`ax_cv_type_ncptl_int_max`
> we can find the maximum '`ncptl_int`' value

`ax_cv_type_ncptl_int_min`
> we can find the minimum '`ncptl_int`' value

`ax_cv_var___cpu_mhz`
> '`__cpu_mhz`' is defined

`ax_cv_var__my_pnid`
> '`_my_pnid`' is defined

Referenced Applications

A number of third-party applications are mentioned throughout this document. For your convenience, the following list states the URL of each application's home page.

a2ps   http://www.gnu.org/software/a2ps/

Dimemas   See Paraver.

dot   See Graphviz.

Emacs   http://www.gnu.org/software/emacs/

Extrae   See Paraver.

Environment Modules
            http://modules.sourceforge.net/

GeSHi   http://qbnz.com/highlighter/

Graphviz   http://www.graphviz.org/

LaTeX   http://www.latex-project.org/

LibSea   http://www.caida.org/tools/visualization/libsea/

Mediawiki   http://www.mediawiki.org/

MPICL   http://www.csm.ornl.gov/picl/

ParaGraph
            http://www.csar.uiuc.edu/software/paragraph/

Paraver   http://www.bsc.es/paraver

PICL   See MPICL.

pkg-config http://pkg-config.freedesktop.org/wiki/

PSTricks   http://www.tug.org/applications/PSTricks/

Python   http://www.python.org/

SLOCCount
            http://www.dwheeler.com/sloccount/

Source-highlight
            http://www.gnu.org/software/src-highlite/

TeX   See LaTeX.

Vim   http://www.vim.org/

Walrus   http://www.caida.org/tools/visualization/walrus/

License

Summary

This is a BSD license with the additional proviso that modified versions of CONCEPTUAL must indicate that they are, in fact, modified.

Index

C

G

H

I

N